arXiv:2403.02951v3 [cs.CL] 18 Mar 2026

SQLBench: A Comprehensive Evaluation for Text-to-SQL Capabilities of Large Language Models

Bin Zhang[1,2,*] **Yuxiao Ye**[5,*] **Guoqing Du**[3] **Xiaoru Hu**[3] **Zhishuai Li**[3] **Chi Harold Liu**[5]
Zhiwei Xu[4] **Guoliang Fan**[1,2] **Rui Zhao**[3] **Ziyue Li**[6] **Hangyu Mao**[7,✉]

[1]Institute of Automation, Chinese Academy of Sciences
[2]School of Artificial Intelligence, University of Chinese Academy of Sciences
[3]SenseTime Research
[4]School of Artificial Intelligence, Shandong University
[5]School of Computer Science and Technology, Beijing Institute of Technology
[6]Technical University of Munich, Heilbronn Data Science Center, Munich Data Science Institute
[7]Institute of Microelectronics, Chinese Academy of Sciences

Abstract

Large Language Models (LLMs) have emerged as a powerful tool in advancing the Text-to-SQL task, significantly outperforming traditional methods. Nevertheless, as a nascent research field, there is still no consensus on the optimal prompt templates and design frameworks. Additionally, existing benchmarks inadequately explore the performance of LLMs across the various sub-tasks of the Text-to-SQL process, which hinders the assessment of LLMs' cognitive capabilities and the optimization of LLM-based solutions. To address the aforementioned issues, we firstly construct a new dataset designed to mitigate the risk of overfitting in LLMs. Then we formulate five evaluation tasks to comprehensively assess the performance of diverse methods across various LLMs throughout the Text-to-SQL process. Our study highlights the performance disparities among LLMs and proposes optimal in-context learning solutions tailored to each task. These findings offer valuable insights for facilitating the development of LLM-based Text-to-SQL systems.

1 Introduction

Text-to-SQL automatically transforms natural language (NL) questions into structured SQL statements, e.g., based on the user's question "*How many musicians do we have?*" and the database schema (e.g., table/column names), the ideal Text-to-SQL module can generate SQL statements "`SELECT count(*) FROM musician`", which will then be executed on the SQL database engine to retrieve the correct answer. Text-to-SQL facilitates seamless user interaction with databases [32]. Previous approaches primarily focus on pattern matching between natural language and SQL statements, utilizing machine learning models to acquire the mapping between the two sequences [53, 20]. However, Large Language Models (LLMs) have brought about a substantial transformation in this field [9, 30] as powerful tools [35, 17, 38], showcasing tremendous potential in comprehending complex NL questions and generating accurate SQL statements. By combining advanced reasoning techniques and in-context learning capabilities, LLMs have significantly pushed the boundaries of the state-of-the-art in this domain, outperforming traditional methods by a considerable margin.

Despite the continuous improvement of LLM-based methods in the two popular datasets, i.e., Spider [49] and BIRD [22], there remains a critical gap in the systematic benchmarking of these

*These authors contributed equally.

Preprint. Under review.

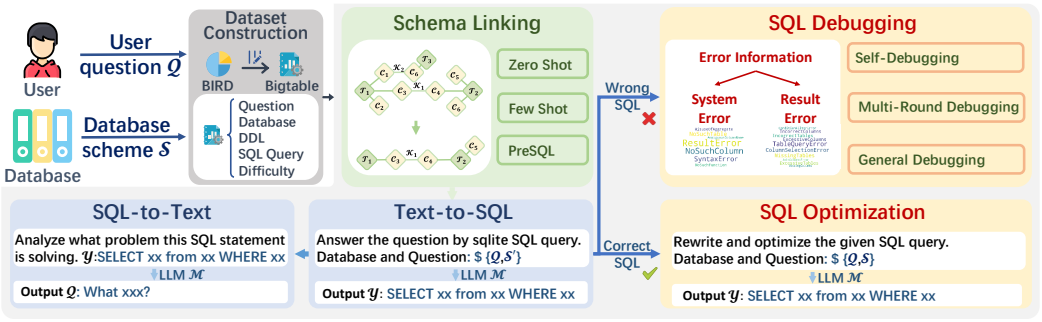

Figure 1: Benchmarking tasks within SQLBench, encompassing various Text-to-SQL pipeline stages. We first construct the Bigtable dataset. Subsequently, different types of LLMs, including general-purpose and coding-specific models, are employed to evaluate various subtasks. These subtasks include end-to-end Text-to-SQL, SQL-to-Text tasks and Schema Linking, which aligns entity references in the question with the corresponding schema tables or columns. For the questions that the LLMs fail to process correctly, a SQL Debugging task is designed with the aim of having the LLMs identify and correct the errors to generate the correct results. A SQL Optimization task is designed for the parts that the LLMs are able to handle correctly, with the goal of improving the execution efficiency of the SQL statements.

solutions [16, 32, 18]. **(1) Are the current solutions overfitting Spider and BIRD?** The risk of overfitting, particularly for LLMs trained on coding tasks and open-source datasets, poses a significant challenge to the generalizability [34]. **(2) What is the optimal prompt?** The optimal prompt engineering strategies play a crucial role in guiding LLMs to generate accurate SQL queries. Although various design schemes are explored [12], there is no consensus on the most effective prompt template. **(3) Is Text-to-SQL the only task that matters?** Furthermore, the current benchmarks [49, 22], while comprehensive in their assessment of end-to-end Text-to-SQL tasks, have not yet provided a detailed exploration of the models' performance across other various sub-tasks and components of the Text-to-SQL process [29, 26]. A detailed exploration of these sub-tasks is crucial for a thorough evaluation of LLMs' cognitive capabilities and their role in facilitating the Text-to-SQL process. Thus, it is essential to develop a more granular benchmarking approach that can accurately reflect the multifaceted nature of Text-to-SQL and inform more effective LLM-based solutions.

To address the aforementioned challenges and fill the gap in the systematic benchmarking of LLMs in Text-to-SQL, we construct a comprehensive testing benchmark called **SQLBench** that provides a holistic assessment of LLM capabilities in this domain. Our approach begins with the construction of a Text-to-SQL dataset, designed to mitigate the risk of overfitting by considering question complexities, database sizes, and prerequisite knowledge. It contains questions involving queries across varying numbers of tables. This dataset, constructed on the foundation of existing open-source data, complemented by existing open-source data, serves as a robust platform for evaluating LLMs across a spectrum of tasks. Furthermore, we devise five distinct tasks: Text-to-SQL, SQL Debugging, SQL Optimization, Schema Linking, and SQL-to-Text, to comprehensively evaluate LLMs' capabilities across the full spectrum of the Text-to-SQL process (see Figure 1). Subsequently, we perform an extensive analysis of various techniques that are essential for improving the in-context learning abilities of LLMs and their precision in generating SQL queries. Specifically, our contributions are summarized as follows:

- To determine the optimal prompt template, we partition the prompt text into distinct components and perform thorough testing of LLMs' performance on end-to-end Text-to-SQL tasks across all possible combination patterns.
- Our benchmarking approach encompasses a range of LLMs, including both general-purpose and coding-specific models with varying parameter sizes. We determine the performance boundaries of these models and identify their performance disparities (see Figure 2).
- For each task, we systematically assess the impact of information granularity on model performance and identify the optimal context learning strategies, such as zero-shot and few-shot, to maximize the performance of the models.

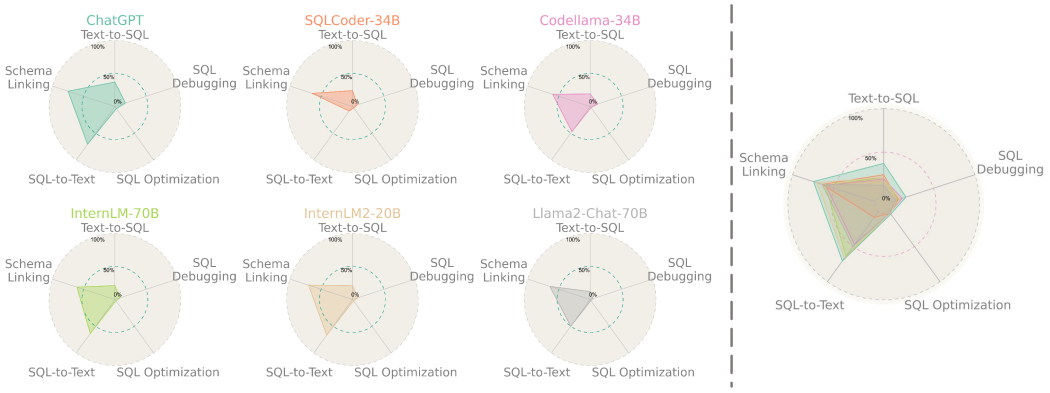

Figure 2: ***Left:*** Six LLMs' performance in five sub-tasks of SQLBench. ***Right:*** The summarized performance of all LLMs.

2 Related Work

Text-to-SQL Methods. Numerous traditional learning-based Text-to-SQL methods existed before the emergence of LLMs [53, 45, 27]. These methods fall into two categories based on their network architecture: non-seq2seq and seq2seq. Non-seq2seq methods [43, 1, 14] first utilize a relation-aware self-attention mechanism for encoding, and then use various decoders for generating SQLs, either as abstract syntax trees [48, 2] or via slot-filling [8, 15, 5]. Seq2seq methods [25, 31] directly translate NL questions into SQL queries through transformer-based [42] architectures in an end-to-end manner. These methods obtain competitive performance by fine-tuning rather than training from scratch, with less effort [33, 21, 20]. However, the highest accuracy achieved by these traditional methods on Spider is 79.9%, which is still far from being a reliable Text-to-SQL parser [16].

Given the powerful zero-shot reasoning and domain generalization abilities of LLMs, LLM-based Text-to-SQL methods successively refresh the record on Spider [23]. C3 [9], leveraging ChatGPT, achieved 82.3% accuracy by refining model inputs, biases, and outputs. DIN-SQL [30] decomposes the task into manageable sub-tasks, reaching 85.3% accuracy, while DAIL-SQL [12] achieves 86.6% through fine-tuning and in-context learning optimizations in few-shot settings. Similarly, other studies investigate the demo selection by synthesising in-domain examples [3] and retrieving question skeletons [13]. Nevertheless, these studies largely rely on OpenAI LLMs without testing other open-source models, nor do they systematically examine sub-task performance of Text-to-SQL.

Datasets and Evaluation Metrics. The rapid advancement of Text-to-SQL datasets has expanded their scope from single-domain [6, 50] to cross-domain [53, 19]. Most of the recent Text-to-SQL works are done on **Spider** [49], recognized as the most challenging cross-domain benchmark featuring multi-table queries, complex SQL clauses, and nested SQLs. Many variants of Spider are designed to test the adaptability and robustness of Text-to-SQL methods [7, 10, 11]. The newest seminal benchmark is **BIRD** [22]. It is the first to be curated for evaluating SQL execution efficiency in large-scale databases. In this paper, we construct a novel dataset built upon BIRD and use it in the evaluation (see Section 4).

We evaluate the accuracy of SQL queries using **Execution Accuracy (EX)**, which measures whether the result of executing the predicted query matches the gold value. Additionally, **Valid Efficiency Score (VES)**, an integrated metric that assesses both the accuracy and efficiency of SQL queries, is employed to evaluate SQL optimization in Section 5.3. Refer to Appendix C.2 for detailed definitions and discussions of evaluation metrics.

3 Settings

Task Formulation. In an LLM-based Text-to-SQL system, LLMs are employed to facilitate the transformation of natural language questions into executable SQL queries. Specifically, Let $\mathcal{Q}$ be a natural language question and $\mathcal{S}$ be the database schema. $\mathcal{S}$ is defined by a tuple $\mathcal{S} = (\mathcal{T}, \mathcal{C}, \mathcal{K})$,

where $\mathcal{T} = \{t_1, ..., t_m\}$ represents multiple tables, $\mathcal{C} = \{c_1, ..., c_n\}$ represents columns, and $\mathcal{K}$ represents foreign key relationships. The goal is to produce a SQL query $\mathcal{Y}$ such that $\mathcal{Y}$ is executable and accurately represents the intent of $\mathcal{Q}$. Given the prompt template $\mathcal{P}(\mathcal{Q}, \mathcal{S})$, the generation process of the SQL query $\mathcal{Y}$ by an LLM $\mathcal{M}$ can be formally defined as a conditional probability distribution:

$$\mathbb{P}_{\mathcal{M}}(\mathcal{Y}|\mathcal{P}(\mathcal{Q}, \mathcal{S})) = \prod_{i=1}^{|\mathcal{Y}|} \mathbb{P}_{\mathcal{M}}(\mathcal{Y}_i|\mathcal{P}(\mathcal{Q}, \mathcal{S}), \mathcal{Y}_{1:i-1}). \tag{1}$$

Here, LLM autoregressively generates each token, $\mathcal{Y}_i$ denotes the i-th token of the SQL query $\mathcal{Y}$, and $|\mathcal{Y}|$ denotes the length of the query $\mathcal{Y}$.

Evaluation Models. Our benchmarking study evaluates the performance of two distinct categories of LLMs with varying parameter sizes: general-purpose and coding-specific. General-purpose LLMs are designed for versatile text generation and comprehension across diverse domains, trained on extensive internet text datasets. Specifically, ChatGPT (gpt-35-turbo-16k), LLaMa2-Chat-70B [41], InternLM-70B [40] and InternLM2-20B are selected as the main baseline models. Coding-specific LLMs are fine-tuned and optimized for programming scenarios, excelling in code generation and technical language understanding. In this paper, the performance analysis of Codellama-34B [34] and SQLCoder-34B [39] are provided.

4 Dataset Construction

Table 1: EX (%) of different LLMs on open source datasets.

	SQLCoder-34B	InternLM-70B	Codellama-34B	LLama2-Chat-70B
Spider Dev	65.00	67.40	**71.60**	54.70
BIRD Dev	**32.07**	29.60	28.29	20.60

We conduct a preliminary assessment of the performance of various LLMs on multiple open-source datasets. As depicted in Table 1, the performance of LLMs varies inconsistently across different datasets. Specifically, on the Spider dataset, Codellama-34B outperforms InternLM-70B and SQLCoder-34B, while on the Bird dataset, SQLCoder-34B surpasses InternLM-70B and Codellama-34B. The performance differences between LLMs may stem from variations in the problem types they are adept at handling. Furthermore, as LLMs are trained on extensive datasets, the discrepancies could be attributed to the use of open-source datasets during fine-tuning of coding-specialized LLMs. This poses challenges in ensuring the reliability of evaluation results obtained on these datasets.

To address the potential overfitting of LLMs, particularly those specialized in coding tasks, and to ensure a reliable and accurate assessment of their capabilities, we construct a novel dataset, termed "BigTable". This dataset is an extension and augmentation of the BIRD dataset. In the process of

Table 2: The data distribution of "BigTable" includes the average number of ground truth (GT) tables and columns involved in the instances.

No. of GT Tables	No. of Instances	Avg. No. of Columns
1	50	23.30
2	50	56.94
3	50	31.73
>3	50	29.04

modifying and expanding the original queries, our primary goal was to introduce a diverse set of linguistic and structural variations that would challenge the model's ability to generalize beyond simple lookup-based responses. Specifically, our construction process involves a systematic analysis of the original BIRD dataset, identifying queries of varying difficulty levels and involving different numbers of tables (1, 2, 3, and more than 3). We modify and expand these queries by altering table and column names, as well as filtering conditions, to create a more diverse set of challenges. In cases where the original dataset lacks sufficient examples with four or more tables (there are only 20 instances in BIRD-Dev dataset), queries that involved three tables are expanded to four. As shown in Table 2, this process generates 50 new instances for each category, resulting in the "BigTable" dataset.

Moreover, each item in the dataset underwent mutual verification by at least two individuals to ensure the accuracy.

5 Evaluation

In this section, we formally evaluate the different sub-tasks of SQLBench within the Text-to-SQL process to determine the performance differences among various LLMs and provide recommendations for addressing specific task requirements. The relationships among the sub-tasks tested are illustrated in Figure 3. The code is available at `https://anonymous.4open.science/r/SQLBench-E83A`.

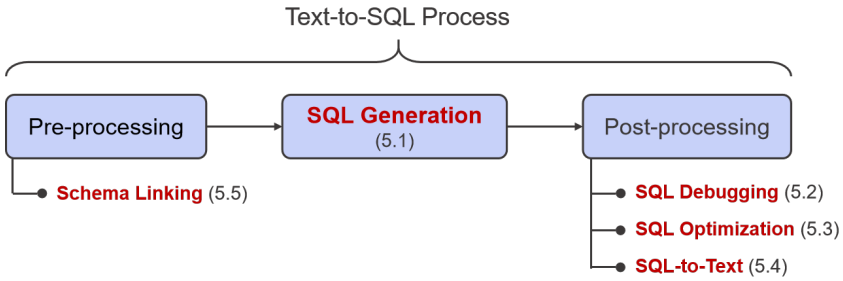

Figure 3: The paradigm for LLM-based Text-to-SQL and the tasks evaluated in SQLBench.

5.1 SQL Generation

Table 3: EX (%) of Text-to-SQL with various zero-shot prompt templates on Spider dev set.

Prompt Template	SQLCoder-34B	Codellama-34B	InternLM-70B	Llama2-Chat-70B
DDL-HTML-Chat	57.8	63.7	**65.0**	49.6
DDL-HTML-Complete	61.8	**65.2**	53.8	50.2
DDL-MD-Chat	63.2	**68.4**	66.3	48.7
DDL-MD-Complete	62.4	**69.8**	64.1	46.8
DDL-Coding-Chat	60.3	**67.1**	66.1	48.4
DDL-Coding-Complete	59.7	**66.9**	62.9	53.4
SimpleDDL-MD-Chat	**65.0**	**71.6**	**67.4**	**54.7**
SimpleDDL-MD-Complete	63.3	**66.0**	61.7	50.2

Unlike previous learning-based studies, the primary challenge in LLM-based Text-to-SQL is the design of an effective prompt template $\mathcal{P}$ (as introduced in Section 3) for LLMs to generate accurate SQL queries, known as prompt engineering. Researchers have evaluated a variety of prompt templates [12]. However, these representations lack uniformity in their structure, making it difficult to find out how a specific feature within a prompt template impacts performance.

To address this issue, we investigate a more unified series of prompt templates. As shown in Listing 1 (a)-(c), these templates differ across three features:

- **DDL/SimpleDDL prefix** affects the representation of the database schema $\mathcal{S}$. "DDL" (Data Definition Language) encompasses the standardized language that includes commands for defining the structure and properties of a database. "SimpleDDL" is simplified by only supplying table and column names.
- **MD/HTML/Coding infix** wraps the entire prompt template with Markdown syntax, HTML snippets and code comment blocks.
- **Complete/Chat postfix** indicates the task of either completing SQL statements based on the "SELECT" clause or directly answering questions.

These features are combined to form a complete prompt template $\mathcal{P}$, and more details of these representations can be found in Appendix D.1. We test these templates on Spider dev set. As shown in Table 3, "SimpleDDL-MD-Chat" in Listing 1 (d) consistently outperforms all other prompts when applied to all 5 backbone LLMs. To this end, we consistently utilize the prompt template "SimpleDDL-MD-Chat" throughout the subsequent evaluations in this paper. A detailed evaluation of "BigTable" can be found in Appendix A.2.1

```
DDL Format:
CREATE TABLE stadium (stadium_id NUMBER PRIMARY KEY, location TEXT,
↳ name TEXT, capacity NUMBER, highest NUMBER, lowest NUMBER, average
↳ NUMBER);
SimpleDDL Format:
stadium(Stadium_ID,Location,Name,Capacity,Highest,Lowest,Average);
```

(a) Examples of "DDL" and "SimpleDDL" formats.

```
MD Format:
### How many singers do we have?
HTML Format:
<Question>How many singers do we have?</Question>
Coding Format:
/* How many singers do we have? */
```

(b) Examples of "MD", "HTML" and "Coding" formats.

```
Complete Format:
Question: How many singers do we have?
SQL: SELECT
Chat Format:
Question: How many singers do we have?
SQL:
```

(c) Examples of "Complete" and "Chat" formats.

MD format

```
### Answer the question by sqlite SQL query only and with no
↳ explanation
### Sqlite SQL tables, with their properties:          SimpleDDL
#
# stadium(Stadium_ID,Location,Name,Capacity,Highest,Lowest,Average);
# singer(Singer_ID,Name,Country,Song_Name,Song_release_year,Age,
↳ Is_male);
# concert(concert_ID,concert_Name,Theme,Stadium_ID,Year);
# singer_in_concert(concert_ID,Singer_ID).
#
### How many singers do we have?                       Chat
### SQL:
```

(d) Proposed prompt template "SimpleDDL-MD-Chat".

Listing 1: (a)-(c): different options for three format features; (d) the proposed optimal prompt templates.

Core Conclusion 1. *The prompt template "SimpleDDL-MD-Chat" achieves optimal performance in Text-to-SQL.*

5.2 SQL Debugging

In recent studies, researchers have demonstrated that LLMs possess self-reflection and self-correction capabilities similar to those of humans [47, 36, 51]. Additionally, previous studies also investigate the potential of LLMs to debug the code they generate [4, 30]. In this section, we provide a comprehensive analysis of the performance of numerous SQL debugging methods across LLMs.

Debugging Dataset. We collect incorrect SQL queries generated by various LLMs in "BigTable" and visualize the distribution of their error information, as shown in Figure 3. The error information can be divided into two categories: (1) **System Error** refers to syntax errors in the SQL statement, and the detailed system error information is generated by the Database Management System (DBMS), e.g., "syntax error" and "no such column"; (2) **Result Error** indicates that the syntax of the SQL statement is correct, but the execution result does not match the ground truth.

The word cloud distribution reveals that "No Such Column" and Result Error are the primary areas of error concentration for all models. However, the concise nature of the error information in the Result Error category significantly hampers the debugging performance. To this end, we categorize these errors based on the logical construction of SQL statements. As shown in Figure 4, this classification is prioritized according to the logical structure within the SQL query. It is delineated in order into the following five subcategories: (1) **Table Query Error** pertains to issues related to the selection of tables in the SQL query. It is further subdivided into three types: Excessive/Missing/Incorrect Tables, which respectively address scenarios where unnecessary tables are included, required tables are omit-

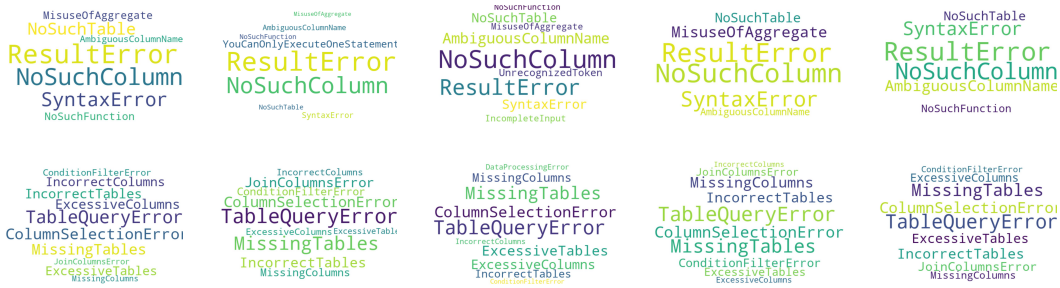

(a) SQLCoder (b) Codellama (c) LLama2 (d) InternLM (e) InternLM2

Figure 3: Word cloud representation of error information for incorrect SQL queries generated by LLMs. ***Top***: System Error and Result Error. ***Bottom***: Detailed classification of Result Error.

ted, or the wrong tables are referenced; (2) **Column Selection Error** focuses on the appropriateness of column selection. Similar to the Table Query Error, it is broken down into Excessive/Missing/Incorrect Columns; (3) **Join Columns Error** examines the errors associated with JOIN operations; (4) **Condition Filter Error** encompasses errors that occur in the conditions used to filter the data, including incorrect comparisons or misapplied filters; (5) **Data Processing Error** pertains to errors in the data processing stage, which includes aggregations, calculations, etc. applied to the data within the SQL query.

Debug Evaluation. To assess the impact of different levels of information granularity on performance, we propose 5 self-debugging strategies, progressively incorporating more details:

- **Regenerate**: Simply regenerate the SQL query with the same prompt in Section 5.1.
- **w/ Wrong SQL**: Let LLMs generate a new SQL query based on the wrong SQL statement.
- **w/ Wrong SQL + System_error_info**: Provide the wrong SQL statement, the corresponding System Error information and the rough Result Error information.
- **w/ Wrong SQL + All_error_info**: Add detailed Result Error information for those SQL queries that are syntactically correct but semantically wrong.
- **w/ Wrong SQL + All_error_info + Comment**: Add manual annotations for all error information. See Appendix D.2 for a detailed prompt template.

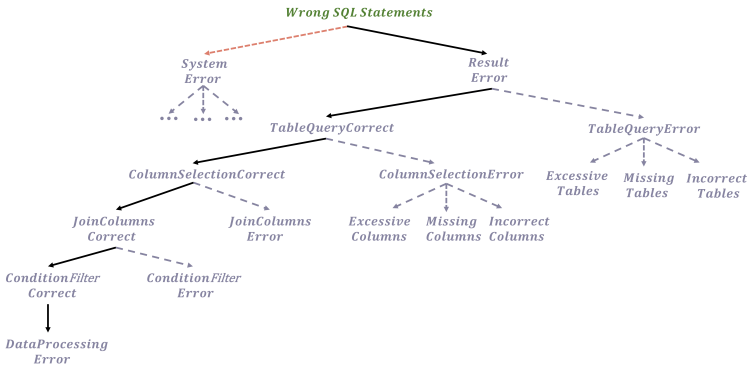

Figure 4: Process for determining error types. System Errors are directly got from the feedback of the SQL compiler. The five error types of Result Error are assessed sequentially based on their logical structure. When evaluating for Data Processing Errors, for example, it indicates that Table Query, Column Selection, and Join Column aspects have been correctly addressed.

What is the most powerful information organization of self debug? As shown in Figure 5, it is evident that the self-debugging performance of LLMs exhibits an upward trend with the introduction of more granular error information. In the absence of additional information, LLM does not possess the capability to regenerate correct answers. However, all models are able to comprehend fine-grained error information, whether it includes comments or not, and rectify their own mistakes.

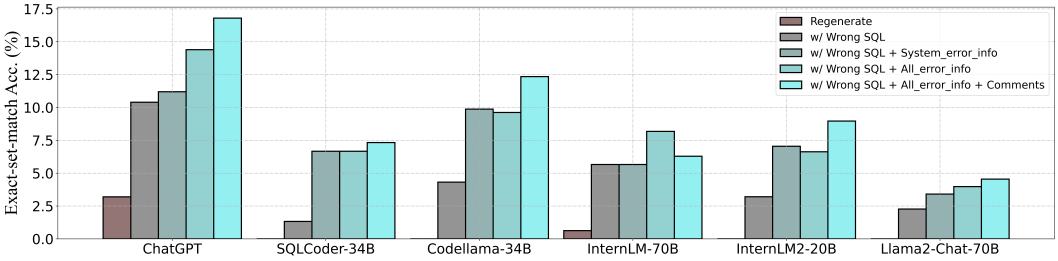

Figure 5: EX (%) improvement brought by self debug.

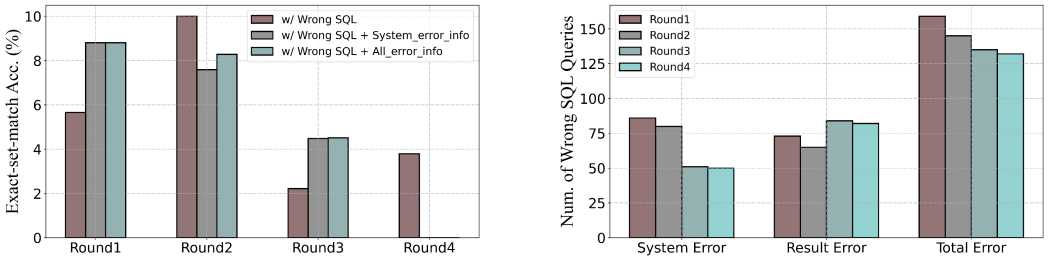

Figure 6: ***Left***: EX (%) improvement brought by multi-round self debug. ***Right***: Numbers of wrong SQL queries of detailed error type during the process of multi-round self debug.

Core Conclusion 2. *Detailed error information and corresponding annotations greatly enhance the capabilities of LLMs, enabling them to effectively correct errors.*

Can LLMs benefit from multi-round self debug? As shown in Figure 6, substantial improvements in EX are achieved in the initial rounds of debugging, yet the performance gain become marginal later. This indicates that conducting 1-2 debugging rounds might strikes a favorable balance between performance improvement and economic efficiency. In addition, we analyze the distribution of detailed error types during the multi-round self-debugging process. As the debugging rounds advance, we observe a reduction in System Error and a slight rise in Result Error. This suggests that although syntax errors are fixed, the regenerated SQL statements may still have semantic errors. With each round of debugging, the generated statements tend to transition from System Error to Result Error before eventually converging towards correct SQL statements.

Core Conclusion 3. *Multi-round self-debugging aids in error correction for LLMs, but there exists a performance boundary, with 1-2 rounds of debugging being the optimal choice.*

Table 4: EX (%) results of general debug. Each column demonstrates the debugging results of SQLCoder or InternLM for the wrong SQL statements generated by the corresponding model.

Dataset $\longrightarrow$	**SQLCoder-34B**	**Codellama-34B**	**InternLM-70B**	**Llama2-Chat-70B**
SQLCoder Debugs All Errors				
Regenerate	**0.00**	12.42	13.21	18.75
w/ Wrong SQL	**0.66**	2.48	6.92	7.95
w/ Wrong SQL + System_error_info	**4.64**	11.18	9.43	11.36
w/ Wrong SQL + All_error_info	**4.64**	11.80	8.81	10.80
InternLM Debugs All Errors				
Regenerate	9.27	11.80	**0.63**	15.34
w/ Wrong SQL	7.95	10.56	**5.66**	9.66
w/ Wrong SQL + System_error_info	10.60	10.56	**8.81**	11.93
w/ Wrong SQL + All_error_info	11.92	11.80	**8.81**	13.07

Can an LLM debug the error incurred by other LLMs (general debugging)? As depicted in Table 4, we select two representative LLMs, SQLCoder and InternLM, to assess their general debugging capabilities. When debugging erroneous SQL statements generated by other LLMs, additional error

information can potentially impair the performance. Conversely, the most naive approach of simply regenerating the SQL statements often yields the best results. This highlights the differences between different LLMs. The debugger LLMs may not encounter errors caused by other LLMs, and these error type informations might confuse them. Consequently, it is unlikely to achieve performance improvement through general debugging. However, the integration of results generated by different LLMs holds promise as a future research direction.

Core Conclusion 4. *The performance of cross-LLM SQL debugging is inferior to the direct regeneration. A multi-agent approach that integrates outputs from different models shows great potential.*

5.3 SQL Optimization

Execution efficiency of SQL queries is a critical aspect, particularly in real-time systems that utilize large-scale databases. In this section, we further explore whether LLMs are able to enhance the execution efficiency of correct SQL queries.

VES (see Equation 8) is commonly employed to evaluate the efficiency of SQL query execution. However, in practice, LLMs can sometimes rewrite a correct SQL query into an incorrect one, making it challenging to figure out if the main reason for the decline in VES is due to these incorrect rewrites or a decrease in the SQL execution efficiency. To this end, we suggest adopting a complementary metric C-VES (Correct-VES), which only assesses the execution efficiency of correct SQL queries:

$$\text{C-VES} = \frac{\sum_{n \in \mathcal{N}_c} R(\mathcal{Y}_n^{\text{o}}, \hat{\mathcal{Y}}_n)}{|\mathcal{N}_c|}, \tag{2}$$

$$\text{where } R(\mathcal{Y}_n^{\text{o}}, \hat{\mathcal{Y}}_n) = \sqrt{\frac{\mathbf{E}(\mathcal{Y}_n^{\text{o}})}{\mathbf{E}(\hat{\mathcal{Y}}_n)}},\ \mathcal{N}_c \doteq \{n | \mathbb{1}(V_n, \hat{V}_n) = 1\}. \tag{3}$$

Here, $\mathcal{Y}_n^{\text{o}}$ and $\hat{\mathcal{Y}}_n$ denote the optimized predicted SQL and GT SQL for the n-th instance, respectively. $\mathbf{E}(\cdot)$ serves as a function that quantifies the absolute execution efficiency (e.g., the reciprocal of execution time) of each SQL within a specified environment. $\mathcal{N}_c$ represent the set of accurate SQLs (see Appendix C.2 for detailed notations). C-VES is designed exclusively to validate the capability of LLMs to generate more efficient SQL queries, regardless of the potential drawback of rewriting correct SQLs into erroneous ones.

Do LLMs have the capability for SQL self-optimization? To the best of our knowledge, we are the first to consider utilizing LLMs for SQL optimization. Specifically, we devise an extensive suite of prompts $\mathcal{P}_{\text{o}}$ curated to SQL optimization:

- **w/** $\mathcal{Y}$: In this basic form, only original SQL statements are provided.
- **w/** $\mathcal{Y} + \mathcal{S} + \mathcal{Q}$: Further incorporates the database schema $\mathcal{S}$ and the user question $\mathcal{Q}$.
- **w/ demo**: Introduce few-shot demonstrations without explanations.
- **w/ demo + comments**: Add an explanation for the few-shot demonstrations.

Table 5: VES and C-VES results of different SQL optimization methods.

Methods	Prompt Template		Metrics	ChatGPT	SQLCoder-34B	Codellama-34B	InternLM-70B	InternLM2-20B
Baseline	**SimpleDDL-MD-Chat**		**VES**	36.90	24.28	18.94	22.63	19.81
			C-VES	102.50	101.17	99.68	110.39	94.33
Two-Stage Generation	**zero-shot**	**with** $\mathcal{Y}$	**VES**	30.73	16.86	17.34	15.33	20.28
			C-VES	102.43	102.19	102.08	102.17	101.42
		w/ $\mathcal{Y} + \mathcal{S} + \mathcal{Q}$	**VES**	32.21	18.90	19.28	18.44	20.77
			C-VES	102.24	102.14	101.47	102.43	101.30
	few-shot	**w/ demo**	**VES**	32.19	18.84	18.34	17.97	20.86
			C-VES	102.18	101.84	101.88	102.68	101.75
		w/ demo + comments	**VES**	32.65	18.28	18.52	17.45	20.92
			C-VES	102.03	101.54	102.86	102.66	102.06
Direct Generation	**SimpleDDL-MD-Chat-Efficiency**		**VES**	**39.26**	**27.77**	**20.75**	**25.23**	**25.93**
			C-VES	**103.31**	**102.84**	**103.75**	**102.98**	**101.70**

These methods focus on the second stage optimization of the SQL queries that have been correctly generated by LLMs during the Text-to-SQL stage. In addition, we also evaluate a direct genaration method **SimpleDDL-MD-Chat-Efficiency**: Require LLMs to directly generate the most efficient SQL query statement based on user query. The effectiveness of these SQL optimization methods

is shown in Table 5. Almost all two-stage methods experience a significant decrease in VES. It can be attributed to the possibility of LLMs optimizing the correct SQL statements into incorrect ones, resulting in a further decrease in accuracy. Even when considering only the correct results, the performance improvement in terms of execution efficiency brought by the optimized SQL statements is almost negligible.

Core Conclusion 5. *In-context learning methods present challenges in achieving effective SQL optimization with LLMs.*

5.4 SQL-to-Text

The goal of SQL-to-Text is to transform the SQL query back into its original natural language question [44, 28, 37]. It can assist researchers in evaluating the semantic comprehension capabilities of different LLMs, thus facilitating the development of more effective Text-to-SQL methodologies.

Table 6: SQL-to-Text performance of different LLMs, including the F1 scores of Rouge and BertScore, as well as the accuracy rate assessed by LLM.

	ChatGPT	SQLCoder-34B	Codellama-34B	InternLM-70B	InternLM2-20B	Llama2-Chat-70B
Rouge-1	**0.428**	0.118	0.359	0.408	0.410	0.368
Rouge-2	**0.384**	0.039	0.167	0.198	0.294	0.159
Rouge-L	**0.545**	0.108	0.320	0.372	0.358	0.325
BertScore	**0.887**	0.779	0.856	0.886	**0.887**	0.885
ChatGPT Evaluator	**70.5%**	8.5%	47.0%	63.5%	66.5%	50.0%
InternLM2 Evaluator	62%	22.5%	49.8%	65.8%	**71.3%**	60.0%

We assess the performance of SQL-to-Text across different LLMs (See Appendix D.4 for prompt templates). The selected metrics for evaluation encompass the F1 values of Rouge-1/2/L and BertScore, along with the application of LLM to assess the semantic coherence between the two texts. The evaluation results are depicted in Table 6. ChatGPT and InternLM2 demonstrate the highest performance, followed by InternLM, while Codellama and SQLCoder exhibit comparatively lower performance. This highlights that even in regards to semantic description of code, general-purpose models exhibit significantly stronger descriptive capabilities compared to coding-specific models.

Core Conclusion 6. *Utilizing a general-purpose model for semantic description of SQL statements is a better choice.*

5.5 Schema Linking

Schema linking is recognized as a crucial prerequisite of generating correct SQL queries [46]. It involves aligning entity references in the question with the corresponding schema tables or columns, requiring the model to understand both structure and value of the database, as well as the the semantics of user questions. In LLM-based Text-to-SQL, prior studies [9, 30] design prompt instructions with in-context learning examples to enable LLMs to retrieve linked tables and columns, which are then used for the downstream Text-to-SQL task. However, there is currently no comprehensive evaluation of the performance of multiple diverse methods across various backbone LLMs. In this section, we aim to bridge these gaps by: (1) introducing a elaborately designed metrics to assess schema linking methods; (2) presenting a novel schema linking method "PreSQL", which demonstrates superior performance; (3) conducting a comprehensive evaluation of a range of schema linking methods across various LLMs.

Table 7: RES results of different Schema Linking methods.

	ChatGPT		SQLCoder-34B		InternLM-70B		Codellama-34B		InternLM2-20B		Llama2-Chat-70B	
	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk
Zero Shot	0.6384	0.6399	0.2278	0.4686	0.5745	0.5652	0.3675	0.4835	0.5687	0.5811	0.4759	0.5566
Few Shot	0.6222	0.6402	0.3657	0.3919	0.5302	0.4961	0.4401	0.4745	0.464	0.4829	0.4123	0.5375
PreSQL	0.6888	0.6610	**0.5661**	**0.6417**	0.4632	0.4881	0.5085	**0.597**	0.5649	0.6267	0.4478	0.5273
Few Shot + PreSQL	**0.7340**	**0.7469**	0.5354	0.6336	**0.5999**	**0.6054**	**0.5576**	0.5646	**0.6594**	**0.7016**	**0.6105**	**0.6418**

Evaluation Metric for Schema Linking. The objective of schema linking is to retrieve all GT tables while avoiding the retrieval of excessive tables (with minimal redundancy). To evaluate this, we

design a comprehensive metric called **Retrieval Efficiency Score (RES)**:

$$\text{RES} = \frac{\sum_{n=1}^{N} \mathbb{1}(T_n, \hat{T}_n) \cdot R(T_n, \hat{T}_n)}{N}, \tag{4}$$

$$\text{where} \quad \mathbb{1}(T_n, \hat{T}_n) = \begin{cases} 1, \text{ if } T_n \subseteq \hat{T}_n \\ 0, \text{ if } T_n \nsubseteq \hat{T}_n \end{cases}, \quad R(T_n, \hat{T}_n) = \sqrt{\frac{|T_n|}{|\hat{T}_n|}}, \tag{5}$$

where T_n and $\hat{T}_n$ denote the set of GT tables and retrieved tables for the n-th instance, respectively. $|\cdot|$ refers to the scale of a set.

Schema Linking Evaluation. We evaluate the most representative **Zero Shot** and **Few Shot** methods, which are proposed by C3 [9] and DIN-SQL [30] respectively. In addition, a **PreSQL** method is introduced. First, we employ the zero-shot Text-to-SQL described in Section 5.1 to generate a preliminary SQL query. From this preliminary SQL, table and column entities are parsed to serve as the retrieval results. Moreover, **Few Shot + PreSQL** takes the union of the retrieval results from both the Few Shot and PreSQL methods, aiming to leverage the strengths of each.

Note that we organize the information of the database schema in a way similar to "SimpleDDL-" in the prompt of all the methods mentioned above, which ignores the information about foreign keys. However, as argued in [43], foreign keys embody features of known schema relations and provide important clues for understanding the database structure. To this end, we conduct experiments under both settings, w/ and w/o foreign keys, to investigate how incorporating foreign keys in the prompt influences the performance of schema linking. Results are demonstrated in Table 7 (refer to Appendix A for results on detailed metrics like Exact Match & Subset Match).

Which method achieved the best performance in schema linking? It can be seen that code-specific models excel in performance when utilizing the PreSQL approach, whereas general-purpose models yield optimal results through the Few-Shot + PreSQL method. This aligns with our expectations, as these two types of models excel in coding tasks (see Section 5.1) and semantic understanding tasks (see Section 5.4), respectively.

Can foreign key information facilitate schema linking? The introduction of foreign key information yield improved performance across all methods and all LLMs. This is evident since a valid JOIN operation in SQL queries is typically based on foreign keys. The foreign key information helps the model retrieve more ground truth tables by indicating all potential table pairs involved in a JOIN operation.

Core Conclusion 7. *Foreign key information is capable of advance the performance of schema linking. PreSQL yields the highest performance on coding-specific models, and integrating the results from Few Shot can further enhance performance on general-purpose models.*

6 Conclusion and Future Work

In this study, we conduct a systematic benchmarking of the various sub-tasks within the Text-to-SQL pipeline, encompassing Text-to-SQL, SQL-Debugging, SQL-Optimization, SQL-to-Text, and Schema-Linking. Our comprehensive evaluation involves six distinct LLMs, spanning both general-purpose and coding-specific models. We focus on determining the optimal prompt templates for each task, assessing performance variations among different approaches, and identifying the distinct capabilities and limitations of each LLM. The results of the study demonstrate notable performance variations across the LLMs, underscoring the significance of careful model selection and prompt engineering in attaining optimal outcomes in text-to-SQL tasks. Our benchmarking provides a meticulous perspective on the pipeline, equipping the research community with strategies to improve the semantic understanding and computational performance of LLMs. This advancement contributes to the development of more reliable Text-to-SQL systems.

As LLMs have seen rapid advancements, their performance continues to be quickly surpassed. It is therefore likely that even more powerful models could achieve superior results on our tasks. However, our work focuses on examining the general insights regarding LLM-based Text-to-SQL systems, such as identifying optimal prompt templates and characterizing the performance differences between code-specific and general-purpose models. These common findings provide valuable guidance for further improving the capabilities of this important class of language-based data management tools.

References

[1] Ruichu Cai, Jinjie Yuan, Boyan Xu, and Zhifeng Hao. Sadga: Structure-aware dual graph aggregation network for text-to-sql. *Advances in Neural Information Processing Systems*, 34:7664–7676, 2021.

[2] Ruisheng Cao, Lu Chen, Zhi Chen, Yanbin Zhao, Su Zhu, and Kai Yu. Lgesql: line graph enhanced text-to-sql model with mixed local and non-local relations. *arXiv preprint arXiv:2106.01093*, 2021.

[3] Shuaichen Chang and Eric Fosler-Lussier. Selective demonstrations for cross-domain text-to-sql. *arXiv preprint arXiv:2310.06302*, 2023.

[4] Xinyun Chen, Maxwell Lin, Nathanael Schärli, and Denny Zhou. Teaching large language models to self-debug. *arXiv preprint arXiv:2304.05128*, 2023.

[5] DongHyun Choi, Myeong Cheol Shin, EungGyun Kim, and Dong Ryeol Shin. Ryansql: Recursively applying sketch-based slot fillings for complex text-to-sql in cross-domain databases. *Computational Linguistics*, 47(2):309–332, 2021.

[6] Deborah A Dahl, Madeleine Bates, Michael K Brown, William M Fisher, Kate Hunicke-Smith, David S Pallett, Christine Pao, Alexander Rudnicky, and Elizabeth Shriberg. Expanding the scope of the atis task: The atis-3 corpus. In *Human Language Technology: Proceedings of a Workshop held at Plainsboro, New Jersey, March 8-11, 1994*, 1994.

[7] Xiang Deng, Ahmed Hassan Awadallah, Christopher Meek, Oleksandr Polozov, Huan Sun, and Matthew Richardson. Structure-grounded pretraining for text-to-sql. *arXiv preprint arXiv:2010.12773*, 2020.

[8] Li Dong and Mirella Lapata. Coarse-to-fine decoding for neural semantic parsing. *arXiv preprint arXiv:1805.04793*, 2018.

[9] Xuemei Dong, Chao Zhang, Yuhang Ge, Yuren Mao, Yunjun Gao, Jinshu Lin, Dongfang Lou, et al. C3: Zero-shot text-to-sql with chatgpt. *arXiv preprint arXiv:2307.07306*, 2023.

[10] Yujian Gan, Xinyun Chen, Qiuping Huang, Matthew Purver, John R Woodward, Jinxia Xie, and Pengsheng Huang. Towards robustness of text-to-sql models against synonym substitution. *arXiv preprint arXiv:2106.01065*, 2021.

[11] Yujian Gan, Xinyun Chen, and Matthew Purver. Exploring underexplored limitations of cross-domain text-to-sql generalization. *arXiv preprint arXiv:2109.05157*, 2021.

[12] Dawei Gao, Haibin Wang, Yaliang Li, Xiuyu Sun, Yichen Qian, Bolin Ding, and Jingren Zhou. Text-to-sql empowered by large language models: A benchmark evaluation. *arXiv preprint arXiv:2308.15363*, 2023.

[13] Chunxi Guo, Zhiliang Tian, Jintao Tang, Pancheng Wang, Zhihua Wen, Kang Yang, and Ting Wang. Prompting gpt-3.5 for text-to-sql with de-semanticization and skeleton retrieval. In *Pacific Rim International Conference on Artificial Intelligence*, pages 262–274. Springer, 2023.

[14] Binyuan Hui, Ruiying Geng, Lihan Wang, Bowen Qin, Bowen Li, Jian Sun, and Yongbin Li. S^2sql: Injecting syntax to question-schema interaction graph encoder for text-to-sql parsers. *arXiv preprint arXiv:2203.06958*, 2022.

[15] Wonseok Hwang, Jinyeong Yim, Seunghyun Park, and Minjoon Seo. A comprehensive exploration on wikisql with table-aware word contextualization. *arXiv preprint arXiv:1902.01069*, 2019.

[16] George Katsogiannis-Meimarakis and Georgia Koutrika. A survey on deep learning approaches for text-to-sql. *The VLDB Journal*, pages 1–32, 2023.

[17] Yilun Kong, Jingqing Ruan, Yihong Chen, Bin Zhang, Tianpeng Bao, Shiwei Shi, Guoqing Du, Xiaoru Hu, Hangyu Mao, Ziyue Li, et al. Tptu-v2: Boosting task planning and tool usage of large language model-based agents in real-world systems. *arXiv preprint arXiv:2311.11315*, 2023.

[18] Ayush Kumar, Parth Nagarkar, Prabhav Nalhe, and Sanjeev Vijayakumar. Deep learning driven natural languages text to sql query conversion: A survey. *arXiv preprint arXiv:2208.04415*, 2022.

[19] Chia-Hsuan Lee, Oleksandr Polozov, and Matthew Richardson. Kaggledbqa: Realistic evaluation of text-to-sql parsers. *arXiv preprint arXiv:2106.11455*, 2021.

[20] Haoyang Li, Jing Zhang, Cuiping Li, and Hong Chen. Resdsql: Decoupling schema linking and skeleton parsing for text-to-sql. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pages 13067–13075, 2023.

[21] Jinyang Li, Binyuan Hui, Reynold Cheng, Bowen Qin, Chenhao Ma, Nan Huo, Fei Huang, Wenyu Du, Luo Si, and Yongbin Li. Graphix-t5: Mixing pre-trained transformers with graph-aware layers for text-to-sql parsing. *arXiv preprint arXiv:2301.07507*, 2023.

[22] Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. *Advances in Neural Information Processing Systems*, 36, 2024.

[23] Zhishuai Li, Xiang Wang, Jingjing Zhao, Sun Yang, Guoqing Du, Xiaoru Hu, Bin Zhang, Yuxiao Ye, Ziyue Li, Rui Zhao, et al. Pet-sql: A prompt-enhanced two-stage text-to-sql framework with cross-consistency. *arXiv preprint arXiv:2403.09732*, 2024.

[24] Chin-Yew Lin. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*, pages 74–81, 2004.

[25] Xi Victoria Lin, Richard Socher, and Caiming Xiong. Bridging textual and tabular data for cross-domain text-to-sql semantic parsing. *arXiv preprint arXiv:2012.12627*, 2020.

[26] Aiwei Liu, Xuming Hu, Lijie Wen, and Philip S Yu. A comprehensive evaluation of chatgpt's zero-shot text-to-sql capability. *arXiv preprint arXiv:2303.13547*, 2023.

[27] Qin Lyu, Kaushik Chakrabarti, Shobhit Hathi, Souvik Kundu, Jianwen Zhang, and Zheng Chen. Hybrid ranking network for text-to-sql. *arXiv preprint arXiv:2008.04759*, 2020.

[28] Da Ma, Xingyu Chen, Ruisheng Cao, Zhi Chen, Lu Chen, and Kai Yu. Relation-aware graph transformer for sql-to-text generation. *Applied Sciences*, 12(1):369, 2021.

[29] Mohammadreza Pourreza and Davood Rafiei. Evaluating cross-domain text-to-sql models and benchmarks. *arXiv preprint arXiv:2310.18538*, 2023.

[30] Mohammadreza Pourreza and Davood Rafiei. Din-sql: Decomposed in-context learning of text-to-sql with self-correction. *Advances in Neural Information Processing Systems*, 36, 2024.

[31] Jiexing Qi, Jingyao Tang, Ziwei He, Xiangpeng Wan, Yu Cheng, Chenghu Zhou, Xinbing Wang, Quanshi Zhang, and Zhouhan Lin. Rasat: Integrating relational structures into pretrained seq2seq model for text-to-sql. *arXiv preprint arXiv:2205.06983*, 2022.

[32] Bowen Qin, Binyuan Hui, Lihan Wang, Min Yang, Jinyang Li, Binhua Li, Ruiying Geng, Rongyu Cao, Jian Sun, Luo Si, et al. A survey on text-to-sql parsing: Concepts, methods, and future directions. *arXiv preprint arXiv:2208.13629*, 2022.

[33] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *The Journal of Machine Learning Research*, 21(1):5485–5551, 2020.

[34] Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.

[35] Jingqing Ruan, Yihong Chen, Bin Zhang, Zhiwei Xu, Tianpeng Bao, Guoqing Du, Shiwei Shi, Hangyu Mao, Xingyu Zeng, and Rui Zhao. Tptu: Task planning and tool usage of large language model-based ai agents. *arXiv preprint arXiv:2308.03427*, 2023.

[36] Noah Shinn, Beck Labash, and Ashwin Gopinath. Reflexion: an autonomous agent with dynamic memory and self-reflection. *arXiv e-prints*, pages arXiv–2303, 2023.

[37] Chang Shu, Yusen Zhang, Xiangyu Dong, Peng Shi, Tao Yu, and Rui Zhang. Logic-consistency text generation from semantic parses. *arXiv preprint arXiv:2108.00577*, 2021.

[38] Guanghu Sui, Zhishuai Li, Ziyue Li, Sun Yang, Jingqing Ruan, Hangyu Mao, and Rui Zhao. Reboost large language model-based text-to-sql, text-to-python, and text-to-function–with real applications in traffic domain. *arXiv preprint arXiv:2310.18752*, 2023.

[39] Defog Team. Sqlcoder, `https://github.com/defog-ai/sqlcoder`, 2023.

[40] InternLM Team. Internlm: A multilingual language model with progressively enhanced capabilities, `https://github.com/InternLM/InternLM`, 2023.

[41] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

[42] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

[43] Bailin Wang, Richard Shin, Xiaodong Liu, Oleksandr Polozov, and Matthew Richardson. Rat-sql: Relation-aware schema encoding and linking for text-to-sql parsers. *arXiv preprint arXiv:1911.04942*, 2019.

[44] Kun Xu, Lingfei Wu, Zhiguo Wang, Yansong Feng, and Vadim Sheinin. Sql-to-text generation with graph-to-sequence model. *arXiv preprint arXiv:1809.05255*, 2018.

[45] Xiaojun Xu, Chang Liu, and Dawn Song. Sqlnet: Generating structured queries from natural language without reinforcement learning. *arXiv preprint arXiv:1711.04436*, 2017.

[46] Sun Yang, Qiong Su, Zhishuai Li, Ziyue Li, Hangyu Mao, Chenxi Liu, and Rui Zhao. Sql-to-schema enhances schema linking in text-to-sql. *arXiv preprint arXiv:2405.09593*, 2024.

[47] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.

[48] Pengcheng Yin and Graham Neubig. A syntactic neural model for general-purpose code generation. *arXiv preprint arXiv:1704.01696*, 2017.

[49] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-sql task. *arXiv preprint arXiv:1809.08887*, 2018.

[50] John M Zelle and Raymond J Mooney. Learning to parse database queries using inductive logic programming. In *Proceedings of the national conference on artificial intelligence*, pages 1050–1055, 1996.

[51] Bin Zhang, Hangyu Mao, Jingqing Ruan, Ying Wen, Yang Li, Shao Zhang, Zhiwei Xu, Dapeng Li, Ziyue Li, Rui Zhao, et al. Controlling large language model-based agents for large-scale decision-making: An actor-critic approach. *arXiv preprint arXiv:2311.13884*, 2023.

[52] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. Bertscore: Evaluating text generation with bert. *arXiv preprint arXiv:1904.09675*, 2019.

[53] Victor Zhong, Caiming Xiong, and Richard Socher. Seq2sql: Generating structured queries from natural language using reinforcement learning. *arXiv preprint arXiv:1709.00103*, 2017.

A Supplemental Evaluation Results

A.1 The Performance of GPT-4

Based on the experiments (as shown in Figure 7), GPT-4 consistently outperformed ChatGPT and other models across all sub-tasks. However, the performance gains were marginal, particularly in more challenging sub-tasks such as SQL optimization. These results underscore the difficulty of our benchmark and its significance in advancing the development of future Text-to-SQL models and approaches. Moreover, we believe that the core insights learned from our broader analysis will remain relevant even as specific model capabilities evolve.

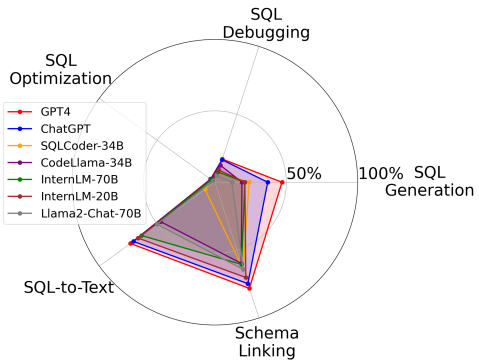

Figure 7: Overall performance of different LLMs in various sub-tasks of SQLBench. We add additional performance of GPT-4-turbo. GPT-4 consistently outperformed ChatGPT and other models across all sub-tasks. However, the performance gains are marginal.

A.2 Additional Experimental Results on "BigTable"

A.2.1 End-to-End Text-to-SQL evaluation.

Table 8: EX (%) of different LLMs on "BigTable". Prompt template is "SimpleDDL-MD-Chat".

No. of GT Tables	ChatGPT	SQLCoder-34B	Codellama-34B	InternLM-70B	Llama2-Chat-70B	InternLM2-20B
1	**60.00**	44.00	32.00	44.00	22.00	40.00
2	**20.00**	8.00	10.00	10.00	10.00	6.00
3	**38.00**	24.00	22.00	20.00	8.00	18.00
>3	**30.00**	20.00	12.00	10.00	8.00	20.00
Total	**37.00**	24.00	19.00	21.00	12.00	21.00

This research conducts an end-to-end evaluation of the Text-to-SQL capabilities of various LLMs using the "SimpleDDL-MD-Chat" prompt template on the "BigTable" dataset, with results depicted in Table 8.

Comparison of Performance Across Different Models. The results demonstrate a clear performance hierarchy among the models, with SQLCoder, CodeLlama, InternLM, and InternLM2 consistently outperforming Llama2-Chat. This finding highlights the effectiveness of coding-specific models, such as SQLCoder and CodeLlama, in the Text-to-SQL domain. Additionally, certain general-purpose models, like InternLM and InternLM2, can achieve performance levels comparable to specialized models, even without fine-tuning for coding tasks.

Difficulty Comparison Across Different Numbers of GT Tables. We examine the query difficulty based on the number of GT tables involved. The results reveal a decrease in EX as the number of GT tables increases. Notably, the EX for models on queries with two GT tables is unexpectedly lower compared to those with three or more GT tables. This observation can be attributed to the fact that queries involving two GT tables have the highest average number of columns (see Table 2).

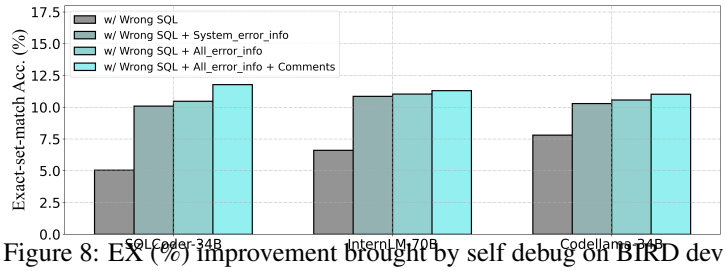

Figure 8: EX (%) improvement brought by self debug on BIRD dev set.

A.2.2 Schema Linking.

For a more detailed analysis of schema linking, we introduce two supplemental metrics, Subset Match and Exact Match. These metrics measure the proportion of instances in which the GT tables are a subset of, or exactly match, retrieved tables, respectively. We observe from Table 9 that the Few Shot + PreSQL and PreSQL method achieves the highest Subset Match and the Exact Match, respectively, which explains the higher RES achieved by these methods.

Table 9: The results of table linking include the metrics of exact match and subset match.

		SQLCoder-34B		**InternLM-70B**		**Codellama-34B**		**InternLM2-20B**		**Llama2-Chat-70B**	
		w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk
Zero Shot	**Subset Match**	0.28	0.55	0.66	0.645	0.495	0.655	0.605	0.67	0.655	0.76
	Exact Match	0.135	0.295	0.355	0.34	0.19	0.27	0.235	0.415	0.12	0.135
Few Shot	**Subset Match**	0.5	0.475	0.595	0.555	0.645	0.7	0.53	0.545	0.57	0.72
	Exact Match	0.155	0.22	0.355	0.34	0.19	0.155	0.33	0.355	0.125	0.18
PreSQL	**Subset Match**	0.6	0.675	0.475	0.505	0.525	0.62	0.6	0.67	0.51	0.65
	Exact Match	**0.465**	**0.55**	**0.43**	**0.435**	**0.45**	**0.525**	**0.5**	**0.525**	**0.3**	**0.23**
Few Shot + PreSQL	**Subset Match**	**0.715**	**0.77**	**0.67**	**0.675**	**0.755**	**0.785**	**0.74**	**0.78**	**0.815**	**0.86**
	Exact Match	0.195	0.345	0.395	0.405	0.2	0.185	0.42	0.46	0.19	0.165

A.3 Additional Experimental Results on BIRD

To further justify the effectiveness and robustness of the core conclusions of this paper, we conduct identical experiments on the BIRD development dataset. The BIRD dev set contains 1,534 instances, which exhibit a larger scale compared to "BigTable". In the experiments, we only choose representative LLMs and methods that are influential in deriving core conclusions. Results for five sub-tasks are shown in Table 10, Figure 8, Tables 11-13, respectively. These experimental results are largely consistent with the core conclusions obtained on "BigTable", supporting the validity of the conclusions.

Table 10: EX (%) of different LLMs on BIRD dev set. prompt template is "SimpleDDL-MD-Chat".

No. of GT Tables	**SQLCoder-34B**	**InternLM-70B**	**Codellama-34B**	**InternLM2-20B**
1	40.33	**43.37**	40.88	40.33
2	**28.42**	25.24	27.57	28.31
3	**21.05**	15.79	16.75	**21.05**
>3	**30.00**	5.00	15.00	**30.00**
Total	**30.25**	27.97	29.07	30.18

B Details for Dataset Construction

In the process of modifying and expanding the original queries, our primary goal is to introduce a diverse set of linguistic and structural variations that would challenge the model's ability to generalize beyond simple lookup-based responses.

Specifically, we categorize the SQL statements into logical components, including Table Queries, Column Selection, Join Columns, Condition Filters, and Data Processing. This categorization is consistent with the approach we used in constructing the dataset for the SQL Debugging task. By

Table 11: VES results of SQL Optimization on BIRD dev set.

Prompt Template	SQLCoder-34B	InternLM-70B	Codellama-34B
with $\mathcal{Y}$	99.68	99.99	99.04
w/ $\mathcal{Y}+\mathcal{S}$	102.65	100.10	101.48
w/ $\mathcal{Y}+\mathcal{S}+\mathcal{Q}$	102.92	100.69	101.78
SimpleDDL-MD-Chat-Efficiency	101.26	100.51	100.34

Table 13: RES results of Schema Linking on BIRD dev set.

	SQLCoder 34B		InternLM-70B		Codellama-34B	
	w/o fk	w/ fk	w/o fk	w/ fk	w/o fk	w/ fk
Zero Shot	0.7034	0.7138	0.7505	0.7739	**0.7130**	0.6564
Few Shot	0.4343	0.4078	0.7365	0.7574	0.5712	0.5822
PreSQL	**0.7274**	**0.7715**	0.5563	0.5958	0.6389	**0.6615**
Few Shot + PreSQL	0.6621	0.7296	**0.7686**	**0.7936**	0.6551	0.6417

systematically modifying these five components, we are able to generate new queries $\mathcal{Q}$ and SQL statements $\mathcal{Y}$ that differed from the original examples in Bird. Furthermore, to ensure a balanced distribution of the number of tables involved in the queries, we expand queries with three tables to include four tables in cases where the original dataset lacked sufficient examples with four or more tables (as there are only 20 such instances in the BIRD-Dev dataset).

Listing 1 provides an example of expanding table queries and column selection by adding condition filters.

```
Original Question: For the client who applied 98832 USD loan in
↳ 1996/1/3, when was his/her birthday?
Original  SQL: SELECT T3.birth_date FROM loan AS T1 INNER JOIN account
↳  AS T2 ON T1.account_id = T2.account_id INNER JOIN client AS T3 ON
↳ T2.district_id = T3.district_id WHERE T1.date = '1996-01-03' AND T1.
↳ amount = 98832
New Question:  When is the birthday of a client who applied 98832 USD
↳ loan in East Bohemia in 1996/1/3?
New SQL: SELECT T3.birth_date FROM loan AS T1 INNER JOIN account AS T2
↳  ON T1.account_id = T2.account_id INNER JOIN client AS T3 ON T2.
↳ district_id = T3.district_id WHERE T1.date = '1996-01-03' AND T1.
↳ amount = 98832 AND T3.district_id = (SELECT district_id FROM
↳ district WHERE A3 = 'east Bohemia' )
```

Listing 1: An example of the construction of a sample in BigTable.

C Details for Sub-tasks

C.1 Task Formulation

The formulations for each sub-task are as follows:

- SQL Debugging: Transforming the initial erroneous SQL queries $\mathcal{Y}^{\mathrm{w}}$ into the correct form $\mathcal{Y}^{\mathrm{r}}$: $\mathcal{Y}^{\mathrm{r}} = f_{\mathcal{M}}(\mathcal{P}_{\mathrm{d}}(\mathcal{Y}^{\mathrm{w}}, \mathcal{I}))$.
- SQL Optimization: Transforming the initial SQL queries $\mathcal{Y}$ into an optimized form $\mathcal{Y}^{\mathrm{o}}$, aiming to improve efficiency while preserving correctness: $\mathcal{Y}^{\mathrm{o}} = f_{\mathcal{M}}(\mathcal{P}_{\mathrm{o}}(\mathcal{Y}, \mathcal{I}))$.
- SQL-to-Text: Converting SQL queries $\mathcal{Y}$ into a natural language description $\mathcal{Q}'$ that effectively conveys the content of the query: $\mathcal{Q}' = f_{\mathcal{M}}(\mathcal{P}_{\mathrm{t}}(\mathcal{Y}, \mathcal{I}))$.

- Schema Linking: Generating a linking representation $\mathcal{L} = (e_i, t_j, c_k)$ based on the natural language query $\mathcal{Q}$ and the database schema $\mathcal{S}$ that maps entities or keywords e_i in $\mathcal{Q}$ to tables t_j and columns c_k in schema $\mathcal{S}$: $\mathcal{L} = f_{\mathcal{M}}(\mathcal{P}_1(\mathcal{Q}, \mathcal{S}))$.

Here, $\mathcal{P}$ and $\mathcal{I}$ are the corresponding prompt template and the additional information used for different tasks, $f_{\mathcal{M}}(\cdot)$ represents the mapping function of the LLM $\mathcal{M}$.

C.2 Details for Evaluation Metrics

In this section, we present the detailed definition and discussion of evaluation metrics in this paper.

Two primary evaluation metrics for assessing the accuracy of SQLs on Spider are Exact Matching(EM) and **Execution Accuracy(EX)**. EM measures whether the predicted query as a whole is equivalent to the gold query. It is possible to encounter false negative evaluations since a question might be solvable by multiple syntactically different but semantically identical SQL statements. EX is a more widely used metric, measures whether the result of executing the predicted query matches the gold value. We use EX to evaluate the accuracy of SQLs in this paper, which is defined as:

$$\text{EX} = \frac{\sum_{n=1}^{N} \mathbb{1}(V_n, \hat{V}_n)}{N}, \tag{6}$$

$$\text{where} \quad \mathbb{1}(V_n, \hat{V}_n) = \begin{cases} 1, \text{ if } V_n = \hat{V}_n \\ 0, \text{ if } V_n \neq \hat{V}_n \end{cases}, \tag{7}$$

where V_n and $\hat{V}_n$ denote execution results of the predicted SQL and the GT SQL for the n-th instance in the evaluation set, respectively.

BIRD further proposes **Valid Efficiency Score (VES)**, an integrated metric assessing both accuracy of execution results (i.e., EX) and the execution efficiency of SQL queries. VES is defined as:

$$\text{VES} = \frac{\sum_{n=1}^{N} \mathbb{1}(V_n, \hat{V}_n) \cdot R(\mathcal{Y}_n^{\text{o}}, \hat{\mathcal{Y}}_n)}{N}, \tag{8}$$

$$\text{where} \quad R(\mathcal{Y}_n^{\text{o}}, \hat{\mathcal{Y}}_n) = \sqrt{\frac{\mathbf{E}(\mathcal{Y}_n^{\text{o}})}{\mathbf{E}(\hat{\mathcal{Y}}_n)}}, \tag{9}$$

where $\mathcal{Y}_n^{\text{o}}$ and $\hat{\mathcal{Y}}_n$ denote the optimized predicted SQL and GT SQL for the n-th instance, respectively. $\mathbf{E}(\cdot)$ serves as a function that quantifies the absolute execution efficiency (e.g., the reciprocal of execution time) of each SQL within a specified environment.

C.3 Details for SQL Debugging

Detailed comments for all Result Error are shown in Table 14. In addition, Figure 9 and Figure 10 provide more detailed information about error types, including the quantity of erroneous data.

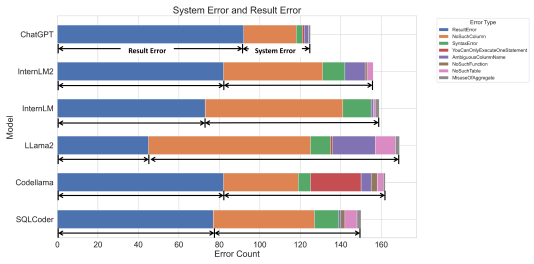

Figure 9: Error information of incorrect SQL statements generated by LLMs. Models with higher performance are less likely to commit basic System Errors (e.g., SQL statements that fail to execute and get errors information from the SQL compiler) and are more prone to Result Errors (i.e., SQL statements that execute correctly but yield incorrect answers). Therefore, it is necessary to further classify Result Errors.

Table 14: Comments for different categories of Result Error used in SQL debugging.

Error Type	Subcategory	comments prompt
Table Query Error	**Excessive Tables**	The tables you inquired about is incorrect, you query too much tables.
	Missing Tables	The tables you inquired about is incorrect, you need to query more tables.
	Incorrect Tables	The tables you inquired about is incorrect.
Column Selection Error	**Excessive Columns**	You have found the correct tables. But you select wrong columns,you select too much Columns.
	Missing Columns	You have found the correct tables. But you select wrong columns,you need to select more Columns.
	Incorrect Columns	You have found the correct tables.But you select wrong columns.
Join Columns Error		You have found the correct tables. You have selected the correct Columns. But you combine wrong rows when JOIN two tables.
Condition Filter Error		You have found the correct tables.You have selected the correct Columns. You have combined (JOIN) the correct tables. But an error occurred in the conditional filter.
Data Processing Error		You have found the correct tables. You have selected the correct Columns. You have combined (JOIN) the correct tables. You have used the correct conditional filtering. But there was an error in your processing of the data.

Figure 10: Result Errors for incorrect SQL statements generated by LLMs are categorized into five types: Table Query Errors, Column Selection Errors, Join Column Errors, Condition Filter Errors, and Data Processing Errors. Specifically, Table Query Errors encompass issues related to excessive, missing, or incorrect tables, while Column Selection Errors include excessive, missing, or incorrect columns.

C.4 Details for SQL Optimization

We suspect the reasons LLMs are challenging to achieve effective SQL optimization are as follows:

- **Lack of in-context learning guidance.** For self-debug, models can benefit from the error messages given by DBMS (see the last two rows of Figure 5). However, SQL optimization tasks lack explicit feedback from the system and guidance through In-context learning, despite the comprehensively designed prompt templates.
- **Lack of pre-training data.** During the pre-training phase, models can access massive correct SQL statements as training data, yet they may lack information on what kind of SQL statements are efficient.
- **System-level optimization is inherently challenging.** The efficiency of a SQL query can only be assessed by executing it within a database system. This is a system-level optimization beyond the scope of syntactic-level and semantic-level language modeling areas in which LLMs excel. It is challenging for models to capture the mapping between SQLs and their execution efficiency.

C.5 Details for SQL-to-Text

Detailed description of evaluation metrics used for SQL-to-Text task are as follows:

- ***Rouge*** [24] focuses on token-level syntactical similarity by breaking down the predicted question and the GT question into N-grams and calculating the recall rate. We use ROUGE-1, ROUGE-2, and ROUGE-L, respectively.
- ***BertScore*** [52] focuses on semantic similarity between predicted and GT questions. Based on the BERT pre-trained model, BertScore uses contextual embeddings to describe sentences and calculates the cosine similarity between them.
- ***Using LLMs as evaluator*** is an intuitive and recently prevalent evaluation approach. This involves presenting two questions to an LLM to determine whether they are semantically identical, with the accuracy rate being recorded. We use InternLM2 as evaluator, which achieves the best performance in the first two metrics. In practice, to mitigate the potential bias introduced by positional preference, we alter the order of predicted and GT questions in the input of InternLM2 and report the average result of two experiments.

D Prompt Template

D.1 Text-to-SQL

Listing 2- 7 demonstrate eight prompt templates evaluated in the end-to-end Text-TO-SQL task, where "SimpleDDL-MD-Chat" obtains the best evaluation performance.

```
Figure out corresponding SQLite SQL Query of Question according to
↳ database.
<Database>
CREATE TABLE stadium (stadium_id NUMBER PRIMARY KEY, ...);
CREATE TABLE singer (singer_id NUMBER PRIMARY KEY, name TEXT, ...);
</Database>
<Question>How many singers do we have?</Question>
(<SQL> SELECT)
```

Listing 2: Prompt template "DDL-HTML-Chat/Complete". The content in parentheses is unique to"Complete".

```
Figure out corresponding SQLite SQL Query of Question according to
↳ database.
<Database>
stadium(Stadium_ID,Location,Name,Capacity,Highest,Lowest,Average);
singer(Singer_ID,Name,Country,Song_Name,...);
</Database>
<Question>How many singers do we have?</Question>
(<SQL> SELECT)
```

Listing 3: Prompt template "SimpleDDL-HTML-Chat/Complete". The content in parentheses is unique to"Complete".

D.2 SQL Debugging

Listing 8 is used for error classification. In the detailed categorization process of Result Error, Table Query Error, Column Selection Error, and Join Column Error can be determined by comparing SQL Query with the ground truth using predefined rules. For Data processing Error and Condition Filter Error, we employ InternLM2-20B for binary classification tasks. Listing 9 is the prompt template used for SQL debugging tasks.

```
### Answer the question by sqlite SQL query only and with no
↳ explanation
### Sqlite SQL tables, with their properties:
CREATE TABLE stadium (stadium_id NUMBER PRIMARY KEY,...);
CREATE TABLE singer (singer_id NUMBER PRIMARY KEY, name TEXT, ...);
### Question: How many singers do we have?
### SQL: (SELECT)
```

Listing 4: Prompt template "DDL-MD-Chat/Complete". The content in parentheses is unique to"Complete".

```
### Answer the question by sqlite SQL query only and with no
↳ explanation
### Sqlite SQL tables, with their properties:
#
# stadium(Stadium_ID,Location,Name,Capacity,Highest,Lowest,Average);
# singer(Singer_ID,Name,Country,Song_Name,...);
#
### How many singers do we have?
### SQL: (SELECT)
```

Listing 5: Prompt template "SimpleDDL-MD-Chat/Complete". The content in parentheses is unique to"Complete".

D.3 SQL Optimization

Listing 10 is the prompt template "w/ demo + comments" used for SQL optimization tasks. Other variations of the template correspond to the removal of relevant content.

D.4 SQL-to-Text

Listing 11 is the prompt template used for SQL-to-Text task. Listing 12 is utilized to assess the consistency between the transformed text description and the original user query semantics. To mitigate the influence of the order of the two text inputs, we employ ChatGPT and InternLM2-20B to conduct two sequential tests and calculate the average of the results.

D.5 Schema Linking

Listing 13 is a schema linking prompt template proposed in C3 [9]. Listing 14 is used for few-shot learning proposed in DIN-SQL [30]. In addition, we conduct tests on the PreSQL method. This method generates an SQL query statement based on the user query and database schema, and extracts the corresponding tables from that statement.

```
/* Given the following database schema: */
CREATE TABLE stadium (stadium_id NUMBER PRIMARY KEY, ...);

CREATE TABLE singer (singer_id NUMBER PRIMARY KEY, name TEXT, ...);

/* Answer the following by SQLite SQL Query according to database:
↳ Show name, country, age for all singers ordered by age from the
↳ oldest to the youngest. */
/* SQL Query here*/
(SELECT)
```

Listing 6: Prompt template "DDL-Coding-Chat/Complete". The content in parentheses is unique to"Complete".

```
/* Given the following database schema: */
stadium(Stadium_ID,Location,Name,Capacity,Highest,Lowest,Average);
singer(Singer_ID,Name,Country,Song_Name,...);

/* Answer the following by SQLite SQL Query according to database:
↳ Show name, country, age for all singers ordered by age from the
↳ oldest to the youngest. */
/* SQL Query here*/
(SELECT)
```

Listing 7: Prompt template "SimpleDDL-Coding-Chat/Complete". The content in parentheses is unique to"Complete".

```
You are an expert in SQL queries. Please provide the error categories
↳ for incorrect SQL queries based on the Question and the correct SQL
↳ query.
Please think step by step and check for the following errors in order:
1. Condition Filter Error: Incorrect filtering of conditions.
2. Data Processing Error:The condition is filtered correctly, but the
↳ data processing is wrong. Note that the premise of this error is
↳ that the conditional filtering is correct.

Question: What is the name and capacity for the stadium with highest
↳ average attendance?
Correct SQL Query: SELECT name , capacity FROM stadium ORDER BY
↳ average DESC LIMIT 1
Wrong SQL Query: SELECT Name, Capacity FROM stadium WHERE Average = (
↳ SELECT MAX(Average) FROM stadium) ORDER BY Highest DESC

Give your Thought and Answer based on the information above.
```

Listing 8: Prompt template for conducting the "Condition Filter Error" and "Data Processing Error" classification tasks.

```
### Write the correct SQLite SQL Query corresponding to the Question
↳ based on the database, the Wrong SQL Query and the cause of the
↳ error.
### Sqlite SQL tables, with their properties:
#
#{SimpleDDL}
#
### Question: Under whose administration does the school with the
↳ highest number of test takers whose total SAT Scores are greater or
↳ equal to 1500 belong to? Indicate his or her full name.
### Wrong SQL Query:
SELECT T1.AdmFName1 , T1.AdmLName1 FROM schools AS T1 JOIN satscores
↳ AS T2 ON T1.CDSCode = T2.cds WHERE T2.NumTstTakr = ( SELECT
↳ NumTstTakr FROM satscores GROUP BY cds HAVING NumGE1500 >= 1500
↳ ORDER BY NumTstTakr DESC LIMIT 1 )
### Error Information:
Executed correctly, but with the wrong result.
You have found the correct tables. But you select wrong columns, you
↳ need to select more Columns.
### Correct SQL:
```

Listing 9: Prompt template "w/ Wrong SQL + All_error_info + Comment" for SQL Debugging.

```
### Rewrite and optimize the given SQL query to improve SQL query efficiency and minimize SQL execution time while ensuring correctness. Only output sql query, do not output any other content. Only output sql query, do not output any other content.
### Here are some reference cases:
#
# Question: List out the age of users who located in Vienna, Austria obtained the badge?
# SQL Query: SELECT Age FROM users WHERE Location = 'Vienna, Austria' AND Id IN (SELECT UserId FROM badges)
# New SQL Query: SELECT u.Age FROM users AS u INNER JOIN badges AS b ON u.Id = b.UserId WHERE u.Location = 'Vienna, Austria'
# Explanation: By applying a JOIN operation instead of a subquery with IN can improve efficiency, as the database may execute the JOIN and filtering processes concurrently in just one operation without the need to store the intermediate results to filter primary query.
#...
#
### Sqlite SQL tables, with their properties:
#
#{SimpleDDLdd}
#
### Question: What is the short name and fifa ID for Queens Park Rangers Football Team?...
### SQL Query:SELECT team_short_name , team_fifa_api_id FROM Team WHERE team_long_name = "Queens Park Rangers"
### New SQL Query:
```

Listing 10: Prompt template "w/ demo + comments" for SQL Optimization.

```
<Instruction>
You are an expert in database analysis and processing of SQL statements.
I will provide an SQL statement and relevant evidence. You need to help me analyze what problem this SQL statement is solving.
Here are some reference cases:
SQL:SELECT list_id FROM lists_users WHERE user_id = 85981819 ORDER BY list_creation_date_utc ASC LIMIT 1
question:What is the list ID that was first created by user 85981819?
SQL:SELECT COUNT(T2.user_id) FROM movies AS T1 INNER JOIN ratings AS T2 ON T1.movie_id = T2.movie_id WHERE T1.movie_title = 'Pavee Lackeen: The Traveller Girl' AND T2.rating_score = 4
question:How many users gave \"Pavee Lackeen: The Traveller Girl\" movie a rating score of 4?
Please answer strictly in the following format and do not change the format arbitrarily:
question:This is a problem description.
</Instruction>
<SQL>SELECT ...</SQL>
```

Listing 11: Prompt template for SQL-to-Text.

```
<Instruction>Determine whether the following two sentences ask the same question and whether their corresponding answers are the same .</Instruction>
<sentence1>How many singers do we have?</sentence1>
<sentence2>How many singers are there in total?</sentence2>
<Question>Just output True or False, do not output anything else</Question>
```

Listing 12: Prompt template for semantic consistency checking.

```
Given the database schema and question, perform the following actions:
1 - Rank all the tables based on the possibility of being used in the SQL according to the question from the most relevant to the least relevant, Table or its column that matches more with the question words is highly relevant and must be placed ahead.
2 - Check whether you consider all the tables.
3 - Output a list object in the order of step 2, Your output should contain all the tables. The format should be like:
[
    "table_1", "table_2", ...
]

Database schemas with their properties:
{SimpleDDL}

Question: What is the short name and fifa ID for Queens Park Rangers Football Team?In the database, short name of the football team refers to team_short_name; Queens Park Rangers refers to team_long_name = 'Queens Park Rangers';fifa ID refers to team_fifa_api_id.
Answer(Only output the list object containing all tables, do not output other content):
```

Listing 13: Prompt template for zero-shot schema linking.

```
Given the database schema and question, perform the following actions:
1 - Evaluate the importance of each table **in relation to the SQL query**, prioritizing tables and columns that closely match the question words. Rank the tables from the most crucial to the least crucial.
2 - Focus on identifying and listing only the most important tables based on the evaluation in step 1.
3 - Output a list object representing the order determined in step 2. The output should include **the most important tables** and follow this format:
[
    "most_important_table_1", "most_important_table_2", ...
]

Schema:
# department(Department_ID,Name,Creation, ...)
# head(head_ID,name,born_state,age)
# management(department_ID,head_ID,temporary_acting)
Foreign key:
management(department_ID) REFERENCES department(Department_ID)\nmanagement(head_ID) REFERENCES head(head_ID)
Question: what are the distinct creation years of the departments managed by a secretary born in state 'Alabama'?
Answer: ["department","management","head"]

Schema:
# Country(id,name)
# League(id,country_id,name)
# Player(id,player_api_id,player_name,player_fifa_api_id,birthday,height,weight)
# Player_Attributes(id,player_fifa_api_id, ...)
# Team(id,team_api_id,team_fifa_api_id,team_long_name,team_short_name)
# Team_Attributes(id,team_fifa_api_id, ...)
# sqlite_sequence(name,seq)
Foreign key:
Player_Attributes(player_api_id) REFERENCES Player(player_api_id)\nLeague(country_id) REFERENCES country(id)\nTeam_Attributes(team_api_id) REFERENCES Team(team_api_id)\nMatch(away_player_11) REFERENCES Player(player_api_id)
Question: List the names of all left-footed players who have overall rating between 85 and 90.
Answer: ["Player","Player_Attributes"]

Database schemas with their properties:
{SimpleDDL}

Question: What is the short name and fifa ID for Queens Park Rangers Football Team? In the database, short name of the football team refers to team_short_name; Queens Park Rangers refers to team_long_name = 'Queens Park Rangers';fifa ID refers to team_fifa_api_id.
Answer:
```

Listing 14: Prompt template for few-shot schema linking.