RELAXED SYSTEM LAB

HEXGEN-FLOW: Optimizing LLM Inference Request Scheduling for Agentic Text-to-SQL

You Peng[1*], **Youhe Jiang**[1*], **Wenqi Jiang**[2], **Chen Wang**[3†], **Binhang Yuan**[1†]

[1]HKUST, [2]ETH Zurich, [3]Tsinghua University

[*]Equal contribution, [†]Corresponding author

arXiv:2505.05286v3 [cs.DB] 8 Mar 2026

Abstract

Recent advances in agentic large language models (LLMs) have substantially improved Text-to-SQL, enabling users without database expertise to query databases intuitively. However, deploying agentic LLM-based Text-to-SQL systems in production remains challenging due to multi-stage dependencies, strict latency requirements, and deployment complexity across heterogeneous GPUs in enterprise clusters. Existing LLM serving frameworks are designed mainly for independent inference tasks, leading to suboptimal performance and frequent service-level objective (SLO) violations for Text-to-SQL workloads. In this paper, we introduce HEXGEN-FLOW, a framework for scheduling and executing agentic multi-stage LLM-based Text-to-SQL workflows on heterogeneous GPU clusters serving multi-tenant requests. HEXGEN-FLOW adopts a hierarchical scheduler that combines global workload-balanced task dispatching with an adaptive local priority queue, guided by a systematic analysis of agentic Text-to-SQL workflows. We also propose a lightweight simulation-based method to tune key scheduling hyperparameters, improving robustness and adaptability. Evaluations on realistic Text-to-SQL benchmarks show that HEXGEN-FLOW significantly outperforms state-of-the-art LLM serving frameworks. Across all traces, HEXGEN-FLOW reduces P95 tail latency by $1.42\sim1.56\times$ and increases throughput by $1.49\sim1.81\times$, demonstrating consistent gains under diverse workloads.

1 Introduction

Recent advances in large language models (LLMs) have substantially improved Text-to-SQL, i.e., translating natural language questions into executable SQL statements [1–4]. This progress can democratize database access by enabling non-experts to query structured data without writing SQL. In enterprise settings, agentic Text-to-SQL can improve analytics productivity and accessibility.

A typical agentic Text-to-SQL workflow (Figure 1) consists of multiple interdependent stages: (1) *schema linking* to map query terms to the database schema; (2) *SQL candidate generation* to propose multiple SQL statements; (3) *self-correction* to execute and iteratively refine erroneous candidates; and (4) *evaluation* to select a final statement using LLM generated unit tests. Compared to non-agentic Text-to-SQL methods with a small, fixed number of LLM calls, this agentic pipeline improves quality via verifier-guided refinement and candidate selection, but introduces a longer, dependency-constrained sequence of LLM calls. Unlike many general agentic workflows, it typically runs without human-in-the-loop, and SQL execution is on the order of milliseconds—negligible relative to end-to-end latency dominated by multi-stage LLM inference.

Deploying agentic LLM-based Text-to-SQL systems in production requires not only strong models [5–9] and inference-time scaling algorithms [10, 11], but also an efficient and scalable serving infrastructure. In practice,

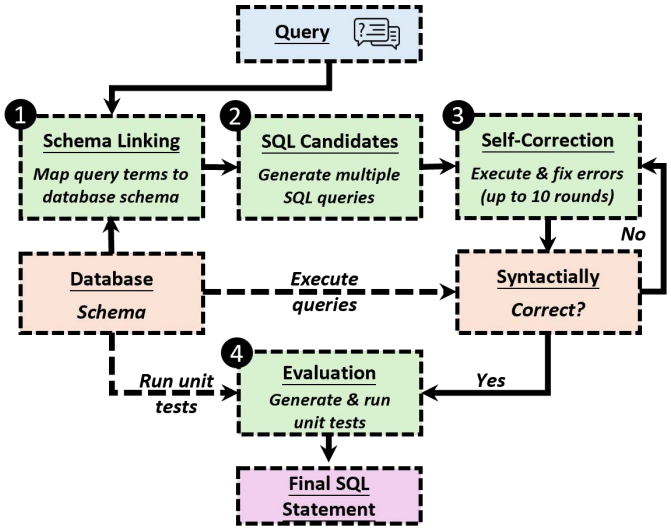

Figure 1 A visual illustration of the agentic multiple-stage Text-to-SQL workflow with inter-stage dependency.

enterprise environments often rely on heterogeneous GPUs for LLM serving[1]. Therefore, the goal is to serve the multi-stage Text-to-SQL pipeline efficiently on heterogeneous clusters, achieving low end-to-end latency and high throughput.

However, efficient serving is challenging due to the workflow dependencies and the heterogeneity in both request demands and GPU serving capacities. Each end-to-end query triggers a sequence of dependent LLM calls that must be coordinated across model instances with different performance and load. The combination of workflow dependencies, resource heterogeneity, and diverse per-query service-level objectives (SLO) complicates scheduling and resource allocation. Concretely, we summarize three core challenges below:

- **C1. LLM inference request dependencies**: Agentic Text-to-SQL workflows involve multiple interdependent stages with different urgency levels. Later-stage tasks cannot commence until the preceding stages are completed. Consequently, delays in early stages affect downstream stages, increasing the risk of end-to-end SLO violations.
- **C2. Heterogeneity in workflow and infrastructure**: The heterogeneity in Text-to-SQL workflow lies in two aspects: (i) The LLM inference requests exhibit substantial variability in different stages, driven by differences in the length of the query prompt and the number of output tokens generated; (ii) Enterprise production environments commonly leverage heterogeneous GPUs with different computational capabilities, so the latency/throughput of the same request can differ across model instances.
- **C3. Varying SLO constraints in serving**: Deployments must handle concurrent Text-to-SQL queries from different users with different SLOs. Ensuring that all LLM inference requests within a query collectively meet its end-to-end SLO requires fine-grained, per-query deadline awareness and adaptive prioritization across stages and model instances.

Due to these challenges, scheduling policies in popular LLM serving systems [16–18] are often inefficient for Text-to-SQL workloads. Existing frameworks largely target independent inference requests and lack explicit coordination across dependent stages or enforcement of end-to-end deadlines. While recent works on adaptive batching [19], priority-aware allocation [20, 21], and GPU load balancing [22] improve general serving, the unique combination of multi-stage pipeline dependencies and GPU resource heterogeneity in agentic Text-to-SQL serving is still largely unexplored.

To address these challenges, we introduce HEXGEN-FLOW, a novel framework that efficiently schedules and executes agentic Text-to-SQL workloads in heterogeneous GPU-serving environments while supporting multi-

[1]In the current exciting era of generative AI, chip vendors typically release new generations of AI chips every 24 months. For example, Nvidia introduced the Turing architecture in 2018 [12], Ampere in 2020 [13], Hopper in 2022 [14], and Blackwell in Q4, 2024 [15]. On the other hand, one particular version of an AI chip often serves for a much longer period.

tenant queries with diverse SLOs. HEXGEN-FLOW is built around a *two-level* scheduling architecture that directly responds to C1-C3:

- To tackle C1, HEXGEN-FLOW explicitly tracks workflow progress and dispatches dependent inference requests only when prerequisites complete. Parallelizable requests are dispatched concurrently to reduce stage completion time.
- To address C2, HEXGEN-FLOW employs a workload-balanced dispatcher at the global coordinator, selecting a model instance by jointly considering (i) the request's computational demand and (ii) each instance's serving capacity and current load. This design ensures that heavy or latency-sensitive tasks run on powerful hardware, while lighter tasks utilize underutilized resources, achieving balanced utilization on a heterogeneous cluster.
- To satisfy C3, each model instance maintains an urgency-guided local priority queue that reorders pending LLM requests at batch boundaries during continuous batching based on remaining budget and estimated execution time. Per-query SLO budgets are continuously recomputed to propagate deadline pressure to downstreams.

Lastly, we summarize our contributions as follows:

- **Contribution 1.** We identify and formalize three core design principles for serving agentic Text-to-SQL workflows at scale: explicit multi-stage dependency management, heterogeneity-aware request allocation, and end-to-end SLO guarantees in multi-tenant settings.
- **Contribution 2.** We propose HEXGEN-FLOW, a serving framework that combines the global workload-balanced dispatching policy with a local urgency-driven priority queue to enable SLO-aware execution on heterogeneous clusters.
- **Contribution 3.** We conduct a comprehensive evaluation demonstrating HEXGEN-FLOW 's performance on agentic Text-to-SQL workloads. Across diverse traces and hardware configurations, HEXGEN-FLOW reduces P95 tail latency by $1.42\sim1.56\times$ and increases throughput by $1.49\sim1.81\times$ compared to state-of-the-art LLM serving systems.

2 Preliminaries

In this section, we first introduce the key concepts underlying our serving system design: §2.1 decomposes the Text-to-SQL workflow into its constituent stages, highlighting the sequential dependencies and multi-stage pattern that necessitate specialized scheduling for the workflow; §2.2 analyzes limitations of the current scheduling and queuing policy in existing LLM serving systems, demonstrating why general-purpose schedulers fail to meet the end-to-end latency requirements of the agentic LLM-based Text-to-SQL workflow.

2.1 Agentic LLM-based Text-to-SQL Workflow

The agentic LLM-based Text-to-SQL workflow involves several key stages to transform a natural language query into an executable SQL statement. As shown in Figure 1, we summarize the key stages in the state-of-the-art agentic Text-to-SQL paradigm (mainly following CHESS [23]) as below:

- **Schema linking**: Given metadata and column descriptions for tables, the LLM identifies and aligns entities mentioned in the user's natural language query to relevant tables and columns within the database schema. This step is essential to accurately ground the user's query in the underlying database structure.
- **SQL candidates generation**: Utilizing schema alignments from the previous step, the LLM generates a fixed number of candidate SQLs from the natural language query. Multiple LLM inference requests are executed concurrently, employing different prompts and illustrative examples, to produce a diverse set of candidates. This parallel approach aims to capture multiple interpretations of the user's intent.
- **Self-correction**: Candidate SQL queries are executed in the database, and any execution errors prompt iterative refinement throughout subsequent LLM invocations. The agent can iteratively refine queries up to a predefined limit (e.g., 10 iterations), systematically improving their correctness.

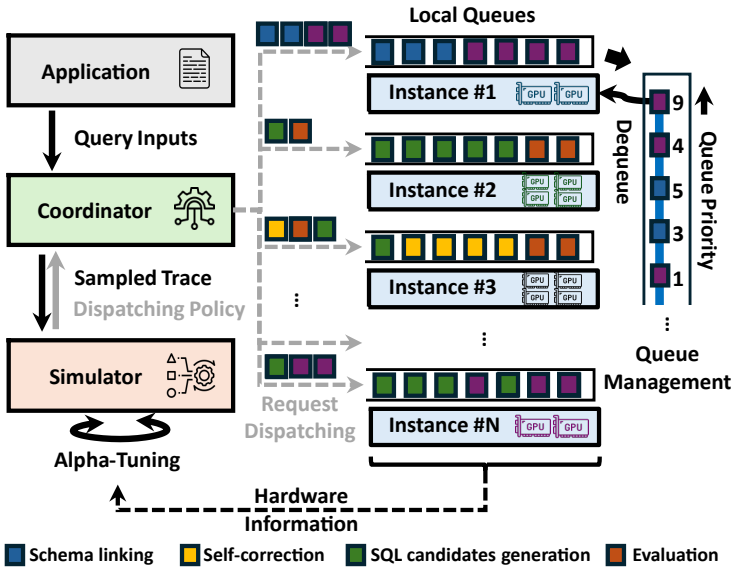

Figure 2 HEXGEN-FLOW system. Incoming LLM inference requests are dispatched by a global coordinator to model instances based on workload balance and task suitability. Each model instance manages its queue using an urgency-guided priority mechanism. Different color chunks represent an LLM inference request from a different stage in the workflow.

- **Evaluation**: After self-correction removes execution errors, the LLM synthesizes a checklist-style set of "unit tests" in natural language (i.e., verifiable criteria derived from the user query). An LLM-as-a-judge scores each finalized SQL candidate against these criteria and selects the highest-scoring query. These tests verify both the semantic accuracy and functional correctness of the SQL candidates, ensuring alignment with the user's original intention.

2.2 Scheduling in LLM Serving

Current LLM serving frameworks, such as vLLM [16], Text Generation Inference (TGI) from Hugging Face [24], and TensorRT-LLM from NVIDIA [25], employ various scheduling and queueing strategies optimized for general-purpose LLM serving. For example, vLLM utilizes continuous batching [26] with a first-come-first-served (FCFS) policy, allowing new LLM inference requests to join ongoing batches during decoding. TGI groups incoming LLM inference requests based on prompt length and generation parameters to optimize batch formation. TensorRT-LLM implements an "in-flight" batching strategy, managing concurrent LLM inference requests through the scheduling policies implemented in Triton [27], including FIFO and priority-based queuing. While effective for standard LLM tasks, these frameworks are not inherently designed to handle the complexities of multi-stage, dependency-aware agentic workflows like Text-to-SQL. Advanced queuing methods, such as Queue Management for LLM serving (QLM) [18] and the Virtual Token Counter (VTC) [17], have been proposed to address specific performance and fairness requirements—QLM introduces priority scheduling to meet SLOs by prioritizing urgent LLM inference requests and employing techniques like preemption and state swapping; VTC ensures equitable resource allocation among users by tracking the number of tokens served and prioritizing those with lower consumption. However, these mechanisms still can not inherently account for the nuances of the Text-to-SQL pipeline, such as varying computational costs across different stages or the impact of query latency on downstream database performance.

3 HEXGEN-FLOW

This section first introduces three system requirement that Text-to-SQL serving systems should meet (§3.1), then demonstrates the framework overview (§3.2), and enumerates each core functionality (§3.3).

3.1 Anatomy of Text-to-SQL Serving System

To eliminate the limitations of current LLM serving systems serving Text-to-SQL workflows, we present three essential requirements for an agentic LLM-based Text-to-SQL system.

- **Explicit multi-stage dependency management.** Agentic Text-to-SQL decomposes an end-to-end query into sequential and parallel stages (schema linking, candidate generation, self-correction, evaluation), each triggering one or more LLM inference requests. The system should explicitly track inter-stage dependencies, dispatching a stage only after prerequisites complete, while executing parallelizable requests concurrently to reduce stage completion time.
- **Heterogeneity-aware LLM inference request allocation.** Production environments commonly feature hardware heterogeneity—mixed GPU deployments with varying computational capabilities, memory capacities, and performance characteristics due to incremental upgrades or cost considerations. Efficient serving therefore requires routing each LLM inference request to a suitable model instance by accounting for both the request's compute demand (e.g., token lengths, memory footprint) and each instance's serving capacity and current load [28–30].
- **End-to-end request SLO guarantees.** Production-level Text-to-SQL serving systems must adhere to end-to-end SLOs to ensure consistent, predictable performance in multitenant environments with diverse workloads. Since a query spans multiple dependent LLM calls, the system should track per-query deadlines and prioritize requests based on remaining budget and expected execution time, adapting priorities as workflow progress consumes slack.

3.2 HEXGEN-FLOW Overview

To satisfy the requirements listed in §3.1, we propose HEXGEN-FLOW, a distributed system designed to efficiently serve multi-stage, LLM-based Text-to-SQL inference workloads in heterogeneous GPU clusters. Figure 2 provides an overview of the system architecture. HEXGEN-FLOW's serving architecture centers on a centralized global coordinator and multiple LLM model instances over heterogeneous GPUs. The system handles multi-stage inference workflows under heterogeneous GPU deployments while meeting strict per-query SLO deadlines through intelligent dispatch and prioritization using a *two-level* scheduling design: (i) the global coordinator assigns each incoming LLM inference request to an appropriate model serving instance, considering the computational requirements of the request and the availability of each model instance; and each model instance maintains its own local priority queue, dynamically ordering pending inference requests by urgency. After dispatch, requests are processed at the selected instance, where pending requests can be continuously reordered as their urgency increases with waiting time. Correspondingly, the global coordinator focuses on making queue-aware decisions for new arrivals rather than migrating in-flight requests across instances, because transferring KV cache and batch state is costly and can negate latency gains.

3.3 System Implementation

We enumerate the core functionalities in HEXGEN-FLOW system here, the corresponding formulation and detailed solution of the scheduling algorithm are introduced in §4.

Multi-stage dependency management. HEXGEN-FLOW treats each end-to-end Text-to-SQL query as a workflow of dependent LLM inference requests. The global coordinator tracks per-query stage completion and dispatches a request only when its prerequisites finish, ensuring correct execution order. For parallel stages (e.g., candidate generation), it dispatches ready requests concurrently across available instances to reduce stage latency. As stages complete, the coordinator updates the query's remaining deadline, which tightens downstream budgets and increases their urgency in local queues.

LLM inference request allocation. HEXGEN-FLOW routes each request to a suitable model instance under heterogeneous GPUs. The global coordinator maintains empirical performance profiles to estimate a request's execution time on each instance and tracks each instance's backlog. It then dispatches the request by balancing (i) estimated computation cost and (ii) queueing delay, avoiding load-agnostic policies such as round-robin. If an instance becomes slow or unavailable, HEXGEN-FLOW relies on standard serving-layer timeout/retry

semantics: the coordinator treats the attempt as failed and re-issues the request to another instance based on the same workload-balanced dispatching policy. This routing improves utilization and tail latency by matching heavy requests to faster instances while filling residual capacity on slower ones.

Adaptive local priority queue scheduling. Each model instance maintains a local priority queue that selects requests at batch boundaries based on urgency. HEXGEN-FLOW derives per-request SLO budgets from the query's remaining end-to-end deadline and estimated per-stage costs, and recomputes budgets as the workflow progresses. Requests that are closer to exhausting their budgets are prioritized, which reduces deadline misses under multi-tenant contention. We formalize budget allocation and the urgency metric in §4.2.

4 Scheduling Algorithm

In this section, we introduce the mathematical formulation and solution of the global coordinator and local priority queue. We first formulate the scheduling problem as below:

Problem formulation. Consider a stream of Text-to-SQL queries $\{Q_1, Q_2, \ldots\}$ where $Q_i \sim \mathbb{P}_Q$. Each query Q_i consists of LLM inference requests $\{q_{i,1}, \ldots, q_{i,n_i}\}$ with an end-to-end SLO T_i^{SLO}. Given N model instances $\mathbf{M} = \{m_1, \ldots, m_N\}$, request $q_{i,j}$ experiences instance-dependent service time $t_{i,j}^m$ (queueing + computation) if executed on m. We seek a dispatching policy ϕ that maps each request to an instance, $\phi(q_{i,j}) \in \mathbf{M}$, to maximize the probability of meeting end-to-end SLOs:

$$\arg\max_{\phi} \mathbb{P}\left(\sum_{j=1}^{n_i} t_{i,j}^{\phi(q_{i,j})} \leq T_i^{\text{SLO}} \mid Q_i \sim \mathbb{P}_Q\right). \tag{1}$$

Finding optimal policies under unknown arrival and service distributions is intractable—The inherent uncertainty in job characteristics necessitates dynamic decision-making without complete information, rendering the scheduling problem NP-hard [31]. We therefore use a practical heuristic that combines global dispatching with local deadline-driven prioritization. Concretely, our design includes the following components:

- **Global coordinator level (§4.1).** The coordinator performs workload-balanced dispatching, assigning each incoming LLM inference request to a model instance by considering both estimated execution time and current queueing delay.
- **Local priority queue level (§4.2).** Each model instance maintains a priority queue and reorders pending LLM inference requests at batch boundaries using a deadline-aware urgency metric derived from per-request SLO budgets.
- **Simulator-guided tuning (§4.3).** A trace-driven simulator tunes the dispatching hyperparameter online to adapt to workload and hardware dynamics.

Our design ensures optimal exploitation of heterogeneous model instance serving capabilities and reliable compliance with strict per-query latency objectives in dynamic, multi-tenant deployments. Note that HEXGEN-FLOW will not affect the original Text-to-SQL results, where the scheduling does not modify model inputs or the inter-stage dependencies within the workflow. We summarize the notations we are going to use in this section in Table 1 for reference.

4.1 Workload-Balanced Dispatching Policy

The workload-balanced dispatching policy in HEXGEN-FLOW assigns each LLM inference request to the most appropriate LLM serving model instance. This policy jointly considers (i) the execution time of each instance for the incoming LLM inference request and (ii) the current queueing time on each model instance. Additionally, a tunable hyperparameter $\alpha \in [0, 1]$ is introduced to dynamically balance the trade-off between these two factors. We discuss how to tune α in §4.3.

Formulate the inference computation cost. For each LLM inference request $q_{i,j}$, we first estimate its output length via a function $\hat{L}_{\text{out}}(q_{i,j})$ derived from its input length—our implementation is based on the prediction method

Symbol	Description
Q_i	The i^{th} end-to-end Text-to-SQL query
$q_{i,j}$	The j^{th} LLM inference requests for query Q_i
m	A LLM model serving instance
ϕ	LLM request allocation strategy
$L_{\text{in}}(q_{i,j})$	Input length of request $q_{i,j}$
$\hat{L}_{\text{out}}(q_{i,j})$	Estimated output length of request $q_{i,j}$
$t_{i,j}^{m}$	Processing time of request $q_{i,j}$ on instance m
${t_{\text{queue}}}_{i,j}^{m}$	Expected queuing time of $q_{i,j}$ on instance m
${t_{\text{comp}}}_{i',j'}^{m}$	Expected compute time of $q_{i,j}$ on instance m
α	Weight balancing queuing and computing
α^*	Optimal value of hyperparameter α
Θ^m	The local queue of requests on instance m
T_i^{SLO}	Pre-set end-to-end SLO for query Q_i
$t_{i,j}^{\text{SLO}}$	SLO budget allocated for $q_{i,j}$
$\tau_{\text{elapsed}}^{i}$	Time elapsed since arrival of Q_i.
$\tau_{i,j}$	Actual queuing delay for $q_{i,j}$
$U_{i,j}$	Urgency of request $q_{i,j}$

Table 1 Summarization of notations.

introduced by Zheng et al. [32]. Based on the estimation, we can get the predicted computational execution cost of $q_{i,j}$ on a particular model instance m, denoted ${t_{\text{comp}}}_{i,j}^{m}$, by the following equation:

$$ {t_{\text{comp}}}_{i,j}^{m} = {t_{\text{prefil}}}^{m}\left(L_{\text{in}}(q_{i,j})\right) + {t_{\text{decode}}}^{m}\left(\hat{L}_{\text{out}}(q_{i,j})\right) \tag{2} $$

where ${t_{\text{prefil}}}^{m}\left(L_{\text{in}}(q_{i,j})\right)$ and ${t_{\text{decode}}}^{m}\left(\hat{L}_{\text{out}}(q_{i,j})\right)$ denote the estimated execution time of the prefill and decoding phase on model instance m based on number of input tokens and the estimated output tokens.

Formulate the (maximal) queueing cost. The expected queuing time cost of $q_{i,j}$ on instance m, denoted by ${t_{\text{queue}}}_{i,j}^{m}$, is estimated by as the sum of the execution costs of all tasks currently in model instance m's queue Θ^m:

$$ {t_{\text{queue}}}_{i,j}^{m} = \sum_{q_{i',j'} \in \Theta^m} {t_{\text{comp}}}_{i',j'}^{m} \tag{3} $$

Note that ${t_{\text{queue}}}_{i,j}^{m}$ captures the potentially longest time task $q_{i,j}$ could wait before execution begins if it is dispatched to model instance m.

Select the serving model instance. Given the estimation of inference computation time and queuing time, an ideal instance has a low estimation for both of them. However, a linear combination of these two factors is problematic—the execution time is relatively predictable given the LLM inference query, while the queuing time can be aggressively adjusted based on the urgency we implement within each local priority queue at each model instance. Thus, we define the following non-linear combination as the heuristic score:

$$ \text{Score}\left(q_{i,j}, m\right) = (1-\alpha) \cdot \frac{\beta}{{t_{\text{queue}}}_{i,j}^{m}} - \alpha \cdot {t_{\text{comp}}}_{i,j}^{m} \tag{4} $$

For the LLM inference request $q_{i,j}$ and model instance m, $q_{i,j}$ will be dispatched to the instance with the highest score. Notice that there are two parameters in this heuristic score (α and β): α is tuned online, while β is a deployment-specific normalization constant that makes the queue-sensitive term and the computation time term numerically comparable. Concretely, we have $\beta \triangleq C_{\text{ref}} \cdot Q_{\text{ref}}$, where C_{ref} is a reference computation time and Q_{ref} is a reference queuing time. With this definition, the queue-related term can be interpreted as a computation time equivalent reward: when $t_{\text{queue}} = Q_{\text{ref}}$, the reward equals C_{ref}; it increases when the target instance's queue is shorter than the reference level and decreases when it is longer. Since C_{ref} and Q_{ref} are primarily determined by stable deployment properties and change only when the deployment configuration materially changes, we calibrate β once per deployment via few warmed-up queries and then keep it fixed. The weighting factor α determines the degree to which dispatching favors fast execution versus load balancing. When $\alpha = 1$, only execution speed is considered; when $\alpha = 0$, only queue-related signal matters. We determine an optimal α via sliding window-based monitoring and simulations (see §4.3).

4.2 Local Priority Queue

To meet the end-to-end SLO, the local priority queue includes two advanced mechanisms:

SLO budget allocation. To determine the priority of each LLM inference request in the local queue, we allocate a per-request SLO budget $t_{i,j}^{\text{SLO}}$ for $q_{i,j}$ based on both execution cost and the remaining end-to-end deadline:

$$t_{i,j}^{\text{SLO}} = \left(T_i^{\text{SLO}} - \tau_{\text{elapsed}}^i\right) \cdot \frac{\bar{t}_{\text{comp}_{i,j}}}{\sum_{k=j}^{n_i} \bar{t}_{\text{comp}_{i,k}}} \tag{5}$$

where τ_{elapsed}^i is the time elapsed since the arrival of Q_i at the global coordinator, and $\bar{t}_{\text{comp}_{i,k}}$ is the estimated execution time of LLM inference $q_{i,j}$ averaged over all of the model instances, i.e.,

$$\bar{t}_{\text{comp}_{i,j}} = \frac{1}{N} \sum_{m \in \mathbf{M}} {t_{\text{comp}}}_{i,j}^{m} \tag{6}$$

This proportional allocation ensures that more time is budgeted for costlier downstream LLM inference requests. It's worth noting that although the workflow contains dynamic iterations, n_i can be determined using configured bounds when available (e.g., a maximum refinement limit) together with the query's observed progress to estimate the remaining work. Additionally, unlike the preset end-to-end SLO T_i^{SLO}, the per-request SLO budget $t_{i,j}^{\text{SLO}}$ is determined right before the inference request being sent. As shown in Figure 3, $t_{i,j}^{\text{SLO}}$ is recomputed in real time as τ_{elapsed}^i is updated upon completion of $q_{i,j-1}$. This design ensures that Text-to-SQL queries behind schedule receive progressively tighter SLO constraints.

Urgency-based priority ranking. Our urgency metric is inspired by Least-Laxity-First (LLF) style deadline-driven scheduling [33]. For an LLM inference request $q_{i,j}$ dispatched to model instance m, we define the urgency $U_{i,j}$ of $q_{i,j}$ as the difference between its execution cost and the remaining SLO margin:

$$U_{i,j} = {t_{\text{comp}}}_{i,j}^{m} - \left(t_{i,j}^{\text{SLO}} - \tau_{i,j}\right) \tag{7}$$

where ${t_{\text{comp}}}_{i,j}^{m}$ is the estimated execution cost of the task on instance m as we defined in Equation 2, and $\tau_{i,j}$ denotes the actual queuing delay tracked by the local priority queue since $q_{i,j}$ entered the local queue. Higher urgency indicates greater risk of SLO violation. The urgency score is time-aware and updated continuously as a request waits. As $\tau_{i,j}$ increases, the request's urgency increases, which prevents indefinite starvation even in the presence of competing high-urgency requests. To handle estimation errors, any unexpected execution cost will be reflected in the per-request SLO budget $t_{i,j+1}^{\text{SLO}}$ of the subsequent inference request $q_{i,j+1}$. Note that the priority is re-evaluated when a new batch is launched for inference, allowing urgent tasks to be admitted ahead of less critical ones. Whenever a request completes, the model instance m always selects the LLM inference request with the highest urgency from the local queue:

$$q^* = \arg \max_{q_{i,j} \in \Theta^m} U_{i,j} \tag{8}$$

where Θ^m is the current set of queued LLM inference requests at the model instance m. This adaptive strategy prioritizes requests at risk of missing their SLO and accounts for instance-specific execution characteristics.

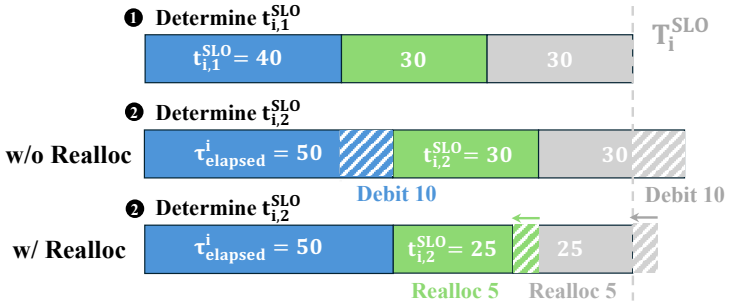

Figure 3 The SLO budget assigned to $q_{i,1}$ is 40ms, but $q_{i,1}$ takes 50ms to finish. The additional 10ms debit is recorded by τ^i_{elapsed} and reallocated to the following inference requests, resulting in a tighter SLO budget for $q_{i,2}$.

4.3 α-Tuning Process

The trade-off determined by α reflects an important trade-off in dispatching. Assigning each request to its most suitable model instance improves per-query efficiency, but can lead to uneven load distribution and queue buildup. On the other hand, balancing load evenly may assign queries to suboptimal instances, increasing latency. The hyperparameter α in the dispatching score function (Equation 4) allows the system to navigate this trade-off between execution cost and queuing cost adaptively. To tune α in response to real-time workload conditions, we implement a lightweight online parameter tuning process based on *tail latency*, i.e., the delays experienced by the slowest requests in a system, typically the 95^{th} to 99^{th} percentile of response time [34].

Initialization. During the cold-start of HEXGEN-FLOW system, the initial α^* is initialized to 0, prioritizing queue length minimization. During the first short period of operation, HEXGEN-FLOW uses this policy to serve incoming Text-to-SQL queries. In parallel, the system collects execution trajectory, including arrival times, queue delays, and stage durations for each LLM inference request. Such information is essential for later simulation-guided auto-tuning.

Sliding window-based monitoring. After initialization, HEXGEN-FLOW assumes workload stationarity over short intervals and continues using the current α^*. The system monitors tail latency in a sliding window. The window length is set to ensure a minimum number of completed end-to-end queries (e.g., 5 samples) for a stable tail-latency estimate, and it adapts to workload changes such that heavier workloads typically result in longer windows. At the end of each sliding window, it computes the 95^{th} percentile latency $\bar{T}_{\text{new}}$ for all queries served in that period and compares it against the baseline latency $\bar{T}_{\text{ref}}$ from the previous window.

While sudden workload shifts (e.g., peak bursts or anomalies) can cause performance degradations, we perform a one-sided two-sample t-test to determine whether latency has degraded significantly:

$$H_0 : \bar{T}_{\text{new}} = \bar{T}_{\text{ref}} \quad \text{vs.} \quad H_1 : \bar{T}_{\text{new}} > \bar{T}_{\text{ref}} \tag{9}$$

If the p-value falls below 0.01, the null hypothesis is rejected, indicating statistically significant latency regression. This triggers a tuning procedure using the most recent trace.

Simulation-guided auto-tuning. Auto-tuning is carried out using a trace-driven simulator that replays historical Text-to-SQL queries and evaluates performance under different α values on the fly. The optimal α^* is determined by minimizing the 95^{th} percentile latency (tail-aligned SLO objective):

$$\alpha^* = \arg\min_{\alpha \in [0,1]} T_{0.95}(\alpha) \tag{10}$$

where $T_{0.95}(\alpha)$ denotes the simulated completion time that 95% queries will be completed under parameter α. The search iterates over $\alpha \in \{0.0, 0.1, 0.2, \ldots, 1.0\}$ in uniform steps.

Discussion of tuning overhead. The simulator executes entirely on the CPU and only incurs negligible overhead cost compared to the actual serving. In practice, tuned α^* values remain stable across adjacent time windows unless there are abrupt changes in workload patterns. This enables robust and low-overhead adaptation to evolving serving conditions, ensuring sustained low-latency performance.

5 Evaluation

In this section, we evaluate the design of HEXGEN-FLOW. We begin by asking the following questions to analyze the end-to-end performance of our framework, as well as each component's contribution towards the overall efficiency:

- *How does the end-to-end performance of* HEXGEN-FLOW *compare to that of state-of-the-art LLM inference serving systems in Text-to-SQL workflows?*
- *How effective is each component during scheduling?*
- *What are the benefits and cost of α-tuning process?*

We state our experiment setup in §5.1, and present the experimental results to answer each question in §5.2, §5.3, and §5.4, respectively.

5.1 Experimental Setup

Runtime. Each end-to-end Text-to-SQL query was processed following a state-of-the-art agentic framework, i.e., CHESS [23] that, at the time of this study, represents the leading approach in Text-to-SQL workflow. We perform evaluations in the following execution environments:

- Hetero-1: This setup consists of two types of GPUs, A100 and A6000, each supporting two model instances, with four GPUs allocated per instance for inference serving.
- Hetero-2: This setup consists of three types of GPUs: A100, L40, and A6000. A100 GPUs are responsible for serving two model instances, while L40 and A6000 GPUs each are responsible for serving one model instance. Each model instance is allocated four GPUs for inference serving.
- Homo: This setup consists solely of A100 GPUs, where the cluster hosts four model instances, each utilizing four GPUs for inference serving.

Note that due to the large number of model parameters, we employ a tensor parallelism degree of eight for serving all model instances using vLLM [16] as the LLM inference engine, and all model instances adopt the continuous batching technique during inference.

Model and dataset. The LLM inference requests following the agentic workflow in CHESS [23] are conducted using the LLAMA3.1-70B model for the main experiments. We follow prior work to generate three workload traces based on the development set from BIRD-bench [35], which is a cross-domain dataset designed specifically for Text-to-SQL evaluation. Our testing traces are sampled from queries related to `financial` and `formula1` databases, incorporating "simple," "moderate", and "challenging" queries defined by BIRD (14%, 72%, 14% respectively). In particular, `Trace 1` samples queries purely from the `financial` database, `Trace 2` samples queries from the `formula1`, and `Trace 3` contains queries from both databases. Depending on the complexity of the query, the workflow may require between zero and ten rounds of revision to refine the SQL query. We count the number of LLM calls per Text-to-SQL query across three trials, observing an average of 19.4, 20.7, and 21.1 calls for `Trace 1`, `Trace 2`, and `Trace 3`, respectively, with corresponding variances of 12.1, 14.7, and 16.3. The variances reflect the different complexity of sampled queries among traces, which widely exists in real-world Text-to-SQL workloads. To emulate the stochastic arrival pattern of users' Text-to-SQL queries, we send queries using a Poisson process with arrival rates of 0.5 query per second and 1.0 query per second (QPS), where each end-to-end query triggers around 20 LLM inference requests on average. This modeling approach captures the inherent randomness and non-stationarity in user interactions and aligns with methodologies employed in prior studies [21]. To further stress the system under heavier workload, we additionally evaluate a 30 QPS setting on `Trace 3` as a stress test. Beyond stationary arrivals, we also construct a fluctuating workload with non-stationary arrival rates (5 QPS to 20 QPS to 5 QPS) to examine how our method adapts to abrupt changes in traffic intensity. We also demonstrate generalizability by employing an alternative agentic workflow, MAC-SQL [36], with the QWEN3-30B model [37]. For this evaluation, we sampled a `Trace 4` from Spider1.0 [38], containing 9 new domain databases including `dog kennels`, `cars`, etc. Finally, to evaluate robustness under hardware sharing, we derive `Trace 5` and `Trace 6`

by augmenting `Trace 3` with 25% and 50% independent single-shot LLM inference requests, each with their own SLOs.

Baseline. To evaluate the end-to-end performance under three deployment setups: Hetero-1, Hetero-2, and Homo, we compare HEXGEN-FLOW with five strong baselines:

- **vLLM**: We compare HEXGEN-FLOW against vanilla vLLM, which is a widely adopted inference serving system that uses first-come-first-served (FCFS) and continuous batching to manage local queues. We dispatch the LLM inference requests to model instances based on the round-robin strategy. This naive approach is commonly utilized in existing inference serving systems [39, 40].
- **VTC**: We include the existing scheduling algorithm Virtual Token Counter (VTC) [17] under all three traces. VTC defines fairness based on a cost function accounting for input/output token counts served for different users. To leverage this design, we view one particular end-to-end Text-to-SQL query as a user in VTC, since each end-to-end Text-to-SQL query is composed of multiple inference requests sent in order. In such a way, VTC can balance the serving for various end-to-end Text-to-SQL queries.
- **QLM**: We also compare HEXGEN-FLOW with a SLO-oriented queue management system QLM [20], which groups inference requests based on their SLOs and uses a request waiting time estimator to estimate the waiting time of each request group in the local queue. If it estimates that the waiting time is too long to violate the group's SLO, all the request groups in the local queue will be reordered to prioritize the group with the least remaining time. Within each request group, inference requests are scheduled according to FCFS.
- **LLF**: We include a Least Laxity First (LLF) [33] baseline that treats each Text-to-SQL query as a "job" with an end-to-end deadline. At runtime, for each stage we schedule inference requests by the job's laxity (i.e., deadline minus elapsed time and remaining processing time), without performing HEXGEN-FLOW's per-request budget recomputation.
- **Ray**: We compare HEXGEN-FLOW with Ray Serve [41], a widely-used general-purpose online serving framework that can express our agentic workflow DAG by chaining deployments. Ray Serve routes requests using power-of-two-choices (Po2C) [42]; however, it does not incorporate cross-stage, end-to-end deadline-aware scheduling strategies.

Additionally, we include two variants of HEXGEN-FLOW in the ablation study to assess the impact of our workload-balanced dispatching (WB) and local priority queue (PQ) independently of other enhancements in HEXGEN-FLOW:

- **RR+PQ**: We consider a truncated version of HEXGEN-FLOW, which replaces the WB mechanism in our global coordinator with a simple round-robin (RR) dispatcher combined with our PQ implementation on each model instance.
- **WB+FCFS**: We also include another ablated version of HEXGEN-FLOW, which reserves the WB mechanism in our global coordinator, while utilizing simple FCFS scheduling on each model instance.

Evaluation metrics. Following the standard setup of existing LLM serving frameworks [43, 44], we evaluate system performance based on SLO attainment and system throughput, where users specify the SLO. For evaluation purposes, we first collect the execution time of each end-to-end Text-to-SQL query in an exclusive environment. During our evaluation, we set the SLO by multiplying the exclusive execution time with various scales (i.e., SLO scale in Figure 4) to assess performance under different levels of operational stringency. Particularly, we are interested in understanding the performance under the setup with minimum SLO scale at which the system achieves 95% SLO attainment (i.e., P95 tail latency). We define the throughput metric as the number of end-to-end Text-to-SQL queries completed per second. We also conduct additional experiments in the multi-tenant scenario where there are two distinct SLOs corresponding to two users.

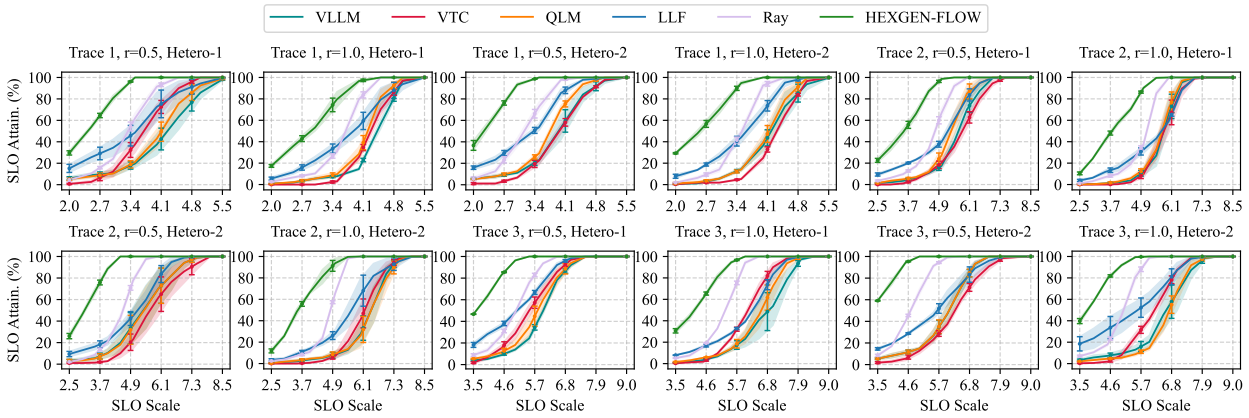

Figure 4 End-to-end SLO attainment comparison. r represents the end-to-end queries arrival rate in queries per second.

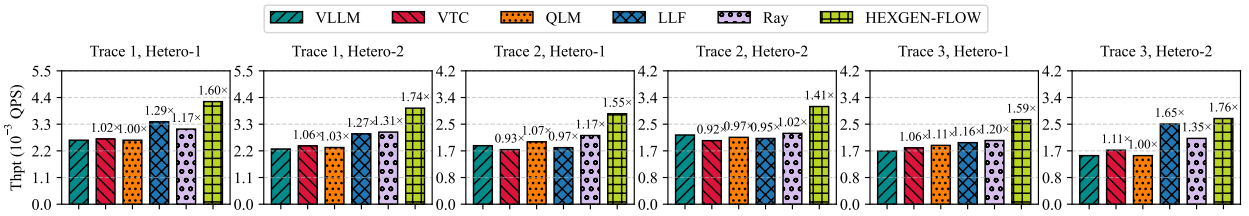

Figure 5 End-to-end system throughput comparison.

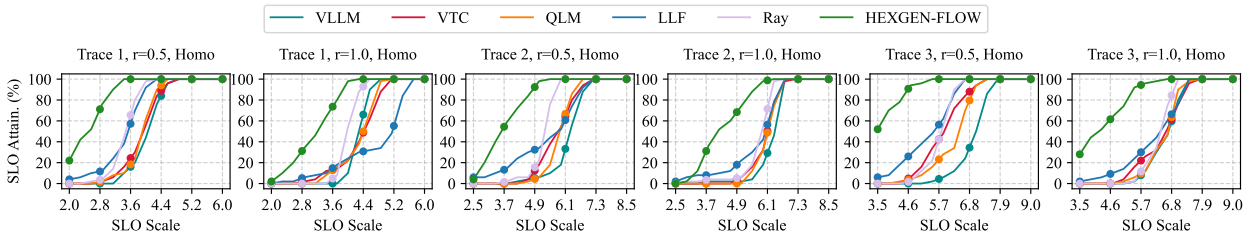

Figure 6 End-to-end experiments of HEXGEN-FLOW under homogeneous settings.

5.2 End-to-End Performance Boost

End-to-end system performance. We evaluate the SLO attainment and throughput of HEXGEN-FLOW across multiple traces, system loads, and heterogeneous hardware configurations, comparing its performance against five baseline systems: vLLM, VTC, QLM, LLF, and Ray. We conduct multiple runs with random seeds $0, 1, 2$, and include error bars in the plot. Figure 4 and Figure 5 demonstrate HEXGEN-FLOW consistently outperforms all baseline systems, concretely:

- Compared with vLLM, HEXGEN-FLOW achieves P95 tail latency reductions ranging from 19.2% to 56.2%, along with throughput improvements ranging from 49.1% to 81.0%. For example, when evaluated on `Trace 1` with Hetero-1 under 0.5 QPS system load, vLLM requires 5.4 SLO scales to achieve 95% attainment, whereas HEXGEN-FLOW requires only 3.5 SLO scales. These results demonstrate that vLLM with naive request dispatching and scheduling methods (FCFS and round robin) fail to provide satisfactory serving performance for complex Text-to-SQL workloads.

- Compared with VTC, HEXGEN-FLOW reduces P95 tail latency by $19.2\%\sim53.8\%$ and increases throughput by $10.3\%\sim39.8\%$. Although VTC provides balanced resource allocation across different Text-to-SQL queries through its fairness-based scheduling mechanism, it fails to account for the heterogeneous query requirements specific to Text-to-SQL workflows, thereby leading to suboptimal system performance. For instance, when evaluated on `Trace 2` with Hetero-2 under 0.5 QPS system load, VTC requires 7.9 SLO scales to achieve 95% attainment, while HEXGEN-FLOW requires only 5.2.

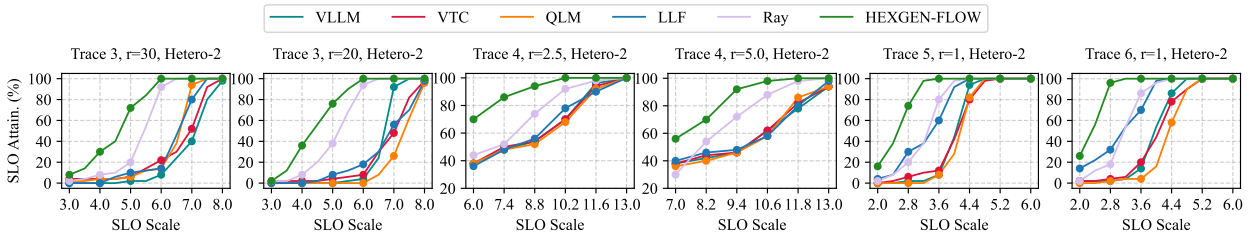

Figure 7 End-to-end SLO attainment comparison. **Left two:** results under a higher QPS and a workload with a sudden shift. **Middle two:** results on the Spider trace. **Right two:** results under 25% and 50% competing workloads. r represents the end-to-end queries arrival rate in queries per second.

- Compared with QLM, HEXGEN-FLOW achieves P95 tail latency reductions between 14.0% and 39.6%, and throughput gains between 17.5% and 36.6%. Although QLM utilizes an estimation-based queue management approach that prioritizes urgent requests to maximize SLO attainment, it applies uniform treatment to all queries without accounting for the diverse computational requirements inherent in Text-to-SQL workflows. This limitation compromises performance optimization. For instance, when evaluated on `Trace 3` with Hetero-2 under 1.0 QPS system load, QLM requires 7.8 SLO scales to achieve 95% attainment, while HEXGEN-FLOW requires only 6.1.
- Compared with LLF, HEXGEN-FLOW achieves P95 tail latency reductions by $12.4\%\sim39.6\%$, and improves throughput by $6.7\%\sim59.8\%$. Although LLF prioritizes ready stage calls by end-to-end job laxity, it does not perform per-request budget recomputation across stages; as a result, it may over- or under-allocate slack to intermediate calls, causing avoidable SLO violations and reduced throughput. For example, on `Trace 3` with Hetero-2 under 1.0 QPS, LLF needs roughly 7.6 SLO scales to reach 95% attainment, while HEXGEN-FLOW reaches 95% at around 5.2.
- Compared with Ray, HEXGEN-FLOW reduces P95 tail latency by $8.6\%\sim23.8\%$ and delivers $30.4\%\sim38.2\%$ higher throughput. Ray Serve's routing strategy neglects the hardware heterogeneity, and its lack of cross-stage deadline-aware scheduling limits the performance. For instance, on `Trace 3` with Hetero-2 under 1.0 QPS, Ray Serve requires approximately 5.9 SLO scales to reach 95% attainment, whereas HEXGEN-FLOW reaches 95% at around 5.2.

Homogeneous environment evaluation. We also evaluate the SLO attainment of HEXGEN-FLOW in homogeneous environments across different traces and system loads. As shown in Figure 6, HEXGEN-FLOW consistently outperforms all baselines in homogeneous configurations. Specifically, across all evaluated scenarios, HEXGEN-FLOW achieves up to 32.6%, 31.4%, 28.6%, 47.3%, and 19.2% lower P95 tail latency compared to vLLM, VTC, QLM, LLF, and Ray respectively. When evaluated on `Trace 3` under 1.0 QPS system load, existing baseline systems require a minimum of 7.2 SLO scales to achieve 95% attainment, whereas HEXGEN-FLOW requires only 5.9 SLO scales. These results demonstrate that HEXGEN-FLOW delivers outstanding system performance even in homogeneous environments, further validating the effectiveness of our approach across diverse hardware configurations.

Evaluation under high and non-stationary arrival rates. As shown in **left two** of Figure 7, under both the heavy (30 QPS throughout) and fluctuating (20 QPS during peak) workloads, HEXGEN-FLOW maintains strong tail-latency control. In particular, HEXGEN-FLOW reduces the P95 tail latency by $5.2\%\sim23.3\%$ compared to all the baselines. Moreover, in the fluctuating workload, α-tuning is triggered twice, matching the two workload shifts and indicating timely adaptation.

Evaluation on alternative workflows and models. We conduct supplementary experiments on MAC-SQL using the QWEN3-30B model. The results demonstrate that our approach maintains consistent improvements across diverse models and tasks, highlighting its potential applicability to other agentic workflows as well. From **middle two** of Figure 7, HEXGEN-FLOW achieves P95 tail latency reductions ranging from 34.0% to 38.3% compared to baseline approaches.

Evaluation in competing workloads scenarios. The **right two** of Figure 7 further evaluate robustness under a

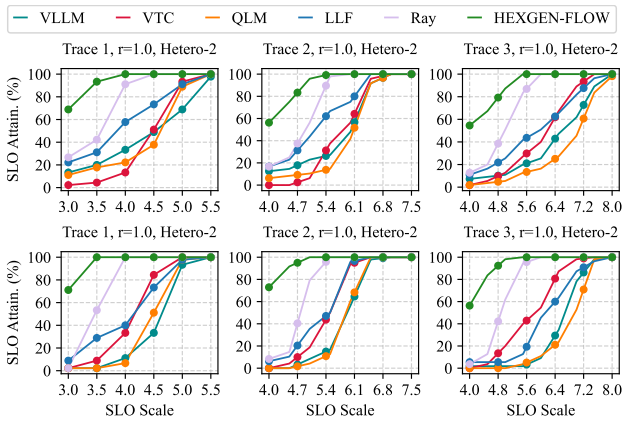

Figure 8 End-to-end SLO attainment comparison for Text-2-SQL queries from two users. The first row corresponds to user 1 and the second row corresponds to user 2. r represents the end-to-end queries arrival rate in queries per second.

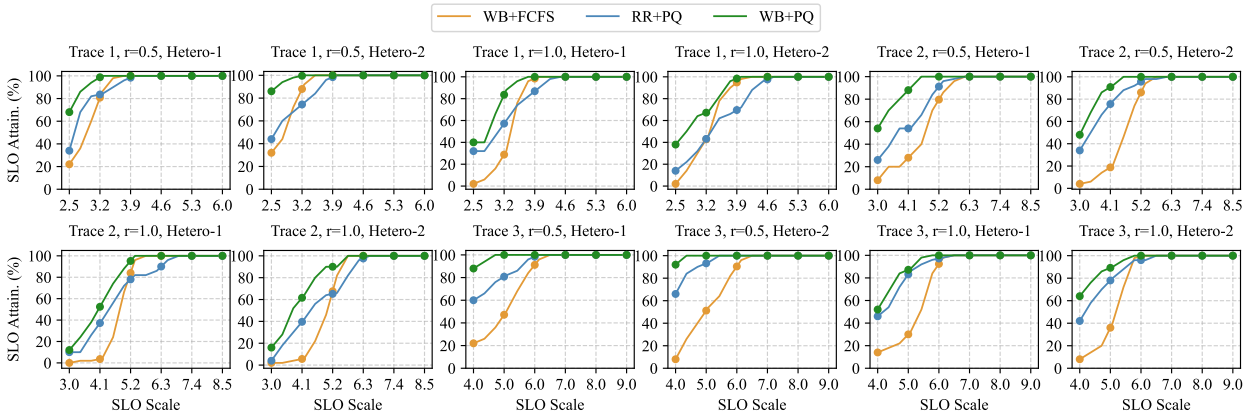

Figure 9 Ablation study of HEXGEN-FLOW's scheduling components, comparing end-to-end SLO attainment rates across: (1) round-robin dispatching + local priority queue, (2) workload-balanced dispatching + First Come First Serve, and (3) workload-balanced dispatching + local priority queue.

broader range of non-Text-to-SQL competing workloads. Across 25% and 50% competing-load mixes, HEXGEN-FLOW reduces P95 tail latency by approximately 18.2% to 36.4% compared to baseline approaches, while sustaining higher SLO attainment across SLO scales.

Evaluation in multi-tenant scenarios. In our design, each LLM inference request's urgency can adapt to various SLOs, allowing our algorithm to integrate seamlessly into multi-tenant scenarios. To evaluate HEXGEN-FLOW in scenarios involving multiple users with heterogeneous SLO requirements, we conduct additional experiments in a setting where two users with distinct SLOs coexist. The results in Figure 8 demonstrate that HEXGEN-FLOW reduces P95 tail latency by $27.9\%\sim35.7\%$ for users with different SLO requirements compared to baseline approaches. Additionally, based on per-tenant SLO attainment rates, we compute Jain's fairness index [45] to be 0.98 across three traces.

5.3 Ablation Study: Effectiveness of Scheduling

To assess the individual contributions of our WB and PQ optimizations, we conduct a systematic assessment by replacing each optimization with a naive baseline implementation (RR or FCFS), and evaluate the resulting SLO attainment. This evaluation spans multiple traces, system loads, and heterogeneous hardware configurations. Note that WB+PQ represents the complete implementation of HEXGEN-FLOW. As shown in Figure 9, the results demonstrate the distinct performance contributions of each optimization component:

- **Effectiveness of WB.** We compare WB+PQ against RR+PQ to evaluate the effectiveness of WB. Experi-

mental results demonstrate that WB consistently outperforms RR, achieving 10.5% to 34.8% lower P95 tail latency. For instance, when evaluated on `Trace 3` with Hetero-1 under 0.5 QPS system load, RR requires 5.8 SLO scales to achieve 95% attainment, while WB requires only 4.3 SLO scales, reducing the tail latency by roughly 35%. These results validate the effectiveness of WB in balancing workloads and resource utilization.

- **Effectiveness of PQ.** We compare WB+PQ against WB+FCFS to evaluate the effectiveness of PQ. Results demonstrate that PQ significantly outperforms FCFS across all evaluated scenarios, achieving up to 48% lower P95 tail latency. For instance, when evaluated on `Trace 2` with Hetero-2 under 0.5 QPS system load, FCFS requires 5.8 SLO scales to achieve 95% attainment, while PQ requires only 4.6. These results further validate the design of PQ, and demonstrate the importance of urgency-based request prioritization in serving Text-to-SQL workloads.

Stage	Dispatcher	I1 (%)	I2 (%)	I3 (%)	I4 (%)
1	RR	20.8	23.1	19.9	26.8
	WB	0.9	0.5	6.1	32.3
2	RR	10.3	7.8	10.6	9.6
	WB	27.9	29.6	16.5	3.1
3	RR	26.5	23.0	25.5	21.7
	WB	47.9	46.6	71.5	8.3
4	RR	42.3	46.0	44.1	41.8
	WB	23.3	23.3	5.8	56.3

Table 2 Impact of dispatching policy on task distributions. Stages 1-4 represent the aforementioned Text-to-SQL stages; I1–I4 represent different GPU instances, where I1 and I2 are A100, I3 is an A6000, and I4 is an L40 instance.

Case study of WB impact on request dispatching. To demonstrate how WB improves request dispatching, we analyze task allocation across model instances with and without WB on `Trace 3` with Hetero-2 under 0.5 QPS system load. As shown in Table 2, unlike RR which distributes tasks evenly across GPU instances, WB strategically allocates computationally intensive stages to more powerful instances (e.g., A100) while assigning lighter stages to less capable ones (e.g., L40). For example, stage 2 and stage 1 requests are evaluated as computationally intensive and light, respectively. Under WB, A100 instances collectively handle 57.5% of stage 2 requests compared to only 18.1% under RR, while the L40 instance processes 3.1% of stage 2 requests versus 9.6% under RR. Conversely, WB allocates 32.3% of stage 1 requests to the L40 instance compared to 26.8% under RR, while allocating only 1.4% of stage 1 requests to A100 instances compared to 43.9% under RR. As a result, WB ensures optimal resource utilization by matching computational requirements with hardware capabilities, thereby enhancing overall performance.

Case study of PQ impact on request processing. To demonstrate how PQ optimizes request processing within local queues, we examine a model instance's local queue state before processing the next LLM call, as shown in Table 3. Under FCFS, the system will process request 1, which arrived earliest at 22.4 seconds. However, this approach disregards the varying urgency levels of requests, potentially leading to SLO violations when early-arriving requests with loose SLO budgets are prioritized over later-arriving requests with stricter SLO requirements. In contrast, PQ prioritizes requests based on urgency values instead of arrival times. In this scenario, PQ will select request 6 to process next, which has the highest urgency value of 26.9 despite arriving at 64.4 seconds. This intelligent prioritization ensures the most time-critical requests are processed first, thereby maximizing the likelihood of meeting SLO requirements across requests in the queue.

Request ID	1	2	3	4	5	6	7
Arrive-at (s)	22.4	46.3	52.3	62.4	62.8	64.4	65.0
Urgency	14.5	13.2	19.0	13.1	19.0	26.9	21.9

Table 3 Snapshot of a local queue before processing next LLM call. Arrive-at represents the timestamp of a LLM call arrival.

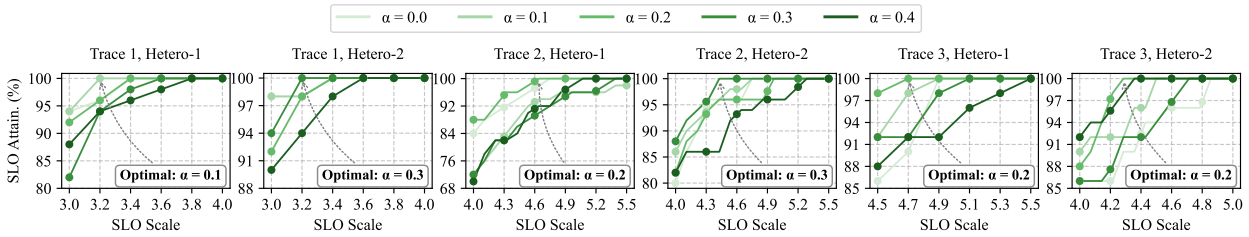

Figure 10 Performance of HEXGEN-FLOW under various α settings.

5.4 Empirical Analysis of α-Tuning

We conduct experiments to estimate how dynamic tuning of the α parameter in HEXGEN-FLOW 's scheduling policy could accommodate hardware heterogeneity and workload variability. Through experiments across all three workloads (`Traces 1,2,3` and two heterogeneous GPU deployments (Hetero-1 and Hetero-2), we evaluate: (i) the impact of different α values on 95% SLO attainment, and (ii) the feasibility of simulation-based tuning during live serving. All experiments are conducted under a fixed query arrival rate of 0.5 QPS, while varying α from 0 (pure workload balancing) to 0.5 (balanced weighting).

Effectiveness of α-tuning. Figure 10 shows how different α values affect 95% SLO attainment across traces and deployment configurations. Concretely, our plots focused on the $[0, 0.5]$ interval for readability purposes, as the optimum lies within that range. For `Trace 1`, $\alpha = 0.1$ yields the best performance under Hetero-1, while $\alpha = 0.3$ performs optically under Hetero-2. Similarly, the results in `Trace 2` indicate optimal values when $\alpha = 0.2$ for Hetero-1, and $\alpha = 0.3$ for Hetero-2. In `Trace 3`, $\alpha = 0.2$ is optimal for Hetero-1, while $\alpha = 0.4$ performs best under Hetero-2, improving 95% SLO attainment latency by 14% over $\alpha = 0$. These results suggest that while workload balancing is a dominant factor, incorporating task suitability (via non-zero α) consistently improves scheduling effectiveness. Moreover, the optimal α varies depending on both hardware configuration and workload composition, indicating that adaptive tuning is essential.

Table 4 Overhead of alpha-tuning.

Setup	Trace 1	Trace 2	Trace 3
(Hetero-1, 0.5QPS)	2.5s	2.4s	2.8s
(Hetero-1, 1.0QPS)	2.7s	2.7s	2.9s
(Hetero-2, 0.5QPS)	1.9s	2.8s	3.1s
(Hetero-2, 1.0QPS)	2.3s	2.7s	3.1s

Overhead of α-tuning. We also measure the overhead of simulation-based α tuning. Under 0.5 and 1.0 QPS, each sliding window lasts for 100 seconds. As shown in Table 4, through parallelized simulations, it takes only $1.9\sim3.1$ seconds per window across different settings. Moreover, tuning is triggered only when a performance degradation is detected in a window; the trigger rate ranged from $28.6\%\sim40.0\%$ across traces. Under stress tests with 20 and 30 QPS load, the window lasts for 300 seconds, while the tuning takes $3.9\sim6.6$ seconds per window, and only $28.6\sim33.3\%$ of windows are triggered. Notice this overhead is minor relative to typical workload drift timescales, which span hours [44], making periodic tuning feasible in practice.

6 Related Work

In this section, we provide a brief summary about the recent advances in LLM-based Text-to-SQL approaches in §6.1 and then discuss some serving system design and implementation for LLM inference requests in §6.2.

6.1 LLM Advances for Text-to-SQL

LLMs have emerged as a revolutionary paradigm for the Text-to-SQL task [1, 46], where the core technique lies in effective SQL generation and schema linking.

SQL generation. Some pioneer studies focus on better *prompting strategies and LLM adaptations* to boost Text-to-SQL performance. Gu et al. [47] propose a structure- and content-based prompt learning method to enhance few-shot Text-to-SQL translation, while Li et al. [48] build an open-source LLM specialized for SQL generation. Other approaches fine-tune or constrain LLM outputs to improve accuracy: Ren et al. [49] propose Purple, which refines a general LLM to make it a more effective SQL writer. Beyond simple LLM adaptations, recent *agentic approaches* leverage multiple LLM inference requests to collaboratively accomplish the Text-to-SQL tasks. Fan et al. [50] combine a small language model with a large one to handle different sub-tasks of Text-to-SQL, thereby improving zero-shot robustness. Similarly, Pourreza et al. [51] explore task decomposition across models: DTS-SQL breaks the problem into stages handled by smaller LLMs sequentially, and their DIN-SQL approach has an LLM refine its own output through iterative self-correction [52] in the prompt. Another line of research enhances the reasoning process of LLMs to produce correct SQL. One strategy is to incorporate intermediate steps or a reasoning framework during inference. Zhang et al. [53] apply the ReAct paradigm [54] to table question answering, which encourages the LLM to generate and reason with intermediate actions (e.g., decomposition or calculations) before finalizing the SQL query. Techniques have also been explored to optimize the context given to the LLM: Talaei et al. [23] present CHESS, which harnesses contextual information efficiently to guide the LLM's SQL synthesis without increasing model size.

Schema linking. Integrating database schema knowledge and domain-specific information into LLM-driven Text-to-SQL is another important procedure. Eyal et al. [55] decompose the natural language question and the corresponding SQL query into semantic sub-units, improving the model's understanding of how question clauses align with schema elements. Dou et al. [56] incorporate external knowledge (e.g., business rules or formulas) into the parsing process to handle queries that require facts beyond the database content. Several works specifically target schema linking challenges in the age of LLMs. Liu et al. [57] propose Solid-SQL, which uses enhanced in-context examples to make an LLM more robust at matching question terms to the schema during generation.

6.2 LLM Inference Request Scheduling

Efficient scheduling of LLM inference requests is crucial in modern AI infrastructure, essential to meeting latency and throughput requirements, particularly under varying system constraints [19, 22, 58–60]. In environments with *consistent hardware and model setups*, scheduling techniques focus on optimizing latency and throughput. Patke et al. [20] introduce QLM, a system that estimates request waiting times to prioritize tasks, ensuring that SLOs are met under load conditions. Gong et al. [61] propose the future scheduler, which uses historical workload patterns and predictive modeling to make informed scheduling decisions, fulfilling SLA guarantees. Fu et al. [58] frame LLM scheduling as a learning-to-rank problem, training models to order queued requests to optimize end-to-end latency and throughput, outperforming traditional heuristics. Agrawal et al. [19] present Sarathi-Serve, a system that adjusts batching and resource allocation to balance throughput and latency, particularly effective for requests of high priority versus low priority. In setups with varying hardware capabilities and model types, recently proposed scheduling strategies adapt to resource *heterogeneity*: Wan et al. [62] develop BROS, a system that differentiates between real-time and best-effort LLM queries, ensuring interactive queries are prioritized without compromising background batch processing. Sun et al. [63] introduce Llumnix, a dynamic scheduling system that adjusts resource allocation for LLM serving in real time as query load patterns change, demonstrating benefits in single-model scenarios. Very recently, Fang et al. [64] investigate the efficiency of multi-LLM application pipelines in an offline setting, using sampling and simulation to optimize inference plans for workflows involving multiple LLMs or sequential model calls. Despite these efforts, there is a notable gap in scheduling strategies that coordinate end-to-end pipelines with multiple dependent LLM request serving stages. Our approach builds upon existing insights, such as urgency-aware prioritization and hardware-sensitive allocation, but effectively extends them to manage complex, multi-stage Text-to-SQL workflows under strict latency requirements and system heterogeneity with a significant performance boost.

7 Conclusion

In this paper, we introduced HEXGEN-FLOW, a novel inference-serving framework specifically designed to efficiently schedule and execute multi-stage agentic Text-to-SQL workloads in heterogeneous GPU clusters

handling multi-tenant requests. By systematically analyzing the unique characteristics of these workloads-including stage dependencies, task variability, resource heterogeneity, and stringent latency constraints—we formulated a hierarchical scheduling solution composed of a global workload-balanced dispatcher and a local urgency-driven queue in each LLM serving model instances. This *two-level* scheduling strategy significantly improves resource utilization and effectively mitigates deadline violations. Comprehensive experiments demonstrated that HEXGEN-FLOW consistently outperforms the existing state-of-the-art LLM serving framework, achieving substantial reductions in query latency and notable improvements in system throughput under realistic production workloads—HEXGEN-FLOW reduces P95 tail latency by $1.42\sim1.56\times$ and increases throughput by $1.49\sim1.81\times$ compared to state-of-the-art systems under diverse, realistic workload conditions. We view extending and validating the framework for other agentic DAG workflows as important future work.

References

[1] Dawei Gao, Haibin Wang, Yaliang Li, Xiuyu Sun, Yichen Qian, Bolin Ding, and Jingren Zhou. Text-to-sql empowered by large language models: A benchmark evaluation. Proceedings of the VLDB Endowment, 17(5):1132–1145, 2024.

[2] Boyan Li, Yuyu Luo, Chengliang Chai, Guoliang Li, and Nan Tang. The dawn of natural language to sql: Are we fully ready? Proceedings of the VLDB Endowment, 17(11):3318–3331, 2024.

[3] Shuaichen Chang, Jun Wang, Mingwen Dong, Lin Pan, Henghui Zhu, Alexander Hanbo Li, Wuwei Lan, Sheng Zhang, Jiarong Jiang, Joseph Lilien, et al. Dr. spider: A diagnostic evaluation benchmark towards text-to-sql robustness. In The Eleventh International Conference on Learning Representations.

[4] Han Fu, Chang Liu, Bin Wu, Feifei Li, Jian Tan, and Jianling Sun. Catsql: Towards real world natural language to sql applications. Proceedings of the VLDB Endowment, 16(6):1534–1547, 2023.

[5] OpenAI. Openai gpt-4o, 2024.

[6] Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy Lillicrap, Jean-baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittwieser, et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. arXiv preprint arXiv:2403.05530, 2024.

[7] Anthropic. The claude 3 model family: Opus, sonnet, haiku, 2024.

[8] Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. arXiv preprint arXiv:2401.04088, 2024.

[9] Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. arXiv preprint arXiv:2307.09288, 2023.

[10] Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling llm test-time compute optimally can be more effective than scaling model parameters. arXiv preprint arXiv:2408.03314, 2024.

[11] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. arXiv preprint arXiv:2407.21787, 2024.

[12] Nvidia. Nvidia reinvents computer graphics with turing architecture, 2018.

[13] Nvidia. Nvidia's new ampere data center gpu in full production, 2020.

[14] Nvidia. Nvidia announces hopper architecture, the next generation of accelerated computing, 2022.

[15] Nvidia. Nvidia blackwell platform arrives to power a new era of computing, 2024.

[16] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In Proceedings of the 29th Symposium on Operating Systems Principles, pages 611–626, 2023.

[17] Ying Sheng, Shiyi Cao, Dacheng Li, Banghua Zhu, Zhuohan Li, Danyang Zhuo, Joseph E Gonzalez, and Ion Stoica. Fairness in serving large language models. In 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24), pages 965–988, 2024.

[18] Archit Patke, Dhemath Reddy, Saurabh Jha, Haoran Qiu, Christian Pinto, Chandra Narayanaswami, Zbigniew Kalbarczyk, and Ravishankar Iyer. Queue management for slo-oriented large language model serving. In Proceedings of the 2024 ACM Symposium on Cloud Computing, pages 18–35, 2024.

[19] Amey Agrawal, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav Gulavani, Alexey Tumanov, and Ramachandran Ramjee. Taming {Throughput-Latency} tradeoff in {LLM} inference with {Sarathi-Serve}. In 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24), pages 117–134, 2024.

[20] Archit Patke, Dhemath Reddy, Saurabh Jha, Haoran Qiu, Christian Pinto, Chandra Narayanaswami, Zbigniew Kalbarczyk, and Ravishankar Iyer. Queue management for slo-oriented large language model serving. In Proceedings of the 2024 ACM Symposium on Cloud Computing, pages 18–35, 2024.

[21] Zhibin Wang, Shipeng Li, Yuhang Zhou, Xue Li, Rong Gu, Nguyen Cam-Tu, Chen Tian, and Sheng Zhong. Revisiting slo and goodput metrics in llm serving. arXiv preprint arXiv:2410.14257, 2024.

[22] Kunal Jain, Anjaly Parayil, Ankur Mallick, Esha Choukse, Xiaoting Qin, Jue Zhang, Íñigo Goiri, Rujia Wang, Chetan Bansal, Victor Rühle, et al. Performance aware llm load balancer for mixed workloads. In Proceedings of the 5th Workshop on Machine Learning and Systems, pages 19–30, 2025.

[23] Shayan Talaei, Mohammadreza Pourreza, Yu-Chen Chang, Azalia Mirhoseini, and Amin Saberi. Chess: Contextual harnessing for efficient sql synthesis. arXiv preprint arXiv:2405.16755, 2024.

[24] HuggingFace. Text generation inference. https://huggingface.co/docs/text-generation-inference/index, 2023.

[25] Nvidia. Tensorrt-llm. https://github.com/NVIDIA/TensorRT-LLM, 2023.

[26] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for {Transformer-Based} generative models. In 16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22), pages 521–538, 2022.

[27] Nvidia. Nvidia triton inference server, 2025.

[28] Yixuan Mei, Yonghao Zhuang, Xupeng Miao, Juncheng Yang, Zhihao Jia, and Rashmi Vinayak. Helix: Serving large language models over heterogeneous gpus and network via max-flow. In Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1, pages 586–602, 2025.

[29] Youhe Jiang, Ran Yan, Xiaozhe Yao, Yang Zhou, Beidi Chen, and Binhang Yuan. Hexgen: Generative inference of large language model over heterogeneous environment. In International Conference on Machine Learning, pages 21946–21961. PMLR, 2024.

[30] Youhe Jiang, Ran Yan, and Binhang Yuan. Hexgen-2: Disaggregated generative inference of llms in heterogeneous environment. arXiv preprint arXiv:2502.07903, 2025.

[31] U Narayan Bhat. An introduction to queueing theory: modeling and analysis in applications, volume 36. Springer, 2008.

[32] Zangwei Zheng, Xiaozhe Ren, Fuzhao Xue, Yang Luo, Xin Jiang, and Yang You. Response length perception and sequence scheduling: An llm-empowered llm inference pipeline. Advances in Neural Information Processing Systems, 36:65517–65530, 2023.

[33] Sung-Heun Oh and Seung-Min Yang. A modified least-laxity-first scheduling algorithm for real-time tasks. In Proceedings Fifth International Conference on Real-Time Computing Systems and Applications (Cat. No.98EX236), pages 31–36. IEEE, 1998.

[34] Khelifa Saif Eddine, Miloud Bagaa, Sihem Ouahouah, Messaoud Ahmed Ouameur, and Adlen Ksentini. Tail-latency aware scheduler for inference workloads. In 2025 International Wireless Communications and Mobile Computing (IWCMC), pages 679–684. IEEE, 2025.

[35] Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. Advances in Neural Information Processing Systems, 36:42330–42357, 2023.

[36] Bing Wang, Changyu Ren, Jian Yang, Xinnian Liang, Jiaqi Bai, LinZheng Chai, Zhao Yan, Qian-Wen Zhang, Di Yin, Xing Sun, and Zhoujun Li. Mac-sql: A multi-agent collaborative framework for text-to-sql. In Proceedings of the

31st International Conference on Computational Linguistics, pages 540–557, Abu Dhabi, UAE, 2025. Association for Computational Linguistics.

[37] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, et al. Qwen3 technical report. arXiv preprint arXiv:2505.09388, 2025.

[38] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, Zilin Zhang, and Dragomir Radev. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-sql task. In Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing, pages 3911–3921, Brussels, Belgium, 2018. Association for Computational Linguistics.

[39] Douglas G. Down and Rong Wu. Multi-layered round robin routing for parallel servers. Queueing Systems, 53(4):177–188, August 2006.

[40] Isaac Grosof, Ziv Scully, and Mor Harchol-Balter. Load balancing guardrails: Keeping your heavy traffic on the road to low response times. arXiv preprint arXiv:1905.03439, 2019.

[41] Ray Project Contributors. Ray Serve: Scalable and Programmable Serving. Online documentation (Ray Serve), 2026. Accessed: 2026-01-26.

[42] Yossi Azar, Andrei Z. Broder, Anna R. Karlin, and Eli Upfal. Balanced allocations. SIAM Journal on Computing, 29(1):180–200, 1999.

[43] Zhuohan Li, Lianmin Zheng, Yinmin Zhong, Vincent Liu, Ying Sheng, Xin Jin, Yanping Huang, Zhifeng Chen, Hao Zhang, Joseph E Gonzalez, et al. {AlpaServe}: Statistical multiplexing with model parallelism for deep learning serving. In 17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23), pages 663–679, 2023.

[44] Yinmin Zhong, Shengyu Liu, Junda Chen, Jianbo Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, and Hao Zhang. {DistServe}: Disaggregating prefill and decoding for goodput-optimized large language model serving. In 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24), pages 193–210, 2024.

[45] Raj Jain, Dah-Ming Chiu, and William Hawe. A quantitative measure of fairness and discrimination for resource allocation in shared computer systems. arXiv:cs/9809099, 1998. Accessed: 2026-01-26.

[46] George Katsogiannis-Meimarakis and Georgia Koutrika. A survey on deep learning approaches for text-to-sql. The VLDB Journal, 32(4):905–936, 2023.

[47] Zihui Gu, Ju Fan, Nan Tang, Lei Cao, Bowen Jia, Sam Madden, and Xiaoyong Du. Few-shot text-to-sql translation using structure and content prompt learning. Proceedings of the ACM on Management of Data, 1(2):1–28, 2023.

[48] Haoyang Li, Jing Zhang, Hanbing Liu, Ju Fan, Xiaokang Zhang, Jun Zhu, Renjie Wei, Hongyan Pan, Cuiping Li, and Hong Chen. Codes: Towards building open-source language models for text-to-sql. Proceedings of the ACM on Management of Data, 2(3):1–28, 2024.

[49] Tonghui Ren, Yuankai Fan, Zhenying He, Ren Huang, Jiaqi Dai, Can Huang, Yinan Jing, Kai Zhang, Yifan Yang, and X Sean Wang. Purple: Making a large language model a better sql writer. In 2024 IEEE 40th International Conference on Data Engineering (ICDE), pages 15–28. IEEE, 2024.

[50] Ju Fan, Zihui Gu, Songyue Zhang, Yuxin Zhang, Zui Chen, Lei Cao, Guoliang Li, Samuel Madden, Xiaoyong Du, and Nan Tang. Combining small language models and large language models for zero-shot nl2sql. Proceedings of the VLDB Endowment, 17(11):2750–2763, 2024.

[51] Mohammadreza Pourreza and Davood Rafiei. Dts-sql: Decomposed text-to-sql with small large language models. In Findings of the Association for Computational Linguistics: EMNLP 2024, pages 8212–8220, 2024.

[52] Mohammadreza Pourreza and Davood Rafiei. Din-sql: Decomposed in-context learning of text-to-sql with self-correction. Advances in Neural Information Processing Systems, 36:36339–36348, 2023.

[53] Yunjia Zhang, Jordan Henkel, Avrilia Floratou, Joyce Cahoon, Shaleen Deep, and Jignesh M Patel. Reactable: Enhancing react for table question answering. Proceedings of the VLDB Endowment, 17(8):1981–1994, 2024.

[54] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik R Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In The Eleventh International Conference on Learning Representations.

[55] Ben Eyal, Amir Bachar, Ophir Haroche, Moran Mahabi, and Michael Elhadad. Semantic decomposition of question and sql for text-to-sql parsing. In 2023 Findings of the Association for Computational Linguistics: EMNLP 2023, pages 13629–13645. Association for Computational Linguistics (ACL), 2023.

[56] Longxu Dou, Yan Gao, Xuqi Liu, Mingyang Pan, Dingzirui Wang, Wanxiang Che, Dechen Zhan, Min-Yen Kan, and Jian-Guang Lou. Towards knowledge-intensive text-to-sql semantic parsing with formulaic knowledge. In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, pages 5240–5253, 2022.

[57] Geling Liu, Yunzhi Tan, Ruichao Zhong, Yuanzhen Xie, Lingchen Zhao, Qian Wang, Bo Hu, and Zang Li. Solid-sql: Enhanced schema-linking based in-context learning for robust text-to-sql. In Proceedings of the 31st International Conference on Computational Linguistics, pages 9793–9803, 2025.

[58] Yichao Fu, Siqi Zhu, Runlong Su, Aurick Qiao, Ion Stoica, and Hao Zhang. Efficient llm scheduling by learning to rank. In The Thirty-eighth Annual Conference on Neural Information Processing Systems.

[59] Shihong Gao, Xin Zhang, Yanyan Shen, and Lei Chen. Apt-serve: Adaptive request scheduling on hybrid cache for scalable llm inference serving. arXiv preprint arXiv:2504.07494, 2025.

[60] Ruicheng Ao, Gan Luo, David Simchi-Levi, and Xinshang Wang. Optimizing llm inference: Fluid-guided online scheduling with memory constraints. arXiv preprint arXiv:2504.11320, 2025.

[61] Ruihao Gong, Shihao Bai, Siyu Wu, Yunqian Fan, Zaijun Wang, Xiuhong Li, Hailong Yang, and Xianglong Liu. Past-future scheduler for llm serving under sla guarantees. In Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2, pages 798–813, 2025.

[62] Wan Borui, Zhao Juntao, Jiang Chenyu, Guo Chuanxiong, and Wu Chuan. Efficient llm serving on hybrid real-time and best-effort requests. arXiv preprint arXiv:2504.09590, 2025.

[63] Biao Sun, Ziming Huang, Hanyu Zhao, Wencong Xiao, Xinyi Zhang, Yong Li, and Wei Lin. Llumnix: Dynamic scheduling for large language model serving. In 18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24), pages 173–191, 2024.

[64] Jingzhi Fang, Yanyan Shen, Yue Wang, and Lei Chen. Improving the end-to-end efficiency of offline inference for multi-llm applications based on sampling and simulation. arXiv preprint arXiv:2503.16893, 2025.