Bridging Business Intent and Data: A Benchmark for Automatic Relational Data Product Generation

Faisal Chowdhury*, **Sola Shirai***, **Sarthak Dash***,
Nandana Mihindukulasooriya, **Horst Samulowitz**
IBM, Yorktown Heights, USA
{mchowdh, sdash, samulowitz}@us.ibm.com, {solashirai, nandana}@ibm.com

arXiv:2512.15798v2 [cs.DB] 25 Aug 2026

Abstract

A data product is designed to address a specific business need by transforming raw data into a curated, usable asset that delivers actionable insights. Despite practical advances in related areas like text-to-SQL and ELT pipelines, there is no comprehensive benchmark for evaluating the end-to-end process of automatically generating such data products from high-level business requests. To fill this gap, we introduce DP-Bench, a first-of-its-kind benchmark for automatic data product creation, built by combining elements from existing ELT and text-to-SQL datasets. We also propose baseline methods using LLMs to generate data products, providing a foundation for future research in bridging natural language business intent and structured data representation. We make the DP-Bench dataset and code available at: https://huggingface.co/datasets/ibm-research/dp-bench.

1 Introduction

The concept of a *data product* was first introduced by Patil (2012) as "a product that facilitates an end goal through the use of data." More recently, it has gained renewed attention as a foundational component of the data mesh paradigm, a decentralized data architecture that organizes data ownership around domain-specific business teams (e.g., marketing, sales, and customer support). In this paradigm, each team is responsible for developing and maintaining high-quality data products that are discoverable, trustworthy, secure, and tailored to specific business needs.

A data product is typically built using a product-centric mindset, emphasizing usability and alignment with user requirements. It represents a curated and structured data asset designed to address a specific analytical objective or business use case, such as predicting customer churn or analyzing sales trends. When implemented effectively, data products can significantly reduce the manual effort involved in data preparation, improve data reliability, and deliver measurable business impact.

Despite significant advances in related areas—such as text-to-SQL, schema linking, and ELT (Extract, Load, Transform) pipeline generation—existing approaches focus on isolated components of the data lifecycle. In contrast, real-world data product creation is inherently holistic: it begins with a high-level business request and culminates in a structured dataset that includes relevant schema elements, derived features, and explicit transformation logic.

However, there is currently no standardized benchmark for evaluating systems that perform this end-to-end transformation. Existing benchmarks primarily target subproblems, such as generating SQL queries from text or constructing data integration pipelines, and therefore fail to capture the broader challenge of translating business intent into reusable data assets.

To address this gap, we introduce **DP-Bench** (Data Product Benchmark), a benchmark specifically designed to evaluate automatic data product generation systems. DP-Bench consists of:

- Natural language descriptions of business use cases, termed *Data Product Requests* (DPRs),
- Corresponding data products, defined as structured subsets of database schemas including both base and derived columns,
- Underlying database schemas, transformation provenance (expressed in SQL), and topic annotations linking DPRs to their associated data products.

Each component of the benchmark is carefully curated and validated by human annotators. The dataset is constructed by integrating and extending existing resources, including ELT-Bench (Jin et al., 2025) and BIRD (Li et al., 2024) benchmarks.

In addition, we propose a number of baseline

*Equal contribution.

approaches for automatic data product generation.

While data products naturally evolve over time, the initial construction phase remains the most complex and resource-intensive. By providing a standardized benchmark and baseline methods, we aim to facilitate research on automating this crucial step and advancing the state of the art in aligning natural language business intent with structured data representations.

We refer to our Limitations section regarding further boundaries of our work and aspects of data product creation beyond this paper's scope.

2 Formal Task Definition

We formalize the task of automatic data product creation as a structured prediction problem over database schemas given a natural language data product request (DPR).

2.1 Data Product Request — DPR

Formally, let $r \in \mathcal{R}$ denote a *Data Product Request (DPR)*. A DPR r is a **natural language description** capturing (1) a business problem, analytical objective, or decision context, (2) the analytical goals or information needs of users, and (3) implicit requirements for relevant data entities and transformations. r serves as the **primary input** to the data product creation process, guiding the selection of relevant tables and columns, derived attributes, transformations, and provenance.

A DPR satisfies the following properties: it is **goal-oriented**, specifying *what* needs to be achieved rather than *how* to compute it; it **abstracts** a set of possible analytical queries or tasks into a unified business intent; and it is **groundable** by mapping to elements of an underlying data schema.

Thus, a DPR can be viewed as a *high-level semantic contract* between business intent and data representation, enabling systems to generate queryable data products aligned with user needs.

2.2 Data Product

Let $\{DB\}$ denote a relational database associated with a schema $\mathcal{S}$ consisting of a set of tables, columns, primary keys, and foreign key relations.

A *data product* DP corresponding to a request r is defined as a tuple:

$$DP = (T, C, C_d, P)$$

where T is a set of selected tables from the schema, C is a set of selected (non-derived) columns from the tables in T, C_d is a set of derived columns constructed from base columns in C, and P is a set of provenance functions. Each $p \in P$ is uniquely associated with one derived column $c_d \in C_d$ and defines the transformation used to compute it from the base columns (e.g., expressed as a SQL query).

2.3 Problem Formulation

Given a data product request r and a database DB (with schema $\mathcal{S}$), the goal is to learn a function:

$$f : (r, \mathcal{S}) \rightarrow \widehat{DP}$$

that produces a predicted data product $\widehat{DP}$.

The predicted data product $\widehat{DP}$ is evaluated against a ground-truth data product DP^* using multiple criteria. **Schema Selection Accuracy** calculates overlap between selected columns and tables in $\widehat{DP}$ versus DP^*. **Derived Column Correctness** is calculated using similarity between generated and ground-truth derived column definitions (e.g., via SQL or AST similarity). Lastly, **Downstream Task Performance** measures the ability of $\widehat{DP}$ to support analytical tasks such as text-to-SQL query answering.

Formally, the objective is to maximize a composite utility function:

$$\max_{f} \quad \mathcal{L}(\widehat{DP}, DP^*)$$

where $\mathcal{L}$ aggregates the evaluation metrics described above.

The above formulation generalizes related problems such as schema linking and text-to-SQL, making it a more expressive and challenging task.

2.4 Constraints

The generated data product must satisfy the following constraints: **Schema Validity**, where all selected tables and columns must exist in the input schema $\mathcal{S}$; **Transformation Validity**, in which derived columns must be computable from available base columns; and **Provenance Availability**, where each derived column must be associated with a valid transformation function.

2.5 Example

Figure 1 contains an input database with multiple tables and a DPR describing analytical needs over mobile applications, such as understanding app performance, reviews, and installation trends.

Database

Playstore

	column_name	column_description
1	App	Name of app
2	Translated_Review	User review (Preprocessed and translated to English)
3	Sentiment	Overall user rating of the app (as when scraped)
4	Sentiment_Polarity	Sentiment polarity score
5	Sentiment_Subjectivity	Sentiment subjectivity score

User Reviews

	column_name	column_description
1	App	Application name
2	Category	Category the app belongs to
3	Rating	Overall user rating of the app (as when scraped)
4	Reviews	Number of user reviews for the app (as when scraped)
5	Size	Size of the app (as when scraped)
6	Installs	Number of user downloads/installs for the app (as when scraped)
7	Type	Paid or Free
8	Price	Price of the app (as when scraped)
9	Content_Rating	Age group the app is targeted at - Children / Mature 21+ / Adult
10	Genres	An app can belong to multiple genres (apart from its main category).

Data Product Request

Collect data that will allow queries on app downloads, installation requirements, sentiments and reviews. It should include information about sentiment ratios, subjectivity scores, and comment reviews, allowing for evaluation of app performance and user attitudes. It should also support assessment of positive and negative sentiments, as well as neutral attitudes, providing insights into app reputation and user feedback.

Figure 1: Example of an input DB and a data product request for this DB.

Data Product

column_name	column_description	provenance
App	Name of the app	
avg_sentiment_subjectivity_score	The average sentiment subjectivity score of the app	SELECT AVG(sentiment_subjectivity) FROM user_reviews
num_negative_comments	Number of negative comments for the app, replace NULL values with 0.	SELECT COUNT(CASE WHEN sentiment = 'Negative' THEN 1 ELSE NULL END) AS negative_comments FROM user_reviews WHERE sentiment = 'Negative' OR sentiment IS NULL;
num_reviews_with_neutral_attitude	The number of reviews with a neutral attitude for the app, replace NULL values with 0.	SELECT COALESCE(SUM(CASE WHEN sentiment = 'Neutral' THEN 1 ELSE 0 END), 0) FROM user_reviews;
Installs	Number of user downloads/installs for the app (as when scraped)	
Reviews	Number of user reviews for the app (as when scraped)	
Category	Category the app belongs to	
Rating	Overall user rating of the app (as when scraped)	
Genres	An app can belong to multiple genres (apart from its main category)	
Sentiment_polarity	Sentiment polarity score	
Translated_review	User review (Preprocessed and translated to English)	
Sentiment_subjectivity	Sentiment subjectivity score	
Sentiment	Overall user rating of the app (as when scraped)	

Figure 2: Example of a DP for the input of Figure 1.

Figure 2 shows a corresponding data product which consists of a single table that integrates relevant information. In this example, base columns such as `App` and `Installs` are directly selected from the source database. In addition, several derived columns (e.g. *"num_negative_comments"*) are constructed from existing attributes using transformation functions (e.g., aggregations or conditional logic). Each derived column is associated with explicit provenance, expressed as SQL, describing how it is computed from base columns.

The resulting data product enables downstream analytical tasks (e.g., answering queries about app popularity or sentiment) in a simplified and structured form. Although the DPR may imply additional attributes (e.g., installation requirements), such columns are excluded if they cannot be derived from the available schema.

3 Data Sources for DP-Bench

There are a number of benchmarks available for various NLP and data integration tasks which use tabular databases. We identified the following two that can be harnessed as sources for our task.

First, **BIRD** (Li et al., 2024) is a cross-domain English text-to-SQL benchmark with large-scale databases. It covers more than 37 domains such as blockchain, healthcare, etc. It contains over 12,751 unique question–SQL pairs for 95 databases.

Second, **ELT-Bench** (Jin et al., 2025) focuses on the ELT (Extract, Load, Transform) task, which is a type of data integration process that moves raw data from a source system to a destination resource, such as a data warehouse. The ELT-Bench benchmark was designed to assess the capabilities of AI agents to build ELT pipelines. It consists of 100 pipelines for 100 databases (i.e. one for each), including 835 source tables and 203 data models across various domains. In the context of ELT-Bench, a data model contains English name and description of the data model, and name description of the columns in the data model.

Examples from the BIRD and ELT-Bench are shown in Appendix A.8.

4 DP-Bench Construction

To construct DP-Bench, we re-purposed ELT-Bench and BIRD to create a new gold standard benchmrak for data product generation. For several phases of DP-Bench construction, we used LLM-based approaches (using Llama-3.3 70B (Grattafiori et al., 2024)) for tasks such column alignment and request generation. LLM outputs were manually vetted by human annotators (all five co-authors of this paper), validating the expected outputs of the benchmark (i.e., the data products)

and then working backwards to annotate the appropriate inputs (i.e., data product requests).

4.1 Align BIRD and ELT-Bench

Out of 100 ELT-Bench pipelines, **78** can be mapped one-on-one to BIRD databases. These 78 BIRD databases together contain **4,306** columns overall and are associated with **10,928** unique question-SQL pairs. The 78 ELT-Bench pipelines have **921** columns within **122** data models.

We used an LLM to align columns within data models of ELT-Bench to the columns within tables of BIRD. This analysis discovered that **634** (out of 921) columns within ELT-Bench's data models are **derived** from columns within BIRD databases. Consequently, the remaining **287** (out of 921) columns are an *exact* match to columns within BIRD. Henceforth, we shall refer to these columns as **non-derived** columns.

4.2 Produce Provenance for Derived Columns

We used the prompt in Appendix A.1 to generate provenance in SQL for the derived columns.

We then manually validated whether the generated SQL provenance for a derived column can be used to populate desired values for that column.

During this step, we had to filter out some derived columns due to erroneous provenance. For example, consider the example shown in Figure 3 in the DB `movielens`. The column description from the corresponding ELT-Bench data model is confusing, and as a result the generated SQL as provenance is actually useless. So, we discarded it. Finally, after the aforementioned filtering, we have **582** derived columns for consideration.

"column name": "is_called_box_office_success_paradox"
"description": "Set to 1 if the director is called the box office success paradox, otherwise, set to 0"
"generated provenance": "SELECT CASE WHEN directorname = 'The Box Office Success Paradox' THEN 1 ELSE 0 END FROM directors;"

Figure 3: Example of confusing derived column.

4.3 Create Preliminary Data Products

We create one data product for each of the 78 databases. We do so by grouping the remaining **582** derived columns and the **287** non-derived columns by table names from the corresponding DB.

To construct the tables for each data product, we followed a systematic procedure. For every data model in an ELT-Bench database, we first identified the corresponding table name aligned with the BIRD database and created a table with that name in the data product. All non-derived columns associated with the ELT-Bench data model were then automatically included in this table. For derived columns, we performed an additional manual review: if a derived column appeared more appropriate for a different table within the BIRD database, we reassigned it to the table with the matching name in the data product.

4.4 Annotate Data Product Requests (DPRs)

4.4.1 Candidate DPR Generation

To build a DPR, as defined in Section 2.1, we use a two stage prompting strategy. The first stage takes in the meta data of the target data product and generates a summary in natural language. The second and final stage transforms the summary in the previous step to a candidate DPR. The prompts for this stage are detailed in Appendices A.2 and A.3 respectively.

For each target data product, we generate *three* summaries, and in turn, *three* candidate DPRs, bringing our total to **234** candidate DPRs for 78 data products.

4.4.2 Manual Curation for Gold DPRs

We did a 2 phase manual curation of the candidate DPRs. As mentioned in Section 4.3, we exploited error-free BIRD question-SQL pairs during data product creation. We randomly chose 5 such BIRD questions for each data product.

In phase 1, we asked each of the 5 human annotators to read 5 randomly chosen question for a data product, and then read each of the 3 candidate DPRs for the data product. Without looking into actual data product, an annotator was asked to decide whether a candidate DPR is phrased in such a way that if a data product is created for the candidate DPR, then this data product could provide insights/information asked by the 5 questions. If the answer is *NO*, the annotator must edit/modify/enhance the business usecase in the candidate DPR accordingly. Each candidate DPR was vetted by 2 different annotators in phase 1.

In phase 2, we asked a 3rd different annotator to look into the edits/modification done by the 2 annotators for a given candidate DPR. If the modifications are not aligned or similar, we asked this 3rd annotator to adjudicate and perform minimal edits and resolve the conflicts in wording in the modifications suggested by the previous 2 annotators.

The adjudicated edit done by the 3rd annotator is considered as the final gold standard DPR.

We found 71% of candidate DPRs (generated by LLM) were deemed by the annotators as sufficient and required no edits. Over the 2 phase of annotations, 67 candidate DPRs across 27 data products received edits.

4.5 Enhance Preliminary Data Products

After SQL error checking and filtering, we found **9,218** (84%) of the previously mentioned 10,928 unique questions–SQL pairs from the aforementioned 78 databases in BIRD are noise free. The column names and table names used in these SQLs can be traced to tables in corresponding databases. We used *sql_metadata*[1] for parsing SQL.

For each data product request, we identify questions (from the 9,218 question–SQL pairs) that have a cosine similarity of at least 0.8 while using the IBM's Granite embeddings. This step yields **836** relevant questions in total.

Finally, for each SQL (for the identified relevant BIRD questions of a BIRD DB), we created tables in the corresponding data product if the tables used in the SQL do not exist in the data product. We added the columns used in the SQL accordingly to the corresponding tables in the data product.

After this step, we have total **1458** non-derived columns and **582** derived columns organized in **383** tables inside the **78** data products. These are our final data products.

In Appendix A.7, Figure 7 shows the distribution of derived and non-derived columns for the data product for each DB in the DP-Bench.

Since some of the databases are small, we designate the 30 DBs (and the corresponding data products) with at least 50 columns as the **Hard** subset of DP-Bench. We report results on both the full benchmark and the *Hard* subset of the benchmark.

4.6 Annotate topics

Topics provide a higher level of abstraction over DPs and their associated DPRs. The goal of annotating topics is threefold: (i) to enable effective indexing and search over DPs, (ii) to support evaluation of semantic alignment between generated and gold DPs, and (iii) to facilitate comparison of abstraction levels between a DPR and its corresponding DP.

[1] https://pypi.org/project/sql-metadata/

To generate topics for each DP, we employed the Llama-3.3 70B model using manually curated three-shot prompting examples. The generated topics were then manually reviewed by annotators to remove irrelevant or redundant entries that did not accurately reflect the tables and columns within the corresponding DP. This process resulted in a total of **564** topics across the 78 data products.

For instance, for a representative DP (see Figure 2), the annotated topics include: *User reviews, Mobile apps, Android apps, App reviews, App ratings, App installation, Sentiment analysis,* and *Play store*.

We applied the same topic generation and validation procedure to DPRs associated with each DP. This resulted in a separate set of **461** topics derived from DPRs.

Continuing the earlier example (Figure 1), the topics generated from the corresponding DPR include: *User reviews, App performance evaluation, App downloads, Sentiment analysis,* and *Mobile app reputation*.

5 Study Design and Evaluation Metrics

We evaluate automatic data product creation systems along three core dimensions: (i) schema selection, (ii) derived column generation, and (iii) downstream task effectiveness.

5.1 Research Questions

Our evaluation is guided by the following 4 research questions: **RQ1:** How accurately can systems select relevant tables and columns for a given DPR? **RQ2:** How well can systems generate derived columns with correct transformations and provenance? **RQ3:** How effectively do generated data products support downstream analytical tasks (e.g., text-to-SQL)? **RQ4:** How does database complexity affect performance?

5.2 Evaluation Overview

Given a data product request (DPR), a ground-truth data product (DP), and a predicted DP, we evaluate systems using multiple metrics.

5.2.1 Column-Level Metrics

We measure the overlap between predicted and ground-truth schema elements using Precision (P), Recall (R), and F1 score over selected columns (and tables implicitly). Formally,

$$P = \frac{|\hat{C} \cap C^*|}{|\hat{C}|}, \quad R = \frac{|\hat{C} \cap C^*|}{|C^*|}, \quad F1 = \frac{2PR}{P+R}$$

where $\hat{C}$ and C^* denote predicted and ground-truth column sets, respectively.

We consider two settings for these metrics. The first is the *non-derived* setting, where we only consider $\hat{C}$ which are base (non-derived) columns. The second is the *inclusive* setting, where $\hat{C}$ also includes the columns which are used to produce derived columns. The columns that make up these two settings can differ because a DP may use columns to produce derived columns that are relevant to the use-case but not select the base columns on their own.

5.3 Derived Column Similarity

To evaluate generated derived columns, we compute similarity between transformation definitions using abstract syntax trees (ASTs). Given two derived columns with ASTs x and y, similarity is defined as:

$$\text{sim}(x, y) = 1 - \frac{\text{diff}(x, y)}{\text{count}(x, y)},$$

where $\text{diff}(x, y)$ is the edit distance between ASTs and $\text{count}(x, y)$ is the total count of comparisons made between the ASTs, which normalizes $\text{sim}(x, y)$ to range from 0 to 1. Derived columns from predicted and ground-truth DPs are then matched to maximize total similarity, and which we use to compute soft Precision, Recall, and F1.

5.4 Downstream Task Evaluation

We evaluate the usefulness of generated data products using a text-to-SQL task. For each DPR, we measure execution accuracy:

$$\text{Execution Accuracy} = \frac{\#\text{correct query results}}{\#\text{queries}},$$

where correctness is determined by comparing outputs of generated SQL queries with ground-truth.

Column-level metrics assess schema selection (RQ1), derived column similarity addresses transformation quality (RQ2), and execution accuracy captures end-to-end utility (RQ3). Evaluation across datasets of varying sizes addresses RQ4.

6 Baselines Approaches for Automatic Data Product Creation

We employ two baseline approaches, the first approach utilizes a hybrid search strategy that combines vector-based and BM25 encodings, whereas, the second uses LLM based prompting to either identify (or generate) existing (or derived) columns.

6.1 Hybrid Search for Data Product Creation

The hybrid search baseline makes use of a hybrid index of dense and sparse embeddings of each table's columns to search for relevant columns. For each table and column in the database, we construct a "sentence" which concatenates the table name and column name together, and then embed this sentence using a sentence transformer model (Reimers and Gurevych, 2019) to produce a dense embedding as well as using a BM25 (Robertson and Zaragoza, 2009) model to produce a sparse embedding. To search for relevant columns, we similarly use a pre-trained sentence transformer[2] and BM25 model to embed the DPR, and then compute the cosine similarity between the DPR's embeddings and each column's embeddings to produce similarity scores. The dense and sparse embedding similarities are weighted (0.8 and 0.2, respectively) and combined to produce a final relevance score, which is then used to rank each column.

For the next step, we use a single cutoff threshold value of 0.75; search results with scores higher than the threshold are chosen as predicted columns. In any case where two tables are included in the predicted DP and have foreign-key relations specified by the database schema, we also augment the predicted DP to include such foreign-key columns.

6.2 LLM-based Data Product Creation

For the second baseline, we use LLMs to process DPRs and output a set of columns as the predicted DP. Given a JSON formatted database schema and a DPR as inputs, we prompt the LLM to produce a chain-of-thought reasoning process to (1) identify relevant columns within the database, and (2) generate relevant derived columns (including SQL transformations). This procedure is significantly more expensive than the hybrid search baseline. However, it also enables the generation of derived columns, which is not possible using only search.

In our experiments, we investigate a wide range of LLM sizes and architectures, including MoE models (GPT OSS 20b and GPT OSS 120b (OpenAI, 2025), Qwen 3 30b (Yang et al., 2025), Llama 4 Maverick (MetaAI, 2025)), a hybrid transformer-mamba architecture (Granite 4 Small (IBM, 2025)), and dense models (Qwen 2.5 72b (Cloud, 2025), Llama 3.3 70b (Grattafiori et al., 2024)).

[2] https://huggingface.co/ibm-granite/granite-embedding-125m-english

7 Experiment Results

7.1 Column Selection Results

First, we highlight key results for column-level metrics for automatic data product creation using our baseline approaches of hybrid search and LLM-based DP creation using 7 different LLMs. Detailed result tables for each experiment can be found in Appendix A.6.

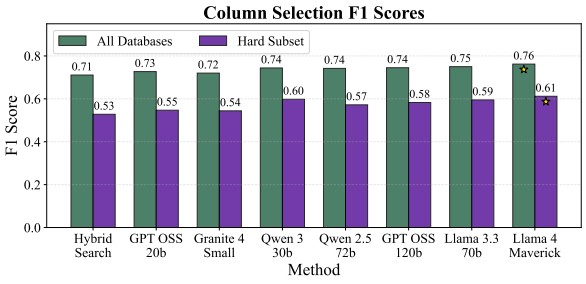

Figure 4: F1 scores for data product column selection across methods, comparing results on all DP-Bench databases and the Hard subset.

Figure 4 shows baseline results over all databases and DPRs in DP-bench, as well as over only the *hard* subset of databases in DP-bench, which consists of 30 databases which contain over 50 total columns. This chart highlights results for only F1 scores and considers all provenance columns for derived columns as also being part of the DP. When looking only at F1, it appears that most of our LLM-based approaches show very similar performance, although in our full results, we can also observe some differences in performance for each model (e.g., both GPT OSS models tended towards having higher recall).

We also can see that the hard subset, which focuses on databases with more columns, shows a decrease in F1 of roughly 0.2 across all models. We generally can see that the absolute scores for the Hard subset are lower across the board, supporting the increased challenge of the Hard subset.

7.1.1 Derived Column Results

Figure 5 shows the computed soft-F1 scores for the derived columns predicted by our baseline approaches. Full result tables can be found in Appendix A.6. Note that hybrid search is not included as a baseline because it does not produce any derived columns.

Compared to the standard column selection, we see greater variation in performance across the different LLMs used for this task. Producing derived

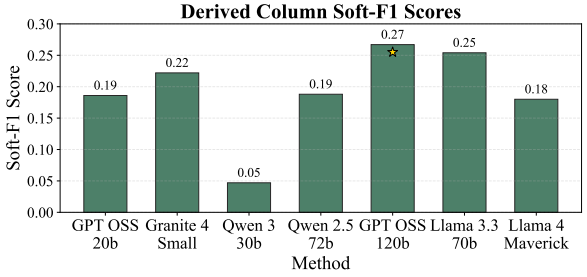

Figure 5: Soft-F1 scores of derived column predictions.

columns which closely match those in the ground-truth DP appears to be a much more challenging task across all databases, as evidenced by the fact that performance on the Hard subset of DP-Bench shows nearly identical results as on all databases.

7.2 Text-to-SQL Evaluation

Lastly, we present results for the downstream text-to-SQL task over DP-Bench. To perform this evaluation, we consider 2 scenarios: first, the *Oracle* setting, where we assume the use of a perfect text-to-SQL translation system, and evaluate the highest possible execution accuracy that could be achieved for a given DP; second, using Llama 4 Maverick to perform text-to-SQL generation.

Execution accuracy for our text-to-SQL task is shown in Table 1, where we perform tests only for questions contained in the Hard subset of DP-Bench (900 total questions across 30 databases). In addition to our baseline DP creation methods, we also present results using "no search" (i.e. use the entire database) and the ground-truth DP as input. Importantly, we observe that for text-to-SQL, using the ground-truth DP leads to better performance than using the entire database, supporting our underlying assumption that selecting a subset of the database as the DP helps to provide a more focused schema for downstream tasks and reduces errors.

Method	Oracle	Llama 4	Input Tokens
No Search	100%	45.6 %	8,432
Ground-Truth	100%	49.8 %	2,192
HybridSearch	63.0%	31.3%	6,186
GPT 20b	80.1%	36.3%	3,630
Granite 4 all	53.7%	26.0%	2,730
Qwen 3 30b	47.0%	24.4%	2,405
Qwen 2.5 72b	63.6%	31.2%	3,031
GPT 120b	84.3%	37.6%	3,587
Llama 3.3 70b	58.4%	28.1%	2,780
Llama 4	43.4%	22.2%	2,218

Table 1: Text-to-SQL execution accuracy over the Hard subset. Oracle indicates the best possible accuracy over the DP, compared to using Llama 4 for Text-to-SQL.

In general, text-to-SQL results largely correspond to Recall performance of each baseline. While results on our ground-truth DP indicate that having a more precise subset of the database is valuable, it is also the case that without high recall, text-to-SQL methods cannot possibly answer many of the questions. However, it is still important to consider that producing DPs with a smaller and more precise schema has benefits. The rightmost column of Table 1 compares the average input token length for each experiment, which helps to highlight how approaches with high recall but poor precision also lead to much higher input cost. While modern LLMs can handle extremely large context windows, it is still the case that unnecessary input context will lead to higher inference costs, which is another benefit of creating a DP rather than simply inserting all tables into a prompt.

8 Key Findings and Insights

We find that column selection can be relatively well handled by baseline approaches, while derived column generation remains a major challenge. Further, we find that recall is critical for down stream task performance, and using the full database can outperform filtered data products – the usefulness of data products thus depends heavily on its precision–recall tradeoff. Further details on our findings and insights can be found in Appendix A.5.

9 Related Work

There is only a very limited published work related to data products. (Hasan and Legner, 2023) gives an overview of data products as well as various insights into different motivations and priorities associated with data products. (Tamilselvan et al., 2025) proposes the concept of autonomous data product improvement through measurable quality contracts and optimization objectives. Prior work does not propose any approach for data product generation or a data product benchmark.

Another recent work is (Zhang et al., 2025), which proposes a data product benchmark from existing text and table QA datasets such as HybridQA by clustering related tables and passages into data products. Compared to this work, DP-Bench includes manual vetting of generated data, goes beyond source datasets through the inclusion of derived columns, produces topic annotations, and proposes baseline approaches for automatic data product generation.

The TARGET benchmark (Ji et al., 2025) considers similar tasks as ours, namely how to retrieve relevant tables and columns for a particular query. While this benchmark is topically relevant and has some overlap in the kinds of insights as our work, the benchmark uses individual factoid questions as queries rather than high-level abstractions, and thus this benchmark does not encapsulate the same kinds of challenges as DP-Bench.

(D'Ambrosio et al., 2025) proposed an approach to find possible duplications in data product catalogs exploiting pre-trained transformer models. This data product catalog does not contain data product requests (i.e. the business requirements for the data products), and thus does not address the task of automatic data product generation.

Schema linking in the context of text-2-sql can be thought of as a related work to the task of data product generation. There are many works (e.g. (Lei et al., 2020)) on this topic. But if a reduced schema w.r.t. a given query in the context of text-2-sql can be thought of a data product, then such reduced schema needs to be identified for every query for this task. Data product mitigates the need for such repetitions by making sure only columns relevant to a business problem (i.e. data product request) are there to serves the questions specific to that business problem. Furthermore, a data product contains derived or inferred columns that reduces the need joining or other expensive operations to satisfy the information need or requirements specified in the corresponding data product request.

10 Conclusion

In this work, we introduced **DP-Bench**, a first-of-a-kind benchmark for evaluating automatic data product creation. The benchmark provides manually validated data product requests (DPRs), corresponding data products with both non-derived and derived features, provenance information, and supporting analytical queries. We also proposed evaluation metrics and baseline approaches to enable systematic study of this task.

Experimental results indicate that while current methods perform reasonably well at selecting relevant schema elements, generating accurate derived features remains a significant challenge. Furthermore, results highlight the importance of balancing precision and recall to support downstream analytical tasks effectively.

DP-Bench establishes a foundation for future

research in automatic data product creation, with promising directions including handling complex schemas, improving transformation generation, enabling cross-database data products, and incorporating agent-based optimization.

Limitations

Boundaries and Assumptions

In this paper, some boundary decision and assumptions regarding the scope of our work lead to limitations regarding the coverage of what aspects of Data Products we address and evaluate.

To enable objective evaluation and tractability of the task, we impose the following boundaries on our work in this paper:

- **Relational focus:** We consider structured relational databases as the primary data source, with optional inclusion of unstructured text.
- **Schema-level outputs:** Data products are represented as collections of tables and columns, rather than full-fledged artifacts such as dashboards, reports, or applications.
- **Single-request setting:** Each data product is generated independently for a single DPR, without iterative refinement or feedback loops.
- **Static generation:** We focus on the initial generation of a data product, rather than its lifecycle management or evolution over time.
- **Non-uniqueness:** Multiple valid data products may satisfy the same DPR; therefore, evaluation is based on similarity rather than exact matching.

Several important aspects of data products creation, usage, and maintenance are outside the scope of this work:

- Visualization layers such as dashboards and reports,
- Cross-database or distributed data product construction,
- Continuous optimization or monitoring of data products,
- Data quality assessment and error remediation.

References

Alibaba Cloud. 2025. Qwen2.5 technical report. *Preprint*, arXiv:2412.15115.

Antonio D'Ambrosio, Paolo Platter, Martina Salis, Francesco Simbola, Diego Reforgiato Recupero, and Daniele Riboni. 2025. Lightweight semantic similarity search of data products. *IEEE Access*, 13:180647–180661.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, and 1 others. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Redwan M Hasan and Christine Legner. 2023. Understanding data products: Motivations, definition, and categorie. In *Proceedings of the 31st European Conference on Information Systems (ECIS 2023)*.

IBM. 2025. Ibm granite 4.0: hyper-efficient, high performance hybrid models for enterprise.

Xingyu Ji, Parker Glenn, Aditya G. Parameswaran, and Madelon Hulsebos. 2025. Target: Benchmarking table retrieval for generative tasks. *Preprint*, arXiv:2505.11545.

Tengjun Jin, Yuxuan Zhu, and Daniel Kang. 2025. Elt-bench: An end-to-end benchmark for evaluating ai agents on elt pipelines. *Preprint*, arXiv:2504.04808.

Wenqiang Lei, Weixin Wang, Zhixin Ma, Tian Gan, Wei Lu, Min-Yen Kan, and Tat-Seng Chua. 2020. Re-examining the role of schema linking in text-to-SQL. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, and 1 others. 2024. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. *Advances in Neural Information Processing Systems*, 36.

MetaAI. 2025. The llama 4 herd: The beginning of a new era of natively multimodal ai innovation.

OpenAI. 2025. gpt-oss-120b & gpt-oss-20b model card. *Preprint*, arXiv:2508.10925.

DJ Patil. 2012. *Data Jujitsu: The Art of Turning Data into Product*. O'Reilly Media, Incorporated.

Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China. Association for Computational Linguistics.

Stephen Robertson and Hugo Zaragoza. 2009. The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends in Information Retrieval*, 3:333–389.

Priyadarshini Tamilselvan, Gregory Bramble, Sola Shirai, Ken C. L. Wong, Faisal Chowdhury, and Horst Samulowitz. 2025. Agentic control center for data product optimization. In *Proceedings of the IEEE International Conference on Data Mining (ICDM) 2025 :: Demo*.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, and 1 others. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Liangliang Zhang, Nandana Mihindukulasooriya, Niharika S. D'Souza, Sola Shirai, Sarthak Dash, Yao Ma, and Horst Samulowitz. 2025. From factoid questions to data product requests: Benchmarking data product discovery over tables and text. *Preprint*, arXiv:2510.21737.

A Appendix

A.1 Prompt to generate provenance in SQL for the derived columns

We used the following prompt to generate provenance in SQL for the derived columns:

```
Generate only SQL to find
{COL_DESC} from the following
database schema.

{DB}
```

Here *"{COL_DESC}"* is the description of the derived columns (obtained from the corresponding ELT-Bench data model) and *"{DB}"* is the complete database schema for the corresponding BIRD database. Such a schema information includes table names, column names, column descriptions, data type of the columns, the primary keys of for each table and the foreign keys for each table.

A.2 Prompt to generate DP summaries

Given the tables of a data product, we used the LLM to generate 3 possible summaries using the following prompt -

```
Summarize the following tables.

{DP_TABLES}

The summary should be limited to
100 words. Try to avoid numbers.
The summarization should start
with something like "This table
contains data that enables
analysis of ..."
```

Here, *"{DP_TABLES}"* is a placeholder for the metadata of the tables in a data product. Such metadata includes table names, column names (both derived and non-derived) and the column descriptions.

A.3 Prompt to generate candidate data product requests

Given a summary, we used the LLM to convert it into a **candidate data product request** using the following prompt containing an example –

```
Transform the input text using
the following example. Do not
add any explanation.

Input Text: This table contains
data that enables analysis of
geographic and demographic
information, including
population characteristics,
economic data, and congressional
representation. It covers
various aspects of residential
areas, such as household income,
education, and employment, as
well as information about the
location itself, like elevation
and time zone. The data also
includes details about the local
government and organization,
allowing for a comprehensive
understanding of the area.

Transformed text: Collect
data that will allow queries on
geographic and demographic
information, including
population characteristics,
economic data, and congressional
representation. It should cover
various aspects of residential
areas, such as household income,
education, and employment, as
well as information about the
location itself, like elevation
and time zone. It should also
```

```
include details about the local
government and organization,
allowing for a comprehensive
understanding of the area.

Input Text: {SUMMARY}

Transformed text:
```

Here *"{SUMMARY}"* refers to a summary of a data product generated earlier.

A.4 DP quality evaluation using topics

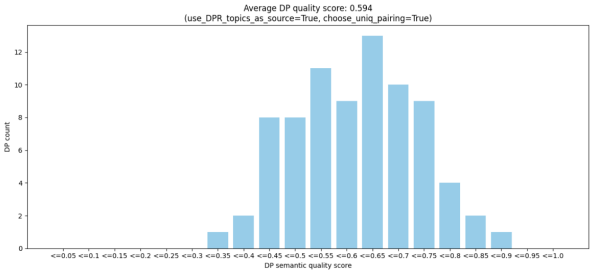

Figure 6: Distribution of semantic quality scores for the generated DPs from the *GPT OSS 120b* based method.

Using the approach described in Section 4.6, we generated topics for the generated DPs. For each topic from a gold DPR, we selected the best possible topic from the topics of the corresponding generated DP by computing a similarity score between the embeddings using the *IBM slate.30m.english.rtrvr* model. The average of the best similarity scores for the selected topics of the corresponding generated DP is used that as the semantic quality score for the generated DP.

Figure 6 shows the distribution of topic-based semantic quality score of the generated DPs (for each of the 78 DBs) from our best baseline method of *GPT OSS 120b*. The average DP quality score is *0.594*. Our analysis of indicates that there is no correlation between the size of the original DB and the eventual topic-based quality score for the corresponding generated DP given the DPR. From our random manual analysis, it appears that usually the more vague or less precise the DPR (i.e. the business use case) is, the lower the generated DP quality score.

A.5 Expanded Key Findings and Insights

In this section, we summarize the key insights obtained from our experimental evaluation of automatic data product creation systems on DP-Bench. These findings highlight both the progress made by current approaches and the remaining challenges in this problem space.

Insight 1: Column selection is relatively well-handled by current approaches. Across all evaluated methods, particularly LLM-based approaches, we observe reasonably strong performance in selecting relevant columns, with F1 scores reaching close to 0.7 on the full benchmark. This indicates that modern models are effective at identifying relevant subsets of the schema corresponding to a given DPR, especially when compared to naive baselines.

Insight 2: Derived column generation remains a major challenge. In contrast to column selection, generating derived columns with correct transformations proves significantly more difficult. All evaluated methods achieve relatively low soft-F1 scores on derived column similarity, suggesting that accurately modelling transformations and feature engineering logic remains an open problem.

Insight 3: Recall is critical for downstream task performance. Our analysis of the text-to-SQL evaluation reveals a strong correlation between recall in column selection and execution accuracy. Systems that fail to include necessary columns in the generated data product cannot support correct query generation, even if other aspects of the data product are accurate.

Insight 4: Using the full database can outperform filtered data products. Surprisingly, we observe that in some cases, using the entire database (i.e., the no-search baseline) yields higher text-to-SQL execution accuracy than generated data products. This finding suggests that modern LLMs are increasingly capable of handling large schemas, and that aggressive schema pruning can negatively impact performance when recall is insufficient.

Insight 5: Data product usefulness depends on the precision–recall trade-off. While full-schema approaches may achieve higher recall, they come at the cost of significantly larger input sizes and increased computational overhead. In contrast, data products provide a more compact and focused schema, reducing prompt size and improving efficiency. This highlights the importance of balancing precision and recall in data product creation.

Insight 6: The quality of DPRs strongly influences outcomes. Our analysis indicates that the clarity and specificity of data product requests play a crucial role in determining the quality of generated data products. Vague or underspecified DPRs

tend to produce lower-quality outputs, as reflected in semantic similarity and topic-based evaluations.

Insight 7: Benchmark difficulty reveals performance gaps. Results on the Hard subset of DP-Bench, consisting of larger and more complex databases, show consistently lower performance across all methods. This demonstrates that current approaches struggle to scale effectively with schema complexity, highlighting an important direction for future work.

Overall, our findings show that while current methods are capable of identifying relevant schema elements, generating high-quality, transformation-rich data products remains challenging. Furthermore, the interplay between schema reduction and downstream task performance suggests that more principled approaches are needed to balance efficiency and effectiveness in automatic data product creation.

A.6 Detailed Results

Here, we show detailed results for experiments surrounding column selection.

A.6.1 Column Selection Precision, Recall, and F1

Table 2 shows baseline results over all databases and DPRs in DP-Bench. We present results in terms of computing precision, recall, and F1 on *non-derived* columns only, where we ignore all derived column content from the predicted and ground-truth DP. Similarly, Table 3 shows results for the *incl. derived* setting, where we consider all provenance columns for derived columns as also being part of the DP, and compute each metric accordingly. Additionally, we indicate precision, recall, and F1 metrics for a *No Search* baseline, where we simply use the entire database as the Data Product and compute metrics.

We can observe that even without performing any search, an F1 score of .648 can be achieved due to our benchmark containing many small databases, where simply selecting all columns in the database results in relatively high precision. However, we do still observe that all baselines out-perform the no search baseline in terms of precision and F1 score with the exception of Hybrid Search in the Incl. Derived setting.

Across the LLM-based approaches, we also can observe some key differences in performance for each metric. Notably, the GPT OSS models both tend towards having have very high recall at the cost of higher precision. Llama 3.3 and Llama 4 show the best performance in terms of F1 score.

Method	P	R	F1
No Search	.522	1.00	.648
Hybrid Search	.575	.872	.652
GPT OSS 20b	.558	.931	.662
Granite 4 Small	.590	.835	.665
Qwen 3 30b	.637	.802	.687
Qwen 2.5 72b	.582	.890	.679
GPT OSS 120b	.563	**.952**	.677
Llama 3.3 70b	.601	.880	.693
Llama 4 Maverick	**.654**	.807	**.704**

Table 2: Precision (P), Recall (R), and F1 scores for DP creation using baseline approaches computed over only non-derived columns.

Method	P	R	F1
No Search	.598	1.00	.715
Hybrid Search	.653	.869	.711
GPT OSS 20b	.638	.929	.727
Granite 4 Small	.670	.829	.720
Qwen 3 30b	.728	.797	.744
Qwen 2.5 72b	.665	.887	.742
GPT OSS 120b	.646	**.950**	.745
Llama 3.3 70b	.688	.871	.750
Llama 4 Maverick	**.749**	.802	**.762**

Table 3: Precision (P), Recall (R), and F1 scores for DP creation using baseline approaches while including the source columns used to produce derived columns.

Table 4 and Table 5 show the same metrics over only the *hard* subset of databases in DP-bench, which consist of 30 databases which contain over 50 total columns. Here, the value of generating data products is greater, as the no search baseline only has a precision of .310, indicating that a much larger number of columns are not useful for downstream tasks. The relative performance among approaches appears to be the same for this subset, but we observe that the best baseline – using Llama 4 Maverick – improves on the no search and hybrid search baselines by nearly 0.1. We generally can see that the absolute scores for the Hard subset are lower across the board, supporting the increased challenge of the Hard subset.

A.6.2 Derived Column Similarity Metrics

Next, we present results for derived column similarity metrics, computed as soft-precision, soft-recall, and soft-F1 scores. Table 6 shows results over all databases in DP-Bench, and Table 7 presents re-

Method	P	R	F1
No Search	.310	1.00	.459
Hybrid Search	.358	.819	.470
GPT OSS 20b	.363	.872	.480
Granite 4 Small	.409	.697	.489
Qwen 3 30b	.476	.675	.534
Qwen 2.5 72b	.396	.779	.504
GPT OSS 120b	.371	**.905**	.507
Llama 3.3 70b	.443	.757	.537
Llama 4 Maverick	**.520**	.649	**.557**

Table 4: Precision (P), Recall (R), and F1 scores for baseline approaches over the Hard subset (30 databases which contain over 50 total columns) of our benchmark, computed only for non-derived columns.

Method	P	R	F1
No Search	.381	1.00	.534
Hybrid Search	.429	.805	.528
GPT OSS 20b	.437	.866	.547
Granite 4 Small	.490	.690	.544
Qwen 3 30b	.578	.670	.598
Qwen 2.5 72b	.479	.777	.572
GPT OSS 120b	.454	**.903**	.583
Llama 3.3 70b	.532	.742	.595
Llama 4 Maverick	**.619**	.641	**.612**

Table 5: Precision (P), Recall (R), and F1 scores for baseline approaches over the Hard subset including derived columns.

sults for only the hard subset of databases. Note that hybrid search is not included as a baseline because it does not produce any derived columns.

Compared to the standard column selection, we see some greater variation in performance across the different LLMs used for this task. For example, while Qwen 3 30b was one of the best performing models for basic column selection, it struggled significantly in producing derived columns which were similar to the ground truth. On the other hand, Granite 4 showed relatively stronger performance at producing derived columns compared to its performance on column selection.

Producing derived columns which closely match those in the ground-truth DP appears to be a much more challenging task across all databases, as evidenced by the fact that performance on the Hard subset of DP-Bench shows nearly identical results as on all databases.

Method	sP	sR	sF1
GPT OSS 20b	.176	.241	.186
Granite 4 Small	.290	.195	.222
Qwen 3 30b	.101	.034	.047
Qwen 2.5 72b	.250	.165	.188
GPT OSS 120b	.270	**.290**	**.267**
Llama 3.3 70b	**.329**	.221	.254
Llama 4 Maverick	.298	.135	.180

Table 6: Soft-precision, -recall, and -F1 scores of the derived columns produced by each baseline method.

Method	sP	sR	sF1
GPT OSS 20b	.177	.235	.185
Granite 4 Small	.305	.193	.223
Qwen 3 30b	.100	.038	.050
Qwen 2.5 72b	.245	.156	.179
GPT OSS 120b	.273	**.282**	**.266**
Llama 3.3 70b	**.338**	.218	.254
Llama 4 Maverick	.299	.136	.180

Table 7: Soft-precision, -recall, and -F1 scores of derived columns over the Hard subset tables.

A.7 Statistical details of the data products in DP-Bench

Figure 7 shows various statistics including number of derived and non-derived columns in the data products in DP-Bench. The first 30 DBs are part of the BP-Bench Hard subset.

A.8 Examples from the BIRD and ELT-Bench benchmarks

Figure 9 shows an example of a data model named "customers" that is created from the database "retails". The data model column "c_custkey" can be directly traced to the a column with the same name inside the table "Customers" in the "retails" database. On the other hand, the data model column "order_date_highest_total_price" is derived by consolidating information from several columns in tables "Customers" and "Orders" in that database.

A.9 License For Artifacts and Intended Use

The license for DP-Bench is CC BY-NC-SA 4.0 (Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International). BIRD dataset has CC BY-SA 4.0 license. ELT-Bench has CC BY-NC-ND 4.0 license.

The licenses of the datasets that we used as source do not require consent for using or curating.

DB index	DB name	Total columns in DB	Total DERIVED columns in the corresponding data product (DP)	Total NON-DERIVED columns in the corresponding DP	Ratio of non-derived columns in DP (%) w.r.t. total columns in DB
1	works_cycles	455	5	31	6.81
2	hockey	300	11	25	8.33
3	european_football_2	199	15	28	14.07
4	professional_basketball	157	15	36	22.93
5	mondial_geo	139	5	29	20.86
6	card_games	115	5	26	22.61
7	legislator	107	5	11	10.28
8	college_completion	102	0	15	14.71
9	formula_1	94	5	32	34.04
10	california_schools	89	5	23	25.84
11	movie_3	89	9	26	29.21
12	soccer_2016	85	4	24	28.24
13	synthea	85	6	25	29.41
14	address	81	5	26	32.1
15	public_review_platform	75	9	19	25.33
16	codebase_community	71	10	33	46.48
17	donor	71	5	23	32.39
18	world_development_indicators	67	5	20	29.85
19	book_publishing_company	64	15	26	40.63
20	thrombosis_prediction	64	5	17	26.56
21	beer_factory	61	11	17	27.87
22	retails	61	9	26	42.62
23	superstore	61	9	29	47.54
24	car_retails	59	16	32	54.24
25	retail_complains	58	6	25	43.1
26	financial	55	10	27	49.09
27	movie_platform	55	14	20	36.36
28	movies_4	53	10	29	54.72
29	chicago_crime	52	6	19	36.54
30	books	50	16	19	38
31	student_club	48	11	26	54.17
32	talkingdata	48	5	16	33.33
33	bike_share_1	46	5	23	50
34	menu	45	5	17	37.78
35	retail_world	42	5	12	28.57
36	simpson_episodes	42	5	22	52.38
37	law_episode	41	10	27	65.85
38	regional_sales	41	11	16	39.02
39	food_inspection_2	40	10	18	45

DP-Bench Hard

DB index	DB name	Total columns in DB	Total DERIVED columns in the corresponding data product (DP)	Total NON-DERIVED columns in the corresponding DP	Ratio of non-derived columns in DP (%) w.r.t. total columns in DB
40	cookbook	39	5	20	51.28
41	software_company	38	5	10	26.32
	ice_hockey_draft	37	6	20	54.05
	coinmarketcap	34	5	14	41.18
	airline	32	13	10	31.25
	olympics	32	11	25	78.13
	shipping	32	10	19	59.38
	superhero	31	6	31	100
	cs_semester	28	15	23	82.14
	movie	27	10	14	51.85
	sales_in_weather	26	8	3	11.54
	codebase_comments	25	5	13	52
	food_inspection	25	5	15	60
	language_corpus	24	7	12	50
	movielens	24	12	21	87.5
	world	24	5	13	54.17
	disney	23	5	13	56.52
	authors	21	5	15	71.43
	debit_card_specializing	21	5	11	52.38
	social_media	21	5	9	42.86
	video_games	21	0	18	85.71
	human_resources	20	10	18	90
	image_and_language	20	5	20	100
	university	20	10	15	75
	shakespeare	19	5	14	73.68
	cars	16	5	13	81.25
	music_platform_2	16	9	13	81.25
	sales	16	5	10	62.5
	app_store	15	5	11	73.33
	genes	15	5	14	93.33
	student_loan	15	16	13	86.67
	computer_student	12	5	8	66.67
	restaurant	12	5	12	100
	trains	12	5	10	83.33
	european_football_1	11	5	10	90.91
	toxicology	11	5	10	90.91
	music_tracker	10	6	9	90
	mental_health_survey	8	5	8	100
	citeseer	6	5	6	100

Figure 7: The number of derived and non-derived columns in the data product for each DB in the DP-Bench.

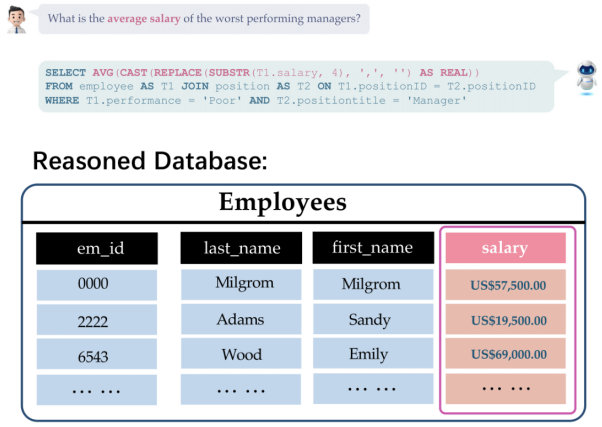

Figure 8: Example of a natural question–SQL pair in BIRD from https://bird-bench.github.io

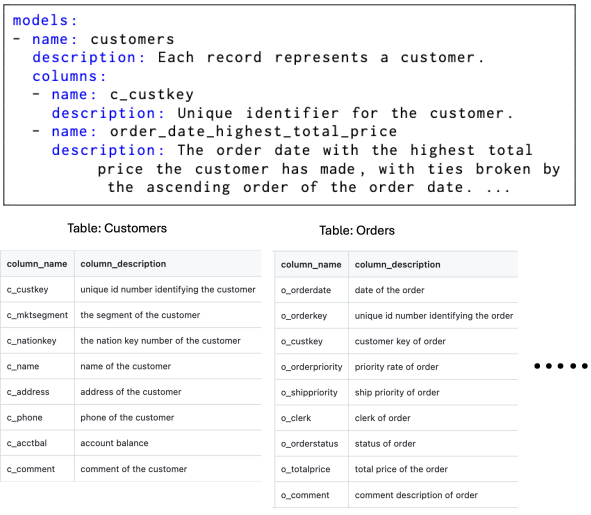

Figure 9: A partial description of the "customers" data model for the "retails" database in ELT-Bench.

It is permitted by the corresponding licenses.

The intended use of DP-Bench is to support research on bridging the gap between natural language intent and structured data representation. It is made publicly and freely available for research purposes at: https://huggingface.co/datasets/ibm-research/dp-bench.

A.10 Ethics Review Board Approval

The usage of BIRD and ELT-Bench datasets and the subsequent annotations for DP-Bench were reviewed and approved by lawyers of our company.