Human-Level Text-to-SQL via Reinforcement Learning on Verified Data, Without Pipeline Engineering

Yuxuan Zhu
University of Illinois
yxx404@illinois.edu

Tengjun Jin
University of Illinois
tengjun2@illinois.edu

Yoojin Choi
University of Illinois
yoojinc3@illinois.edu

Daniel Kang
University of Illinois
Bridgewater AIA Labs
ddkang@illinois.edu

arXiv:2603.20004v4 [cs.DB] 21 Aug 2026

Abstract

Translating natural language questions to SQL queries (Text-to-SQL) is a long-standing problem in database research. Recent efforts have focused on improving Text-to-SQL accuracy by building increasingly complex *multi-stage large language model (LLM) pipelines*, layering task decomposition, schema linking, and model-based query selection on top of an LLM. Despite this growing pipeline complexity, a substantial gap (>10%) between such systems and human experts persists on benchmarks, suggesting that pipeline engineering alone has hit a ceiling.

We show that human-level Text-to-SQL performance is achievable by fine-tuning an LLM using reinforcement learning with verifiable rewards (RLVR) on clean data, without any pipeline components. In this paper, we identified the dominant bottleneck for RLVR on Text-to-SQL: existing training data contains pervasive annotation errors that mislead optimization. To address this, we developed a multi-round, expert-driven verification pipeline and used it to curate BIRD-Platinum, a dataset of 2.5k verified instances sampled from BIRD Train, with errors corrected in 61% of instances. We show that fine-tuning Qwen3-235B-A22B on BIRD-Platinum yields consistent improvements (11–16%) over BIRD Train on Arcwise-Plat (an expert-verified version of BIRD Mini-Dev) and Spider2, outperforming state-of-the-art open-source systems by 0.6–16%. Furthermore, we diagnosed two failure modes of standard RLVR on Text-to-SQL. We find that (1) result-based rewards have non-trivial false positive rates (32.8% on average), and (2) models systematically ignore the external knowledge in BIRD-style problems. To address these failure modes, we propose ReViSQL-BIRD, a specialized reward shaping method that combines result-based rewards with SQL equivalence verification and leverages process rewards for incentivizing external-knowledge use. We fine-tuned Kimi-K2.6 with ReViSQL-BIRD. On Arcwise-Plat, ReViSQL-BIRD-K2.6 is the first method to achieve human-level accuracy (92.96%), outperforming the top five open-source systems on the leaderboard by 10–22%.

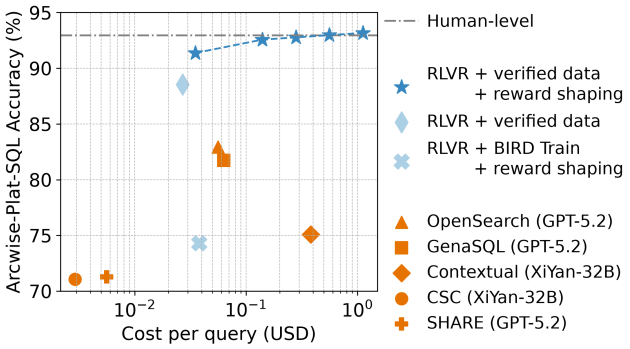

Figure 1: On Arcwise-Plat-SQL, an expert-verified BIRD Mini-Dev set constructed by prior work [4, 22], ReViSQL-BIRD achieves human-level performance [52] using RLVR (Kimi-K2.6) with verified data and reward shaping. ReViSQL-BIRD obviates the need for the complex pipelines [2, 14, 43, 47, 67] with 10–22% higher accuracy.

1 Introduction

Translating natural language questions to SQL queries (Text-to-SQL) is a foundational capability for democratizing data access [1, 3, 27, 45]. Recently, LLMs with strong reasoning and coding capabilities have emerged as its core engine [23, 32]. However, existing LLM-based methods still fall significantly short of human-level performance on the BIRD Leaderboard [52], one of the most used benchmarks. Closing this gap is a precondition for deploying LLMs in production analytics, where SQL errors carry real cost [10].

To improve the accuracy of out-of-the-box LLMs on Text-to-SQL, existing work has built increasingly complex multi-stage pipelines that compensate for the failure modes of LLMs. These LLM pipelines layer task decomposition, schema linking, in-context example retrieval, and model-based candidate selection (Table 1) around the underlying LLM [11, 14, 18, 25, 26, 28, 29, 34, 36, 39, 42, 43, 47–49, 51, 62, 67, 71]. Despite this growing complexity, the best such systems still underperform human experts by over 10% on BIRD [52], suggesting that pipeline engineering alone has hit a ceiling.

We show that human-level Text-to-SQL performance (92.96% [30]) is achievable by fine-tuning an LLM with RLVR on clean data, without any of the pipeline components. We propose two training ingredients that together close the gap to human-level performance on Text-to-SQL: a verification pipeline that produces clean training data for RLVR (Section 3), and reward shaping methods tailored to Text-to-SQL problems (Section 4). Our empirical evidence shows that clean training data is a dominant factor that brings RLVR within 3–4% of human parity, while domain-specific reward shaping closes the residual gap by providing accurate learning signals. Once both

are applied, the pipeline machinery built around the LLM becomes unnecessary for reaching human-level accuracy.

BIRD-Platinum: a verified dataset for RLVR on Text-to-SQL. Previous data-centric studies have identified pervasive annotation errors in existing Text-to-SQL datasets, including BIRD and Spider2 [22, 33, 66], where over 50% of their public instances contain errors. These errors lead to incorrect RLVR reward signals that cause fine-tuned models to make similar mistakes, limiting the effectiveness of prior RLVR efforts [12, 35, 68]. To enable effective RLVR for Text-to-SQL, we designed a verification pipeline involving 2–4 independent manual correction rounds per data instance by SQL experts. Using this pipeline, we corrected 2.5k instances sampled from BIRD Train uniformly at random [30]. We identified and corrected annotation errors across 52% of SQL queries, 26% of natural language questions, and 18% of external knowledge entries. We analyze the error patterns and release the guidelines of our verification pipeline in Section 3 for future Text-to-SQL data curation.

ReViSQL: RLVR on verified data. We introduce ReViSQL, a recipe that combines RLVR with BIRD-Platinum. Despite its simplicity, ReViSQL significantly improves accuracy and achieves generalizability across benchmarks. We evaluated on two verified versions [22] of BIRD Mini-Dev, Arcwise-Plat-Full (errors in questions, knowledge, and gold SQL fixed) and Arcwise-Plat-SQL (only gold SQL fixed; noisy questions and knowledge preserved). Fine-tuning Qwen3-235B-A22B with ReViSQL (ReViSQL-235B) yields 11% and 16% higher accuracy than RLVR with BIRD Train on Arcwise-Plat-Full and -SQL. Furthermore, we evaluated on Spider2-SQLite and Spider2-Snow [24], where ReViSQL-235B yields 12% and 14% higher accuracy than RLVR on BIRD Train. Beyond demonstrating the value of verified training data, ReViSQL-235B also matches or outperforms prior open-source pipelines at lower inference costs. Compared to the best open-source LLM pipelines (Table 4), ReViSQL-235B achieves 4.4%, 0.6%, and 2.2% higher greedy-decoding accuracy on Arcwise-Plat-SQL, -Full, and Spider2-SQLite, with 7–8× lower costs. For Spider2-Snow, ReViSQL-235B with self-consistency[1] over 128 candidates outperforms open-source LLM agents that use larger models such as DeepSeek-V3 and Qwen3-Coder.

ReViSQL-BIRD: achieving human parity on BIRD. Although ReViSQL significantly improves RLVR, fine-tuning one of the best open-source LLMs, Kimi-K2.6, using ReViSQL still underperforms human parity by 3–4% on Arcwise-Plat datasets. To close the residual gap, we diagnosed two failure modes of ReViSQL and developed a reward shaping method, ReViSQL-BIRD, to address them (Section 4). First, we find that the result-based reward, the widely used reward in Text-to-SQL RLVR [12, 35, 68], leads to an empirical false positive rate of 32.8%. This is because a semantically incorrect generated SQL query can yield the same results as the gold query on a concrete database. We mitigate such false positive rewards by integrating verification-based grading [20] with result-based grading [31, 35]. Second, we identify that fine-tuned models systematically ignore the external-knowledge entries provided in BIRD instances, defaulting to their pretraining priors over domain semantics. To address this failure, we incentivize the model to effectively leverage the external knowledge provided in a Text-to-SQL problem by assigning rule-based process rewards in addition to the widely used outcome rewards (Section 4.2).

We fine-tuned Kimi-K2.6 with ReViSQL-BIRD on BIRD-Platinum and evaluated the resulting model, ReViSQL-BIRD-K2.6, on Arcwise-Plat-Full and Arcwise-Plat-SQL [22]. With greedy decoding, ReViSQL-BIRD-K2.6 achieves 93.8% accuracy on Arcwise-Plat-Full and 91.4% accuracy on Arcwise-Plat-SQL. With self-consistency over 16 candidates, ReViSQL-BIRD-K2.6 reaches 94.2% and 93.0% respectively. On both Arcwise-Plat benchmarks, our pipeline-free results match the proxy human-level performance (92.96%) measured on BIRD Test [30]. Furthermore, ReViSQL-BIRD-K2.6 outperforms OpenSearch [67], the strongest prior open-source pipeline powered by GPT-5.2, by 8.4% on Arcwise-Plat-SQL and 5.6% on Arcwise-Plat-Full.

We summarize our contributions as follows:

(1) **A paradigm shift for Text-to-SQL**. To our knowledge, we are the first to demonstrate that human-level Text-to-SQL is achievable through clean data and RLVR alone, without multi-stage pipelines, suggesting a paradigm shift from system design to high-quality data curation for Text-to-SQL.
(2) **BIRD-Platinum and ReViSQL**. We curated BIRD-Platinum, a verified dataset with 2.5k high-quality Text-to-SQL instances. We show that fine-tuning LLMs with ReViSQL (RLVR on BIRD-Platinum) achieves consistent improvements on Arcwise-Plat (verified BIRD Mini-Dev) [22] and Spider2 [24], matching or outperforming state-of-the-art LLM pipelines with lower costs.
(3) **ReViSQL-BIRD**. We diagnosed two failure modes of RLVR on Text-to-SQL: false-positive result-based rewards and systematic neglect of external knowledge. To address them, we propose ReViSQL-BIRD, a domain-specialized reward shaping method leveraging verification-aware rewards and process rewards. We show that fine-tuning Kimi-K2.6 using ReViSQL-BIRD achieves human-level performance on Arcwise-Plat for the first time.

2 Related Work

In this section, we review the landscape of existing Text-to-SQL methods (Section 2.1) and examine the critical limitations of current Text-to-SQL datasets (Section 2.2).

2.1 Text-to-SQL Methods

We categorize prior work into three distinct paradigms: multi-stage LLM pipelines, training with massive synthetic data, and RLVR.

Multi-stage LLM pipelines. Advancements in Text-to-SQL, particularly on benchmarks including BIRD [30] and Spider2 [24], have been heavily driven by the development of complex multi-stage LLM pipelines [14, 26, 29, 34, 39, 43, 47–49, 51, 62, 67]. As shown in Table 1, top-ranking methods on BIRD rely on hand-engineered auxiliary modules around the LLM at inference time. We summarize prior work by the failure mode each component addresses, and argue that RLVR fine-tuning targets the same failure modes directly in the model, rather than working around them externally.

General-purpose LLMs struggle to map natural-language entities to the right tables and columns when the data schema is complex, leading to invalid column references and incorrect joins. Schema linking and dynamic prompts patch this externally and are

[1]We use the standard self-consistency procedure for inference-time scaling [15, 17, 47, 63, 67]. See Section 5.2 for full specification.

Table 1: Comparison of the pipeline complexity among top-30 BIRD Leaderboard methods that have disclosed technical details. Each filled dot marks a component the method requires at inference in addition to SQL query generation. ReViSQL-BIRD exceeds these methods on accuracy while requiring none of them.

Method	Dynamic pipeline	Dynamic prompts	Schema linking	Iterative revision	Selection
JoyDataAgent [5]	●	●	●	●	○
SiriusAI [21]	●	○	●	●	○
Agentar-Scale [62]	○	●	○	●	●
SHARE [43]	○	●	●	●	○
OpenSearch [67]	○	●	●	●	○
Distillery [36]	○	●	●	●	○
CHASE-SQL [39]	○	○	●	●	●
ReasoningSQL [42]	○	○	●	●	●
AskData [49]	○	●	●	○	○
GenaSQL [14]	○	○	●	○	●
XiYan-SQL [34]	○	○	●	○	●
CSC-SQL [47]	○	○	○	○	●
Contextual [2]	○	○	○	○	●
ReViSQL-BIRD	○	○	○	○	○

widely used in prior systems, including AskData [49], SHARE [43], OpenSearch [67], and Distillery [36]. They pre-filter the schema, rewrite questions, and retrieve few-shot examples to narrow the LLM's search space before generation. RLVR addresses the same failure mode internally. By rewarding the exploration trajectories that lead to correct queries, RLVR teaches the model to explore the database and ground its outputs in the schema, eliminating the need to externally link a question to the schema.

Furthermore, LLMs frequently produce queries that fail to execute or return incorrect results on the first attempt due to syntax errors, type mismatches, or subtle semantic mistakes. Iterative revision loops patch this externally. Prior methods, such as JoyDataAgent-SQL [5], Agentar-Scale-SQL [62], and CHASE-SQL [39], execute candidate queries against the target database and feed execution errors back to the LLM for iterative debugging. RLVR addresses the same failure mode by penalizing incorrect queries that would trigger revision at inference, pushing the model toward queries that execute correctly the first time.

Even with refinement, individual generations remain unreliable. Hence, many systems trade compute for accuracy via model-based, post-generation selection. GenaSQL [14], XiYan-SQL [34], and CSC-SQL [47] sample multiple candidates and use secondary reward models to score, filter, or vote among them. JoyDataAgent-SQL [5] further uses an LLM to plan and route across this multi-stage pipeline. These methods not only add inference complexity but also incur extra cost to train reward models. RLVR addresses this failure mode by shifting the distribution of generations themselves. For example, GRPO [46], a typical RLVR objective, maximizes the expected accuracy of a group of generations per question, training a model to generate correct queries with fewer candidates. Empirically, we show that, on a verified benchmark without noisy questions (Arcwise-Plat-Full) [22], greedy decoding from an RLVR model matches human-level accuracy without selection mechanisms. On a benchmark with noisy questions (Arcwise-Plat-SQL), an RLVR model achieves human parity using model-free, self-consistency generation with 16 candidates per question (Section 5.2).

While each module provides distinct utility in patching a specific LLM failure mode, their composition introduces inference-time overhead and cascading error propagation [7, 62]. For example, a wrong schema-linking result corrupts every downstream stage. More fundamentally, these pipelines remain limited by the underlying model's intrinsic SQL capability. ReViSQL-BIRD instead trains the model itself, using RLVR on clean data with tailored reward shaping, eliminating the external machinery that exists to compensate for gaps in the underlying model's reasoning.

Training on synthetic data. In parallel to building LLM pipelines, prior work has also explored data-centric approaches, developing large-scale synthetic Text-to-SQL data for training LLMs [28, 41, 65]. SynSQL-2.5M [28] developed a pipeline to generate 2.5 million synthetic Text-to-SQL instances, aiming to cover every conceivable SQL pattern. SQL-GEN [41] addresses the specific challenge of dialect adaptation, covering common dialects including PostgreSQL, BigQuery, and SQLite. SQaLe [65] takes a different angle, focusing on schema realism. It generates queries based on 135,875 real-world schemas, exposing models to "messy" production databases (e.g., thousands of tables, obscure column names).

While synthetic data enables large-scale training [28, 41] and improves robustness [65], it fails to address the inherent gap between human data and synthetic logic. These datasets typically generate natural language questions with clear, consistent logic, which does not reflect real-world deployment. Our research suggests that rigorous curation is more important than large-scale generation.

RLVR for Text-to-SQL. RLVR has emerged as a promising alternative to multi-stage pipelines for Text-to-SQL, shifting investment from inference-time orchestration to training-time optimization of the underlying LLM. Prior work has applied RLVR to Text-to-SQL tasks [35, 38, 42, 56, 68]. For example, Arctic-Text2SQL-R1 [68] fine-tuned open-weight models using result-based rewards on BIRD Train. Concurrent work from Databricks [12] and SQL-R1 [35] adopted similar execution-grading approaches. Moreover, ReasoningSQL [42] extended this with partial reward terms targeting approximate matching of SQL queries.

However, prior RLVR efforts still fail to outperform complex LLM pipelines [52]. We argue this stems from several limitations that no prior work has addressed. First, all of these efforts train, at least in part, on the official BIRD Train set whose data quality has not yet been fully verified. Second, prior efforts rely on standard result-based grading as the reward signal, which we show in Section 4 is unreliable on Text-to-SQL. In addition, domain-specific problems include external knowledge that the model must integrate (e.g., evidence entries in BIRD [30] and external knowledge in Spider2 [24]), but standard outcome-based rewards provide no signal about whether the model actually used it. Finally, prior work establishes the effectiveness of RLVR only for relatively small base models (7B–32B) evaluated on unverified benchmarks. In this work, we address these limitations simultaneously (Section 3, 4, and 5).

2.2 Noise in Text-to-SQL Datasets

Pervasive annotation errors hinder the rigorous evaluation and training of Text-to-SQL models. In BIRD [30] and Spider2 [24], two

Table 2: Correction statistics of two expert-verified datasets based on BIRD mini-Dev.

Correction	Arcwise-Plat-Full	Arcwise-Plat-SQL
SQL query	132 (26.5%)	80 (16.1%)
Question	88 (17.7%)	0
External knowledge	97 (19.4%)	0

widely used datasets, recent audits have uncovered errors in up to 52% of questions, SQL queries, and external knowledge across both evaluation and training sets [4, 22, 33, 40, 66], causing concerns in both the development and evaluation of Text-to-SQL methods.

Errors in evaluation sets. The annotation errors in the BIRD Dev set have been heavily investigated. Jin et al. [22] identified annotation errors in 52% of instances. Similarly, Wretblad et al. [66] analyzed the financial domain subset, finding that ambiguous questions and incorrect gold SQL queries introduced noise into 15–49% of instances per database. While Liu et al. [33] proposed labeling these errors to test error-detection capabilities, they did not correct them. To enable reliable evaluation, Arcwise [4] and subsequently Jin et al. [22] released two verified variants of the BIRD Mini-Dev that we used as our evaluation benchmarks: Arcwise-Plat-Full and Arcwise-Plat-SQL. These two benchmarks differ in which sources of error are corrected (Table 2; Appendix A walks through a representative instance). Arcwise-Plat-Full delivers a fully cleaned version with errors corrected in all three sources: 26.5% of gold SQL queries, 17.7% of natural language questions, and 19.4% of external knowledge entries. Arcwise-Plat-SQL corrects only the gold SQL queries (16.1% of instances), while *deliberately preserving* the factual errors and ambiguities in questions and external knowledge [22]. The two benchmarks thus evaluate different capabilities. Arcwise-Plat-Full measures accuracy under clean inputs, while Arcwise-Plat-SQL measures robustness to the noisy questions and external knowledge that are closer to real-world Text-to-SQL deployments [32].

Errors in training sets. While prior work has verified the BIRD Mini-Dev evaluation set, Text-to-SQL training data remains largely unverified. Pourreza and Rafiei [40] partially categorized errors in the training set into three types: sorting failures (ties in `ORDER BY`), schema matching ambiguity, and incorrect content assumptions. They detected errors in 18.2% of a subset of problems, which likely underestimates the true noise level. To date, there are no large-scale, publicly available corrections for the BIRD Train set. This lack of clean training data represents a fundamental barrier to training models via reinforcement learning.

3 Curating BIRD-Platinum: A Verified Dataset for Text-to-SQL RLVR

We constructed BIRD-Platinum, a verified dataset of 2.5k instances sampled from BIRD Train and corrected through a hybrid LLM–human-expert pipeline. In this section, we describe the design principles behind our correction process (§3.1), the five-stage pipeline that implements them (§3.2), a taxonomy of the errors we encountered with worked examples (§3.3), and summary statistics for both the construction process and the resulting dataset (§3.4).

3.1 Design Principles and Verification Checklist

Soundness over completeness. For a verified dataset to improve RLVR training, two properties of data cleaning are essential.

(1) Every correction must target a *genuine* annotation error.
(2) Corrections themselves must not introduce new errors.

We optimize the pipeline for these two soundness properties, even at the cost of completeness. We do not claim to find every error in every instance. Instead, we ensure that every change we make is principled, defensible, and verified by an independent expert.

Checklist. To achieve soundness, we develop a checklist that human experts apply to every instance. The checklist enumerates the specific failure modes we observed during pilot annotation and serves as both a guide for correction and verification. We organize it around the components of a Text-to-SQL instance: the natural language question, the external knowledge, and the gold SQL query.

For each natural language question, we check:

(1) Is the question answerable from the given database?
(2) Does the question admit a single interpretation?
(3) Is the expected output, including whether duplicates should be removed, unambiguously specified?
(4) Does the question yield a non-empty answer?
(5) Do all filter conditions in the question constrain the result?

For each external knowledge entry, we check:

(1) Is the external knowledge relevant to answering the question?
(2) Is the external knowledge consistent with the question?

For each gold SQL query, we check:

(1) Does the SQL query output all and only the required columns?
(2) Does the SQL query handle ties correctly?
(3) Does the SQL query handle duplicates correctly?
(4) Does the SQL query handle missing values correctly?
(5) Does the SQL query use the correct tables via appropriate joins and set operations?
(6) Are all filter conditions from the question correctly translated?
(7) Are aggregations computed over the correct partitions?
(8) Does string-value extraction generalize across the full database, not just the sampled rows?

3.2 Construction Pipeline

During data construction, each instance passes through five stages designed to combine the breadth of LLM-based error detection with the precision of human expert judgment (Figure 2). The pipeline incorporates independent correction with independent verification to avoid propagating mistakes of an expert to the final dataset.

Stage 1: LLM-assisted error flagging. We applied an LLM reviewer [22] based on OpenAI's o3 to detect annotation errors and provide correction suggestions. The agent can execute SQL queries on the example database via function calls, with a budget of 30 executions per instance. The agent generates an error report describing errors in the question, external knowledge, and SQL query, as well as a suggested correction of the SQL query.

Stage 2: Independent expert correction. A human expert independently investigates the instance using the checklist (§3.1), *before* reading the LLM agent's report. The order ensures the expert

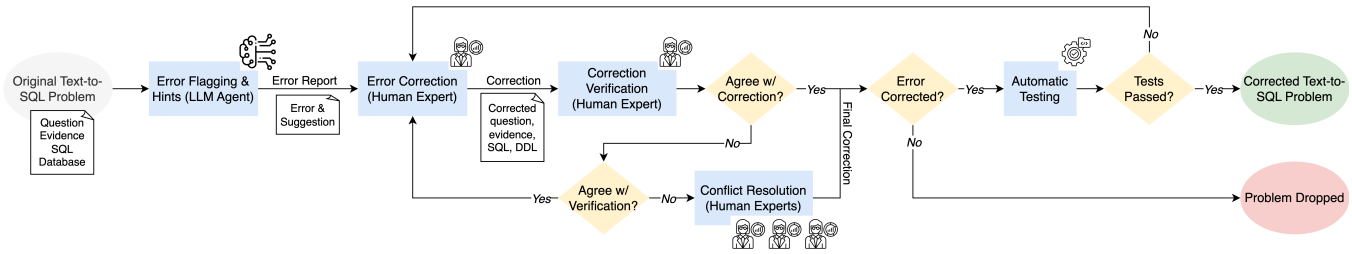

Figure 2: Our hybrid LLM–human-expert pipeline for correcting text-to-SQL annotation errors in BIRD.

detects errors the LLM missed rather than completely relying on the LLM's findings. The expert then incorporates the LLM report and produces a draft correction covering the question, external knowledge, and SQL query.

Stage 3: Independent verification. A second expert independently re-investigates the instance and the draft correction, applying the same checklist. The verifier confirms each correction is sound and searches for errors the first expert may have missed. Disagreement between the correction and verification experts triggers revision by a third expert. If the third expert disagrees with the verification result, we proceed to conflict resolution.

Stage 4: Conflict resolution. When correction and verification experts disagree on the substance of a correction, three experts convene to reach consensus. In our process, every conflict was resolved at this stage. The one outcome that we treat specially is when the experts agree that no correction is possible. This is typically because the question cannot be answered from the given database (e.g., §3.3.1, item 1). In these cases, we drop the instance.

Stage 5: Automatic testing. We then run the following two automated checks on the corrected instance. If either check fails, we route the instance back to manual revision and verification.

(1) The corrected SQL query outputs a non-empty result.
(2) The output of the corrected SQL query must differ from the output of the original SQL query.

3.3 Taxonomy of Annotation Errors

The corrections produced by our pipeline form a more comprehensive taxonomy of annotation errors than those investigated in prior work [22, 33, 66]. We describe the taxonomy to explain what BIRD-Platinum corrects and hopefully to inspire future Text-to-SQL data curation. We group errors into four categories: question and external knowledge errors, gold SQL query errors, schema errors, and grading-method errors.

3.3.1 Question and External Knowledge Correction. We identified and corrected the following three types of errors in the natural language questions and evidence:

(1) Unanswerable from schema. Some questions cannot be answered from the provided database. We dropped such instances. In the following example, the question presumes a sale-status attribute that the database does not contain:

> *Original Question*: Among the products, how many of them were discontinued in production?
>
> *Action*: (Instance dropped)

(2) Empty result. We should not use filter conditions that yield empty answer sets as RLVR targets, because a model can achieve maximum reward on these instances with trivial queries such as `SELECT 1 WHERE False`, a classic reward-hacking pathology [50]. We revise filter conditions to ensure non-empty results.
(3) External knowledge-question mismatch. When external knowledge entries contradict a question, we align the external knowledge to the question. For example, we found a question asking for "at least 200 movies" that has external knowledge "> 200". We then updated the external knowledge to ">=200".

3.3.2 Gold SQL Query Errors.

(1) Mishandled missing data. In practice, `NULL` values, sentinel values, or indicator columns often mark missing data. Aggregation or subsequent data operations should exclude them to produce meaningful results when necessary. In the following example, missing values in `weight` and `height` corrupt the `MAX` and `MIN` aggregations of BMI:

> **Question:** What are the basketball players' BMI ranges?
>
> **Original SQL:**
>
> ```
> SELECT
>     MIN(weight * 1.0 / (height * height) * 703),
>     MAX(weight * 1.0 / (height * height) * 703)
> FROM players
> ```
>
> **Corrected SQL:**
>
> ```
> SELECT ... FROM ... -- same as the original
> WHERE weight > 0 AND height > 0;
> ```

(2) Unhandled ties. When a question asks for an extreme value, `ORDER BY ... LIMIT 1` silently discards tied rows. We corrected such queries to return all rows matching the extreme value:

> *Question*: Name the movies with the most ratings.
>
> *Original SQL Query*:
>
> ```
> SELECT movie_title FROM movies
> GROUP BY movie_title
> ORDER BY COUNT(movie_title) DESC LIMIT 1
> ```
>
> *Corrected SQL Query*:

```
WITH cte AS (
    SELECT m.movie_title title, count(*) cnt
    FROM movies AS m JOIN ratings AS r
    ON m.movie_id = r.movie_id
    GROUP BY m.movie_id, m.movie_title )
SELECT title FROM cte
WHERE cnt = (SELECT MAX(cnt) FROM cte)
```

(3) Wrong or missing joins. The example above also illustrates this category. The original SQL query miscounts movies without ratings, while the corrected query joins `movies` and `ratings`.
(4) Wrong or missing predicates. In the example below, the question specifies "over 1,000", while the original query uses "> 10":

Question: Among all the users that have posted a tweet with over 1000 likes, how many of them are male?

Original SQL Query:

```
SELECT COUNT(T1.TweetID)
FROM twitter AS T1 INNER JOIN user AS T2
ON T1.UserID = T2.UserID
WHERE T1.Likes > 10 AND T2.Gender = 'Male'
```

Corrected SQL Query:

```
SELECT COUNT(DISTINCT T1.UserID)
FROM ... JOIN ... -- same as the original
WHERE T1.Likes > 1000 AND T2.Gender = 'Male'
```

(5) Unhandled duplicates. The example above also illustrates the failure to handle duplicated values when aggregating. After joining `twitter` and `user`, the column `UserID` is no longer unique. Hence, counting users requires `DISTINCT`. In our experiments, this error leads to significant consequences for RLVR. In an ablation, a model trained on uncorrected data applied `DISTINCT` correctly in only 5 of the 30 problems that require it, indicating that the model picked up the failure modes in the training data.
(6) Wrong or missing output columns. We corrected queries whose `SELECT` clauses omit columns requested by the question or include columns the question does not request.
(7) Wrong or missing group keys. We corrected `GROUP BY` clauses that aggregate over incorrect partitions.
(8) Wrong or missing tables. We added tables that the original query failed to join or set-combine, using the data schema to determine the correct relation.
(9) Wrong ranking criteria. When `ORDER BY` ranks by expression inconsistent with the question, we corrected it:

Question: List the product ID of the top five products, by descending order, the number of quantities ordered.

Original SQL Query:

```
SELECT ProductID FROM Products
ORDER BY UnitsInStock DESC LIMIT 5
```

Corrected SQL Query:

```
SELECT `ProductID` FROM `OrderDetails`
GROUP BY `ProductID`
ORDER BY SUM(`Quantity`) DESC LIMIT 5
```

(10) Execution Errors: We fixed syntax and reference errors that prevent the gold query from executing. The example above also demonstrates this. The original query references a column `UnitsInStock` that does not exist in the table `Products`.

3.3.3 Schema Errors. We corrected the schema information for 9 (1.7%) tables to maintain consistency with the example databases. Our corrections fall into the following two types.

(1) Removing references to tables or columns that do not exist in the provided example databases.
(2) Correcting value descriptions that contradict the actual data. In the example below, applying the original description would imply that a person can be both a professor and a student at the same time:

A sample of the table person:

```
p_id|professor|student|...
   3|        0|      1|...
   4|        0|      1|...
```

Original Value Explanation:

```
column, value_description
professor, [...] 0: professor. 1: student.
student, [...] 0: professor. 1: student.
```

Corrected Value Explanation:

```
column, value_description
professor, [...] 1: professor. 0: student.
student, [...] 0: professor. 1: student.
```

3.3.4 Grading-Method Errors. We find that the standard set-based grading method used in BIRD is insufficient for two classes of questions. In addition to set-based grading, we annotated instances with the following alternative grading methods:

(1) Subset-based. For any-k questions, such as "Please list the titles of any two papers that Jundu has written", we modified the gold query to return the full set without size limitation. Then, we check whether the result set of a generated SQL query is a subset of that of the gold SQL query with a correct size.
(2) List-based. For questions requiring listing rows in a specific order, such as "List the product ID of the top five products, by descending order, the number of quantities ordered", we check whether the result table of a generated SQL query has the same values as that of the gold query at the same table cell position.

3.4 Construction and Dataset Statistics

We then describe the summary statistics for our data curation process and our curated dataset BIRD-Platinum.

Construction process. In Figure 3, we highlight the effectiveness of the LLM reviewer, the necessity of the involvement of SQL experts, and the importance of multi-round correction and verification. As shown in Figure 3a, the LLM reviewer achieves a high precision of 90.6% in terms of detecting annotation errors, indicating that the errors identified by the LLM reviewer are highly reliable. However, the LLM reviewer achieves a recall of only 24.5%, missing a significant number of errors, which emphasizes the importance of expert involvement in data correction.

As shown in Figure 3b, we find that 16.2% of the first-round proposed corrections fail the subsequent verification, indicating the necessity of multi-round correction and verification. In this

Table 3: Structural complexity comparison between the original BIRD Train and BIRD-Platinum. Values for the datasets denote the average count of each structural component per query; the difference is the relative change from BIRD Train to BIRD-Platinum. We find that gold SQL queries exhibit a significant increase in complexity across every structural metric.

Per-query stats	# tables	# joins	# functions	# aggregations	# set operations	# subqueries	# CTEs	# window func.
BIRD Train	2.08	1.02	1.67	0.62	0.0024	0.088	0.0	0.0012
BIRD-Platinum	2.36	1.06	2.10	0.80	0.024	0.26	0.11	0.0012
Difference	+13.5%	+3.9%	+25.7%	+29.0%	+900%	+195.5%	—*	+2.0%

Differences are computed from unrounded values.
* The original BIRD Train contains no common table expressions, so the relative change is undefined; BIRD-Platinum introduces 0.11 CTEs per query.

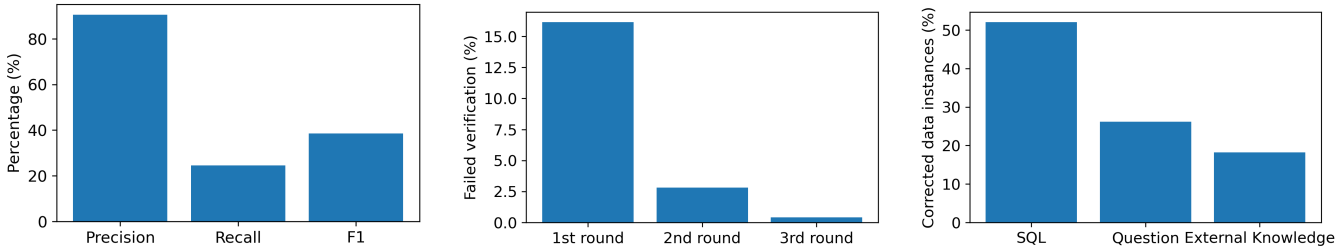

(a) LLM reviewer demonstrates high precision but critically low recall.

(b) Iterative expert correction and verification required up to four rounds.

(c) Pervasive annotation errors. Over 52% of original gold SQL queries are incorrect.

Figure 3: Quantitative analysis of the BIRD-Platinum expert curation process. Relying solely on automated LLMs for data correction is insufficient due to poor recall (Fig. 3a). Consequently, BIRD-Platinum uses an expert verification pipeline that requires up to four iterative rounds to fully resolve errors (Fig. 3b). We identified and corrected errors across 52.1% of SQL queries, 26.2% of natural language questions, and 18.2% of external knowledge contexts (Fig. 3c).

process, we encountered five cases (0.2%) where conflicts occurred between correction and verification, which were all resolved at the conflict-resolution stage. Eventually, all the proposed corrections passed verification after up to four rounds.

Dataset characteristics. In Figure 3c and Table 3, we summarize the statistics of our corrections and a comparison between the original BIRD Train and our BIRD-Platinum. As shown in Figure 3c, our process corrected gold SQL queries in 52.1% of instances, natural language questions in 26.2%, and external knowledge in 18.2%. Additionally, we discarded 1.5% of instances whose questions are unanswerable given the database schema. In total, 61.1% of the evaluated instances contained at least one type of annotation error.

In Table 3, we show the average number of tables, joins, functions, aggregations, set operations, subqueries, common table expressions (CTEs), and window functions per query. We find that our corrected queries have higher average counts in all categories, indicating that the corrected gold queries exhibit greater structural complexity, reflecting a more rigorous adherence to domain constraints.

4 ReVISQL-BIRD: Shaping RLVR for Effective Text-to-SQL Learning

We propose ReVISQL-BIRD, a reward shaping method with two techniques tailored to Text-to-SQL problems. In this section, we identify two failure modes of standard result-based RLVR rewards on Text-to-SQL (§4.1) and then propose new reward shaping technique that addresses them (§4.2).

4.1 Diagnosing the Failure Modes of Result-Based RLVR Rewards

We discuss the problem setting and the RLVR formulation our method builds on and then diagnose where standard rewards fail.

BIRD-style Text-to-SQL problems. We target the Text-to-SQL setting introduced by BIRD [30]. A *BIRD-style Text-to-SQL problem* provides (i) a natural language question, (ii) a relational database, given by its schema and contents, and (iii) a set of *external knowledge* entries. The goal is to produce a SQL query that, when executed against the database, returns the answer to the question. The external knowledge, termed *evidence* in BIRD, consists of short natural-language annotations that are "required to map the natural language instructions into counterpart database values" [30]. Because a question is often under-specified without this evidence, a Text-to-SQL model is expected to read the external knowledge and reflect it in the generated query.

RLVR background. We train models for this setting with RLVR, the dominant method for post-training LLMs [19, 37]. Group relative policy optimization (GRPO) [46] is one of the most widely used algorithmic instantiations of RLVR. In each training iteration, for a question-answer pair (q, a), GRPO samples a group of rollouts $\{\tau_1, \ldots, \tau_G\}$ from the model π_θ and estimates the advantage of output τ_i as follows:

$$\hat{A}_i = \frac{R(\tau_i, a) - \text{mean}\left(\left\{R(\tau_j, a)\right\}_{j=1}^{G}\right)}{\text{std}\left(\left\{R(\tau_j, a)\right\}_{j=1}^{G}\right)} \tag{1}$$

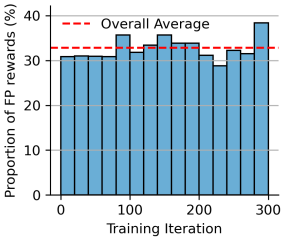

Figure 4: False positive rewards persist across the entire training process of Kimi-K2.6.

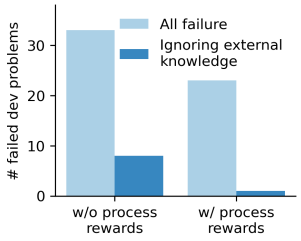

Figure 5: Process rewards reduce the failure cases caused by ignoring external knowledge entries.

where R is a task-specific reward function. GRPO then maximizes a clipped surrogate objective over these advantages. The choice of R thus fully determines what the algorithm optimizes for.

Prior RLVR work for Text-to-SQL [12, 35, 68] uses a binary result-based reward, setting $R(\tau, a) = 1$ if executing the generated SQL against the example database produces the same result as the gold SQL query, and 0 otherwise. While this formulation is cheap to compute during training, we show that such reward signals corrupt Text-to-SQL learning in two systematic ways that limit the achievable accuracy of RLVR.

Failure mode 1: false positive rewards from result coincidence. It is known that execution equivalence on a specific database does not guarantee semantic equivalence between two SQL queries [9, 20]. A query with an incorrect join, a missing predicate, or an accidentally compensating constant can produce the gold result on an example database at training time. If such queries receive positive rewards, the model is pushed toward systematic logical errors that happen to be consistent with the example database.

We measured the proportion of false positive rewards across the training process on Kimi-K2.6 by applying VeriEQL [20] which verifies semantic SQL equivalence under bounded inputs (Appendix D). A reward is false positive if result-based grading accepts the query, while VeriEQL refutes its semantic equivalence to the gold. During our pilot experiments (300 training steps), VeriEQL identifies 28.8–38.4% of positive rewards as false positives, with an average rate of 32.8% (Figure 4). Furthermore, this rate does not decrease as training progresses and the model improves. The example below illustrates a typical case. The gold query averages over order totals, while the generated query averages over individual line-item prices. The two queries compute different quantities in general, but coincide on the example database because the customer happens to have orders containing a single line item each.

Question: How much money on average does Lucas Wyldbore spend on book orders?

Gold SQL Query:

```
SELECT AVG(order_total)
FROM (
    SELECT T1.order_id, SUM(T1.price) AS order_total
    FROM order_line AS T1 JOIN cust_order As T2 ...
    GROUP BY T1.order_id
);
```

Incorrect Generated SQL Query getting positive rewards:

```
SELECT AVG(price)
FROM order_line AS T1 JOIN cust_order As T2 ...
-- same as the gold except for no GROUP BY
```

Failure mode 2: outcome rewards do not enforce external-knowledge use. BIRD-style problems provide an evidence field as domain-specific external knowledge. Standard outcome rewards based only on final execution results provide no signal about how the model arrived at its answer. Therefore, a model trained with outcome-only rewards may learn to ignore the external knowledge, reverting to its pre-training priors over domain semantics.

To assess whether this failure mode appears empirically, we conducted a pilot manual analysis of the failures in a validation set for a model (Kimi-K2.6) trained with standard outcome-based RLVR. We treat a failed instance as "ignoring external knowledge" when the generated query is consistent with a plausible alternative interpretation of the question that the external knowledge is meant to rule out. Among 33 failures, 8 instances (24.2%) meet this criterion (Figure 5, left bars). The example below illustrates a representative case. The external knowledge specifies "sodium-free refers to `sodium=0`," while the model chooses the threshold of `sodium < 5`.

Question: Among the recipes from The California Tree Fruit Agreement, calculate the percentage of sodium-free recipes.

External knowledge: The California Tree Fruit Agreement is a source; calculation = `MULTIPLY(DIVIDE(COUNT(sodium = 0 THEN recipe_id), COUNT(recipe_id)), 100)`

Gold SQL Query:

```
SELECT CAST(SUM(CASE WHEN T2.sodium = 0 THEN 1 ELSE 0
    END) AS REAL) * 100 / COUNT(*) FROM ...
```

Generated SQL Query that fails due to ignoring external knowledge:

```
SELECT CAST(SUM(CASE WHEN n.sodium < 5 THEN 1 ELSE 0
    END) AS REAL) * 100 / COUNT(*)FROM ...
```

Unfortunately, outcome rewards are structurally incapable of fixing the failures due to ignoring rules in the external knowledge. By construction, outcome rewards depend only on the final execution result, not on the reasoning that produced it. An attempt that ignores the external knowledge but coincidentally returns the correct result and an attempt that uses the external knowledge correctly are indistinguishable to an outcome reward. This observation motivates the process reward component of ReVISQL-BIRD, which we will describe in the next section.

4.2 Reward Shaping in ReVISQL-BIRD

In ReVISQL-BIRD, we develop a composite reward that complements the result-based reward and targets the two failure modes diagnosed above. For a generated rollout τ with gold answer a, the ReVISQL-BIRD reward is

$$R_{\text{ReVISQL-BIRD}}(\tau, a) = \underbrace{\rho(\tau, a, \beta)}_{\substack{\text{shaped} \\ \text{outcome reward}}} - \underbrace{\sum_{k \in \mathcal{K}} \lambda_k \cdot \mathbb{1}[\text{violation}_k(\tau)]}_{\substack{\text{external knowledge} \\ \text{process penalties}}} \quad (2)$$

where ρ is a verification-aware outcome reward that mitigates false positives with penalty β, and $\mathcal{K}$ is a set of rules for external-knowledge use whose penalties λ_k are deducted when violated. We describe each component as follows.

Verification-aware outcome reward. To address Failure Mode 1, we integrate VeriEQL [20] with result-based grading as a second-stage check. Let $g(\tau, a) \in \{\text{True}, \text{False}\}$ denote the result-based verdict and $v(\tau, a) \in \{\text{True}, \text{False}, \perp\}$ denote VeriEQL's verdict, where $\perp$ indicates a VeriEQL timeout or a SQL query unsupported by VeriEQL. The shaped outcome reward is

$$\rho(\tau, a, \beta) = \begin{cases} 0, & g = \text{False}, \\ 1, & g = \text{True} \ \wedge \ v \in \{\text{True}, \perp\}, \\ 1 - \beta, & g = \text{True} \ \wedge \ v = \text{False}. \end{cases} \tag{3}$$

We downweight rather than zero out rewards that VeriEQL refutes because VeriEQL itself can fail due to timeouts or unsupported SQL operators (e.g., window functions) [20]. Therefore, soft downweighting preserves the equivalence signal where VeriEQL is reliable while bounding the damage where it is not.

Process supervision for external knowledge use. To address Failure Mode 2, we require the model to explicitly translate and verify each external knowledge entry as part of its reasoning trace. We curate a system prompt that requires two enumerated blocks indexed by the position of an external knowledge entry:

(1) a "requirement" block that translates each external knowledge entry into a concrete query constraint.
(2) a "verification" block right before the final answer that audits whether the generated SQL query satisfies each external knowledge entry.

We deduct penalties λ_k when the model generates a solution without complying with these requirements. We conducted a pilot evaluation on the dev set to validate our reward shaping intervention. With process rewards added, only 1 of 23 failed instances (4.3%) is attributable to ignoring external knowledge. We defer the full ablation results on benchmarks to Section 5.2.

4.3 Training Algorithm and Implementation

Given the reward specification, we fine-tune models using the state-of-the-art (SOTA) RLVR settings [31] and loss function [8, 69].

Prompt. Our input prompt comprises two components: a universal system prompt shared by all Text-to-SQL problems (Appendix E), and the database schema of each problem. We render the schema in data definition language and, following prior work [28, 31], annotate each column with the column description provided by BIRD. When a column description is unavailable, we instead include three randomly sampled values from that column.

Rollout. During training, the model performs multiple rollouts to generate multiple SQL queries. Within a rollout during training and inference, the model is allowed to issue intermediate SQL queries using the provided SQL tool to explore data, test queries, and refine answers. We configure a maximum number of turns in which the model can issue intermediate queries during a rollout. To aid the model's planning, we explicitly append a turn counter to the environment's response after each intermediate execution. We defer a detailed specification of a model rollout to Appendix E.

Loss function. We use CISPO [8], a SOTA loss function of RLVR algorithms that stabilizes training and improves sample efficiency. Compared to GRPO, CISPO introduces a constraint that clips gradients asymmetrically to prevent token-level instability and mitigate the extreme variance caused by off-policy distribution shifts [8].

Hyperparameters. We use a consistent set of hyperparameters as follows. We split BIRD-Platinum with an 85:15 ratio for training and validation. For training, we use a group size of 16, a LoRA rank of 32, a batch size of 64, and a learning rate of 5×10^{-5}. During rollouts, we allow up to five turns of interaction with the database for the model to issue intermediate queries, given that our pilot experiments show that the average number of turns fluctuates between 3 and 4. For each turn, we follow prior work to allow the model to output at most 3,072 tokens [31]. For the composite reward computation, we set $\beta = 0.2$ for verification-aware rewards and $\lambda_k = 0.1$ for process supervision. We determine the training duration based on the convergence of the validation accuracy and select the checkpoint with the highest validation accuracy for the evaluation on benchmarks.

5 Evaluation

We assess the contributions of BIRD-Platinum and ReViSQL-BIRD across four benchmarks, addressing three questions. First, does ReViSQL-BIRD achieve human-level accuracy on BIRD, and how does it compare to the strongest prior open-source pipelines and models (Section 5.2)? Second, how do verified data and reward shaping each contribute to this result (Section 5.2)? Third, does the benefit of BIRD-Platinum generalize beyond BIRD-style problems (Section 5.3)? We first describe the experimental setup (Section 5.1).

5.1 Experimental Setup

We introduce our experimental settings, including benchmarks, baselines, metrics, and implementation details.

Benchmarks. We evaluate on four benchmarks spanning two classes of Text-to-SQL problems:

(1) *Arcwise-Plat-Full*: As the original BIRD benchmark contains substantial annotation errors [22, 40, 53, 66], the Arcwise company introduced a manually corrected version of the BIRD Mini-Dev set (498 problems) [4]. In addition, prior work [22] leveraged SQL experts and LLMs to fix extra errors in the Arcwise-corrected version and introduced Arcwise-Plat-Full.
(2) *Arcwise-Plat-SQL*: In addition to Arcwise-Plat-Full, prior work [22] also introduced Arcwise-Plat-SQL, a variant that only corrects errors in the gold SQL while intentionally keeping the noise in natural language questions and external knowledge unchanged. Arcwise-Plat-SQL reflects the noise that often appears in real-world deployment [22].
(3) *Spider2-SQLite*: As a successor of the widely used Spider1 benchmark [70], Spider2 consists of 547 challenging Text-to-SQL problems [24]. Following prior work [28], we use 135 problems of Spider2 that are based on SQLite. Gold SQL queries in Spider2

are more complex than those in BIRD, containing 5.2× more tokens on average [24].
(4) *Spider2-Snow*: As a variant of Spider2, Spider2-Snow rewrites all 547 gold queries using the Snowflake dialect and requires the model to write Snowflake queries for evaluation [24].

Baselines. We compare against the five strongest open-source pipelines on the BIRD leaderboard as of May 1, 2026, four open-weight Text-to-SQL models, and one proprietary reasoning model. We configured each pipeline with its leaderboard-optimal setting. If the leaderboard-optimal setting is not published, we use the default settings described in their technical report. For single-model baselines, we use the prompt described in their technical report. If they do not specify a prompt template, we apply the same prompts as ReViSQL-BIRD (Appendix E.2). We summarize the settings of baselines as follows.

(1) *Contextual-SQL* relies on a dual-model architecture. It uses Qwen2.5-Coder-32B-Instruct for candidate generation, and a fine-tuned 32B reward model (trained on the entire 9.5k-example BIRD Train set) to perform the final candidate selection [2].
(2) *CSC-SQL* uses a fine-tuned model, XiYanSQL-QwenCoder-32B-2412, to generate candidate SQL queries and uses Qwen2.5-Coder-7B-Instruct to select a final query for each problem [47].
(3) *GenaSQL* requires a multi-API orchestration pipeline spanning GPT-4o, Gemini 1.5, and Cohere to handle schema linking, retrieval, and generation tasks. To standardize evaluation, we used the same complex workflow and applied OpenAI's text-embedding-3-small for retrieval and GPT-5.2 for all subsequent orchestration and generation steps [14].
(4) *OpenSearch* proposes multi-stage retrieval and refinement, using bge-large-en-v1.5 for few-shot example retrieval and GPT-4o for schema extraction, query generation, and refinement [67]. In our experiment, we used the same retrieval model and replaced GPT-4o with GPT-5.2.
(5) *SHARE* takes a baseline SQL query as input and refines it across an ensemble of three distinct, distilled LLMs dedicated to reasoning, schema linking, and iterative query fixing. In our evaluation setup, this complex external refinement pipeline is applied to baseline queries generated by GPT-5.2 [43].
(6) *OmniSQL-32B* is a fine-tuned model based on Qwen2.5-Coder-32B-Instruct using a combination of synthetic data (2.5 million examples), the Spider Train set, and the BIRD Train set [28].
(7) *XiYanSQL-QwenCoder-32B* is a fine-tuned model based on Qwen2.5-Coder-32B over a closed-source dataset [34].
(8) *Infly-RL-SQL-32B* is a fine-tuned model without publicly disclosed information about its fine-tuning process and usage [56].[2] We used the same inference setup as OmniSQL-32B.
(9) *Arctic-R1-7B* is an RLVR model based on Qwen2.5-Coder-7B-Instruct over 28k examples from BIRD Train set, Spider Train set, Spider Dev set, and synthetic data [68].
(10) *GPT-5.2* is a frontier reasoning LLM. We execute it using the same inference settings as our models.

Metrics. We use two metrics. First, following prior work [24, 30], we measure the accuracy of generated SQL queries using their execution accuracy with set-based grading. Second, we measure the inference cost for each SQL query. For methods based on proprietary models, we calculate the costs based on the official pricing. For methods involving open-source models, we estimate their token costs using the cheapest pricing of their corresponding base models provided by Together [60], Groq [55], Sail [44], or Fireworks [54]. We used pricing data from May 1, 2026. In addition, we compute confidence intervals and run statistical significance tests when comparing a method against the human-level proxy and defer the technical details of these statistical procedures to Appendix C.

Implementation details. For the training stage of ReViSQL-BIRD, we used Tinker [61] as the infrastructure. Each intermediate query issued by the model is subject to a 30-second timeout. We conducted our training experiments on a machine with an Intel 8568Y CPU, 512 GB memory, and ran open-source models on one H100.

5.2 ReViSQL-BIRD Reaches Human-Level Accuracy on BIRD

We instantiated ReViSQL-BIRD on Kimi-K2.6 [57] and evaluated the resulting model, ReViSQL-BIRD-K2.6, on Arcwise-Plat-SQL and Arcwise-Plat-Full under two decoding methods: (1) greedy decoding and (2) self-consistency over n candidates. For self-consistency, we follow the standard procedure used in prior work [15, 17, 47, 63, 67]. We sample n candidates with temperature 1 and top-p 1, group them by execution results, and randomly return a query from the largest group, breaking ties uniformly at random. We repeat this process five times and report the average accuracy.

ReViSQL-BIRD-K2.6 is the first method to achieve human-level accuracy. As shown in Figure 6b, with greedy decoding, ReViSQL-BIRD-K2.6 achieves 93.8% (95% CI [91.3, 95.6]) on Arcwise-Plat-Full. As shown in Figure 6a, with self-consistency over 16 candidates, ReViSQL-BIRD-K2.6 achieves 93.0% (95% CI [90.4, 94.9]) on Arcwise-Plat-SQL. For both results, the human-level accuracy of 92.96% [52] lies within the 95% confidence interval. In contrast, every prior system we evaluate falls significantly below the human proxy (e.g., OpenSearch at 82.9%/88.2%, $p < 10^{-4}$). To our knowledge, ReViSQL-BIRD-K2.6 is the first Text-to-SQL method to achieve human parity on expert-verified BIRD datasets.

ReViSQL-BIRD-K2.6 dominates complex pipelines. As shown in Figures 6a and 6b, the prior best open-source pipeline, OpenSearch (based on GPT-5.2), achieves 82.9% on Arcwise-Plat-SQL and 88.2% on Arcwise-Plat-Full at a cost of $0.056 per query. ReViSQL-BIRD-K2.6 with greedy decoding outperforms OpenSearch by 8.4% on Arcwise-Plat-SQL and 5.6% on Arcwise-Plat-Full, *while costing 37% less per query*. Compared to other pipelines, ReViSQL-BIRD-K2.6 with greedy decoding achieves 9.6–20.3% and 8.8–18.1% higher accuracy, respectively. With self-consistency over up to 32 candidates, ReViSQL-BIRD-K2.6 outperforms all the complex pipelines by 9.6–22.1% and 5.8–18.5% on Arcwise-Plat-SQL and -Full, respectively.

ReViSQL-BIRD-K2.6 outperforms open-weight and frontier models. In Figures 6c and 6d, we compare ReViSQL-BIRD-K2.6 against the best open-weight Text-to-SQL models and GPT-5.2, as the number of generated candidates varies from 4 to 32. Across all candidate budgets and both benchmarks, ReViSQL-BIRD-K2.6

[2]We contacted the authors to request instructions for using the model. However, we have not received a response yet.

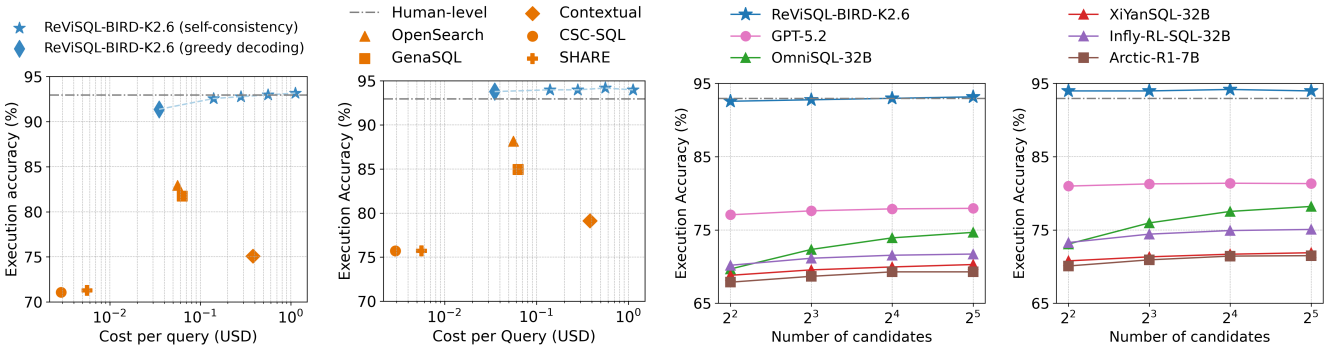

(a) Arcwise-Plat-SQL. (b) Arcwise-Plat-Full. (c) Arcwise-Plat-SQL. (d) Arcwise-Plat-Full.

Figure 6: ReViSQL-BIRD-K2.6 achieves human-level accuracy on BIRD while dominating prior methods. (a, b) Accuracy vs. cost on Arcwise-Plat-SQL and Arcwise-Plat-Full, comparing ReViSQL-BIRD-K2.6 against the top 5 open-source methods on BIRD Leaderboard. (c, d) Accuracy vs. candidate count, comparing ReViSQL-BIRD-K2.6 against single-model baselines.

achieves 12.7–24.7% higher accuracy. Furthermore, ReViSQL-BIRD-K2.6 with 4 candidates exceeds the performance of all baselines even when the latter are given a budget of 32 candidates. This result indicates that test-time compute fails to compensate for weak intrinsic model capabilities on Text-to-SQL.

Verified data is a dominant factor while reward shaping bridges the gap to human parity. To isolate the contribution of each ingredient, we ablate ReViSQL-BIRD-K2.6 under greedy decoding on the Arcwise-Plat benchmarks (Figure 7), comparing ReViSQL-BIRD against three configurations: (1) ReViSQL (no reward shaping), (2) ReViSQL-BIRD on BIRD Train, and (3) the base model with the same prompts as ReViSQL-BIRD. All configurations share the same hyperparameters and checkpoint-selection method discussed in Section 4.3.

We find that verified data is the dominant driver. Although BIRD Train contains 3.8× more instances than BIRD-Platinum, training on BIRD-Platinum improves the base model by 7.2% and 5.0% on Arcwise-Plat-SQL and -Full, whereas training on BIRD Train lowers it by 7.0% and 5.0%. We find that this is because rewarding a model against erroneous gold queries teaches it to reproduce the annotation errors (e.g., the `DISTINCT` omission discussed in §3.3.2). Furthermore, reward shaping bridges the residual gap, improving the accuracy by 2.8% and 3.8%. While these absolute gains appear modest, they are decisive at this accuracy level. Reward shaping removes 24.6% and 38.0% of the errors that verified data alone leaves behind. It is what lifts ReViSQL-BIRD-K2.6 to human parity.

The residual gap is dominated by question ambiguity. To characterize the residual gap, we manually inspected every ReViSQL-BIRD-K2.6 self-consistency failure across both Arcwise-Plat benchmarks, classifying each as (1) question ambiguity, where the model produces a valid reading of a natural question that differs from the gold's, or (2) a genuine model failure. We find that 69% of the cases are attributable to question ambiguity, and 31% to model defects. Specifically, Arcwise-Plat-SQL carries more ambiguity-driven failures than Arcwise-Plat-Full (24 vs. 20), which is consistent with its preserved noisy questions. Among genuine model failures, 60% of the cases are due to omitting a schema-induced filter (e.g., filtering

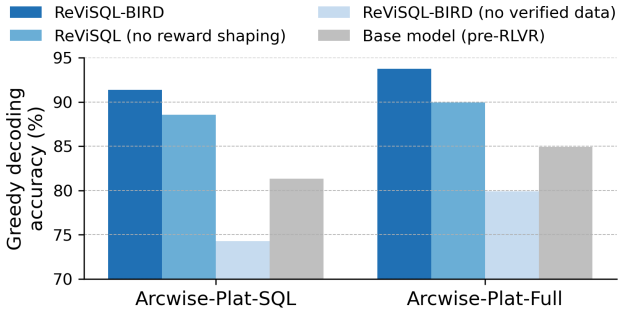

Figure 7: Ablation of ReViSQL-BIRD-K2.6 on Arcwise-Plat-SQL and -Full under greedy decoding. Training on BIRD-Platinum rather than the original BIRD Train improves accuracy by 13.9–17.1%. Reward shaping then further improves accuracy by 2.8–3.8%, reducing error rates by 24.6–38.0%.

`NULL` values). We provide the full taxonomy, per-type analysis, and case studies in Appendix B.

5.3 BIRD-Platinum Generalizes Across Benchmarks

We evaluate whether our curated training data, BIRD-Platinum, and our ReViSQL recipe (RLVR on BIRD-Platinum) generalize beyond BIRD. To isolate the data-quality contribution from any algorithmic effect, we trained Qwen3-235B-A22B using ReViSQL without the reward shaping methods discussed in Section 4. We then evaluated ReViSQL-235B on all four benchmarks.

ReViSQL-235B is trained on BIRD-Platinum, where all SQL queries are in SQLite syntax. Therefore, evaluating it on Spider2-Snow, which only accepts Snowflake-dialect queries, requires two inference-time adaptations without re-training. First, we translated the intermediate exploratory SQL queries issued by the model from SQLite to Snowflake using SQLGlot [59]. When SQLGlot fails, we fall back to a minimal automatic conversion that quotes identifiers with double quotes. Second, because Spider2-Snow databases have schemas and

Table 4: ReViSQL-235B (RLVR on BIRD-Platinum) outperforms the strongest prior open-source pipelines across two BIRD-derived benchmarks (Arcwise-Plat-SQL, Arcwise-Plat-Full) and the out-of-distribution Spider 2-SQLite benchmark, at lower per-query cost, demonstrating the generalizability of BIRD-Platinum.

Method	Arcwise-Plat-SQL (%)	Arcwise-Plat-Full (%)	Spider 2-SQLite (%)	Per-query cost (USD)
SHARE (GPT-5.2)	71.29	75.70	21.48	5.6×10^{-3}
CSC-SQL (XiYan-32B)	71.08	75.70	10.37	2.9×10^{-3}
Contextual-SQL (XiYan-32B)	75.10	79.12	11.11	3.8×10^{-1}
GenaSQL (GPT-5.2)	81.73	84.94	34.81	6.3×10^{-2}
OpenSearch (GPT-5.2)	82.93	88.15	34.07*	5.6×10^{-2}
ReViSQL-235B (greedy)	**87.34**	**88.75**	**37.04**	7.7×10^{-3}

* OpenSearch [67] relies on training data for dynamic few-shot prompting, which Spider 2 does not provide. We tested with substituting the BIRD and Spider 1 training sets, as well as disabling the module entirely. We report the performance with dynamic few-shot disabled, as it achieved the highest performance.

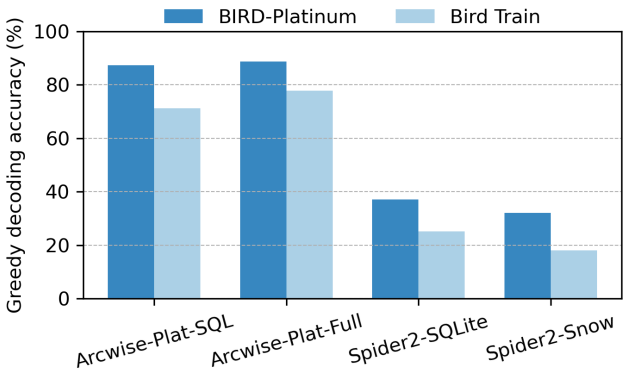

Figure 8: Training Qwen3-235B-A22B with RLVR on BIRD-Platinum (ReViSQL) yields consistent accuracy improvements (11.0–16.1%) over training on the original BIRD Train, across Arcwise-Plat and Spider2 benchmarks.

questions whose combined length exceeds our fine-tuned model's context window, we applied the column- and table-filtering procedure from ReFoRCE [13] to select schema elements relevant to each question. All other inference settings match our BIRD evaluation.

ReViSQL improves accuracy on all benchmarks. In Figure 8, we compare ReViSQL-235B trained on BIRD-Platinum against the same base model trained on the original BIRD Train, under greedy decoding. On the BIRD-derived Arcwise-Plat benchmarks, ReViSQL-235B outperforms its BIRD Train counterpart by 16.1% on Arcwise-Plat-SQL and 11.0% on Arcwise-Plat-Full. Furthermore, the improvement persists on Spider2 benchmarks. ReViSQL-235B achieves 11.8% and 14.0% higher accuracy on Spider2-SQLite and Spider2-Snow compared to training on BIRD Train. The improvements across benchmarks (different SQL query complexities and dialects) demonstrate that the data-quality effect is not specific to BIRD. RLVR with BIRD-Platinum produces Text-to-SQL models that generalize more broadly.

ReViSQL-235B outperforms the open-source pipelines on Spider2-SQLite. In Table 4, we compare ReViSQL-235B against the top 5 open-source pipelines on the BIRD Leaderboard. On Spider2-SQLite, with greedy decoding, ReViSQL-235B achieves 37.04% accuracy at $0.0077 per query, outperforming the previous best pipeline (GenaSQL) by 2.2% at 8× lower cost, and outperforming OpenSearch by 3% at 7× lower cost. ReViSQL-235B also leads on Arcwise-Plat-SQL and Arcwise-Plat-Full, outperforming baselines by 0.6–16.3%.

BIRD-Platinum generalizes to Snowflake dialects. On the more challenging Spider2-Snow benchmark, which requires Snowflake-dialect SQL, ReViSQL-235B with self-consistency over 128 candidates achieves 55.58% accuracy, outperforming agentic pipelines built on substantially larger models, including ReFoRCE and AutoLink on DeepSeek (671B) and Spider-Agent on Qwen3-Coder (480B), by 17.55%, 0.74%, and 24.5% [58]. Since Spider2-Snow is held out from ReViSQL's BIRD-Platinum training data, this result provides evidence that ReViSQL's improvements transfer to new SQL dialects.

6 Conclusion

We show that human-level Text-to-SQL accuracy is achievable by fine-tuning a single LLM with RLVR on clean data, without the complex pipelines that dominate top-ranking LLM-based methods. We introduce two advances in RLVR for Text-to-SQL. First, we develop BIRD-Platinum, a 2.5k-instance verified dataset constructed through a hybrid LLM–human verification pipeline. We repair the annotation errors in 61% of the sampled instances from BIRD Train, yielding 11–16% accuracy improvements over training on BIRD Train under the same RLVR algorithm. Second, we introduce ReViSQL-BIRD, a specialized reward shaping method designed for the failure modes of RLVR in Text-to-SQL tasks. We propose a verification-aware outcome reward to mitigate the false positive rewards in standard result-based rewards, and leverage rule-based process rewards to prevent models from ignoring external knowledge. We trained ReViSQL-BIRD-K2.6, the first method to match human-level accuracy (92.96%) on expert-verified variants of BIRD, outperforming top open-source methods by 10–22%.

Acknowledgments

This research was supported in part by a sponsorship from Bridgewater AIA Labs. We are grateful to the CloudLab for providing computing resources for experiments [16].

References

[1] Katrin Affolter, Kurt Stockinger, and Abraham Bernstein. 2019. A comparative survey of recent natural language interfaces for databases. *The VLDB Journal* 28, 5 (2019), 793–819.

[2] Sheshansh Agrawal and Thien Nguyen. 2025. Open-Sourcing the Best Local Text-to-SQL System. https://contextual.ai/blog/open-sourcing-the-best-local-text-to-sql-system/

[3] Ion Androutsopoulos, Graeme D Ritchie, and Peter Thanisch. 1995. Natural language interfaces to databases–an introduction. *Natural language engineering* 1, 1 (1995), 29–81.

[4] Arcwise. 2025. BIRD minidev - corrections. https://docs.google.com/spreadsheets/d/1IGm9Otruey60ujUnl8AOkepY3qgWHdFJHnX7hQGUeCw Accessed: 2025-09-15.

[5] Agent Team at JDCHO. 2025. *JoyAgent-JDGenie*. https://github.com/jd-opensource/joyagent-jdgenie

[6] Lawrence D Brown, T Tony Cai, and Anirban DasGupta. 2001. Interval estimation for a binomial proportion. *Statistical science* 16, 2 (2001), 101–133.

[7] Mert Cemri, Melissa Z Pan, Shuyi Yang, Lakshya A Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, et al. 2025. Why do multi-agent llm systems fail? *arXiv preprint arXiv:2503.13657* (2025).

[8] Aili Chen, Aonian Li, Bangwei Gong, Binyang Jiang, Bo Fei, Bo Yang, Boji Shan, Changqing Yu, Chao Wang, Cheng Zhu, et al. 2025. Minimax-m1: Scaling test-time compute efficiently with lightning attention. *arXiv preprint arXiv:2506.13585* (2025).

[9] Chunyu Chen, Zhengjie Miao, Yong Zhang, and Jiannan Wang. 2025. ParSEval: Plan-aware Test Database Generation for SQL Equivalence Evaluation. *Proceedings of the VLDB Endowment* 18, 11 (2025), 4750–4762.

[10] Kaiwen Chen, Yueting Chen, Nick Koudas, and Xiaohui Yu. 2025. Reliable Text-to-SQL with Adaptive Abstention. *Proceedings of the ACM on Management of Data* 3, 1 (2025), 1–30.

[11] Team Cohere, Arash Ahmadian, Marwan Ahmed, Jay Alammar, Milad Alizadeh, Yazeed Alnumay, Sophia Althammer, Arkady Arkhangorodsky, Viraat Aryabumi, Dennis Aumiller, et al. 2025. Command a: An enterprise-ready large language model. *arXiv preprint arXiv:2504.00698* (2025).

[12] Databricks. 2025. *The Power of RLVR: Training a Leading SQL Reasoning Model on Databricks*. https://www.databricks.com/blog/power-rlvr-training-leading-sql-reasoning-model-databricks Accessed: 2025-10-10.

[13] Minghang Deng, Ashwin Ramachandran, Canwen Xu, Lanxiang Hu, Zhewei Yao, Anupam Datta, and Hao Zhang. 2025. ReFoRCE: a text-to-SQL agent with self-refinement, consensus enforcement, and column exploration. *arXiv preprint arXiv:2502.00675* (2025).

[14] Yusuf Denizay Dönder, Derek Hommel, Andrea W Wen-Yi, David Mimno, and Unso Eun Seo Jo. 2025. Cheaper, Better, Faster, Stronger: Robust Text-to-SQL without Chain-of-Thought or Fine-Tuning. *arXiv preprint arXiv:2505.14174* (2025).

[15] Xuemei Dong, Chao Zhang, Yuhang Ge, Yuren Mao, Yunjun Gao, Jinshu Lin, Dongfang Lou, et al. 2023. C3: Zero-shot text-to-sql with chatgpt. *arXiv preprint arXiv:2307.07306* (2023).

[16] Dmitry Duplyakin, Robert Ricci, Aleksander Maricq, Gary Wong, Jonathon Duerig, Eric Eide, Leigh Stoller, Mike Hibler, David Johnson, Kirk Webb, Aditya Akella, Kuangching Wang, Glenn Ricart, Larry Landweber, Chip Elliott, Michael Zink, Emmanuel Cecchet, Snigdhaswin Kar, and Prabodh Mishra. 2019. The Design and Operation of CloudLab. In *Proceedings of the USENIX Annual Technical Conference (ATC)*. 1–14. https://www.flux.utah.edu/paper/duplyakin-atc19

[17] Dawei Gao, Haibin Wang, Yaliang Li, Xiuyu Sun, Yichen Qian, Bolin Ding, and Jingren Zhou. 2023. Text-to-sql empowered by large language models: A benchmark evaluation. *arXiv preprint arXiv:2308.15363* (2023).

[18] Dawei Gao, Haibin Wang, Yaliang Li, Xiuyu Sun, Yichen Qian, Bolin Ding, and Jingren Zhou. 2024. Text-to-SQL Empowered by Large Language Models: A Benchmark Evaluation. *Proceedings of the VLDB Endowment* 17, 5 (2024), 1132–1145.

[19] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948* (2025).

[20] Yang He, Pinhan Zhao, Xinyu Wang, and Yuepeng Wang. 2024. VeriEQL: Bounded equivalence verification for complex SQL queries with integrity constraints. *Proceedings of the ACM on Programming Languages* 8, OOPSLA1 (2024), 1071–1099.

[21] Jie Jiang, Haining Xie, Siqi Shen, Yu Shen, Zihan Zhang, Meng Lei, Yifeng Zheng, Yang Li, Chunyou Li, Danqing Huang, et al. 2024. Siriusbi: A comprehensive llm-powered solution for data analytics in business intelligence. *arXiv preprint arXiv:2411.06102* (2024).

[22] Tengjun Jin, Yoojin Choi, Yuxuan Zhu, and Daniel Kang. 2026. Pervasive Annotation Errors Break Text-to-SQL Benchmarks and Leaderboards. *VLDB* (2026).

[23] George Katsogiannis-Meimarakis and Georgia Koutrika. 2023. A survey on deep learning approaches for text-to-SQL. *The VLDB Journal* 32, 4 (2023), 905–936.

[24] Fangyu Lei, Jixuan Chen, Yuxiao Ye, Ruisheng Cao, Dongchan Shin, Hongjin Su, Zhaoqing Suo, Hongcheng Gao, Wenjing Hu, Pengcheng Yin, et al. 2024. Spider 2.0: Evaluating language models on real-world enterprise text-to-sql workflows. *arXiv preprint arXiv:2411.07763* (2024).

[25] Boyan Li, Yuyu Luo, Chengliang Chai, Guoliang Li, and Nan Tang. 2024. The Dawn of Natural Language to SQL: Are We Fully Ready? *Proceedings of the VLDB Endowment* 17, 11 (2024), 3318–3331.

[26] Boyan Li, Jiayi Zhang, Ju Fan, Yanwei Xu, Chong Chen, Nan Tang, and Yuyu Luo. [n.d.]. Alpha-SQL: Zero-Shot Text-to-SQL using Monte Carlo Tree Search. In *Forty-second International Conference on Machine Learning*.

[27] Fei Li and Hosagrahar V Jagadish. 2014. Constructing an interactive natural language interface for relational databases. *Proceedings of the VLDB Endowment* 8, 1 (2014), 73–84.

[28] Haoyang Li, Shang Wu, Xiaokang Zhang, Xinmei Huang, Jing Zhang, Fuxin Jiang, Shuai Wang, Tieying Zhang, Jianjun Chen, Rui Shi, et al. 2025. Omnisql: Synthesizing high-quality text-to-sql data at scale. *arXiv preprint arXiv:2503.02240* (2025).

[29] Haoyang Li, Jing Zhang, Hanbing Liu, Ju Fan, Xiaokang Zhang, Jun Zhu, Renjie Wei, Hongyan Pan, Cuiping Li, and Hong Chen. 2024. Codes: Towards building open-source language models for text-to-sql. *Proceedings of the ACM on Management of Data* 2, 3 (2024), 1–28.

[30] Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. 2023. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. *Advances in Neural Information Processing Systems* 36 (2023), 42330–42357.

[31] Shu Liu, Sumanth Hegde, Shiyi Cao, Alan Zhu, Dacheng Li, Tyler Griggs, Eric Tang, Akshay Malik, Kourosh Hakhamaneshi, Richard Liaw, Philipp Moritz, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. 2025. SkyRL-SQL: Matching GPT-4o and o4-mini on Text2SQL with Multi-Turn RL. Notion Blog.

[32] Xinyu Liu, Shuyu Shen, Boyan Li, Peixian Ma, Runzhi Jiang, Yuxin Zhang, Ju Fan, Guoliang Li, Nan Tang, and Yuyu Luo. 2025. A Survey of Text-to-SQL in the Era of LLMs: Where are we, and where are we going? *IEEE Transactions on Knowledge and Data Engineering* (2025).

[33] Xinyu Liu, Shuyu Shen, Boyan Li, Nan Tang, and Yuyu Luo. 2025. Nl2sql-bugs: A benchmark for detecting semantic errors in nl2sql translation. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*. 5662–5673.

[34] Yifu Liu, Yin Zhu, Yingqi Gao, Zhiling Luo, Xiaoxia Li, Xiaorong Shi, Yuntao Hong, Jinyang Gao, Yu Li, Bolin Ding, et al. 2025. Xiyan-sql: A novel multi-generator framework for text-to-sql. *arXiv preprint arXiv:2507.04701* (2025).

[35] Peixian Ma, Xialie Zhuang, Chengjin Xu, Xuhui Jiang, Ran Chen, and Jian Guo. 2025. Sql-r1: Training natural language to sql reasoning model by reinforcement learning. *arXiv preprint arXiv:2504.08600* (2025).

[36] Karime Maamari, Fadhil Abubaker, Daniel Jaroslawicz, and Amine Mhedhbi. [n.d.]. The Death of Schema Linking? Text-to-SQL in the Age of Well-Reasoned Language Models. In *NeurIPS 2024 Third Table Representation Learning Workshop*.

[37] OpenAI. 2025. *Introducing OpenAI o3 and o4-mini*. https://openai.com/index/introducing-o3-and-o4-mini/ Accessed Dec. 10, 2025.

[38] Simone Papicchio, Simone Rossi, Luca Cagliero, and Paolo Papotti. 2025. Think2sql: Reinforce llm reasoning capabilities for text2sql. *arXiv preprint arXiv:2504.15077* (2025).

[39] Mohammadreza Pourreza, Hailong Li, Ruoxi Sun, Yeounoh Chung, Shayan Talaei, Gaurav Tarlok Kakkar, Yu Gan, Amin Saberi, Fatma Ozcan, and Sercan O Arik. [n.d.]. CHASE-SQL: Multi-Path Reasoning and Preference Optimized Candidate Selection in Text-to-SQL. In *The Thirteenth International Conference on Learning Representations*.

[40] Mohammadreza Pourreza and Davood Rafiei. 2023. Evaluating Cross-Domain Text-to-SQL Models and Benchmarks. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. 1601–1611.

[41] Mohammadreza Pourreza, Ruoxi Sun, Hailong Li, Lesly Miculicich, Tomas Pfister, and Sercan O Arik. 2024. Sql-gen: Bridging the dialect gap for text-to-sql via synthetic data and model merging. *arXiv preprint arXiv:2408.12733* (2024).

[42] Mohammadreza Pourreza, Shayan Talaei, Ruoxi Sun, Xingchen Wan, Hailong Li, Azalia Mirhoseini, Amin Saberi, Sercan Arik, et al. 2025. Reasoning-sql: Reinforcement learning with sql tailored partial rewards for reasoning-enhanced text-to-sql. *arXiv preprint arXiv:2503.23157* (2025).

[43] Ge Qu, Jinyang Li, Bowen Qin, Xiaolong Li, Nan Huo, Chenhao Ma, and Reynold Cheng. 2025. SHARE: An SLM-based Hierarchical Action CorREction Assistant for Text-to-SQL. *arXiv preprint arXiv:2506.00391* (2025).

[44] Sail Research. 2025. Sail Models and Pricing Page. https://docs.sailresearch.com/pricing Accessed: 2026-06-12.

[45] Jaydeep Sen, Fatma Ozcan, Abdul Quamar, Greg Stager, Ashish Mittal, Manasa Jammi, Chuan Lei, Diptikalyan Saha, and Karthik Sankaranarayanan. 2019. Natural language querying of complex business intelligence queries. In *Proceedings of the 2019 international conference on management of data*. 1997–2000.

[46] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Yang Wu, et al. 2024. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint*

arXiv:2402.03300 (2024).

[47] Lei Sheng and Shuai-Shuai Xu. 2025. CSC-SQL: Corrective Self-Consistency in Text-to-SQL via Reinforcement Learning. *arXiv preprint arXiv:2505.13271* (2025).

[48] Lei Sheng and Shuai-Shuai Xu. 2025. SLM-SQL: An Exploration of Small Language Models for Text-to-SQL. *arXiv preprint arXiv:2507.22478* (2025).

[49] Vladislav Shkapenyuk, Divesh Srivastava, Theodore Johnson, and Parisa Ghane. 2025. Automatic Metadata Extraction for Text-to-SQL. *arXiv preprint arXiv:2505.19988* (2025).

[50] Joar Skalse, Nikolaus HR Howe, Dmitrii Krasheninnikov, and David Krueger. 2022. Defining and characterizing reward hacking. In *Proceedings of the 36th International Conference on Neural Information Processing Systems*. 9460–9471.

[51] Shayan Talaei, Mohammadreza Pourreza, Yu-Chen Chang, Azalia Mirhoseini, and Amin Saberi. 2024. Chess: Contextual harnessing for efficient sql synthesis. *arXiv preprint arXiv:2405.16755* (2024).

[52] BIRD Team. 2024. *BIRD-bench*. https://bird-bench.github.io/ Accessed: 2025-02-10.

[53] BIRD Team. 2025. Panel of BIRD Annotation Issues. https://github.com/AlibabaResearch/DAMO-ConvAI/issues/39 Accessed: 2025-10-15.

[54] Fireworks Team. 2025. Fireworks Pricing Page. https://fireworks.ai/pricing Accessed: 2026-02-10.

[55] Groq Team. 2025. Groq Pricing Page. https://groq.com/pricing Accessed: 2025-10-10.

[56] Infly Team. 2025. infly/inf-rl-qwen-coder-32b-2746. https://huggingface.co/infly/inf-rl-qwen-coder-32b-2746

[57] Kimi Team. 2026. *Kimi K2.6: Advancing Open-Source Coding*. https://www.kimi.com/blog/kimi-k2-6 Accessed: 2026-05-10.

[58] Spider 2.0 Team. 2025. *Spider 2.0*. https://spider2-sql.github.io/ Accessed: 2026-08-09.

[59] SQLGlot Team. 2025. *SQLGlot: Python SQL Parser and Transpiler*. https://github.com/tobymao/sqlglot Accessed: 2025-10-10.

[60] Together AI Team. 2025. Together AI Pricing Page. https://together.ai/pricing Accessed: 2025-10-10.

[61] Tinker. 2025. Tinker: a training API for researchers and developers. https://tinker-docs.thinkingmachines.ai/ Accessed: 2025-10-10.

[62] Pengfei Wang, Baolin Sun, Xuemei Dong, Yaxun Dai, Hongwei Yuan, Mengdie Chu, Yingqi Gao, Xiang Qi, Peng Zhang, and Ying Yan. 2025. Agentar-Scale-SQL: Advancing Text-to-SQL through Orchestrated Test-Time Scaling. *arXiv preprint arXiv:2509.24403* (2025).

[63] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2022. Self-consistency improves chain of thought reasoning in language models. *arXiv preprint arXiv:2203.11171* (2022).

[64] Edwin B Wilson. 1927. Probable inference, the law of succession, and statistical inference. *J. Amer. Statist. Assoc.* 22, 158 (1927), 209–212.

[65] Cornelius Wolff, Daniel Gomm, and Madelon Hulsebos. [n.d.]. SQaLe: A large text-to-SQL corpus grounded in real schemas. In *EurIPS 2025 Workshop: AI for Tabular Data*.

[66] Niklas Wretblad, Fredrik Riseby, Rahul Biswas, Amin Ahmadi, and Oskar Holmström. 2024. Understanding the Effects of Noise in Text-to-SQL: An Examination of the BIRD-Bench Benchmark. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*. 356–369.

[67] Xiangjin Xie, Guangwei Xu, Lingyan Zhao, and Ruijie Guo. 2025. Opensearch-sql: Enhancing text-to-sql with dynamic few-shot and consistency alignment. *Proceedings of the ACM on Management of Data* 3, 3 (2025), 1–24.

[68] Zhewei Yao, Guoheng Sun, Lukasz Borchmann, Zheyu Shen, Minghang Deng, Bohan Zhai, Hao Zhang, Ang Li, and Yuxiong He. 2025. Arctic-Text2SQL-R1: Simple Rewards, Strong Reasoning in Text-to-SQL. *arXiv preprint arXiv:2505.20315* (2025).

[69] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, et al. 2025. Dapo: An open-source llm reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476* (2025).

[70] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, et al. 2018. Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*. 3911–3921.

[71] Bohan Zhai, Canwen Xu, Yuxiong He, and Zhewei Yao. 2025. ExCoT: Optimizing reasoning for text-to-SQL with execution feedback. *arXiv preprint arXiv:2503.19988* (2025).

A Case Study: The Two Expert-Verified Arcwise-Plat Benchmarks

We evaluate on two expert-verified variants of the BIRD Mini-Dev set released by prior work [4, 22]: Arcwise-Plat-Full and Arcwise-Plat-SQL. The two are built from the same instances but differ in which sources of annotation error are corrected (Table 2). Arcwise-Plat-Full corrects all three sources: the natural language question, the external knowledge, and the gold SQL query. Therefore, it measures Text-to-SQL accuracy under clean inputs. Arcwise-Plat-SQL corrects only the gold SQL query and *deliberately preserves* the original, often noisy, questions and external knowledge. Hence, it measures robustness to the ambiguous inputs that are typical of real-world Text-to-SQL deployments [32]. We walk through a representative instance to make the distinction concrete.

A representative instance. In Figure 9, we show one BIRD Mini-Dev instance under all three versions. The original instance (Figure 9a) contains an error in each of the three sources. The question asks about schools serving "grades 1 through 12," an incorrect phrasing of the intended "K-12" criterion; the external knowledge includes a redundant entry; and the gold SQL query omits a `NOT NULL` filter, so its result wrongly includes rows about districts in addition to schools.

Arcwise-Plat-Full repairs all three (Figure 9b): it rephrases the question to "K-12," removes the redundant external-knowledge entry, and augments the gold SQL with the missing filter. A model evaluated here sees an unambiguous question and is graded against a correct gold query, so the benchmark isolates the model's SQL-reasoning ability.

Arcwise-Plat-SQL repairs only the gold SQL (Figure 9c): the missing `NOT NULL` filter is added, but the noisy "grades 1 through 12" question and the redundant external knowledge are left untouched. A model must now recover the intended "K-12" meaning from an imperfect question without the benefit of a sanitized prompt. This is the harder, deployment-realistic condition: the input is exactly as noisy as the original benchmark, yet the gold answer is now correct, so a model is neither rewarded for matching an erroneous gold query nor rescued by a cleaned question.

Implication. Contrasting the two benchmarks isolates two distinct capabilities. A gap between a method's Arcwise-Plat-Full and Arcwise-Plat-SQL accuracy reflects its sensitivity to noisy natural-language inputs rather than its raw SQL competence. This is consistent with the headline result in Section 5.2: ReViSQL-BIRD-K2.6 reaches human-level accuracy on the clean Arcwise-Plat-Full with greedy decoding, but recovers the same level on the noisier Arcwise-Plat-SQL only with self-consistency over multiple candidates, indicating that inference-time scaling chiefly compensates for question ambiguity rather than for limited SQL reasoning.

B Detailed Analysis of Residual Failures

In this appendix, we report a case-by-case analysis of the failures that remain after ReViSQL-BIRD-K2.6 reaches human-level accuracy (Section 5.2). We aim to answer the question: of the small residual gap, how much reflects genuine limitations of the model's SQL generation, and how much reflects inherent ambiguity in the natural-language questions? We treat the expert-verified Arcwise-Plat gold labels as correct throughout and attribute every disagreement either to question ambiguity or to a model defect.

Table 5: Classification of the 64 residual ReViSQL-BIRD-K2.6 failures (the most frequent incorrect self-consistency query on each problem where the model disagrees with the gold, across both Arcwise-Plat benchmarks). Most failures trace to ambiguity in the question rather than to defects in the model's SQL generation.

Attribution	Mechanism	# cases
Question ambiguity	Entity / column / value mapping	9
	Output specification	20
	Aggregation / filter scope	15
	Subtotal	**44 (69%)**
Model defect	Schema / value grounding	12
	Ignored external knowledge	3
	Join / output / aggregation	5
	Subtotal	**20 (31%)**
Total		**64**

B.1 Scope and Method

Failure set. We collected, for every problem on which ReViSQL-BIRD-K2.6's self-consistency answer disagreed with the gold, the single most frequent incorrect query the model produced. This yields 64 cases spanning 43 distinct problems over 11 databases, 30 from Arcwise-Plat-Full and 34 from Arcwise-Plat-SQL (a problem may appear in both benchmarks). For each case we examined the question, the external knowledge, the gold query, the model's query, and the database schema.

Two-axis taxonomy. We classify each case along two axes. The first is the source of the disagreement: *question ambiguity*, where the question admits more than one reasonable reading and the model realizes a valid interpretation other than the one the gold encodes; or a *model defect*, where the model's query is genuinely incorrect. The second axis is the *mechanism*, refined into three categories on each side (Table 5). Two of the authors conduct the classification manually and reach agreement via discussion.

B.2 Findings

The residual gap is dominated by question ambiguity, not SQL defects. As summarized in Table 5, 44 of the 64 cases (69%) arise from question ambiguity: the model selects a defensible reading of an under-specified question that differs from the gold's. Only 20 cases (31%) are genuine model defects. This indicates that ReViSQL-BIRD-K2.6's SQL-generation competence is largely saturated on Arcwise-Plat, and that the residual measured error is dominated by the interpretation of ambiguous natural language rather than by the model's ability to write correct SQL.

Model defects are stable across input noise, while ambiguity is not. The genuine model defects are evenly split between the

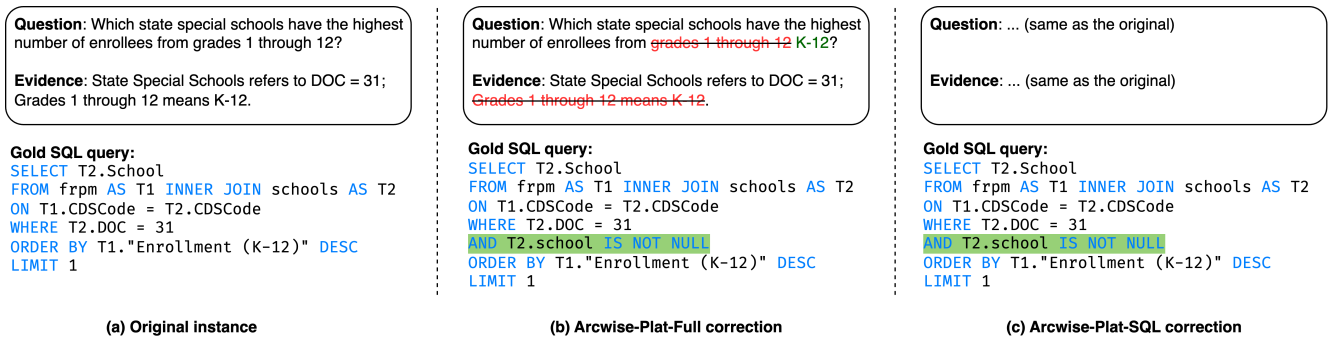

Figure 9: A representative instance of two expert-verified variants of BIRD Mini-Dev. (a) The original instance contains an incorrect phrasing in the question ("grades 1 through 12" as opposed to "K-12") and a gold SQL query that omits a `NOT NULL` filter, causing it to include data about districts and schools. (b) Arcwise-Plat-Full corrects all three error sources: rephrasing the question to "K-12," removing the redundant external-knowledge entry, and augmenting gold SQL. (c) Arcwise-Plat-SQL corrects only the gold SQL, preserving the original noisy question and external knowledge.

two benchmarks (10 on Arcwise-Plat-Full and 10 on Arcwise-Plat-SQL), whereas Arcwise-Plat-SQL carries more ambiguity-driven failures than Arcwise-Plat-Full (24 vs. 20). Because Arcwise-Plat-SQL preserves the original noisy questions and external knowledge while Arcwise-Plat-Full rewrites them, this is exactly the pattern expected if the model's intrinsic SQL competence is fixed and the additional Plat-SQL failures come from unresolved question ambiguity. It also explains why self-consistency, which aggregates over interpretations, recovers human-level accuracy on the noisier Arcwise-Plat-SQL (Section 5.2).

When the model fails, it under-constrains. Of the 20 model defects, 12 are *schema/value-grounding* errors in which the model omits an existence or validity filter the schema requires, leaving 3 cases of ignoring otherwise-usable external knowledge and 5 join/output/aggregation errors. The model rarely produces logically scrambled SQL. In contrast, when the model fails, it typically writes a locally reasonable query that fails to encode a schema-specific constraint.

B.3 Question Ambiguity (44 cases)

Entity, column, or value mapping (9 cases). The question does not uniquely determine which schema element it refers to. In Q39, "Fresno schools" is graded against `County = 'Fresno'` while the model filters `City = 'Fresno'`. Fresno is both a city and a county. In Q892, "the driver with the most points scored" is computed by the gold over the per-race `results.points` column and by the model over the cumulative `driverStandings.points` column. The schema exposes two columns named `points` with different semantics.

Output specification (20 cases). The question leaves the exact output under-determined, such as which columns to return, in what form, or how a value should be matched or scaled. In some cases, the model's choice differs from the gold's. In Q587, "posts tagged humor" is matched by the gold with exact string equality (`Tags = '<humor>'`) and by the model with substring membership (`Tags LIKE '%<humor>%'`). The two differ on posts that carry additional tags. In Q894, "list the recorded lap time" returns the human-readable `time` string in the model's query and the raw `milliseconds` integer in the gold. In Q1149, the question asks for a *ratio* of male in-patients to outpatients. The gold reports it scaled to a percentage while the model returns the unscaled ratio.

Aggregation or filter scope (15 cases). The question leaves the set being aggregated, or the handling of ties and multiplicities, open to reading. In Q197, "single-bonded molecules" is read by the gold as molecules with *at least one* single bond and by the model as molecules whose bonds are *all* single. In Q1080, "players who prefer the left foot … and have at least one record indicating low work rate" is encoded by the gold as both conditions holding in the *same* attribute record, while the model, following the "at least one record" phrasing, matches the two conditions across different records of the same player.

B.4 Model Defects (20 cases)

Schema/value grounding (12 cases). The model omits a constraint the schema requires. In Q11, "codes of the schools" is graded against a query that filters `School IS NOT NULL` because the `frpm` table contains non-school rows; the model omits the filter and over-returns—the same omission the case study in Appendix A illustrates.

```
-- gold: restricts frpm to real schools
... WHERE ("Enrollment_(K-12)" + "Enrollment_(Ages_5-17)") > 500
    AND School IS NOT NULL
-- model: omits the existence filter, over-returns non-school rows
SELECT CDSCode FROM frpm
WHERE ("Enrollment_(K-12)" + "Enrollment_(Ages_5-17)") > 500;
```

In Q1531, the model generates a query that computes average price per item as `SUM(Price)/SUM(Amount)`, not recognizing that `Price` is a per-unit value, so the schema-correct expression is `SUM(Price*Amount)/SUM(Amount)`.

Other join/output/aggregation logic (5 cases). These include an incorrect inner join that drops valid rows where a left join is required (Q773) and a query that omits a column the question requests (Q1235).

Ignored external knowledge (3 cases). In Q407, the evidence maps "all types of cards" to the `subtypes` and `supertypes` columns, but the model instead returns the separate `type` column, disregarding an explicit and usable hint.

B.5 The Full–SQL Contrast

In this section, we present two problems to highlight concretely why Arcwise-Plat-SQL is the harder benchmark, by changing character between the two variants of the *same* problem. In Q48, the Arcwise-Plat-Full evidence states the ratio's denominator explicitly. A model that nonetheless narrows that denominator is committing a genuine ignored-evidence defect. In Arcwise-Plat-SQL the same problem's evidence omits that formula, so the denominator's scope is genuinely ambiguous and the model's reading becomes defensible. In Q1486, Arcwise-Plat-Full asks "how much more do SMEs pay," whereas Arcwise-Plat-SQL preserves the original "is it true that more SMEs pay…if so, how many more," which invites a differently shaped answer. In both, the additional difficulty of Arcwise-Plat-SQL is attributable to preserved question ambiguity.

C Statistical Methodology

This appendix details how we compute the confidence intervals and significance tests reported in Section 5.2. All procedures operate on the per-instance binary correctness of each method under set-based execution accuracy, which we release together with our code. Each Arcwise-Plat benchmark contains $n = 498$ instances, and a method's accuracy is the mean of n independent Bernoulli outcomes.

Confidence intervals. Because the accuracies we report are high ($\hat{p} \gtrsim 0.9$), the normal (Wald) approximation has poor coverage near the boundary at this sample size. We therefore report the Wilson score interval [6, 64], which is well calibrated for proportions close to 1 and moderate n. For an observed accuracy $\hat{p} = x/n$ at confidence level $1 - \alpha$ with $z = z_{1-\alpha/2}$ (here $\alpha = 0.05$, $z \approx 1.96$), the interval is

$$\frac{\hat{p} + \dfrac{z^2}{2n} \;\pm\; z\sqrt{\dfrac{\hat{p}(1-\hat{p})}{n} + \dfrac{z^2}{4n^2}}}{1 + \dfrac{z^2}{n}}. \tag{4}$$

For example, the greedy-decoding accuracy of 93.8% on Arcwise-Plat-Full yields the 95% interval $[91.3, 95.6]$.

Comparison against the human proxy. We treat BIRD's reported human execution accuracy, $p_H = 92.96\%$ [52], as a fixed reference and use a one-sided exact binomial test. For each method we test $H_0 \colon p \geq p_H$ against $H_1 \colon p < p_H$ and call the method *significantly below human parity* when we reject H_0 at $\alpha = 0.05$. Every prior system we evaluate is significantly below p_H ($p < 10^{-4}$); the closest, OpenSearch on Arcwise-Plat-Full at 88.2%, gives $p \approx 1.4 \times 10^{-5}$. For ReViSQL-BIRD-K2.6, the same test fails to reject H_0 in every configuration ($p > 0.05$): its accuracy is not significantly below p_H, and equivalently p_H lies within the two-sided 95% Wilson interval of Eq. 4. We therefore report ReViSQL-BIRD-K2.6 as statistically indistinguishable from the human proxy, and as the first Text-to-SQL system for which this holds while all prior systems remain significantly below it.

D VeriEQL Configuration

We use VeriEQL [20], a bounded SQL equivalence verifier, in two places: to measure the false-positive rate of result-based rewards during diagnosis (§4.1), and as the verification-aware component of the ReViSQL-BIRD reward during training (§4.2). In both, VeriEQL is a *secondary* check layered on top of result-based grading, and it is disabled at inference time. This appendix specifies its configuration.

When VeriEQL is invoked. During training, VeriEQL is called only when a candidate query's execution result already matches the gold result. Its role is to test whether the two queries are *formally* equivalent, catching cases where the results coincide on the particular database instance without the queries being semantically equivalent (Failure Mode 1, §4.1).

Solver settings. We run VeriEQL with a bounded-verification bound of 2 tuples per table. Each check runs with a 30-second timeout during diagnosis and training.

Schema coarsening. We extract each table's schema automatically from the SQLite database file and map every column to one of VeriEQL's two supported types: `INT`, and `VARCHAR`. This coarsening, together with the bounded verification, means VeriEQL may return an inconclusive verdict, which we handle conservatively (§4.2).

E Prompts and Rollout Specification

This appendix specifies the exact prompts and rollout mechanics used for ReViSQL and ReViSQL-BIRD, during both RLVR training and inference. ReViSQL uses the base system prompt of §E.1, while ReViSQL-BIRD extends it with the two external-knowledge use directives of §E.2. The ReViSQL-BIRD's prompt is also what we give the pre-RLVR base model and any single-model baseline that does not publish a system prompt of its own (Section 5.1).

E.1 Base System Prompt (ReViSQL)

We use a universal system prompt for all Text-to-SQL problems, across all benchmarks and both SQL dialects. It defines the task, the input format, the output contract, and the multi-turn protocol.

```
Task Overview:
You are a data science expert. Below, you are provided with a database
       schema and a natural language question. Your task is to understand
       the schema and generate a valid SQL query to answer the question
       within limited turns. You should breakdown the problem, draft your
       reasoning process, and generate the solution.

The task you will receive should have the following format:

Database Engine:
{engine}

Database Schema:
<<db_details>>
This schema describes the database's structure, including tables, columns,
       primary keys, foreign keys, and any relevant relationships or
       constraints.

External Knowledge:
<<external_knowledge>>

Question:
<<question>>
```

```
Instructions:
- Make sure you only output the information that is asked in the question.
      If the question asks for a specific column, make sure to only include
       that column in the SELECT clause, nothing more.
- The generated query should return all of the information asked in the
      question without any missing or extra information.
- If the external knowledge is not empty, you must use all the information
      provided in the external knowledge to help you generate the SQL query
      .
- Before generating the final SQL query, please think through the steps of
      how to write the query. It should include detailed considerations
      such as analyzing questions, summarizing relevant findings,
      brainstorming new ideas, verifying the accuracy of the current steps,
       refining any errors, thinking of how to call SQL tools, and
      revisiting previous steps.

Format:
- Conduct thinking every time you get new observation or information.
- You can use the SQL tool to explore data or verify result. You will
      receive the execution results or error information of the SQL from a
      user. Based on this information, you can think again and refine.
- The returned dataframe will be truncated in 50 rows if observation is too
       long.
- Only if you find no further exploration is needed or reach max turns, you
       directly provide the final SQL query solution inside <solution>...</
      solution>.
- Do not request a SQL tool execution and provide a solution in the same
      response.
```

E.2 Evidence-Structured Prompt (ReViSQL-BIRD)

ReViSQL-BIRD uses the same system prompt with two additional directives in the `Format` section, which operationalize the process supervision discussed in §4.2. Every other part of the prompt is identical to §E.1.

```
- In the first or second turn, you must perform a translation from the
      external knowledge to requirements of the final SQL query. The
      external knowledge may contain multiple pieces of information
      separated by semicolons. You must list the requirement for each
      external knowledge piece the following format: "Requirement of
      external knowledge {i} ({raw external knowledge text}): ...". This
      will help you better understand the question and the external
      knowledge, and guide your thinking in the subsequent turns. If the
      external knowledge is empty, skip this translation step.

- In the response where you produce the final solution inside <solution
      >...</solution>, you must first verify the external knowledge one by
      one and conduct a final check to ensure all the external knowledge is
       correctly used in the generated SQL query. You must use the
      following format to verify each piece of external knowledge: "
      Verification of external knowledge {i} ({raw external knowledge text
      }): ...". If you find any issues during the verification, you must
      update your draft before finalizing the query inside <solution>. If
      the external knowledge is empty, skip this verification step.
```

Relation to process supervision. The *requirement* directive forces the model to translate each external-knowledge entry into an explicit query constraint before it writes SQL. The *verification* directive forces it to audit the generated query against each entry, and to revise the draft if the audit fails, in the same response that commits the answer. Because both blocks are emitted in a fixed, indexed format, we check the compliance by a string-level test on the reasoning trace, with no learned reward model in the loop. §E.5 specifies the resulting rule set $\mathcal{K}$ and its penalties.

E.3 Input Format

Each problem is rendered into the four fields declared by the system prompt.

(1) **Database Engine.** The target dialect (`SQLite` or `Snowflake`), which fixes the syntax the model must emit.
(2) **Database Schema.** The full schema in data definition language, including tables, columns, primary keys, and foreign keys. Following prior work [28, 31], we annotate each column with the column description provided by BIRD; when no description is available, we instead attach three randomly sampled values from that column, which lets the model infer the value format without querying the database.
(3) **External Knowledge.** The evidence entries associated with the problem, verbatim. Multiple entries are separated by semicolons, and this ordering defines the index i used by the requirement and verification blocks. Problems that carry no external knowledge leave this field empty.
(4) **Question.** The natural language question, verbatim.

E.4 Rollout Specification

A rollout is a multi-turn interaction between the model and a database environment. The model is given a tool, `execute_sql_query`, which executes a SQL query against the problem's database and returns the result.

Intermediate turns. On each intermediate turn the model emits its reasoning followed by a tool call. The environment executes the query with a 30-second timeout and replies with one of the following state.

(1) *Success.* The result table, rendered as text. To bound the context, the table is truncated to its first 50 rows and individual cells are truncated to 200 characters. We explicitly tell the model when truncation occurred.
(2) *Execution error.* The database error message, so that the model can diagnose and repair its query.
(3) *Malformed action.* If the tool call cannot be parsed, the environment returns an instruction to re-emit a tool call.

To aid planning, the environment appends a turn counter to every successful response: "`You have {k} turns left to complete the task.`" The model may not request a tool execution and submit a final solution in the same response.

Terminal turn. The model terminates a rollout by emitting its final query inside `<solution>...</solution>`, which it is instructed to do only when further exploration is unnecessary or the turn budget is exhausted. A response that violates the output contract receives a reward of -1.

E.5 Reward Specifications

Table 6 summarizes which prompt and which reward terms each configuration uses. ReViSQL is trained with the base prompt and a purely result-based reward; the verification-aware stage and the process penalties are specific to ReViSQL-BIRD. The reward for a rollout is computed as follows.

(1) **Format check.** A solution that does not satisfy the output contract receives -1.

Table 6: Prompt and reward configuration of each method. The pre-RLVR base model receives the same evidence-structured prompt as ReVISQL-BIRD. ReVISQL, our general recipe, uses the base prompt and a purely result-based reward. ReVISQL-BIRD couples the evidence-structured prompt with the penalties that enforce it.

	System prompt	Verification-aware reward (β)	Process penalties (λ_k)
Base model	Evidence	*not trained*	
ReVISQL	Base	○	○
ReVISQL-BIRD	Evidence	●	●

(2) **Result-based grading.** We execute the generated query and the gold query against the database and compare their results using the grading method annotated for that instance—set-based, subset-based, or list-based (§3.3). A mismatch yields reward 0. This is the only reward term used by ReVISQL.

(3) **Verification-aware grading (ReVISQL-BIRD only).** If result-based grading accepts the query, we invoke VeriEQL [20] with a 30-second timeout to test semantic equivalence to the gold query. If VeriEQL refutes equivalence, the reward is down-weighted to $1 - \beta = 0.8$; if VeriEQL confirms equivalence, times out, or does not support the query, the reward is 1. This stage is applied to SQLite problems, the dialect VeriEQL supports.

(4) **Process penalties (ReVISQL-BIRD only).** We deduct $\lambda_k = 0.1$ for each violated rule $k \in \mathcal{K}$, as specified below.

Evidence-supervision rules. The rule set $\mathcal{K}$ of Equation 2 contains two rules, each carrying $\lambda_k = 0.1$. Both are evaluated by string-level checks on the reasoning trace.

(1) **Requirement rule.** The requirement block must appear by the second turn. We evaluate this rule once per rollout, at turn two, and deduct 0.1 if the block is absent. Because it is checked before the model commits an answer, the penalty applies whether or not the rollout ultimately produces a solution.

(2) **Verification rule.** The verification block must appear in the same response as the final `<solution>`. We evaluate this rule when the model commits a solution and deduct 0.1 if the block is absent.

We relax both rules in a special case when the model answers on the first turn, which happens on easy problems that require no exploration. In this case, demanding a separate translate-then-verify pass is unnecessary, so we deduct 0.1 only if *both* blocks are missing, and nothing if either is present. A rollout therefore incurs at most 0.2 in process penalties, and at most 0.1 when it terminates on its first turn. Problems that carry no external knowledge require neither block.