arXiv:2608.28276v2 [cs.LG] 13 Oct 2026

Parser Transitions Already Reveal Structural Responsibility: KV Persistence for Structured Generation

Linze Wu **Xinrui Chen***
Hangzhou Institute for Advanced Study,
University of Chinese Academy of Sciences, Hangzhou, China
{wulinze26, chenxinrui25}@mails.ucas.ac.cn

Abstract

LLM agents use structured generation for JSON, SQL, and function calls. Constrained decoding ensures valid formats, but KV compression can discard history needed for correct function and argument selection. Structural responsibility can diverge from model saliency: importance scores can overlook generated tokens' structural roles. Parser transitions already computed for legality encode these roles, including required arguments and structural boundaries, providing a free retention signal. We introduce **PASK** (**P**arser-**A**ware **S**tructural **K**V Persistence), which reuses this signal for KV persistence by layer group without online importance estimation. Task-error sensitivity sets protection floors, while attention-output distortion ranks upgrades under the remaining budget. Offline calibration produces a frozen table for structural lookup and budget accounting. PASK outperforms budget-controlled baselines on function calling and text-to-SQL, improving BFCL accuracy by up to 8.67 pp over the strongest comparator. Against Full KV, it cuts peak GPU memory by 47.5% and achieves up to 3.3× TPOT speedup.

1 Introduction

LLM agents access tools and databases through structured outputs: JSON, SQL (Li et al., 2023), function calls (Patil et al., 2025), and other machine-executable formats (Geng et al., 2025). Constrained decoding ensures format validity by allowing only tokens consistent with a schema, grammar, or function signature (Scholak et al., 2021; Geng et al., 2023; Dong et al., 2025). Compressing the key–value (KV) cache of reusable token representations improves serving efficiency but can discard essential history: ***grammatically valid calls may still select incorrect functions or arguments***.

Existing KV policies guide retention using attention saliency (Zhang et al., 2023; Li et al., 2024), quantization error (Liu et al., 2024b; Hooper et al., 2024), recency (Xiao et al., 2024), or generic token importance (Zhou et al., 2025; Wu et al., 2025). Structure-aware approaches use coarse roles in schema-dense inputs (Mandal, 2026). These signals can overlook generated tokens' structural roles, creating a responsibility–saliency mismatch between model importance and retention needs for task correctness. Under our shared constrained-decoding protocol at approximately one-third KV budget, attention-guided baselines lose up to 70.0 accuracy points relative to Full KV on BFCL non-Live parallel-multiple calls. Fig. 1 motivates retention guided by structural responsibility.

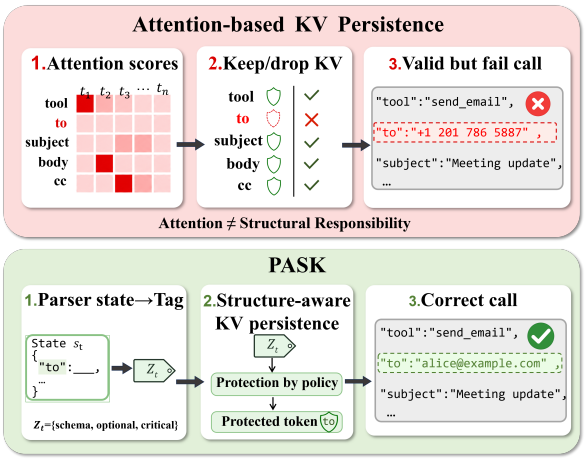

Figure 1: Motivation: parser transitions expose structural roles complementing model-side saliency.

Our insight is that constrained decoding already provides structural information for KV retention. As each token is accepted, the parser updates its state to track progress through the output structure. These transitions reveal roles such as required arguments, optional fields, and structural boundaries. ***Already computed for legality enforcement, they provide a freely available signal for memory management.*** We repurpose this signal to connect a token's structural role to its KV persistence, giving the constraint engine a second role in protecting history for subsequent generation.

*Corresponding author.

We introduce **PASK** (**P**arser-**A**ware **S**tructural **K**V Persistence), which assigns release, low-precision retention, or high-precision retention by parser-derived structural tag and Transformer layer group, allowing each token's protection to vary with depth. ***Its risk-constrained allocation separates task protection from residual spending***: offline interventions measure task-error sensitivity to releasing or quantizing each structure–layer group to set protection floors, while offline attention-output distortion ranks upgrades by distortion reduction per unit of additional memory. After development-set selection of the lowest-cost candidate satisfying a task-error tolerance, PASK compiles actions, floors, and upgrade priorities into a frozen table. During decoding, PASK retrieves them by structural tag and layer group and adjusts retention to the budget without recomputing importance or distortion. With fixed layer groups, per-token control work is independent of retained-history length. A lightweight prefill companion compresses prompt KV, enabling end-to-end compression of prompt and generated history.

PASK achieves SOTA accuracy among the evaluated budget-controlled KV compression methods on BFCL and BIRD. At roughly one-third total KV budget, it stays within 2.5 points of Full KV on BFCL non-Live parallel-multiple calls across three models. Relative to Full KV, it ***reduces peak GPU memory by*** 47.5% ***and achieves up to*** $3.3\times$ ***TPOT speedup*** in the primary serving matrix. Ablations and transfer experiments support the design and calibration reuse.

Our contributions are:

- We identify a responsibility–saliency mismatch in compressed structured generation and establish existing parser transitions as a structural signal for generated-KV retention.
- We formulate risk-constrained KV persistence: task-error sensitivity sets protection floors; numerical distortion allocates remaining memory.
- We develop PASK, compiling structural and layer-dependent retention decisions into a frozen policy with lightweight online control.
- We demonstrate accuracy and serving gains on structured-generation benchmarks, with ablations and transfer experiments supporting the mechanism and policy reuse.

2 Related Work

Structured Generation. Structured generation enforces schemas, grammars, or function signatures so that outputs remain parseable and executable (Scholak et al., 2021; Geng et al., 2023). Systems compile them into guided decoding procedures (Willard and Louf, 2023; Zheng et al., 2024; Beurer-Kellner et al., 2024; Ugare et al., 2025), explicit parser states (Scholak et al., 2021; Dong et al., 2025; Li et al., 2026b), or token masks (Dong et al., 2025; Li et al., 2026b); evaluations emphasize coverage, correctness (Li et al., 2023; Geng et al., 2025; Patil et al., 2025; Chen et al., 2025), and efficiency. PASK extends parser state usage from legality enforcement to KV retention.

KV Compression. KV compression reduces inference memory through low-bit quantization (Liu et al., 2024b; Hooper et al., 2024; Lin et al., 2025; Sharma et al., 2025; Boroujeni et al., 2026) to store KV states with reduced numerical precision, attention- or distortion-based token selection (Liu et al., 2023; Zhang et al., 2023; Xiao et al., 2024; Li et al., 2024; Ge et al., 2024; Behnam et al., 2025; Goel et al., 2025; Feng et al., 2026; Zhang et al., 2026b; Jo et al., 2026; Chen et al., 2026c) to preserve history estimated to matter for subsequent decoding, and system-level cache organization (Kwon et al., 2023) to reduce fragmentation and improve memory access efficiency. Complementary work exploits inter-layer redundancy (Liu et al., 2024a), non-uniform layer budgets (Cai et al., 2025b), task- and context-adaptive allocation (Zhou et al., 2025; Liu et al., 2026; Chen et al., 2026b), distillation (Jiang and Jin, 2026), reasoning-stage signals (Cai et al., 2025a; Ramachandran et al., 2026), and agent-level cache directives (Ma et al., 2026). KV utility varies across tokens, layers, and tasks, but retention signals remain model-internal statistics or task-level heuristics that do not directly encode whether accepted generated tokens realize required keys, arguments, enum values, or schema boundaries.

Structure-Aware Memory. PASK's closest antecedents incorporate temporal (Shen et al., 2026), semantic (Basu, 2026; Liu et al., 2025; Li et al., 2026a), program-context (Chen et al., 2026a), and structural-role signals (Mandal, 2026) into KV allocation or retention. Prior work exploits structural cues from attention patterns, program context, execution traces, or prompt organization. PASK instead repurposes token-level parser transitions

already computed for legality during constrained decoding. This signal exposes generated tokens' structural responsibility and guides their KV retention without online importance estimation.

3 Method

PASK reuses parser transitions computed for constrained decoding for layer-group KV persistence. Offline calibration separates task protection from distortion-based allocation, compiling a frozen table for online lookup and budget accounting.

3.1 Preliminaries

KV cache. For one Transformer layer and head, causal attention at sequence position i computes

$$\mathbf{o}_i = \mathbf{V}_i \,\mathrm{softmax}\Big(\mathbf{K}_i^\top \mathbf{q}_i / \sqrt{d_k}\Big), \qquad (1)$$

where $\mathbf{q}_i$ is the query, $\mathbf{K}_i = [\mathbf{k}_1, \ldots, \mathbf{k}_i]$ and $\mathbf{V}_i = [\mathbf{v}_1, \ldots, \mathbf{v}_i]$ contain keys and values, and d_k is the key dimension. KV caching reuses past keys and values. Eviction renormalizes attention over retained entries; quantization changes their representations. Both can alter attention outputs.

Constrained decoding. Given request $r = (x, G)$ with prompt x and constraint G, parser state s_{t-1} summarizes $y_{1:t-1}$. Decoding selects $y_t \in \mathcal{A}_{\text{legal}}(s_{t-1})$ and updates $s_t = \delta_G(s_{t-1}, y_t)$ before the next forward pass computes the non-stop token's KV. Format validity does not ensure correct functions, arguments, or field values.

3.2 Risk-Constrained KV Persistence

For generated KV (prompt compression: Section 3.6), PASK partitions Transformer layers into groups $\mathcal{C}$ and assigns each non-stop token $a_{t,c} \in \{R, L, H\}$ per group c: release, INT4 retention, or BF16 retention, ordered by increasing protection and cost. Recent tokens remain high-precision; releases occur after window exit. Scheduled long-range payload $m(a, c)$ includes group size and quantization scales, with $m(R, c) = 0$. Recent-window overhead and physical GPU memory are accounted for separately (Appendix B.3). Define $C_{\text{full}} = \sum_c m(H, c)$ and

$$\begin{aligned} R_{\text{task}}(\pi) &= \mathbb{E}_{r\sim\mathcal{D}}\big[\mathbb{I}[\text{Fail}(y_r^\pi)]\big], \\ C_{\text{KV}}(\pi) &= \frac{\mathbb{E}_{r\sim\mathcal{D}}\big[\sum_{t=1}^{T_r^\pi} \sum_c m(a_{r,t,c}^\pi, c)\big]}{\mathbb{E}_{r\sim\mathcal{D}}[T_r^\pi] C_{\text{full}}}, \end{aligned} \qquad (2)$$

where T_r^π counts generated non-stop tokens and Fail follows the benchmark task oracle, including incomplete outputs. The denominator uses the same generated lengths. Relative to the all-high reference π_{ref}, our objective is

$$\begin{aligned} \pi_\alpha^\star \in \arg\min_{\pi\in\Pi_{\text{struct}}} & C_{\text{KV}}(\pi) \\ \text{s.t.} \quad & R_{\text{task}}(\pi) - R_{\text{task}}(\pi_{\text{ref}}) \le \alpha. \end{aligned} \qquad (3)$$

Π_{struct} contains policies using structural tags, layer groups, and request-local budget state, and α bounds the absolute increase in task failure probability. PASK searches a finite family, enforcing the tolerance empirically on development data.

3.3 Parser-Derived Persistence Buckets

A deterministic grammar adapter maps each accepted transition, already computed for legality, to a structural tag and pairs it with a layer group:

$$\begin{aligned} z_t &= \phi_G(s_{t-1}, y_t, s_t) \\ &= (\text{class}_t, \text{role}_t, \text{state}_t, \text{next}_t), \\ q_{t,c} &= (z_t, c), \qquad c \in \mathcal{C}. \end{aligned} \qquad (4)$$

The components identify token class, schema responsibility, grammar status, and an approaching critical region. Required arguments and field boundaries receive corresponding roles; next_t uses the current grammar frontier, not future tokens. These labels identify structural responsibilities, while offline task feedback determines their protection requirements. Buckets share calibrated decisions across token identities without online importance estimation. Unresolved tags default to high protection. Appendix A specifies the vocabulary, and Appendix B describes parser integration. Lower, middle, and upper layer groups allow identical roles to receive different protection across depth (Liu et al., 2024a; Cai et al., 2025b; Zhou et al., 2025). For example, the same argument token may receive low-precision retention in early layers and high-precision retention in later layers.

3.4 Offline Calibration and Compilation

Task-error sensitivity determines minimum protection; attention-output distortion guides upgrades. Both offline measurements share decoding settings, prompt policy, and recent-window rules.

Task-error sensitivity. Policy $\pi_{q,a}$ applies $a \in \{R, L\}$ to every occurrence of bucket q during complete autoregressive generation, keeping other buckets at H. Occurrences follow the intervened trajectory, allowing the measurement to capture downstream changes in generation. On the same

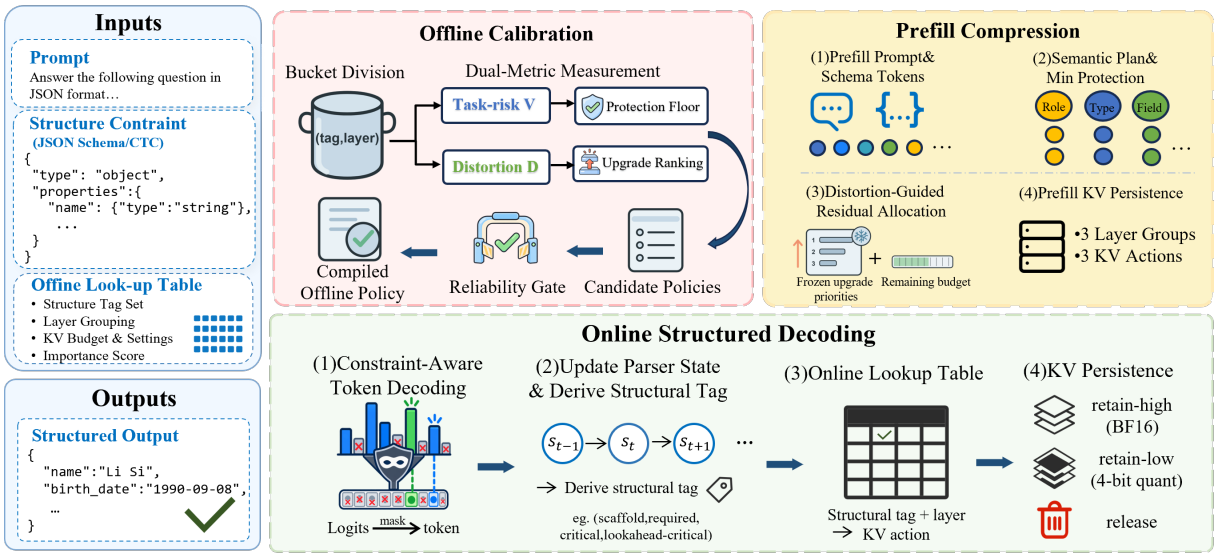

Figure 2: Overview of PASK. Parser transitions and layer groups index an offline-calibrated table that assigns high-precision retention, low-precision retention, or release to generated KV states.

calibration requests,

$$
\begin{aligned}
&V(q,a) = \widehat{R}_{\mathrm{cal}}(\pi_{q,a}) - \widehat{R}_{\mathrm{cal}}(\pi_{\mathrm{ref}}),\\
&\bar{V}(q,H) = 0, \qquad \bar{V}(q,L) = \max\{0, V(q,L)\},\\
&\bar{V}(q,R) = \max\{\bar{V}(q,L), V(q,R)\}.
\end{aligned} \tag{5}
$$

The envelope conservatively orders empirical sensitivities by protection. For $0 \le \tau_{\mathrm{low}} \le \tau_{\mathrm{high}}$,

$$
a^{\tau}_{\min}(q) = \begin{cases} H, & \bar{V}(q,L) > \tau_{\mathrm{high}},\\ R, & \bar{V}(q,R) \le \tau_{\mathrm{low}},\\ L, & \text{otherwise.} \end{cases} \tag{6}
$$

A bucket with modest attention distortion can therefore remain protected if releasing it exceeds the task-error threshold. Bucket thresholds construct local floors; α governs complete-policy selection, which checks joint effects.

Attention-output distortion. Using fixed reference queries and K/V tensors, we locally recompute attention after releasing or quantizing eligible entries in q. Reference tokens and queries remain unchanged, isolating numerical deviation from the task-level effects measured by V. Let $\mathcal{S}(q)$ contain affected reference query positions (r,u), shared across actions. Define accumulated distortion per calibration request as

$$
D(q,a) = \frac{1}{|\mathcal{D}_{\mathrm{cal}}|} \sum_{(r,u)\in\mathcal{S}(q)} \sum_{\ell\in c} \max_{h} \left\| o^{\mathrm{ref}}_{\ell,r,u,h} - o^{a(q)}_{\ell,r,u,h} \right\|_2^2, \tag{7}
$$

where h indexes query heads and $D(q,H) = 0$. For adjacent upgrades $R \to L$ or $L \to H$,

$$
\rho_D(q, a \to a') = \frac{[D(q,a) - D(q,a')]_+}{\bar{N}_q[m(a',c) - m(a,c)]}, \tag{8}
$$

where $\bar{N}_q$ is the mean bucket count per reference request. Because a bucket upgrade affects every occurrence, the denominator measures its mean total incremental payload per request. Calibration details appear in Appendix A.4.

Policy selection. For each threshold pair $\tau \in \mathcal{T}$, PASK starts from its floors and repeatedly applies the eligible upgrade with largest ρ_D, generating a nested candidate path. Each candidate retains its floor map and takes its nominal decode target B_{D} from its reference-trajectory payload ratio. We construct its reserve and evaluate the complete controller (Section 3.5) autoregressively on development requests, capturing joint task effects and budget-driven action changes:

$$
\begin{aligned}
&i^{\star} \in \arg\min_{i\in\mathcal{I}_\alpha} \widehat{C}_{\mathrm{KV,dev}}(\pi_i),\\
&\mathcal{I}_\alpha = \{i : \widehat{R}_{\mathrm{dev}}(\pi_i) - \widehat{R}_{\mathrm{dev}}(\pi_{\mathrm{ref}}) \le \alpha\}.
\end{aligned} \tag{9}
$$

Costs use each candidate's actual generated lengths in Eq. 2; the all-high reference provides a feasible fallback. We freeze selected actions, floors, and priorities in Λ, alongside reserve Γ and budget B_{D}.

3.5 Structure-Conditioned Online Control

Each accepted non-stop token adds allowance $b = B_{\mathrm{D}} C_{\mathrm{full}}$. The balance R_t includes the current allowance, with $R_1 = b$. To anticipate upcoming

Algorithm 1: PASK Offline Calibration

Require: $\mathcal{D}_{\text{cal}}, \mathcal{D}_{\text{dev}}, \mathcal{C}, \mathcal{T}, \alpha$
Ensure: $\Lambda, \Gamma, B_{\text{D}}$
$(\bar{V}, D, \bar{N}) \leftarrow \text{Calibrate}(\mathcal{D}_{\text{cal}}, \mathcal{C})$
$\rho_D \leftarrow \text{Priorities}(D, \bar{N}, m)$
$\mathcal{P} \leftarrow \{\pi_{\text{ref}}\}$
for all $\tau \in \mathcal{T}$ **do**
  $f \leftarrow \text{Floor}(\bar{V}; \tau)$
  for all $A \in \text{GreedyPath}(f, \rho_D)$ **do**
    $B \leftarrow \text{ReferenceCostRatio}(A)$
    $\Gamma \leftarrow \text{Reserve}(\mathcal{D}_{\text{cal}}, f, B)$
    Add $\text{Controller}(A, f, \rho_D, B, \Gamma)$ to $\mathcal{P}$
Evaluate complete controllers on $\mathcal{D}_{\text{dev}}$
return $\text{Pack}(\pi_{i^\star})$ *using Eq. 9*

protection costs, the controller looks up a reserve table constructed offline:

$$\Gamma(s; B_{\text{D}}) = Q_\beta\big(\{S_{r,u}^{(K)}(B_{\text{D}}) : s_{r,u}^{\text{ref}} = s\}\big). \quad (10)$$

Here, $S_{r,u}^{(K)}$ is the largest positive cumulative floor cost minus newly accrued allowance over prefixes of the next K reference tokens. We use $K = 8$, $\beta = 0.8$, and a pooled quantile for unseen states. Only the current parser state is needed online; Appendix B.6 details the offline construction.

For the current structural tag, lookup returns base actions $a_{t,c}^0$, floors $a_{t,c}^-$, and frozen priorities for each Transformer layer group. PASK initializes

$$\begin{aligned} a_{t,c} &= \max_{\preceq}(a_{t,c}^0, a_{t,c}^-), \\ B_t &= [R_t - \Gamma(s_t; B_{\text{D}})]_+. \end{aligned} \quad (11)$$

Writing $M(\mathbf{a}) = \sum_c m(a_c, c)$, it demotes eligible actions in ascending stored priority until $M(\mathbf{a}_t) \le B_t$ or only floors remain. For critical or lookahead-critical tags, it promotes in descending priority only while total cost fits B_t. Floors take precedence over the target budget; negative balances carry forward. Adjustments affect only the current token's pending actions, never reconstructing previously discarded KV. Thus, the policy stays fixed while its current-token allocation adapts to accumulated budget.

Given resolved tags and constant-time lookup, control costs $O(|\mathcal{C}| \log |\mathcal{C}|)$ per token, independent of history length for fixed groups. This excludes parser adaptation, KV materialization, and attention computation. Appendix B.6 specifies budget accounting and speculative-decoding rollback.

3.6 Lightweight Prefill Companion

A separate adapter groups prompt spans by semantic label and layer group. Table 6 lists protected content, function names, parameter keys, constraints, and schema scaffold, rather than output-parser transitions. Independently calibrated floors and distortion priorities determine a frozen retention plan before the prompt forward; its $R/L/H$ representations are materialized after high-precision prefill. Prefill and decode use separate tolerances and budgets; their composed accuracy and total KV cost are evaluated end to end. Appendix D.8 reports stage-wise contributions.

Algorithm 2: Online Structural Control

Require: $z_t, s_t, R_t, B_{\text{D}}, \mathcal{C}, \Lambda, \Gamma$
Ensure: $\mathbf{a}_t, R_{t+1}$
$(\mathbf{a}^0, \mathbf{a}^-, \boldsymbol{\rho}) \leftarrow \text{Lookup}(\Lambda, z_t, \mathcal{C})$
$\mathbf{a}_t \leftarrow \max_{\preceq}(\mathbf{a}^0, \mathbf{a}^-)$
$B_t \leftarrow [R_t - \Gamma(s_t; B_{\text{D}})]_+$
$\mathbf{a}_t \leftarrow \text{Demote}(\mathbf{a}_t, \mathbf{a}^-, B_t, \boldsymbol{\rho})$
if $\text{Crit}(z_t)$ **then**
  $\mathbf{a}_t \leftarrow \text{Promote}(\mathbf{a}_t, B_t, \boldsymbol{\rho})$
Schedule $\mathbf{a}_t$ for KV materialization
$R_{t+1} \leftarrow R_t - M(\mathbf{a}_t) + B_{\text{D}} C_{\text{full}}$
return $\mathbf{a}_t, R_{t+1}$

4 Experiments

We directly test the central claim: whether parser-derived structural state can sustain structured-generation correctness at substantially reduced KV cost. Evaluation covers task accuracy, serving efficiency, ablations, and sensitivity analyses.

4.1 Experimental Setup

Models and Datasets. We evaluate PASK with Qwen3-4B-Instruct-2507, Qwen3-14B (Yang et al., 2025), and LLaMA-3.1-8B-Instruct (Grattafiori et al., 2024), abbreviated as Qwen3-4B, Qwen3-14B, and LLaMA-3.1-8B, respectively. Our benchmarks cover BFCL function calling constrained by tool specifications (Patil et al., 2025), including both non-live and Live subsets, and BIRD text-to-SQL generation of executable, schema-grounded SQLite queries (Li et al., 2023). BIRD experiments use Qwen3-4B and the official Dev package; Table 7 reports dataset statistics.

Baselines. We compare against eight baselines covering attention-based retention, mixed-precision storage, recurrence-aware eviction, and structured cache policies: H2O (Zhang et al., 2023), SnapKV (Li et al., 2024), MiKV (Yang et al., 2024), RocketKV (Behnam et al., 2025), CriticalKV (Feng et al., 2026), LazyEviction (Zhang et al., 2026a),

Model	Method	BFCL non-live					BFCL live				
		Simple	Multiple	Parallel	Par.-multiple	Overall	Simple	Multiple	Parallel	Par.-multiple	Overall
Qwen3-4B	Full KV	95.00	95.00	90.00	84.17	91.83	78.57	74.57	60.00	78.57	75.22
	H2O NeurIPS 2023	75.83	81.67	31.67	19.17	56.83	45.45	48.18	0.00	21.43	46.61
	SnapKV NeurIPS 2024	85.00	87.50	42.50	35.00	67.00	55.19	59.08	0.00	14.29	56.84
	MiKV arXiv 2024	87.50	90.83	76.67	76.67	83.83	64.94	64.45	30.00	**50.00**	63.87
	RocketKV ICML 2025	75.42	65.83	33.33	25.00	55.00	48.05	47.24	0.00	28.57	46.49
	CriticalKV ICML 2026	87.92	91.67	59.17	40.00	73.33	59.74	66.19	30.00	28.57	63.87
	LazyEviction ACL 2026	46.67	81.67	20.83	14.17	42.00	33.12	60.03	10.00	28.57	53.76
	TriAxialKV arXiv 2026	90.42	85.00	70.83	56.67	78.67	62.99	59.08	40.00	0.00	58.57
	TransAttn arXiv 2026	88.75	64.17	54.17	43.33	67.83	45.45	27.49	0.00	0.00	30.09
	PASK (Ours)	**92.08**	**93.33**	**80.83**	**81.67**	**88.00**	**74.03**	**72.20**	**60.00**	**50.00**	**72.01**
LLaMA-3.1-8B	Full KV	94.58	95.00	91.67	80.83	91.33	75.32	60.98	40.00	57.14	63.38
	H2O NeurIPS 2023	75.42	90.00	15.83	12.50	53.83	46.75	36.65	0.00	14.29	37.73
	SnapKV NeurIPS 2024	87.92	94.17	43.33	35.00	69.67	51.30	53.55	10.00	35.71	52.28
	MiKV arXiv 2024	91.67	95.00	66.67	51.67	79.33	64.94	48.50	20.00	28.57	50.92
	RocketKV ICML 2025	82.92	90.83	25.00	21.67	60.67	51.30	47.39	0.00	21.43	47.10
	CriticalKV ICML 2026	86.67	93.33	35.83	33.33	67.17	62.34	53.08	10.00	35.71	54.01
	LazyEviction ACL 2026	56.67	88.33	25.83	16.67	48.83	38.96	45.66	10.00	28.57	43.65
	TriAxialKV arXiv 2026	92.08	88.33	85.83	55.83	82.83	66.23	38.23	**50.00**	28.57	43.53
	TransAttn arXiv 2026	90.00	67.50	71.67	38.33	71.50	60.39	32.07	30.00	0.00	36.87
	PASK (Ours)	**93.75**	**96.67**	**90.83**	**80.83**	**91.17**	**73.38**	**58.45**	40.00	**50.00**	**60.91**
Qwen3-14B	Full KV	96.25	98.33	92.50	81.67	93.00	79.87	75.20	70.00	64.29	75.83
	H2O NeurIPS 2023	72.08	85.00	22.50	16.67	53.67	54.55	51.82	0.00	0.00	50.80
	SnapKV NeurIPS 2024	82.50	91.67	60.00	55.83	74.50	61.69	66.98	0.00	14.29	64.24
	MiKV arXiv 2024	83.75	88.33	79.17	65.83	80.17	66.23	63.03	**40.00**	42.86	63.01
	RocketKV ICML 2025	78.33	90.00	30.00	41.67	63.67	55.19	55.13	30.00	42.86	54.62
	CriticalKV ICML 2026	80.42	87.50	54.17	50.83	70.67	61.69	65.72	**40.00**	**57.14**	64.49
	LazyEviction ACL 2026	49.17	91.67	28.33	25.00	48.67	33.77	64.14	20.00	**57.14**	57.71
	TriAxialKV arXiv 2026	87.92	79.17	77.50	64.17	79.33	68.18	61.14	30.00	0.00	61.04
	TransAttn arXiv 2026	88.33	83.33	70.00	49.17	75.83	49.35	31.12	0.00	0.00	33.66
	PASK (Ours)	**92.08**	**94.17**	**86.67**	**79.17**	**88.83**	**73.38**	**73.62**	30.00	**57.14**	**72.75**

Table 1: BFCL accuracy (%) across models and the non-live and Live subsets. Red and orange shading indicate the **best** and second-best accuracy, respectively; Full KV is shown as a reference and excluded from ranking.

TriAxialKV (Shen et al., 2026), and Transactional Attention (Basu, 2026). All share prompts, decoding configurations, and the constrained-decoding backend (Dong et al., 2025).

Evaluation & Implementation. Calibration, selection, and evaluation use disjoint calibration, development, and frozen test splits partitioned by normalized tool-call schemas (BFCL) or databases (BIRD). Appendix C.3 details baseline-specific tuning. Qwen3-4B serving on NVIDIA L40S and A100 uses matched software and decoding settings. B denotes measured, workload-weighted total KV cost (prefill and decode), normalized by measured Full KV ($B = 1.00$). PASK achieves $B \approx 0.33$ in the main comparison; flexible-budget baselines are configured to match this cost, while fixed-profile methods retain native configurations with measured costs reported separately.

4.2 Main Results

Task accuracy. ❶ **Consistent accuracy gains.** PASK leads all evaluated compressed methods across six BFCL Overall settings in Table 1 and all four BIRD budgets in Table 2. On Live BFCL, Qwen3-4B and Qwen3-14B improve over the strongest compressed baselines by 8.14 and 8.26 points, respectively. Fig. 3 shows its BFCL lead across budgets. ❷ **Near-Full-KV compositional accuracy.** On LLaMA-3.1-8B non-Live, PASK nearly matches Full KV in Overall accuracy (91.17% vs. 91.33%) and matches it on parallel-multiple calls (80.83%), outperforming the strongest compressed baseline by 25.0 points. Across all three models, parallel-multiple accuracy remains within 2.5 points of Full KV. ❸ **Accuracy under tighter KV budgets.** On BIRD, PASK achieves 52.71% execution accuracy at $B = 0.275$, exceeding the strongest same-budget baseline (49.48%) and even the best baseline at $B = 0.40$ (51.83%). At $B = 0.60$, it reaches 55.93%, only 0.35 points below Full KV. Detailed BFCL budget results appear in Appendix D.1.

Method	KV budget B				
	100%	60%	50%	40%	27.5%
H2O	56.28	41.10	35.51	32.55	24.08
SnapKV	56.28	49.83	46.77	44.24	36.56
MiKV	56.28	54.01	52.27	50.26	48.25
RocketKV	56.28	54.36	53.66	50.09	41.10
CriticalKV	56.28	50.96	44.68	41.80	37.52
LazyEviction	56.28	55.67	52.88	50.96	49.48
TransAttn	56.28	54.80	53.05	51.83	40.14
PASK (Ours)	56.28	**55.93**	**53.84**	**53.75**	**52.71**

Table 2: Execution accuracy (EX, %) on BIRD with Qwen3-4B. Gray entries show the Full KV reference.

Memory. At a fixed logical capacity of 229,896 tokens, PASK achieves the lowest peak GPU memory across both evaluated GPUs in Table 3, reducing Full-KV memory by 47.5% and using 13.8%

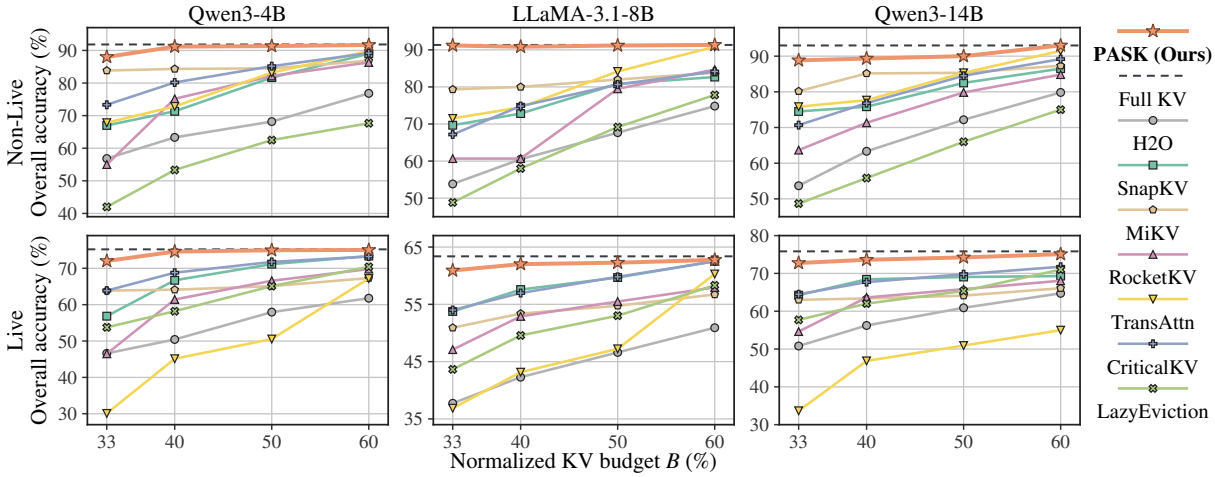

Figure 3: BFCL Overall accuracy versus measured relative total KV cost B across three models and both subsets.

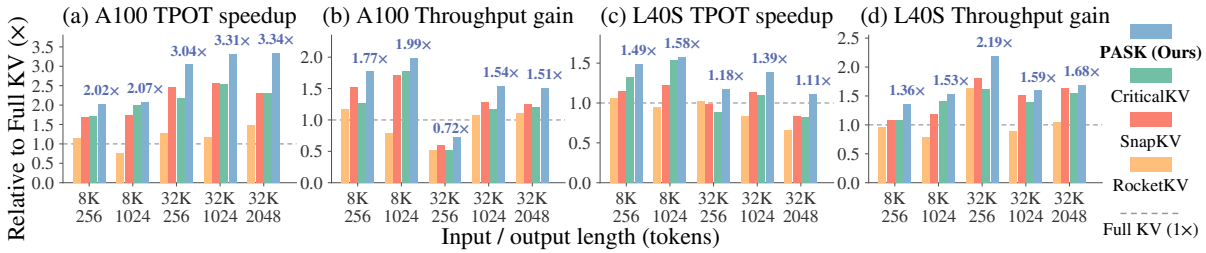

Figure 4: Relative serving efficiency for Qwen3-4B at $c = 8$ and approximate total KV budget $B = 0.33$. (a,c) TPOT speedup, defined as $\mathrm{TPOT}_{\mathrm{Full}}/\mathrm{TPOT}_{\mathrm{method}}$. (b,d) Throughput gain, defined as $\mathrm{TPS}_{\mathrm{method}}/\mathrm{TPS}_{\mathrm{Full}}$.

Method	NVIDIA L40S		NVIDIA A100	
	Peak (MiB) ↓	Saving ↑	Peak (MiB) ↓	Saving ↑
Full KV	40986	—	41008	—
SnapKV	25090	38.78%	25086	38.83%
RocketKV	24976	39.06%	24976	39.09%
CriticalKV	24976	39.06%	24978	39.09%
PASK (Ours)	**21528**	**47.47%**	**21536**	**47.48%**

Table 3: Peak GPU memory for Qwen3-4B at a fixed logical capacity of 229,896 tokens.

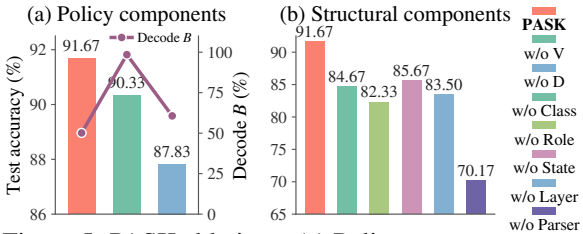

Figure 5: PASK ablations. (a) Policy components: accuracy (bars) and Decode budget (line). (b) Structural components: accuracy at an aligned Decode budget.

less than RocketKV, the strongest compressed comparator in memory.

Speed. PASK achieves the lowest TPOT and highest throughput among compressed baselines across all ten GPU/workload settings in Fig. 4, reaching 3.34× TPOT speedup on A100 (32K/2048) and 2.19× throughput on L40S (32K/256) relative to Full KV. Against the strongest compressed comparator per setting, it reduces TPOT by up to 31.1% and improves throughput by up to 25.9%. At the A100 32K/256 setting, PASK's throughput remains below that of Full KV despite a 3.04× TPOT speedup. This is consistent with packed INT4 storage reducing KV traffic and frozen parser-conditioned lookup avoiding online importance estimation.

4.3 Ablation Study

❶ **Parser-derived signals improve retention.** In Fig. 5, at matched Decode budgets, removing structural roles, layer grouping, token classes, or parser state and lookahead reduces accuracy by 9.33, 8.17, 7.00, and 6.00 points, respectively. Removing parser-transition information entirely causes a larger 21.50-point drop, supporting its importance for KV retention. Appendix C details ablation protocols. ❷ **Task protection reduces budgets.** Removing V-based protection floors increases Decode budget from 50.15% to 98.49% while lowering accuracy from 91.67% to 90.33% in Fig. 5. Thus, distortion ranking alone requires substantially more memory to meet the development reliability criterion. ❸ **Distortion ranking improves residual allocation.** Keeping protection floors, replacing D-ranking with fixed-order, layer-group round-robin allocation increases Decode budget by 10.52 points and reduces accuracy by 3.83 points in Fig. 5, confirming that distortion ranking complements task-

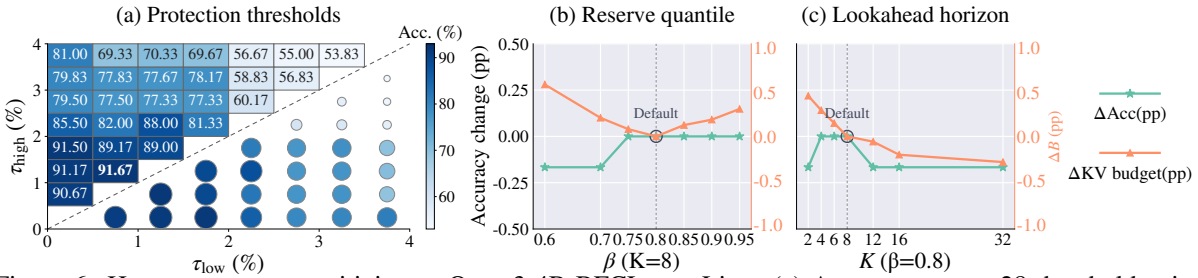

Figure 6: Hyperparameter sensitivity on Qwen3-4B BFCL non-Live. (a) Accuracy across 28 threshold pairs; bubbles mirror the heatmap, and bold marks the selected pair. (b,c) Accuracy and Decode-budget changes (pp) for β and K, relative to the default $(0.8, 8)$.

error sensitivity in residual allocation.

4.4 Hyperparameter Analysis

Fig. 6(a) evaluates 28 protection-threshold pairs for Qwen3-4B on BFCL non-Live at approximately 50% decode KV budget. Pairs with $\tau_{\text{high}} \leq 1\%$ achieve 90.67–91.67% accuracy, versus 53.83–60.17% for $\tau_{\text{low}} \geq 2\%$, highlighting the cost of relaxed protection. The selected $(0.5\%, 1\%)$ pair achieves 91.67% accuracy at realized decode budget $B_{\text{D}} = 0.5015$.

The reserve quantile β and lookahead horizon K control budget reservation for upcoming protected tokens. Fig 6(b,c) varies one parameter while fixing the other at its corresponding default value and reports accuracy and realized Decode KV-budget changes relative to $\beta = 0.8$ and $K = 8$. Stage-specific reliability tolerances α_{D} and α_{P} govern development-set policy selection; their quality–cost trade-offs are reported in Appendix D.5.

4.5 Additional Analysis

Responsibility–distortion mismatch. Table 16 in Appendix D.7 provides evidence for separating task protection from numerical fidelity. Across 27 Qwen3-4B BFCL non-Live calibration buckets, task-error sensitivity and distortion yield disjoint top-20% sets for both release and low-precision retention. Of the most task-sensitive buckets, 100% and 83.3% fall in the respective bottom halves of distortion rankings. Thus, low numerical distortion need not imply low task risk.

Calibration cost and reuse. Table 15 reports about 3.3 GPU-hours for combined Runtime-V/D calibration on one A100-80GB (Qwen3-4B, 200 BFCL non-Live cases per candidate). Table 4 compares unchanged transfer (trans) and per-target recalibration (recalib) at $B \in \{0.33, 0.50\}$; Best baseline denotes each setting's strongest compressed comparator. Transfer outperforms this baseline across all three targets and both budgets,

B	Method	Qwen3-4B Live	Qwen3-14B non-Live	LLaMA-3.1-8B non-Live
0.33	Best baseline	63.87	80.17	82.83
	PASK (trans)	69.67	88.67	87.33
	PASK (recalib)	72.01	88.83	91.17
0.50	Best baseline	71.76	85.33	84.17
	PASK (trans)	73.98	90.83	89.50
	PASK (recalib)	74.97	90.00	91.17

Table 4: BFCL accuracy (%) under calibration transfer.

even without target-specific policy recalibration. On Qwen3-14B non-Live, it differs from recalibration by at most 0.83 points, outperforming it by 0.83 at $B = 0.50$. Recalibration gains on LLaMA-3.1-8B non-Live and Qwen3-4B Live fall from 3.84/2.34 points at $B = 0.33$ to 1.67/0.99 at $B = 0.50$, respectively, showing effective transfer across evaluated targets and greater recalibration benefits under tighter KV budgets.

High-concurrency scaling. At concurrency 32 on A100 (32K inputs, 1024-token checkpoint), PASK achieves 2.17× throughput and 4.84× TPOT speedups over Full KV in Table 14, while maintaining 68–69 ms TPOT across concurrency 8,16,32 as throughput rises from 51.02 to 72.23 tokens/s, demonstrating scalability under increasing concurrency in long-context serving.

5 Conclusion

PASK connects output legality with memory retention by reusing parser transitions already computed during constrained decoding. Calibrated against task errors, this structural signal guides KV protection across Transformer layers, separating task protection from numerical fidelity. Function-calling and text-to-SQL experiments demonstrate stronger accuracy under constrained budgets, complemented by serving efficiency and policy reuse. Future work will explore adaptive recalibration and agents alternating between free-form reasoning and constrained outputs. More broadly, PASK illustrates a principle for efficient generation: reuse

available structural information to allocate memory according to the consequences of losing history, aligning compression with task correctness.

Limitations

PASK relies on offline calibration to construct a structure-conditioned policy table. In the current design, this calibration is performed per model and grammar family before deployment; transfer across model families and to the Live split concedes one to four points to per-target recalibration on some targets (Section 4.5). The implementation also discretizes depth into three layer groups and persistence into three actions; finer partitions or additional precision levels may offer greater flexibility but are not evaluated here. Finally, the method requires a formal constraint engine that exposes parser state during decoding; extending PASK to requests that alternate free-form language with structured calls is future work discussed in Appendix F.

References

Joshua Ainslie, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebron, and Sumit Sanghai. 2023. GQA: Training generalized multi-query transformer models from multi-head checkpoints. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 4895–4901, Singapore. Association for Computational Linguistics.

Abhinaba Basu. 2026. Transactional attention: Semantic sponsorship for KV-cache retention. *Preprint*, arXiv:2604.11288.

Payman Behnam, Yaosheng Fu, Ritchie Zhao, Po-An Tsai, Zhiding Yu, and Alexey Tumanov. 2025. RocketKV: Accelerating long-context LLM inference via two-stage KV cache compression. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 3358–3392. PMLR.

Luca Beurer-Kellner, Marc Fischer, and Martin Vechev. 2024. Guiding LLMs the right way: Fast, non-invasive constrained generation. *Preprint*, arXiv:2403.06988.

Sayed Pedram Haeri Boroujeni, Niloufar Mehrabi, Patrick Woods, Gabriel Hillesheim, and Abolfazl Razi. 2026. Don't waste bits! adaptive KV-cache quantization for lightweight on-device LLMs. *Preprint*, arXiv:2604.04722.

Zefan Cai, Wen Xiao, Hanshi Sun, cheng Luo, Yikai Zhang, Ke Wan, Yucheng Li, Yeyang Zhou, Li-Wen Chang, Jiuxiang Gu, Zhen Dong, Animashree Anandkumar, Abedelkadir Asi, and Junjie Hu. 2025a. R-KV: Redundancy-aware KV cache compression for reasoning models. In *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pages 60980–61005. Curran Associates, Inc.

Zefan Cai, Yichi Zhang, Bofei Gao, Yuliang Liu, Yucheng Li, Tianyu Liu, Keming Lu, Wayne Xiong, Yue Dong, Junjie Hu, and Wen Xiao. 2025b. PyramidKV: Dynamic KV cache compression based on pyramidal information funneling. *Preprint*, arXiv:2406.02069.

Chen Chen, Xinlong Hao, Weiwen Liu, Xu Huang, Xingshan Zeng, Shuai Yu, Dexun Li, Yuefeng Huang, Xiangcheng Liu, Wang Xinzhi, and Wu Liu. 2025. ACEBench: A comprehensive evaluation of LLM tool usage. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 12970–12998, Suzhou, China. Association for Computational Linguistics.

Qiujiang Chen, Jing Xiong, Chenyang Zhao, Sidi Yang, and Ngai Wong. 2026a. CodeComp: Structural KV cache compression for agentic coding. *Preprint*, arXiv:2604.10235.

Yeqiu Chen, Ziyan Liu, Zhenxin Huang, Runquan Gui, Hong Wang, and Lei Liu. 2026b. ArborKV: Structure-aware KV cache management for scaling tree-based LLM reasoning. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 15560–15573. PMLR.

Zhirui Chen, Peiyang Liu, and Ling Shao. 2026c. StructKV: Preserving the structural skeleton for scalable long-context inference. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 12784–12797, San Diego, California, United States. Association for Computational Linguistics.

Yixin Dong, Charlie F. Ruan, Yaxing Cai, Ziyi Xu, Yilong Zhao, Ruihang Lai, and Tianqi Chen. 2025. XGrammar: Flexible and efficient structured generation engine for large language models. In *Proceedings of Machine Learning and Systems*, volume 7. MLSys.

Yuan Feng, Junlin Lv, Haoyu Guo, Yukun Cao, S Kevin Zhou, and Xike Xie. 2026. CriticalKV: Optimizing KV cache eviction from an output perturbation perspective. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 30221–30246. PMLR.

Suyu Ge, Yunan Zhang, Liyuan Liu, Minjia Zhang, Jiawei Han, and Jianfeng Gao. 2024. Model tells you what to discard: Adaptive KV cache compression for LLMs. In *International Conference on Learning Representations*, volume 2024, pages 22975–22988.

Saibo Geng, Hudson Cooper, Michał Moskal, Samuel Jenkins, Julian Berman, Nathan Ranchin, Robert West, Eric Horvitz, and Harsha Nori. 2025. JSONSchemaBench: A rigorous benchmark of structured outputs for language models. *Preprint*, arXiv:2501.10868.

Saibo Geng, Martin Josifoski, Maxime Peyrard, and Robert West. 2023. Grammar-constrained decoding for structured NLP tasks without finetuning. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 10932–10952, Singapore. Association for Computational Linguistics.

Raghavv Goel, Junyoung Park, Mukul Gagrani, Dalton Jones, Matthew Morse, Harper Langston, Mingu Lee, and Chris Lott. 2025. CAOTE: KV cache selection for LLMs via attention output error-based token eviction. *Preprint*, arXiv:2504.14051.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, and 542 others. 2024. The Llama 3 herd of models. *Preprint*, arXiv:2407.21783.

Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. 2024. KVQuant: Towards 10 million context length LLM inference with KV cache quantization. In *Advances in Neural Information Processing Systems*, volume 37, pages 1270–1303. Curran Associates, Inc.

Bo Jiang and Sian Jin. 2026. KVSculpt: KV cache compression as distillation. *Preprint*, arXiv:2603.27819.

Dongwon Jo, Jiwon Song, Yulhwa Kim, and Jae-Joon Kim. 2026. FastKV: Decoupling of context reduction and KV cache compression for prefill-decoding acceleration. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 32167–32186, San Diego, California, United States. Association for Computational Linguistics.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with PagedAttention. *Preprint*, arXiv:2309.06180.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, Xuanhe Zhou, Chenhao Ma, Guoliang Li, Kevin Chang, Fei Huang, Reynold Cheng, and Yongbin Li. 2023. Can LLM already serve as a database interface? a BIg bench for large-scale database grounded text-to-SQLs. In *Advances in Neural Information Processing Systems*, volume 36, pages 42330–42357. Curran Associates, Inc.

Junjie Li, Jiong Lou, and Jie Li. 2026a. IntentKV: Cross-turn intent-aware KV cache pruning for agent inference. *Preprint*, arXiv:2606.09916.

Linzhang Li, Yixin Dong, Guanjie Wang, Ziyi Xu, Alexander Jiang, and Tianqi Chen. 2026b. XGrammar-2: Dynamic and efficient structured generation engine for agentic LLMs. *Preprint*, arXiv:2601.04426.

Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. 2024. SnapKV: LLM knows what you are looking for before generation. In *Advances in Neural Information Processing Systems*, volume 37, pages 22947–22970. Curran Associates, Inc.

Yujun Lin, Haotian Tang, Shang Yang, Zhekai Zhang, Guangxuan Xiao, Chuang Gan, and Song Han. 2025. QServe: W4A8KV4 quantization and system co-design for efficient LLM serving. In *Proceedings of Machine Learning and Systems*, volume 7. MLSys.

Akide Liu, Jing Liu, Zizheng Pan, Yefei He, Gholamreza Haffari, and Bohan Zhuang. 2024a. MiniCache: KV cache compression in depth dimension for large language models. In *Advances in Neural Information Processing Systems*, volume 37, pages 139997–140031. Curran Associates, Inc.

Taowen Tony Liu, Jeffrey T. H. Wong, Can Xiao, Bowen Yang, Hao Mark Chen, and Yiren Zhao. 2026. AgentKV: Phase-aware KV eviction for agentic LLMs. In *AdaptFM: Resource-Adaptive Foundation Model Inference*.

Xiang Liu, Zhenheng Tang, Peijie Dong, Zeyu Li, Yue Liu, Bo Li, Xuming Hu, and Xiaowen Chu. 2025. ChunkKV: Semantic-preserving KV cache compression for efficient long-context LLM inference. In *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pages 28728–28778. Curran Associates, Inc.

Zichang Liu, Aditya Desai, Fangshuo Liao, Weitao Wang, Victor Xie, Zhaozhuo Xu, Anastasios Kyrillidis, and Anshumali Shrivastava. 2023. Scissorhands: Exploiting the persistence of importance hypothesis for LLM KV cache compression at test time. In *Advances in Neural Information Processing Systems*, volume 36, pages 52342–52364. Curran Associates, Inc.

Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. 2024b. KIVI: A tuning-free asymmetric 2bit quantization for KV cache. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 32332–32344. PMLR.

Bole Ma, Jan Eitzinger, and Harald Koestler. 2026. Leyline: KV cache directives for agentic inference. *Preprint*, arXiv:2606.01065.

Soumil Mandal. 2026. Adaptive filtering of the KV cache: Diagnosing and correcting structural-role bias in LLM inference. *Preprint*, arXiv:2607.13205.

Shishir G Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and

Joseph E. Gonzalez. 2025. The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392. PMLR.

Akshat Ramachandran, Marina Neseem, Charbel Sakr, Rangharajan Venkatesan, Brucek Khailany, and Tushar Krishna. 2026. ThinKV: Thought-adaptive KV cache compression for efficient reasoning models. In *International Conference on Learning Representations*, volume 2026, pages 110072–110106.

Torsten Scholak, Nathan Schucher, and Dzmitry Bahdanau. 2021. PICARD: Parsing incrementally for constrained auto-regressive decoding from language models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 9895–9901, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.

Akshat Sharma, Hangliang Ding, Jianping Li, Neel Dani, and Minjia Zhang. 2025. MiniKV: Pushing the limits of 2-bit KV cache via compression and system co-design for efficient long context inference. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 18506–18523, Vienna, Austria. Association for Computational Linguistics.

Hanzhang Shen, Haoran Wu, Yiren Zhao, and Robert Mullins. 2026. TriAxialKV: Toward extreme low-precision KV-cache quantization for agentic inference tasks. *Preprint*, arXiv:2605.17170.

Shubham Dipak Ugare, Rohan Gumaste, Tarun Suresh, Gagandeep Singh, and Sasa Misailovic. 2025. IterGen: Iterative semantic-aware structured LLM generation with backtracking. In *International Conference on Learning Representations*, volume 2025, pages 24369–24394.

Brandon T. Willard and Rémi Louf. 2023. Efficient guided generation for large language models. *Preprint*, arXiv:2307.09702.

Jialong Wu, Zhenglin Wang, Linhai Zhang, Yilong Lai, Yulan He, and Deyu Zhou. 2025. SCOPE: Optimizing key-value cache compression in long-context generation. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10775–10790, Vienna, Austria. Association for Computational Linguistics.

Guangxuan Xiao, Yuandong Tian, Beidi Chen, Song Han, and Mike Lewis. 2024. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations*, volume 2024, pages 21875–21895.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, and 41 others. 2025. Qwen3 technical report. *Preprint*, arXiv:2505.09388.

June Yong Yang, Byeongwook Kim, Jeongin Bae, Beomseok Kwon, Gunho Park, Eunho Yang, Se Jung Kwon, and Dongsoo Lee. 2024. No token left behind: Reliable KV cache compression via importance-aware mixed precision quantization. *Preprint*, arXiv:2402.18096.

Haoyue Zhang, Hualei Zhang, Xiaosong Ma, Jie Zhang, and Song Guo. 2026a. LazyEviction: Lagged KV eviction with attention pattern observation for efficient long reasoning. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 36335–36352, San Diego, California, United States. Association for Computational Linguistics.

Junkai Zhang, Hang Guo, Luca Benini, and Yawei Li. 2026b. RDKV: Rate-distortion bit allocation for joint eviction and quantization of the KV cache. *Preprint*, arXiv:2605.08317.

Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuandong Tian, Christopher Ré, Clark Barrett, Zhangyang "Atlas" Wang, and Beidi Chen. 2023. H2O: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems*, volume 36, pages 34661–34710. Curran Associates, Inc.

Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. 2024. SGLang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37, pages 62557–62583. Curran Associates, Inc.

Xiabin Zhou, Wenbin Wang, Minyan Zeng, Jiaxian Guo, Xuebo Liu, Li Shen, Min Zhang, and Liang Ding. 2025. DynamicKV: Task-aware adaptive KV cache compression for long context LLMs. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 8042–8057, Suzhou, China. Association for Computational Linguistics.

A Structural Bucket Construction

A deterministic adapter converts grammar information already computed for constrained decoding into lookup keys; offline task feedback determines their protection.

A.1 Decode tags and structural attribution

Table 5 lists the task-specific vocabularies for the structural tags defined in Eq. 4. BFCL assigns roles according to function dispatch and required or optional arguments, whereas BIRD uses SQL productions and schema context.

Ambiguous transitions and tokens spanning structural boundaries retain candidate tags for conservative action and floor resolution (Appendix B). Unseen schemas use context and parent fallback; supporting a new grammar family requires an adapter implementing the same four-field interface.

A.2 Layer groups and lookup keys

For L layers, let $b_1 = \lfloor (L+2)/3 \rfloor$ and $b_2 = \lfloor (2L+2)/3 \rfloor$. The lower, middle, and upper groups are $[0, b_1)$, $[b_1, b_2)$, and $[b_2, L)$, respectively. Their zero-indexed ranges are 0–11/12–23/24–35 for Qwen3-4B, 0–13/14–26/27–39 for Qwen3-14B, and 0–10/11–21/22–31 for LLaMA-3.1-8B. Payload costs use the actual dimensions of each group.

Decode and prefill lookup keys are $q^{\mathrm{D}}_{t,c} = (z_t, c)$ and $q^{\mathrm{P}}_{t,c} = (\mathrm{label}_t, c)$, respectively. Each calibrated entry stores its base action, protection floor, and adjacent-upgrade scores.

A.3 Prompt-span labels

Prompt labels come from deterministic span matching. All text starts as `protected_content`. Validated BFCL tool-schema JSON and BIRD SQLite DDL spans become `schema_scaffold`, with finer labels assigned by serialization-specific patterns (Table 6); invalid or unrecognized spans retain the default. No human annotation or general-purpose AST parser is used; new serializations require new patterns.

Span offsets map labels to tokens. Cross-boundary tokens conservatively combine the corresponding actions and floors (Appendix B), without imposing a universal risk order on labels. `protected_content` is a category, not necessarily an H action; sink and recent-window masks independently enforce high retention.

A.4 Calibration and frozen-policy records

Calibration follows Section 3.4 and Algorithm 1. Unobserved buckets use fallback mappings. Affected reference query positions are counted once per bucket and shared across distortion actions.

Within each threshold pair, adjacent upgrades use a fixed serialized bucket/action tie order; zero-score upgrades remain eligible until the all-high endpoint. Candidates retain their floor maps, so identical action maps with different floors remain distinct controllers. Nominal reference-trajectory cost is nondecreasing within each upgrade path, but autoregressive accuracy and realized cost need not be. The all-high fallback controller uses $B_{\mathrm{D}} = 1$, all-high floors, and $\Gamma = 0$.

The selected frozen policy records adapter mappings, tag vocabularies, layer boundaries, quantizer and payload-cost conventions, sink and recent-window settings, thresholds, stage budgets, reserve horizon and quantile, fallback maps, and tie order. These records identify the configuration used for evaluation and policy transfer.

B KV-Cache Backend

Fig. 7 summarizes parser-conditioned storage and attention. Scheduled cost, live KV payload, and physical GPU memory are accounted for separately.

B.1 Parser hand-off and resolution

SGLang (Zheng et al., 2024) receives XGrammar (Dong et al., 2025) frontiers before and after each accepted non-stop token. Each frontier is a sorted state set with fields *rule_id*, *sequence_id*, *element_id*, and *sub_element_id*. The serialized $F_{t-1} \to F_t$, schema key, and context resolve z_t before the next forward pass produces the token's KV. Lookup uses $z_t :: \{\mathrm{lower}, \mathrm{middle}, \mathrm{upper}\}$. Missing or ambiguous mappings use schema context, then

$$(\mathrm{class}, \mathrm{role}, \mathrm{state}, \mathrm{next}) \to (\mathrm{class}, \mathrm{role}, \mathrm{state}) \to (\mathrm{class}, \mathrm{role}) \to \mathrm{class}.$$

For candidate descendants $\mathcal{U}$, group c receives

$$\widetilde{a}^{0}_{c} = \widetilde{a}^{-}_{c} = \max_{\preceq} \{a^{0}_{u,c}, a^{-}_{u,c} : u \in \mathcal{U}\}. \qquad (12)$$

Thus budget adjustment cannot undo conservative resolution. Optional promotions require a unique calibrated score; no compatible entry defaults to H for both action and floor. Cross-label prompt tokens follow the same rule. Parent aggregates are

Component	Function calling (BFCL)	Text-to-SQL (BIRD)
Class	`pure_scaffold`, `schema_bearing`, `free_generation`	`pure_scaffold`, `schema_bearing`, `free_generation`
Role	`scaffold_separator`, `required`, `optional`, `dispatch_key`	`projection`, `relation`, `predicate`, `outside_schema`
State	`boundary`, `critical`	`stable`, `boundary`, `critical`
Next	`lookahead_critical`, `no_lookahead`	`lookahead_critical`, `no_lookahead`

Table 5: Task-specific vocabularies for the four structural-tag components.

Semantic label	Assignment rule
Protected content	Prompt text outside validated schema spans: instructions, requests, and chat context.
Schema scaffold	JSON/DDL structure and unmatched tokens within validated schema spans.
Function name	BFCL tool/function names; BIRD table identifiers.
Parameter key	BFCL property names and `required` entries; BIRD column identifiers.
Type constraint	BFCL `type` values; BIRD declared column types.
Enum constraint	BFCL `enum` values; recognized BIRD constraints, e.g., `PRIMARY KEY`, `REFERENCES`, and `NOT NULL`.
Schema description	BFCL `description` text.

Table 6: Semantic labels for prefill buckets in BFCL and BIRD.

precomputed; frontier adaptation is separate from constant-time lookup. Local conservatism does not guarantee unseen-schema accuracy.

B.2 Mixed-precision pools and protection

Layers share group-level actions but maintain separate KV tensors and slot maps. Current KV remains available during its forward pass, and sink/recent-window masks enforce H until protection expires. Scheduled L/R actions then trigger quantization or release. Subsequent promotion cannot reconstruct discarded or pre-quantization values. Slot maps preserve absolute positions across released entries.

B.3 Payload cost and physical accounting

For KV-head count H_ℓ^{KV} and head dimension d_ℓ, group-wise K/V element and scale counts are

$$n_c = 2\sum_{\ell \in c} H_\ell^{\mathrm{KV}} d_\ell, \quad k_c = 2\sum_{\ell \in c} H_\ell^{\mathrm{KV}}. \tag{13}$$

For even dimensions, scheduled bit costs are

$$\begin{aligned} m(R,c) &= 0, \qquad m(H,c) = 16n_c, \\ m(L,c) &= 4n_c + 16k_c. \end{aligned} \tag{14}$$

Odd dimensions add $4p_c$ padding bits to $m(L,c)$, where $p_c = 2\sum_{\ell \in c} H_\ell^{\mathrm{KV}}(d_\ell \bmod 2)$. Symmetric quantization needs no zero points. Group-specific costs give $C_{\mathrm{full}} = \sum_c m(H,c)$.

The ledger charges long-range schedules, so R may still occupy protected-window storage. Live payload counts prompt/generated KV, protection, and scales without duplicating aliases. Physical memory additionally includes indices, alignment, pages, allocator overhead, and transient buffers. Peak GPU memory is measured across prefill, migration, and decode, not inferred from C_{KV}.

B.4 Prefill planning, migration, and admission

For n prompt tokens, the prefill target is $T_{\mathrm{P}} = B_{\mathrm{P}} n C_{\mathrm{full}}$, subject to sink/recent-window protection. The floor-constrained allocation uses the frozen priorities described in Section 3.6; ties follow token position and then layer-group order. Floor excess is recorded rather than overriding protection.

The selected actions determine persistent counts

$$H_c = \sum_{t=1}^{n} \mathbb{I}[a_{t,c}^{\mathrm{P}} = H], \qquad L_c = \sum_{t=1}^{n} \mathbb{I}[a_{t,c}^{\mathrm{P}} = L],$$

and payload $H_c m(H,c) + L_c m(L,c)$, including mask-protected positions. Reservations add slot/page rounding and metadata. Prefill first uses all-BF16 scratch storage, then retains H, quantizes L, and releases unprotected R.

Admission checks each pool against persistent demand, scratch storage, and overlapping migration buffers, deduplicating only allocator-supported aliases. Capacity shortage or allocation failure defers or retries a request.

B.5 Packed attention and quantization

The kernel reads disjoint high/low maps, separate K/V scales, current KV, and logical positions. GQA maps query heads to shared KV heads (Ainslie et al., 2023), supporting one-head, GQA-2, GQA-4, and split-KV routes. Each visible key appears once across persistent, recent, and current storage; absolute positions and causal masks survive storage gaps.

Each token–KV-head K/V row has an independent symmetric INT4 scale: maximum absolute

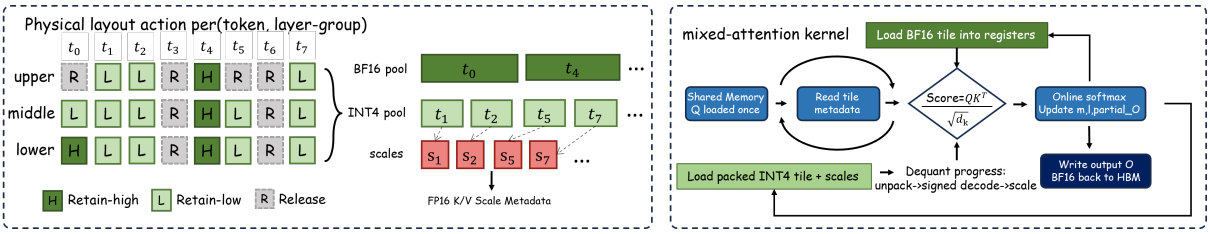

Figure 7: PASK backend. Token–layer KV is retained in BF16/INT4 or released after recent protection. Disjoint index streams, on-demand dequantization, and shared FP32 online softmax support mixed-precision attention.

value divided by 7, rounded to FP16 and bounded below by its smallest positive representable value. P/s^P is rounded to the nearest integer and clipped to $[-7, 7]$, using the same stored scale for dequantization. Zero rows use zero codes and unit scale. Two signed nibbles occupy one byte; odd dimensions add masked zero padding. Scales have shape $(\text{slot}, \text{KV head}, 1)$.

Dequantization occurs inside attention. Query arithmetic, scores, softmax statistics, and value accumulation use FP32; output is BF16. Nonempty disjoint partitions with score maximum v_j, exponential sum l_j, and unnormalized value sum u_j share normalization:

$$v = \max_j v_j, \quad o = \frac{\sum_j e^{v_j - v} u_j}{\sum_j e^{v_j - v} l_j}. \tag{15}$$

Empty partitions contribute zero; independently normalized outputs are not averaged.

B.6 Decode budget lifecycle and reserve

The per-token allowance is $b = B_{\mathrm{D}} C_{\mathrm{full}}$ and retains fractional bits. The recurrence in Algorithm 2 equivalently gives $R_t = tb - \sum_{i<t} M(\mathbf{a}_i)$, with $M(\mathbf{a}) = \sum_c m(a_c, c)$. This balance is a soft scheduled-payload target, not a physical-memory cap. Floor-induced debt carries forward. The quantity $[M(\mathbf{a}_t) - R_t]_+$ measures the excess of the current action cost over the pre-action balance, which may already include debt from previous tokens.

For reference floor cost $f_{r,v} = \sum_c m(a_{\min}(q^{\mathrm{ref}}_{r,v,c}), c)$, the offline reserve statistic is

$$S^{(K)}_{r,u}(B_{\mathrm{D}}) = \max_{0 \le h \le \min(K, T^{\mathrm{ref}}_r - u)} \left[\sum_{j=1}^{h} f_{r,u+j} - hb \right]_+ . \tag{16}$$

The prefix maximum captures temporary deficits that subsequent allowance could offset. Equation 10 uses $K = 8$ and $\beta = 0.8$; the empirical quantile is the value at rank $\lceil \beta N \rceil$ among N sorted samples, with a pooled fallback for unseen states. Future reference tokens enter only this offline construction. Calibration-distribution coverage does not imply coverage on unseen trajectories.

B.7 Termination, rollback, and control complexity

Stop tokens incur no new charge. The controller applies through the generation limit; termination reclaims request storage, and the task oracle scores incomplete outputs. Speculative proposals leave the committed ledger unchanged. Rollback restores parser state, budget, maps, and pending actions, undoing both allowances and charges. Destructive KV updates are deferred or snapshotted until commitment; restoring accounting alone cannot recover discarded KV.

After tag resolution, lookup and parent fallback use precomputed entries. At most two adjacent changes per group in either direction yield $O(|\mathcal{C}| \log |\mathcal{C}|)$ control work and $O(|\mathcal{C}|)$ state per token, independent of history length for fixed groups. This excludes frontier adaptation, quantization, attention, storage, and rollback history. Physical-capacity shortages defer decoding rather than violate floors.

C Experimental Details

C.1 Datasets

BFCL uses a seeded 1:1:3 calibration/development/test split within each non-Live and Live category, assigning complete normalized tool-call schemas to a single partition. For BIRD, we use the official 2024-06-27 Dev package, partitioned by database: `debit_card_specializing` and `european_football_2` for calibration (193 questions), `california_schools` and `financial` for development (195), and the remaining seven databases for held-out evaluation (1,146). Table 7 reports the category-level composition.

Benchmark	Subset	Cases	Test	Description
BFCL non-live (cases 1,000)	Simple	400	240	A single tool call matching one target function schema.
	Multiple	200	120	One tool call selected among several candidate functions.
	Parallel	200	120	Several independent tool calls emitted in a single response.
	Parallel-multiple	200	120	Several calls drawn from several candidate functions.
	All subsets	1,000	600	200 calibration and 200 development cases.
BFCL Live (cases 1,351)	Simple	258	154	Real-user queries with a single tool call.
	Multiple	1,053	633	Real-user queries requiring function selection among candidates.
	Parallel	16	10	Real-user queries with several independent calls.
	Parallel-multiple	24	14	Real-user queries with several calls over candidate functions.
	All subsets	1,351	811	270 calibration and 270 development cases.
BIRD Dev (cases 1,534)	Simple	925	701	Straightforward SQL over the given database schema.
	Moderate	464	330	Queries involving joins, aggregations, or nested conditions.
	Challenging	145	115	Queries requiring complex reasoning over schemas and values.
	All databases	1,534	1,146	Seven held-out SQLite databases: toxicology (145), card_games (191), codebase_community (186), superhero (129), formula_1 (174), thrombosis_prediction (163), and student_club (158).

Table 7: Benchmark composition. **Cases** denotes the full benchmark and **Test** the held-out evaluation subset.

C.2 Task-Accuracy Evaluation

BFCL predictions are scored by the official AST checker, with Overall accuracy defined as the instance-weighted fraction of correct cases. BIRD execution accuracy (EX) uses the official `mini_dev` evaluator. All methods share the same prompts, SGLang/XGrammar constrained-decoding backend, greedy decoding (temperature 0), and task scorer. Results are deterministic single-run point estimates.

PASK uses these task oracles for the calibration interventions defined in Section 3.4, whereas baseline task feedback is used for offline hyperparameter selection. No method receives task labels or task-error feedback during inference. Baseline implementations and their budget conventions are detailed in Appendix C.3.[1]

C.3 Baseline tuning and implementation

Tuning protocol. Each baseline is tuned separately for every model–dataset pair: BFCL non-Live and Live with Qwen3-4B, Qwen3-14B, and LLaMA-3.1-8B, and BIRD with Qwen3-4B. Candidate configurations are selected on the development split using the official task metric and frozen before test evaluation.

Budget-flexible methods are adjusted to the target KV budget. Fixed-profile methods that cannot exactly match the target retain their native allocation and precision profiles, with measured KV payload ratios reported separately. Table 8 lists representative high-impact search dimensions; secondary implementation-specific parameters are omitted.

H2O. H2O accumulates attention received by past tokens across decoding steps and retains the highest-scoring history together with protected sink and recent tokens.

SnapKV. SnapKV uses attention from the final prompt queries to score prefix tokens and pools neighboring scores before selecting the retained prefix.

MiKV. MiKV assigns importance-ranked KV entries to a high-precision tier and stores the remaining entries at lower precision rather than dropping them. We use H2O-style cumulative-attention importance. No author implementation was available when this reproduction was built. The high-precision ratio is tuned separately at each target budget.

RocketKV. RocketKV combines pooled attention-based prompt selection with coarse-to-fine sparse attention during decoding.

We use a batched SGLang implementation of its two-stage algorithm. Prefill selection scores prefix tokens with the observation-window scorer (odd max-pool kernel 63) and retains $\lceil r_{\text{prompt}} \cdot L \rceil$ tokens, where r_{prompt} is the configured prompt-retention ratio under the same shared-token-mask convention as SnapKV. During decoding, newly generated KV remains in BF16, while sparse atten-

[1] Official implementations: H2O (FMInference/H2O), SnapKV (FasterDecoding/SnapKV), RocketKV (NVlabs/RocketKV), LazyEviction (Halo-949/LazyEviction), and CriticalKV (via the authors' FFY0/AdaKV). No standalone official public repositories were available at the time of writing for MiKV, TriAxialKV, or Transactional Attention.

Baseline	Representative tuned knobs	Candidate values
H2O	Sink size; recent window	Sink $\in \{2, 4, 8\}$; recent window $\in \{16, 32, 64, 128\}$.
SnapKV	Observation window; pooling; pooling kernel	Observation window $\in \{16, 32, 64, 128\}$; pooling $\in \{\text{avg}, \max\}$; kernel $\in \{3, 5, 7\}$.
MiKV	High-precision ratio; low-bit width; group size	High-precision ratio $\in \{0.10, 0.20, 0.30, 0.40\}$; low-bit width $\in \{2, 4\}$; group size $\in \{32, 64, 128\}$.
RocketKV	Observation window; pooling kernel; prefill ratio	Observation window $\in \{16, 32, 64, 128\}$; kernel $\in \{31, 63, 127\}$; prefill ratio $\in \{0.20, 0.25, 0.30, 0.35\}$.
CriticalKV	First-stage ratio; stabilizer ϵ; scorer window	First-stage ratio $\in \{0.25, 0.40, 0.50, 0.70\}$; $\epsilon \in \{10^{-5}, 10^{-4}, 10^{-3}\}$; scorer window $\in \{16, 32, 64\}$.
LazyEviction	Observation/recent window; attention threshold α	Observation/recent window $\in \{8, 16, 32, 64\}$; $\alpha \in \{10^{-5}, 10^{-4}, 3\times10^{-4}, 10^{-3}\}$.
TriAxialKV	Protected prefix; precision profile; recent window	Protected prefix $\in \{4, 8, 16, 32\}$; (low, high) bits $\in \{(2, 4), (3, 4), (4, 4), (2, 8)\}$; recent window $\in \{32, 128, 256\}$.
Transactional Attention	Anchor mode; sponsor span; protected fraction	Anchor mode $\in \{\text{regex}, \text{semantic}\}$; sponsor span $\in \{4, 6, 8, 12\}$; protected fraction $\in \{0.15, 0.25, 0.35\}$.

Table 8: Representative baseline hyperparameter search spaces.

tion reads $\lceil 0.5 \cdot \text{prompt budget} \rceil$ selected tokens per request and KV head using the chunked max-key sketch with head- and sequence-dimension reductions. The official fused kernels are integrated with TensorRT-LLM and are not directly composable with our constrained-generation stack; therefore, the same-stack comparison uses this batched implementation.

CriticalKV. CriticalKV first constructs a candidate set using a SnapKV-style attention score and then reranks candidates according to the output perturbation induced by removing individual value vectors.

LazyEviction. LazyEviction tracks recurring attention over an observation window and delays eviction so that tokens that regain importance can remain cached.

TriAxialKV. Our text-only reproduction follows the per-tag attention-output distortion calibration and budget-constrained INT2/INT4 allocation procedure of TriAxialKV (Shen et al., 2026). It assigns precision according to semantic tags, protecting the prefix and recent tail while using lower precision for older prompt tokens. We freeze the calibrated allocation before test evaluation and report its measured KV cost separately, since the evaluated configuration does not match the target budget. Under this configuration, TriAxialKV uses an element-weighted average precision of 3.84 bits for stored K/V values and has a measured relative total KV cost of $B = 0.24$.

Transactional Attention. TA-Fast-Semantic detects semantic anchors and assigns decaying sponsorship vouchers to nearby tokens, which contribute to the token-retention utility score. Additional sponsorship and utility parameters are tuned under the same held-out selection protocol but are omitted from Table 8 for compactness.

C.4 System experiment

The serving stack uses SGLang 0.4.6.post5, PyTorch 2.6.0+cu124 and XGrammar 0.1.21. Radix caching, overlap scheduling, and chunked prefill are disabled in the primary serving matrix.

Each GPU configuration is measured over three successful repetitions, with medians reported. Concurrency c denotes requests submitted concurrently in each client-side wave, including those queued for server admission; checkpoint metrics are collected at the specified generated-token position using matched request-level workloads.

Resource licenses. Our code release contains original PASK code and modified SGLang source files. The modified SGLang files are redistributed under Apache License 2.0, with original license and attribution notices retained and modifications explicitly marked. Qwen3-4B-Instruct-2507 and Qwen3-14B are distributed under Apache License 2.0. LLaMA-3.1-8B-Instruct is distributed under the Llama 3.1 Community License Agreement and is subject to its Acceptable Use Policy. BFCL is distributed under Apache License 2.0, while BIRD's published data license is Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0). We use these models and datasets for offline research evaluation and do not redistribute benchmark datasets or pretrained model weights.

C.5 Structural ablation protocols.

The canonical policy key combines token class, structural role, parser state, lookahead, and layer group. The no-class and no-role variants remove their respective components, while no-state removes both parser state and lookahead. The no-layer variant ties actions across the three layer groups without removing any physical layers.

Method	B	BFCL non-live					BFCL live				
		Simple	Multiple	Parallel	Par.-multiple	Overall	Simple	Multiple	Parallel	Par.-multiple	Overall
Full KV	1.00	95.00	95.00	90.00	84.17	91.83	78.57	74.57	60.00	78.57	75.22
H2O	0.33	75.83	81.67	31.67	19.17	56.83	45.45	48.18	0.00	21.43	46.61
SnapKV	0.33	85.00	87.50	42.50	35.00	67.00	55.19	59.08	0.00	14.29	56.84
MiKV	0.33	87.50	90.83	76.67	76.67	83.83	64.94	64.45	30.00	50.00	63.87
RocketKV	0.33	75.42	65.83	33.33	25.00	55.00	48.05	47.24	0.00	28.57	46.49
TransAttn	0.33	88.75	64.17	54.17	43.33	67.83	45.45	27.49	0.00	0.00	30.09
CriticalKV	0.33	87.92	91.67	59.17	40.00	73.33	59.74	66.19	30.00	28.57	63.87
LazyEviction	0.33	46.67	81.67	20.83	14.17	42.00	33.12	60.03	10.00	28.57	53.76
TriAxialKV	fixed	90.42	85.00	70.83	56.67	78.67	62.99	59.08	40.00	0.00	58.57
PASK (Ours)	**0.33**	**92.08**	**93.33**	**80.83**	**81.67**	**88.00**	**74.03**	**72.20**	**60.00**	**50.00**	**72.01**
H2O	0.40	77.92	84.17	42.50	34.17	63.33	50.00	51.97	0.00	21.43	50.43
SnapKV	0.40	89.17	94.17	45.00	39.17	71.33	59.74	69.98	20.00	28.57	66.71
MiKV	0.40	88.33	90.83	79.17	75.00	84.33	66.23	64.61	30.00	42.86	64.12
RocketKV	0.40	86.67	95.83	65.00	41.67	75.17	59.74	63.03	20.00	35.71	61.41
TransAttn	0.40	90.83	73.33	61.67	47.50	72.83	52.60	44.71	20.00	0.00	45.13
CriticalKV	0.40	91.25	92.50	61.67	64.17	80.17	60.39	72.04	40.00	35.71	68.80
LazyEviction	0.40	58.75	76.67	34.17	38.33	53.33	47.40	61.77	10.00	50.00	58.20
PASK (Ours)	**0.40**	**95.83**	**97.50**	**86.67**	**80.00**	**91.17**	**77.27**	**74.88**	**50.00**	**50.00**	**74.60**
H2O	0.50	81.25	85.83	52.50	40.00	68.17	52.60	60.66	10.00	28.57	57.95
SnapKV	0.50	92.08	95.83	66.67	61.67	81.67	62.34	74.57	50.00	28.57	71.15
MiKV	0.50	89.17	90.83	80.00	73.33	84.50	69.48	65.09	30.00	42.86	65.10
RocketKV	0.50	90.00	94.17	75.83	60.83	82.17	58.44	69.67	30.00	42.86	66.58
TransAttn	0.50	94.58	78.33	85.83	62.50	83.17	68.18	46.76	60.00	21.43	50.55
CriticalKV	0.50	92.92	95.00	75.00	70.00	85.17	64.29	74.57	50.00	42.86	71.76
LazyEviction	0.50	75.42	90.83	35.00	35.83	62.50	55.19	68.56	10.00	57.14	65.10
PASK (Ours)	**0.50**	**95.00**	**97.50**	**87.50**	**81.67**	**91.33**	**79.22**	**74.72**	**50.00**	**57.14**	**74.97**
H2O	0.60	85.42	90.00	65.00	58.33	76.83	57.14	64.14	20.00	35.71	61.78
SnapKV	0.60	93.33	96.67	83.33	78.33	89.00	69.48	75.36	50.00	42.86	73.37
MiKV	0.60	92.08	91.67	83.33	75.00	86.83	71.43	66.51	50.00	71.43	67.32
RocketKV	0.60	90.42	94.17	80.83	75.83	86.33	62.99	72.83	30.00	35.71	69.79
TransAttn	0.60	95.42	95.00	89.17	74.17	89.83	74.68	65.72	60.00	57.14	67.20
CriticalKV	0.60	94.58	95.00	80.00	81.67	89.17	69.48	74.88	50.00	57.14	73.24
LazyEviction	0.60	87.08	83.33	35.83	45.00	67.67	52.60	76.15	0.00	57.14	70.41
PASK (Ours)	**0.60**	**95.42**	**97.50**	**85.83**	**84.17**	**91.67**	**77.92**	**74.88**	**60.00**	**64.29**	**75.09**

Table 9: Accuracy (%) at multiple KV budgets on Qwen3-4B.

The no-parser-transition variant replaces structural tags with request-local position bins, retaining layer groups and the $R/L/H$ action set. We construct ten contiguous bins from Full-KV calibration completion lengths, requiring each bin to be reached by at least 75% of calibration requests in every BFCL category; the final bin absorbs the tail. The parser enforces token legality but does not inform KV-action lookup.

For each variant, we reconstruct its policy buckets and independently recalibrate Runtime-V task-error sensitivity and Runtime-D attention-output distortion using the same calibration split and single-bucket intervention protocol. Each variant uses its measurements and policy.

D Additional Evidence

D.1 Multi-budget BFCL results

Tables 9–11 provide category-level BFCL results underlying Fig. 3.

D.2 BIRD detailed results

Table 12 decomposes BIRD execution accuracy by difficulty tier across the evaluated KV budgets.

D.3 Detailed efficiency results

Table 13 reports TPOT and throughput across input/output lengths on A100 and L40S.

D.4 High-concurrency scaling

Table 14 extends the main system results to higher concurrency at the 1024-token checkpoint. It focuses on 32K contexts, where PASK is designed to reduce persistent KV traffic during concurrent decode.

Across $c \in \{8, 16, 32\}$, PASK increases throughput and reduces TPOT relative to Full KV on both GPUs.

D.5 Reliability-tolerance sensitivity

For each stage, we select the lowest-budget candidate satisfying the development reliability gate.

Method	B	BFCL non-live					BFCL live				
		Simple	Multiple	Parallel	Par.-multiple	Overall	Simple	Multiple	Parallel	Par.-multiple	Overall
Full KV	1.00	96.25	98.33	92.50	81.67	93.00	79.87	75.20	70.00	64.29	75.83
H2O	0.33	72.08	85.00	22.50	16.67	53.67	54.55	51.82	0.00	0.00	50.80
SnapKV	0.33	82.50	91.67	60.00	55.83	74.50	61.69	66.98	0.00	14.29	64.24
MiKV	0.33	83.75	88.33	79.17	65.83	80.17	66.23	63.03	40.00	42.86	63.01
RocketKV	0.33	78.33	90.00	30.00	41.67	63.67	55.19	55.13	30.00	42.86	54.62
TransAttn	0.33	88.33	83.33	70.00	49.17	75.83	49.35	31.12	0.00	0.00	33.66
CriticalKV	0.33	80.42	87.50	54.17	50.83	70.67	61.69	65.72	40.00	57.14	64.49
LazyEviction	0.33	49.17	91.67	28.33	25.00	48.67	33.77	64.14	20.00	57.14	57.71
TriAxialKV	fixed	87.92	79.17	77.50	64.17	79.33	68.18	61.14	30.00	0.00	61.04
PASK (Ours)	**0.33**	**92.08**	**94.17**	**86.67**	**79.17**	**88.83**	**73.38**	**73.62**	**30.00**	**57.14**	**72.75**
H2O	0.40	77.08	87.50	40.00	35.00	63.33	57.79	57.35	10.00	21.43	56.23
SnapKV	0.40	82.08	89.17	65.83	60.00	75.83	64.94	70.30	30.00	50.00	68.43
MiKV	0.40	88.75	92.50	85.00	70.83	85.17	65.58	63.67	40.00	42.86	63.38
RocketKV	0.40	84.58	94.17	51.67	41.67	71.33	57.14	65.88	30.00	57.14	63.63
TransAttn	0.40	91.67	85.00	73.33	46.67	77.67	55.84	45.34	40.00	21.43	46.86
CriticalKV	0.40	81.67	90.00	60.00	70.83	76.83	62.34	69.67	40.00	64.29	67.82
LazyEviction	0.40	62.50	86.67	30.83	36.67	55.83	50.00	66.19	0.00	50.00	62.02
PASK (Ours)	**0.40**	**93.33**	**94.17**	**87.50**	**78.33**	**89.33**	**75.32**	**74.09**	**40.00**	**57.14**	**73.61**
H2O	0.50	82.92	89.17	57.50	48.33	72.17	59.09	62.72	20.00	28.57	60.91
SnapKV	0.50	85.83	93.33	75.83	71.67	82.50	64.29	71.09	40.00	50.00	69.05
MiKV	0.50	88.75	91.67	85.83	71.67	85.33	64.94	64.77	40.00	42.86	64.12
RocketKV	0.50	86.25	94.17	73.33	59.17	79.83	60.39	67.77	40.00	57.14	65.84
TransAttn	0.50	95.42	88.33	89.17	58.33	85.33	64.29	48.50	30.00	28.57	50.92
CriticalKV	0.50	85.00	95.00	83.33	74.17	84.50	66.23	71.09	40.00	71.43	69.79
LazyEviction	0.50	70.00	94.17	49.17	46.67	66.00	52.60	69.35	20.00	57.14	65.35
PASK (Ours)	**0.50**	**93.33**	**94.17**	**88.33**	**80.83**	**90.00**	**76.62**	**74.41**	**50.00**	**57.14**	**74.23**
H2O	0.60	85.83	90.83	74.17	62.50	79.83	61.69	66.35	30.00	50.00	64.73
SnapKV	0.60	86.67	95.00	85.00	79.17	86.50	62.99	71.56	40.00	57.14	69.30
MiKV	0.60	89.58	92.50	88.33	76.67	87.33	65.58	66.51	70.00	50.00	66.09
RocketKV	0.60	89.17	93.33	82.50	70.00	84.83	63.64	69.67	40.00	64.29	68.06
TransAttn	0.60	95.83	98.33	91.67	75.83	91.50	72.73	51.66	30.00	28.57	54.99
CriticalKV	0.60	90.00	95.83	85.83	84.17	89.17	72.73	72.20	60.00	50.00	71.76
LazyEviction	0.60	90.83	89.17	51.67	52.50	75.00	63.64	74.57	10.00	35.71	71.02
PASK (Ours)	**0.60**	**95.00**	**97.50**	**94.17**	**83.33**	**93.00**	**78.57**	**74.88**	**60.00**	**57.14**	**75.09**

Table 10: Accuracy (%) at multiple KV budgets on Qwen3-14B.

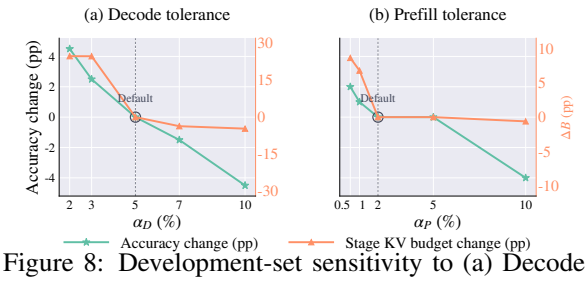

Figure 8: Development-set sensitivity to (a) Decode tolerance α_{D} and (b) Prefill tolerance α_{P}. Accuracy and stage-budget changes are measured in percentage points relative to the respective defaults 5% and 2%.

Fig 8 shows the resulting accuracy–budget trade-off around the defaults $\alpha_{\mathrm{D}} = 5\%$ and $\alpha_{\mathrm{P}} = 2\%$.

D.6 Calibration Cost

Table 15 reports the measured offline calibration cost on a single NVIDIA A100-80GB GPU (Qwen3-4B, the BFCL non-Live calibration split of 200 cases per candidate, the same side table and runtime contract that produced the deployed policy). Runtime-V perturbs one structure–layer bucket at a time and Runtime-D measures the fixed-trajectory distortion of each candidate; median candidate runtimes are 218.8 s for Runtime-V and 222.1 s for Runtime-D. With two concurrent workers sharing the same GPU, the full 54-candidate scans take 1.64 and 1.67 GPU-hours, respectively. The combined Runtime-V and Runtime-D calibration takes approximately 3.3 GPU-hours on one A100. The calibration is therefore a one-time cost of a few GPU-hours, amortized over all downstream use of the frozen lookup tables.

D.7 Direct Evidence of Responsibility–Distortion Mismatch

We compare task responsibility and model-side numerical sensitivity on the calibration buckets used to construct the frozen policy. For each action $a \in \{R, L\}$, task responsibility is measured by the effective sensitivity $\bar{V}(q, a)$ from Eq. 5, while model-side sensitivity is measured by the corre-

Method	B	BFCL non-live					BFCL live				
		Simple	Multiple	Parallel	Par.-multiple	Overall	Simple	Multiple	Parallel	Par.-multiple	Overall
Full KV	1.00	94.58	95.00	91.67	80.83	91.33	75.32	60.98	40.00	57.14	63.38
H2O	0.33	75.42	90.00	15.83	12.50	53.83	46.75	36.65	0.00	14.29	37.73
SnapKV	0.33	87.92	94.17	43.33	35.00	69.67	51.30	53.55	10.00	35.71	52.28
MiKV	0.33	91.67	95.00	66.67	51.67	79.33	64.94	48.50	20.00	28.57	50.92
RocketKV	0.33	82.92	90.83	25.00	21.67	60.67	51.30	47.39	0.00	21.43	47.10
TransAttn	0.33	90.00	67.50	71.67	38.33	71.50	60.39	32.07	30.00	0.00	36.87
CriticalKV	0.33	86.67	93.33	35.83	33.33	67.17	62.34	53.08	10.00	35.71	54.01
LazyEviction	0.33	56.67	88.33	25.83	16.67	48.83	38.96	45.66	10.00	28.57	43.65
TriAxialKV	fixed	92.08	88.33	85.83	55.83	82.83	66.23	38.23	50.00	28.57	43.53
PASK (Ours)	**0.33**	**93.75**	**96.67**	**90.83**	**80.83**	**91.17**	**73.38**	**58.45**	**40.00**	**50.00**	**60.91**
H2O	0.40	78.75	90.83	28.33	25.83	60.50	49.35	41.86	0.00	14.29	42.29
SnapKV	0.40	90.83	94.17	50.00	38.33	72.83	62.99	57.35	20.00	35.71	57.58
MiKV	0.40	91.67	94.17	69.17	53.33	80.00	64.94	51.50	20.00	35.71	53.39
RocketKV	0.40	82.92	90.83	25.00	21.67	60.67	63.64	51.50	10.00	28.57	52.90
TransAttn	0.40	93.33	69.17	80.00	36.67	74.50	59.09	40.44	30.00	0.00	43.16
CriticalKV	0.40	87.92	94.17	65.00	39.17	74.83	64.94	56.08	20.00	35.71	56.97
LazyEviction	0.40	66.25	87.50	33.33	36.67	58.00	37.01	53.08	20.00	50.00	49.57
PASK (Ours)	**0.40**	**93.75**	**95.00**	**90.83**	**80.83**	**90.83**	**73.38**	**59.72**	**40.00**	**57.14**	**62.02**
H2O	0.50	84.58	92.50	43.33	33.33	67.67	53.90	46.29	0.00	14.29	46.61
SnapKV	0.50	92.50	95.83	70.00	53.33	80.83	66.23	59.24	30.00	28.57	59.68
MiKV	0.50	91.67	93.33	70.00	63.33	82.00	64.94	53.40	10.00	35.71	54.75
RocketKV	0.50	89.17	95.00	75.00	49.17	79.50	62.34	54.34	30.00	50.00	55.49
TransAttn	0.50	95.00	90.00	88.33	52.50	84.17	69.48	42.81	40.00	7.14	47.23
CriticalKV	0.50	92.50	98.33	67.50	52.50	80.67	67.53	59.08	30.00	28.57	59.80
LazyEviction	0.50	80.83	94.17	48.33	41.67	69.17	54.55	53.71	0.00	42.86	53.02
PASK (Ours)	**0.50**	**94.58**	**95.00**	**90.83**	**80.83**	**91.17**	**75.32**	**59.56**	**40.00**	**57.14**	**62.27**
H2O	0.60	89.58	93.33	53.33	48.33	74.83	58.44	50.39	10.00	21.43	50.92
SnapKV	0.60	94.17	95.83	70.00	59.17	82.67	68.83	62.09	30.00	35.71	62.52
MiKV	0.60	92.08	95.00	70.83	69.17	83.83	65.58	54.66	50.00	57.14	56.72
RocketKV	0.60	91.25	95.83	81.67	63.33	84.67	63.64	57.82	20.00	28.57	57.95
TransAttn	0.60	95.83	91.67	90.83	80.00	90.83	74.03	57.66	40.00	42.86	60.30
CriticalKV	0.60	95.00	95.83	70.00	65.00	84.17	74.68	60.19	30.00	57.14	62.52
LazyEviction	0.60	94.17	95.00	50.00	55.83	77.83	41.56	63.35	0.00	57.14	58.32
PASK (Ours)	**0.60**	**93.75**	**96.67**	**91.67**	**80.83**	**91.33**	**74.03**	**60.66**	**40.00**	**50.00**	**62.76**

Table 11: Accuracy (%) at multiple KV budgets on LLaMA-3.1-8B.

sponding distortion $D(q, a)$ from Eq. 7. For compactness, let

$$\mathcal{T}^{(a)} = \mathrm{Top}_{20\%}\big(\bar{V}(\cdot, a)\big)\,.$$

$$\mathrm{Recall}^{(a)}_{20\%}(S) = \frac{\big|\mathcal{T}^{(a)} \cap \mathrm{Top}_{20\%}(S(\cdot, a))\big|}{\big|\mathcal{T}^{(a)}\big|}. \tag{17}$$

$$\mathrm{Mismatch}^{(a)}_{20,50}(S) = \frac{\big|\mathcal{T}^{(a)} \cap \mathrm{Bottom}_{50\%}(S(\cdot, a))\big|}{\big|\mathcal{T}^{(a)}\big|}. \tag{18}$$

A high mismatch value means that the model-side ranking places many task-critical buckets in its lower half. The attention-saliency quantity $A(q)$ requires a separate reference-attention pass and is not inferred from $D(q, a)$.

Across the 27 calibration buckets, the top-20% sets ranked by task-error sensitivity $\bar{V}(q, a)$ and distortion $D(q, a)$ do not overlap for either action (Table 16).

D.8 Prefill–Decode Decomposition

We decompose the policy into four prefill–decode combinations under a single strict JSON-schema BFCL non-live test protocol. The prefill policy was selected on development data using the $\alpha_{\mathrm{P}} = 2\%$ gate with decode fixed to all-high; the decode policy was frozen to the default $\tau = \{0.5\%, 1\%\}$. This stage decomposition measures each policy's contribution and their composition.

Protocol and budget accounting. All four rows use BFCL non-live test cases at concurrency 1 with `max_new_tokens=192`. The Full KV and Decode-only rows reuse the matching strict-protocol controls; Prefill-only and Combined use the development-frozen minimum-protection prefill policy. Following Section 4.1, we report measured KV costs normalized by the corresponding Full-KV reference. $B_{\mathrm{P}}^{\mathrm{meas}}$ and $B_{\mathrm{D}}^{\mathrm{meas}}$ denote the

Method	Budget	Simple EX	Moderate EX	Challenging EX	Overall EX
Full KV	1.000	62.77	46.67	44.35	56.28
H2O	0.275	32.81	10.61	9.57	24.08
SnapKV	0.275	44.79	25.15	19.13	36.56
MiKV	0.275	55.63	37.27	34.78	48.25
RocketKV	0.275	50.36	29.70	17.39	41.10
CriticalKV	0.275	45.22	28.18	17.39	37.52
LazyEviction	0.275	57.92	39.09	27.83	49.48
TransAttn	0.275	50.07	29.09	11.30	40.14
PASK (Ours)	**0.275**	**60.77**	**41.52**	**35.65**	**52.71**
H2O	0.4	42.80	15.45	19.13	32.55
SnapKV	0.4	51.64	35.76	23.48	44.24
MiKV	0.4	58.35	38.18	35.65	50.26
RocketKV	0.4	57.77	39.39	33.91	50.09
CriticalKV	0.4	50.50	30.30	21.74	41.80
LazyEviction	0.4	59.77	39.70	29.57	50.96
TransAttn	0.4	58.63	43.94	33.04	51.83
PASK (Ours)	**0.4**	**61.34**	**42.42**	**40.00**	**53.75**
H2O	0.5	44.37	21.21	22.61	35.51
SnapKV	0.5	55.78	37.27	19.13	46.77
MiKV	0.5	60.20	40.30	38.26	52.27
RocketKV	0.5	59.63	46.36	38.26	53.66
CriticalKV	0.5	52.35	35.15	25.22	44.68
LazyEviction	0.5	60.49	43.03	34.78	52.88
TransAttn	0.5	60.06	43.03	39.13	53.05
PASK (Ours)	**0.5**	**61.20**	**44.24**	**36.52**	**53.84**
H2O	0.6	50.50	27.27	23.48	41.10
SnapKV	0.6	57.06	40.91	31.30	49.83
MiKV	0.6	61.77	41.52	42.61	54.01
RocketKV	0.6	60.77	46.06	39.13	54.36
CriticalKV	0.6	58.92	41.52	29.57	50.96
LazyEviction	0.6	60.91	49.09	42.61	55.67
TransAttn	0.6	61.63	44.55	42.61	54.80
PASK (Ours)	**0.6**	**62.48**	**46.06**	**44.35**	**55.93**

Table 12: Detailed BIRD database-disjoint held-out execution accuracy (EX, %) for Qwen3-4B.

prefill and decode cost ratios, respectively. Total B is computed as the sum of measured prefill and decode KV payloads divided by the sum of their corresponding Full-KV reference payloads.

Result. Prefill-only retained 91.17% accuracy at total $B = 0.489$, whereas decode-only retained 90.67% at total $B = 0.841$. Activating both stages yielded 88.83% accuracy at total $B = 0.330$, a 4.17-percentage-point difference from Full KV. These operating points are not matched in total budget, so the table reports stage contributions under the frozen policies.

E Theoretical Properties

We formalize properties of the floor-constrained controller, greedy calibration path, online complexity, and empirical reserve construction.

E.1 Protection-floor preservation

Proposition 1 (Protection floors are invariant)**.** *Assume* DEMOTE *in Algorithm 2 only applies an adjacent demotion when the resulting action remains no lower than the stored floor $a^-_{t,c}$, as specified in Section 3.5. Then for every accepted token t and layer group c, the persisted action satisfies*

$$a_{t,c} \succeq a^-_{t,c}. \tag{19}$$

Proof. Algorithm 2 initializes

$$a_{t,c} = \max_{\preceq}(a^0_{t,c}, a^-_{t,c}),$$

which satisfies the invariant. By assumption, DEMOTE cannot cross the floor, while PROMOTE only moves upward in the order $R \prec L \prec H$. Both operations therefore preserve the invariant, which holds for the final persisted action. □

This invariant guarantees floor preservation, not task correctness; the latter depends on the empirical calibration in Eq. 5.

Input / Output	Method	NVIDIA A100		NVIDIA L40S	
		TPOT (ms) ↓	Throughput (T/S) ↑	TPOT (ms) ↓	Throughput (T/S) ↑
8k / 256	Full KV	$92.73_{(0.0\%)}$	$79.14_{(0.0\%)}$	$65.09_{(0.0\%)}$	$107.26_{(0.0\%)}$
	SnapKV	$54.66_{(41.1\%)}$	$120.00_{(51.6\%)}$	$56.80_{(12.7\%)}$	$116.04_{(8.2\%)}$
	RocketKV	$81.72_{(11.9\%)}$	$93.04_{(17.6\%)}$	$61.05_{(6.2\%)}$	$103.41_{(-3.6\%)}$
	CriticalKV	$54.34_{(41.4\%)}$	$99.68_{(26.0\%)}$	$49.09_{(24.6\%)}$	$115.07_{(7.3\%)}$
	PASK (Ours)	$\mathbf{45.89}_{(50.5\%)}$	$\mathbf{140.16}_{(77.1\%)}$	$\mathbf{43.62}_{(33.0\%)}$	$\mathbf{146.15}_{(36.3\%)}$
8k / 1024	Full KV	$89.49_{(0.0\%)}$	$87.26_{(0.0\%)}$	$63.45_{(0.0\%)}$	$121.67_{(0.0\%)}$
	SnapKV	$51.29_{(42.7\%)}$	$150.04_{(71.9\%)}$	$52.12_{(17.9\%)}$	$145.01_{(19.2\%)}$
	RocketKV	$116.15_{(-29.8\%)}$	$68.21_{(-21.8\%)}$	$66.75_{(-5.2\%)}$	$95.77_{(-21.3\%)}$
	CriticalKV	$45.03_{(49.7\%)}$	$155.22_{(77.9\%)}$	$41.35_{(34.8\%)}$	$172.30_{(41.6\%)}$
	PASK (Ours)	$\mathbf{43.28}_{(51.6\%)}$	$\mathbf{173.58}_{(98.9\%)}$	$\mathbf{40.15}_{(36.7\%)}$	$\mathbf{186.45}_{(53.2\%)}$
32k / 256	Full KV	$256.04_{(0.0\%)}$	$\mathbf{25.77}_{(0.0\%)}$	$87.85_{(0.0\%)}$	$9.72_{(0.0\%)}$
	SnapKV	$103.72_{(59.5\%)}$	$15.32_{(-40.6\%)}$	$89.28_{(-1.6\%)}$	$17.49_{(79.9\%)}$
	RocketKV	$201.13_{(21.4\%)}$	$13.52_{(-47.5\%)}$	$85.42_{(2.8\%)}$	$15.81_{(62.7\%)}$
	CriticalKV	$117.36_{(54.2\%)}$	$13.34_{(-48.2\%)}$	$98.83_{(-12.5\%)}$	$15.68_{(61.3\%)}$
	PASK (Ours)	$\mathbf{84.17}_{(67.1\%)}$	$18.57_{(-27.9\%)}$	$\mathbf{74.52}_{(15.2\%)}$	$\mathbf{21.26}_{(118.7\%)}$
32k / 1024	Full KV	$227.52_{(0.0\%)}$	$33.22_{(0.0\%)}$	$83.94_{(0.0\%)}$	$31.92_{(0.0\%)}$
	SnapKV	$89.09_{(60.8\%)}$	$42.44_{(27.8\%)}$	$74.27_{(11.5\%)}$	$48.15_{(50.8\%)}$
	RocketKV	$195.98_{(13.9\%)}$	$35.49_{(6.8\%)}$	$100.90_{(-20.2\%)}$	$28.62_{(-10.3\%)}$
	CriticalKV	$89.80_{(60.5\%)}$	$38.74_{(16.6\%)}$	$76.53_{(8.8\%)}$	$44.60_{(39.7\%)}$
	PASK (Ours)	$\mathbf{68.69}_{(69.8\%)}$	$\mathbf{51.02}_{(53.6\%)}$	$\mathbf{60.58}_{(27.8\%)}$	$\mathbf{50.84}_{(59.3\%)}$
32k / 2048	Full KV	$224.95_{(0.0\%)}$	$34.53_{(0.0\%)}$	$61.09_{(0.0\%)}$	$31.49_{(0.0\%)}$
	SnapKV	$97.58_{(56.6\%)}$	$43.37_{(25.6\%)}$	$73.59_{(-20.5\%)}$	$51.29_{(62.9\%)}$
	RocketKV	$150.87_{(32.9\%)}$	$38.13_{(10.4\%)}$	$91.90_{(-50.4\%)}$	$33.17_{(5.3\%)}$
	CriticalKV	$97.94_{(56.5\%)}$	$41.64_{(20.6\%)}$	$74.84_{(-22.5\%)}$	$48.95_{(55.4\%)}$
	PASK (Ours)	$\mathbf{67.28}_{(70.1\%)}$	$\mathbf{52.21}_{(51.2\%)}$	$\mathbf{54.87}_{(10.2\%)}$	$\mathbf{52.98}_{(68.2\%)}$

Table 13: Serving efficiency for Qwen3-4B at concurrency $c = 8$ and approximate total KV budget $B = 0.33$ across input/output lengths. Parenthetical percentages are relative to Full KV on the same GPU: positive TPOT values indicate reductions, positive throughput values indicate increases, and negative values indicate regressions.

E.2 Nested residual-allocation frontier

Proposition 2 (Nested greedy path)**.** *Fix a threshold pair τ and its floor policy $a^{\tau}_{\min}$. Let $\Theta_0, \Theta_1, \ldots, \Theta_J$ denote the bucket-action maps produced by* GREEDYPATH*, where each step applies one eligible adjacent upgrade $R \to L$ or $L \to H$ without demotion.*

For fixed reference trajectories $\mathcal{D}_{\mathrm{ref}}$, define the nominal reference-trajectory KV cost as

$$\widetilde{C}^{\mathrm{ref}}_{\mathrm{KV}}(\Theta) = \frac{\sum_{r\in\mathcal{D}_{\mathrm{ref}}} \sum_{t=1}^{T^{\mathrm{ref}}_r} \sum_{c\in\mathcal{C}} m\big(\Theta(q^{\mathrm{ref}}_{r,t,c}), c\big)}{C_{\mathrm{full}} \sum_{r\in\mathcal{D}_{\mathrm{ref}}} T^{\mathrm{ref}}_r}.$$

Here, $q^{\mathrm{ref}}_{r,t,c}$ denotes the structure–layer bucket of a token on the fixed reference trajectory.

Then the action maps are componentwise ordered:

$$\Theta_0 \preceq \Theta_1 \preceq \cdots \preceq \Theta_J. \tag{20}$$

Their nominal reference-trajectory costs are nondecreasing:

$$\widetilde{C}^{\mathrm{ref}}_{\mathrm{KV}}(\Theta_0) \le \widetilde{C}^{\mathrm{ref}}_{\mathrm{KV}}(\Theta_1) \le \cdots \le \widetilde{C}^{\mathrm{ref}}_{\mathrm{KV}}(\Theta_J). \tag{21}$$

Proof. By construction, each greedy step upgrades one bucket action without demoting any other action. Thus, $\Theta_i \preceq \Theta_{i+1}$ componentwise. Since the reference trajectories and their bucket assignments are fixed, and $m(a', c) > m(a, c)$ for each adjacent upgrade $a \prec a'$, the nominal payload numerator cannot decrease. The normalization is also fixed, establishing the cost inequality. □

The proposition applies to each fixed-τ greedy path. Candidate paths generated from different floor thresholds in Algorithm 1 need not be totally ordered, and development selection evaluates them explicitly using Eq. 9. The monotonicity result concerns nominal costs evaluated on fixed reference trajectories, rather than the realized autoregressive cost defined in Eq. 2. During deployment, policy-dependent generation trajectories and budget-controller adjustments may change realized costs, for which monotonicity is not guaranteed.

E.3 Online computational complexity

Proposition 3 (History-length-independent online control)**.** *Let $|\mathcal{C}|$ be the number of layer groups and assume constant-time lookup of one frozen table entry. For each accepted token, Algorithm 2 requires $O(|\mathcal{C}|)$ lookup/storage work and at most $O(|\mathcal{C}| \log |\mathcal{C}|)$ work to rank the $O(|\mathcal{C}|)$ eligible pro-*

GPU	Input	c	Throughput (tokens/s) ↑			TPOT (ms) ↓		
			Full KV	PASK	SnapKV	Full KV	PASK	SnapKV
A100	32K	8	$33.22_{(0.0\%)}$	$\mathbf{51.02}_{(53.6\%)}$	$42.44_{(27.8\%)}$	$227.52_{(0.0\%)}$	$\mathbf{68.69}_{(69.8\%)}$	$89.09_{(60.8\%)}$
A100	32K	16	$33.26_{(0.0\%)}$	$\mathbf{63.23}_{(90.1\%)}$	$51.84_{(55.9\%)}$	$331.35_{(0.0\%)}$	$\mathbf{68.00}_{(79.5\%)}$	$140.50_{(57.6\%)}$
A100	32K	32	$33.34_{(0.0\%)}$	$\mathbf{72.23}_{(116.6\%)}$	$51.66_{(54.9\%)}$	$333.99_{(0.0\%)}$	$\mathbf{68.96}_{(79.4\%)}$	$141.81_{(57.5\%)}$
L40S	32K	8	$31.92_{(0.0\%)}$	$\mathbf{50.84}_{(59.3\%)}$	$48.15_{(50.8\%)}$	$83.94_{(0.0\%)}$	$\mathbf{60.58}_{(27.8\%)}$	$74.27_{(11.5\%)}$
L40S	32K	16	$32.02_{(0.0\%)}$	$\mathbf{60.38}_{(88.6\%)}$	$48.29_{(50.8\%)}$	$85.54_{(0.0\%)}$	$\mathbf{60.74}_{(29.0\%)}$	$73.96_{(13.5\%)}$
L40S	32K	32	$32.50_{(0.0\%)}$	$\mathbf{71.37}_{(119.6\%)}$	$48.09_{(48.0\%)}$	$85.41_{(0.0\%)}$	$\mathbf{61.78}_{(27.7\%)}$	$74.32_{(13.0\%)}$

Table 14: High-concurrency system results for 32K inputs at the 1024-token checkpoint. Each cell is the median over three successful repetitions. Parenthetical percentages are relative to Full KV on the same GPU: positive throughput values indicate increases and positive TPOT values indicate reductions.

	Runtime-V	Runtime-D
All-high control	220.4 s	260.5 s
Candidate time, median	218.8 s	222.1 s
Candidate time, p90	221.1 s	226.5 s
GPU util, mean / peak	62% / 100%	
Peak memory	75.8 GiB	
Full 54-candidate scan	1.64 h	1.67 h

Table 15: Measured calibration cost on one NVIDIA A100-80GB (Qwen3-4B, BFCL non-Live calibration split, 200 cases per candidate).

Action	Model-side score	Recall @20%	Mismatch $_{20,50}$
Release R	$D(q,R)$	0.0%	100.0%
Retain-low L	$D(q,L)$	0.0%	83.3%

Table 16: Direct comparison between effective task responsibility and Runtime-D distortion on 27 Qwen3-4B BFCL non-live calibration buckets. The effective responsibility values use the monotone envelope in Eq. 5.

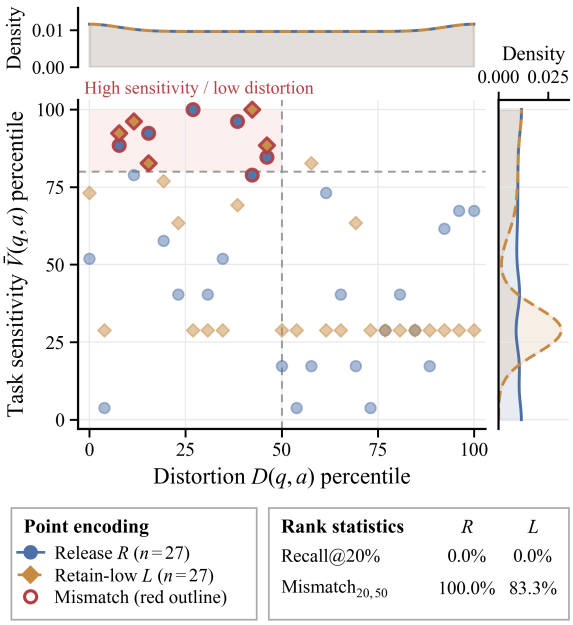

Figure 9: Responsibility–distortion mismatch across 27 calibrated structure–layer buckets. Blue circles represent release (R), and orange diamonds represent retain-low (L). Axes show percentile ranks of distortion $D(q,a)$ and task sensitivity $\bar{V}(q,a)$. Red outlines highlight top-20% sensitivity buckets ranked in the bottom 50% by distortion. Top and right curves show marginal kernel density estimates.

motions or demotions. Its request-local controller state is $O(|\mathcal{C}|)$. Consequently, for the fixed three-group implementation used in this work, the online control overhead is $O(1)$ with respect to sequence length and retained-history length.

Proof. The lookup key is formed once per layer group, giving $|\mathcal{C}|$ table accesses and an action vector of length $|\mathcal{C}|$. Each group has only the three ordered actions $R \prec L \prec H$, so at a token there are only $O(|\mathcal{C}|)$ adjacent promotion/demotion candidates. Ranking these candidates by their stored ρ_D scores is bounded by $O(|\mathcal{C}| \log |\mathcal{C}|)$ using a comparison sort, and applying the selected adjustments is linear in $|\mathcal{C}|$. The controller stores only the action/floor/score vectors and scalar budget state. None of these quantities grows with the number of previously generated tokens. □

This proposition concerns the policy-controller work only; the Transformer attention computation still depends on the retained KV representation and the serving kernel.

E.4 Empirical reserve coverage

Proposition 4 (Calibration-sample coverage of reserve deficits)**.** *For a parser state s with $n_s > 0$, let*

$$\mathcal{I}_s = \{(r,u) : s^{\mathrm{ref}}_{r,u} = s\}, \qquad n_s = |\mathcal{I}_s|, \quad (22)$$

Variant	Prefill policy	Decode policy	$B_{\mathrm{P}}^{\mathrm{meas}} \downarrow$	$B_{\mathrm{D}}^{\mathrm{meas}} \downarrow$	Total $B \downarrow$	BFCL non-live Acc. ↑
Full KV	Full	Full	1.000	1.000	1.000	93.00
Prefill-only	PASK	Full	0.250	1.000	0.489	91.17
Decode-only	Full	PASK	1.000	0.501	0.841	90.67
Combined PASK	PASK	PASK	0.250	0.501	0.330	88.83

Table 17: Prefill–decode stage decomposition for Qwen3-14B on BFCL non-live under the strict JSON-schema protocol.

and define the empirical CDF of the calibration-sample deficits in Eq. 16 as

$$\widehat{F}_{s;K,B_{\mathrm{D}}}(x) = \frac{1}{n_s} \sum_{(r,u)\in\mathcal{I}_s} \mathbb{I}\left[S_{r,u}^{(K)}(B_{\mathrm{D}}) \le x\right]. \tag{23}$$

If the state-conditioned reserve in Eq. 10 uses the left empirical β-quantile,

$$\Gamma(s; B_{\mathrm{D}}) = \inf\left\{x : \widehat{F}_{s;K,B_{\mathrm{D}}}(x) \ge \beta\right\}, \tag{24}$$

then

$$\widehat{F}_{s;K,B_{\mathrm{D}}}(\Gamma(s; B_{\mathrm{D}})) \ge \beta, \tag{25}$$

or equivalently, the empirical fraction of calibration deficits exceeding the reserve is at most $1-\beta$.

Proof. By definition, $\Gamma(s; B_{\mathrm{D}})$ is the left empirical β-quantile of the observed values $S_{r,u}^{(K)}(B_{\mathrm{D}})$ for state s. Therefore, the empirical CDF at this quantile is at least β. The empirical fraction of deficits exceeding the reserve is $1 - \widehat{F}_{s;K,B_{\mathrm{D}}}(\Gamma(s; B_{\mathrm{D}}))$, which is at most $1-\beta$. □

This result describes empirical calibration-sample coverage of the deficit statistic used to construct the reserve. It does not guarantee runtime budget feasibility or coverage on unseen or shifted trajectories. For an unseen state, the controller uses the pooled calibration quantile; the same empirical statement applies to that pooled calibration sample.

The default policies use $\beta = 0.8$ with future horizon $K = 8$.

F Sample-level Compression Case Studies

We present one BIRD SQL request and one BFCL parallel-multiple request to illustrate how PASK assigns persistence by token role and layer group.

BIRD: percentage of carcinogenic molecules

Question (ID 227). What is the percentage of molecules that are carcinogenic? Please provide your answer as a percentage with three decimal places.

PASK output. `SELECT CAST(COUNT(CASE WHEN 'label' = '+' THEN 1 END) AS FLOAT) * 100.0 / COUNT(*) AS percent FROM 'molecule';`

Execution. On the BIRD database, this query returns 44.3149%; the reference query rounds the same result to 44.315%.

Compression trace. The policy mapping assigns 320 RETAIN-LOW and 84 RELEASE actions to the 404 prompt tokens in each layer group, with no RETAIN-HIGH actions. For the 37 non-stop SQL tokens, the reported decode trace gives the following split:

Stage	Layer	H	L	R
Prefill	Lower	0	320	84
	Middle	0	320	84
	Upper	0	320	84
Decode	Lower	37	0	0
	Middle	24	6	7
	Upper	3	34	0

The seven middle-group releases are pieces of the `'+'` predicate literal, the `1` case result and the `100.0` multiplier. Their lower-layer KV remains BF16 and upper-layer KV remains INT4, so each expression is retained in other groups.

BFCL: parallel-multiple tool calls (parallel_multiple_183)

Input. Could you first calculate the probability of drawing a heart from a deck of 52 cards where there are 13 hearts, and then calculate the probability of drawing a queen from the same deck where there are 4 queens? After that, could you retrieve the most recent artwork by the artist named 'Pablo Picasso' with a detailed description? Finally, could you locate the most popular sculpture exhibitions in New York, NY that are happening in the month of December and have high user ratings?

PASK output. Four calls match the reference functions and arguments:

1. `card_game_probability.calculate`: total_cards=52, desired_cards=13, cards_drawn=1.
2. `card_game_probability.calculate`: total_cards=52, desired_cards=4, cards_drawn=1.
3. `get_sculpture_info`:

`artist_name="Pablo Picasso", detail=true.`

4. `find_exhibition: location="New York, NY", art_form=sculpture, month="December", user_ratings="high".`

Prefill trace. The matcher labels the 595 prompt tokens as protected content (181), schema scaffold (130), function names (19), descriptions (167), types (36), parameter keys (53) and enums (9). The layer-group actions are:

Stage	Layer	H	L	R
Prefill	Lower	81	181	333
	Middle	81	348	166
	Upper	81	348	166
Decode	Lower	143	0	0
	Middle	0	84	59
	Upper	30	112	1

The 59 middle-group decode releases comprise 25 function-name pieces, 15 optional-field-key pieces, 14 structural separators and 5 optional-value pieces. The function names remain BF16 in the lower and upper groups; released keys and separators remain in the lower group and mostly at INT4 in the upper group. The four-call output exactly matches the reference on this one case.

Mixed-Mode Generation The cases above illustrate PASK within a single constrained response. Extending this control to agents that alternate between free-form language and structured calls could treat mode boundaries as structural events. The serving stack exposes when a grammar is active, allowing dispatch tokens after free-form text to receive lookahead-critical protection and free-form spans to use a dedicated mode tag with calibrated floors and upgrade priorities. The request-local budget ledger could continue across mode boundaries; in later turns, earlier free-form content becomes prompt text handled by the prefill companion (Section 3.6). Cross-mode calibration would measure how interventions on free-form history affect subsequent structured calls, using jointly intervened rollouts or paired trajectories. Evaluation covering both modes would establish the accuracy–cost trade-off of this extension.