arXiv:2604.01601v2 [cs.LG] 27 Sep 2026

Contrast Matters: Understanding Robustness of In-Context Fine-Tuning to Target-Context Relatedness

Deeptanshu Malu[1*] **Deevyanshu Malu**[1*] **Aditya Nemiwal**[1] **Sunita Sarawagi**[1]
[*]Equal Contribution [1]IIT Bombay
{deeptanshumalu, deevyanshumalu}@cse.iitb.ac.in

Abstract

In-context fine-tuning (IC-Train), training an LLM with labeled examples in-context, is increasingly used in place of standard fine-tuning for domain adaptation and continual absorption of labeled data. We study the robustness of the in-context learning ability that emerges from such training: does the fine-tuned model perform well across test inputs whose in-context examples range from unrelated to nearly identical? Across 32 configurations spanning four open-source LLMs and eight test sets over machine translation, Text-to-SQL, and multilingual semantic parsing, we show that robustness hinges on an overlooked design choice: how in-context examples are selected relative to the target during training. The two prevailing strategies turn out to be accurate over complementary parts of this spectrum: random contexts yield a model that gains little from related examples even when they are placed in its context, while retrieved similar contexts weaken accuracy on targets lacking close neighbors and raise the propensity to copy labels from context. Probes tracking in-weights learning, in-context learning, and copying trace these failures to distinct training dynamics, and show that introducing contrast in target-context similarity both within a context and across batches, restores robustness across the entire spectrum.

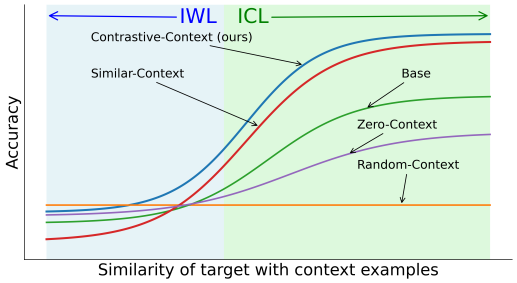

Figure 1: Schematic summary of main findings of the paper. On the X-axis is target to context similarity — IWL important on left side and ICL on right side. Fine-tuning with in-context example (IC-Train) is sensitive to similarity of target to context examples chosen during training: random context leads to sharp drop in ICL, similar context does not develop IWL and instead is prone to blind copying, contrast in the context recovers both limitations.

1 Introduction

Our goal is to fine-tune an LLM on a task with limited labeled data, while retaining the ability to continuously incorporate feedback in the form of new examples at test time, without further training. Such adaptation should boost accuracy on inputs resembling the new examples, while retaining generalization to inputs without close neighbors. In-context learning (ICL) offers a natural mechanism for flexibly absorbing labeled data by placing test inputs alongside labeled examples. Pre-trained LLMs (Brown et al., 2020) show strong ICL ability, yet prior studies (Alves et al., 2023; Duan et al., 2024; Wang et al., 2026; Goyal et al., 2025), as well as our experiments, find that fine-tuning often diminishes it. A natural alternative is to fine-tune in ICL mode (IC-Train) (Min et al., 2022; Alves et al., 2023; Bornschein et al., 2025; Zheng et al., 2024; Zebaze et al., 2025), where the model is trained to predict a target given a few labeled examples in its context.

IC-Train is thus positioned to meet both goals at once: a model adapted to the task, and one that still knows what to do with a related example when it appears in context. In this paper, we show that a critical but overlooked factor that determines how far IC-Train achieves these goals is how the in-context examples are selected relative to the target during training. Most prior work treated example selection as implementation details, and sampled them at random and independent of the target (Alves et al., 2023; Bornschein et al., 2025), including in multi-task meta-training settings (Min

et al., 2022), a strategy we call Random-Context. A more recent line retrieves the top-k most similar examples for each target (Zebaze et al., 2025; Zheng et al., 2024), which we call Similar-Context. Prior studies compared these methods on aggregate accuracy, and showed modest differences.

We study the robustness of a model trained with IC-Train under varying grades of target-context relatedness, to cover the continuous spectrum of useful context available to a random test example during deployment. Through several context distributions, across 32 (model, language-pair, test-set) configurations for machine translation, eleven Text-to-SQL databases, and three languages of multilingual semantic parsing, we show that the two prevailing strategies are each accurate over only part of the spectrum, and they cover complementary parts of it.

1. Random-Context produces a model whose accuracy is nearly flat in target-context similarity: it gains little from a highly related example placed in its context, and in the high-similarity region falls below standard zero-shot fine-tuning.
2. Similar-Context is competitive when close neighbors are available, but relatively weak when they are not, in most of our settings. Figure 1 summarizes these findings.

To understand why, we design probes that track over the course of training, how much a model relies on in-weights learning (IWL) versus genuine input-conditioned use of context, together with a third diagnostic for the propensity to copy context labels irrespective of their relevance to the target. The probes reveal that: under Random-Context, ICL capability decays steadily with training, so the model comes to rely almost entirely on its weights; under Similar-Context, IWL develops more weakly and the propensity to blind copying mildly rises. Table 3 (in Appendix C) presents an example of such blind copy.

Both diagnoses point to the importance of contrast in target-context similarity. A model never shown a related example in context has no reason to learn to use one. A model always trained with related examples learns to copy from context instead of acquiring the task in its weights, and when every context example is related, it has no reason to discriminate among them by relevance. We show that introducing contrast both across training batches and within a single context yields a model that is robust across the full similarity spectrum, retaining both IWL and ICL without resorting to copying. We refer to this recipe as Contrastive-Context. Our claim is not that any single ingredient is novel (mixtures of retrieval strategies are not new), but show that the similarity structure of IC-Train data determines how robust the resulting model is to varying degrees of target-context relatedness.

Contributions. 1. We highlight the importance of a fine-grained evaluation of models trained under IC-Train to diverse degree of relatedness of contexts to target examples. 2. We show that combining across-batch and within-context contrast in similarity yields robustness across the entire spectrum. 3. We evaluate IC-Train across grades of target-context relatedness on 32 MT configurations plus Text-to-SQL and multilingual semantic parsing, and show that the two prevailing strategies are robust over complementary parts of the spectrum. 4. We introduce probes for in-weights learning, in–context learning, and blind copying that trace these profiles to distinct training dynamics, and that can be applied to characterize any IC-trained model

2 Problem Setup and Context Selection Strategies

Let P_θ denote a pre-trained LLM to be finetuned on a task (e.g. low-resource translation). We are given a labeled dataset D of N pairs of inputs-outputs $\{(\mathrm{x}_1, \mathrm{y}_1), \ldots, (\mathrm{x}_N, \mathrm{y}_N)\}$ drawn from the underlying joint distribution $P(X, Y)$ of this task. Our goal is to train P_θ using D so that its performance remains robust under these deployment settings: (1) As more labeled pairs are appended to D (e.g., via user feedback), the model's accuracy should improve on test inputs with highly similar cases in D, *without requiring further parameter updates.* (2) For test inputs lacking similar examples in D, accuracy should be no worse than that of a model trained in the standard zero-shot setting. A natural candidate to meet these goals is IC-Train. In standard fine-tuning, we maximize likelihood on each example independently as $\mathbb{E}_{(\mathrm{x},\mathrm{y})\sim D} \log P_\theta(\mathrm{y}|\mathrm{x})$. IC-Train samples $k+1$ labeled examples from D, place the first k pairs as in-context examples while maximizing likelihood of target as:

$$\mathbb{E}_{(\mathrm{x}_*,\mathrm{y}_*)\sim D}\, \mathbb{E}_{C=\{(\mathrm{x}_i,\mathrm{y}_i):1:k\}\sim D} \log P_\theta(\mathrm{y}_*|C, \mathrm{x}_*) \quad (1)$$

The above training explicitly trains the model to leverage in-context examples, thereby better preparing it to absorb any labeled data during deployment. This paper shows that a crucial but often overlooked factor in IC-Train is the inter-example similarity among the $k+1$ samples. To make this explicit, we rewrite the IC-Train objective in terms of the strategy $\texttt{Choose}(D, \mathrm{x})$, for choosing the k examples from D to accompany a target input x_*:

$$\mathbb{E}_{(\mathrm{x}_*,\mathrm{y}_*)}\, \mathbb{E}_{C\sim\texttt{Choose}(D,\mathrm{x}_*)} \log P_\theta(\mathrm{y}_*|C, \mathrm{x}_*) \quad (2)$$

As we will show, $\texttt{Choose}(D, \mathrm{x}_*)$ critically influences whether IC-Train strengthens or erodes our desired robustness properties. Most prior work studied IC-Train with k context examples chosen at random independent of x_*. We refer this method as **Random-Context**. A second strategy is to select as C the top-k examples from D most similar to x_*, which we call **Similar-Context**. Section 3.1 shows that each of these is accurate over only part of the spectrum of target-context relatedness, and they cover complementary parts of it.

Contrastive-Context: Both regimes are missing contrast in target-context similarity, and we therefore construct training contexts that create contrast both among the examples within a context and across contexts. Contrast across batches comes from drawing a fraction $1-p$ of instances with random context, and fraction $\frac{p}{2}$ with similar context. The remaining $\frac{p}{2}$ creates contrast within a context: we sample one example weighed by similarity to the target, and the rest uniformly randomly. When the labeled pool is small, no example may be close enough to the target during fine-tuning, so the model could lose the capability of harnessing highly related context examples. We therefore augment ϵ fraction of the training instances (Random-Context and Similar-Context) with synthetic highly similar example by small perturbation of x_*. For NLP tasks, the perturbation is obtained by getting a paraphrase of x_*. A pseudocode of the overall method appears in Algorithm 1.

Algorithm 1 CONTRASTIVE-CONTEXT: sampling one IC-Train training instance

Require: Labeled pool D; target $(\mathrm{x}_*, \mathrm{y}_*) \sim D$; context size k; contrast fraction p; augmentation fraction ϵ; similarity $K(\cdot,\cdot)$
Ensure: Context C of k labeled pairs

$u \sim \mathrm{Unif}[0,1]$
if $u > p$ **then** ▷ prob. $1-p$
  $m \leftarrow$ RANDOM-CONTEXT
  $C \leftarrow k$ pairs drawn uniformly from D
else if $u > p/2$ **then** ▷ prob. $p/2$
  $m \leftarrow$ SIMILAR-CONTEXT
  $C \leftarrow \mathrm{TOPK}(D, \mathrm{x}_*, k)$ under K
else ▷ prob. $p/2$
  $m \leftarrow$ INTRA-CONTEXT CONTRAST
  $(\mathrm{x}_s, \mathrm{y}_s) \sim D$ with prob. $\propto K(\mathrm{x}_*, \mathrm{x}_s)$
  $C \leftarrow \{(\mathrm{x}_s, \mathrm{y}_s)\} \cup \{k-1$ pairs drawn uniformly$\}$
end if
if $m \in \{$RANDOM-CONTEXT, SIMILAR-CONTEXT$\}$ **and** $\mathrm{Unif}[0,1] < \epsilon$ **then** ▷ synthetic highly-similar example
  $\tilde{\mathrm{x}} \leftarrow$ PARAPHRASE(x_*)
  replace a random element of C with $(\tilde{\mathrm{x}}, \mathrm{y}_*)$
end if
shuffle C **and return** C

3 Evaluating IC-Train Across Target-Context Relatedness

We fine-tune four open source models on four machine translation tasks, eleven Text-to-SQL task, three multilingual semantic parsing tasks, and in Appendix B.4 also include two synthetic alignment reasoning tasks.

Evaluation Setup. To study the effect of different grades of relatedness of the input to the in-context examples, we evaluate under four different contexts: (1) Randomly selected k examples from D, (2) Select k examples most similar to x_* using BM25 (3) $k-1$ random examples and 1 example most similar to x_* using BM25, (4) $k-1$ random examples and one closely related example obtained to be a paraphrase of x_* with $\mathrm{y}_i = \mathrm{y}_*$. This mode checks the scenario of whether a model can harness closely related feedback from users after training. Over the union of all (context,test) instances, we plot accuracy as a function of the maximum similarity of the input x_* to an example in its context.

Models. We use four open-source LLMs: 1. Llama 3.2 1B base, 2. Llama 3.1 8B base (Grattafiori et al., 2024) 3. Qwen 2.5 7B (Yang et al., 2024). 4. Mistral 7B v0.3 (Jiang et al., 2023)

3.1 Machine Translation (MT) Tasks

We translate from English to four different languages: Lithuanian, Tamil, Hindi, and German. We train on standard benchmarks, and evaluate under two test-sets, for a total of 32 (model,task,test-set) configurations. For testing, we consider two types: a generic dataset from Flores (ID), and to

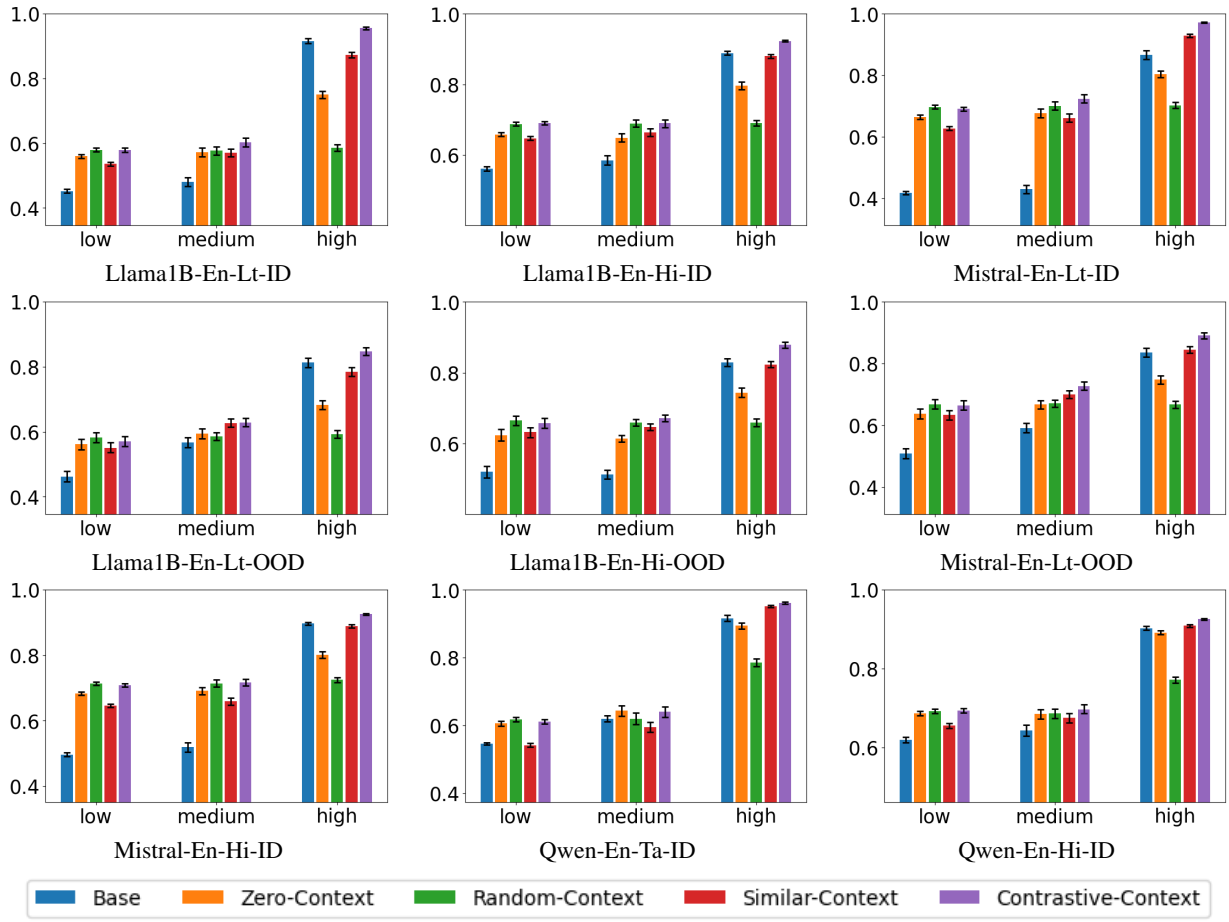

Figure 2: Effect of fine-tuning a base model with different strategies (Zero-Context, and IC-Train under Random-Context, Similar-Context, and Contrastive-Context) on accuracy over varying grades of similarity to in-context examples for nine of 32 different models, language-pairs, and test-sets (remaining plots in Appendix E). X-axis: Level of maximum similarity of target with in-context examples. The similarity ranges here are - Low: $0 - 0.33$, Medium: $0.33 - 0.67$, High: $0.67 - 1$. Y-axis: Accuracy (COMET score).

evaluate domain adaptation, an OOD dataset from a specialized domain such as judiciary and religion. Appendix B.2 provides more information. The prompt is given in Appendix A.1.

Methods Compared. We compare Contrastive-Context against the following baseline methods of fine-tuning LLMs for MT. (1) Zero-Context (Vieira et al., 2024): Standard zero-shot fine–tuning frequently used for domain-adapting MT models. (2) Random-Context: Alves et al. (2023) is one of the earliest works on harnessing IC–Train for MT, and proposes to select examples randomly. (3) Similar-Context: Recently, Zebaze et al. (2025) evaluated IC-Train with similarity based example selection where BM25 emerged as a competitive option, so we use this as a baseline. As reference we also show accuracy for the base untuned models. For all training setups (Random-Context, Similar-Context, Contrastive-Context) we choose k randomly from $[0 \ldots 5]$ and for testing k is set to 5. We use COMET-22 to measure accuracy.

In Figure 2 we plot accuracy against the maximum target-context similarity on combinations of four models, four tasks, two test settings, and five training methods: Zero-Context, IC-Train with Random-Context, Similar-Context, Contrastive-Context, and the original untuned base model. For clarity, we group similarity values into three bins of Low, Medium, and High similarity. The error bars show the 95% confidence interval with 1000 bootstrap resampling of the test data. Raw plots and over all 32 task-model combinations appear in Appendix E. Based on these plots, we make the following observations:

Random-Context worse than Zero-Context in harnessing related in-context examples. IC-Train with Random-Context (**green** bar) provides

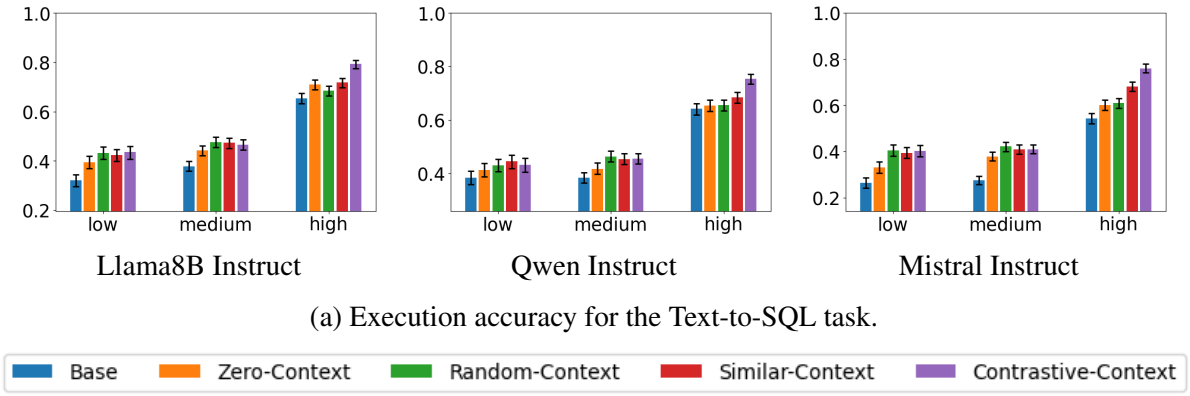

(a) Execution accuracy for the Text-to-SQL task.

Figure 3: Accuracy (Y-axis) against three levels of similarity (X-axis). The similarity ranges here are - (a) Low: $0-33$, Medium: $33-67$, High: $67-100$, (b) Low: $0-0.33$, Medium: $0.33-0.67$, High: $0.67-1$. Error bars show 95% confidence intervals.

almost a flat accuracy even as we increase the similarity of context to examples in almost all 32 model-task-test settings. Its performance is worse than even Zero-Context when presented with similar examples in context! Already Zero-Context (**orange** bar) fine-tuned is worse than the base untuned model (**blue** bar) for high target similarity, and IC-Train was meant to preserve ICL capability, but we observe that with random context, context usage actually deteriorates.

Similar-Context weak in the low similarity region. IC-Train with Similar-Context (**red** bar) suffers in the low similarity region, and is worse than all other forms of fine-tuning in that region. However, it responds to increasing target-context similarity with increasing accuracy.

Contrastive-Context robust on the full target-context similarity spectrum. Contrastive-Context(**violet** bar), is among the highest performing methods. Compared to Similar-Context, the second best performer, Contrastive-Context scores when highly related examples are present in context, and is competitive in low similarity ranges.

3.2 Text-to-SQL and Semantic Parsing

Text to SQL Task We experiment on Text-to-SQL as an instance of a Text-to-code generation task where the need for online adaptation to private databases is compelling. We use the BIRD dataset (Li et al., 2024b), with its official train split for fine-tuning and dev split spanning eleven distinct databases. More details in Appendix B.3. The results shown in Figure 3(a) further add evidence to the robustness of Contrastive-Context to handling contexts at varying similarity levels. Observe how IC-Train with Random-Context is worse than even Zero-Context in the high similarity range, and with Contrastive-Context we obtain the best accuracy across all levels. In this task Similar-Context does not suffer in the low-similarity range because the labeled pool is small, and Similar-Context almost reduces to random for many instances.

Multilingual Semantic Parsing task. We experiment with the cross-lingual semantic parsing task MTOP (Zhang et al., 2023b) on three languages — Spanish, German, and French. The results in Appendix Figure 6(b) show that Similar-Context suffers in the low and medium similarity regions compared to Random-Context and Contrastive-Context. In the high similarity region, Random-Context suffers compared to Similar-Context and Contrastive-Context. In this task too Contrastive-Context performs competitively or better across all similarity levels.

3.3 Probing In-Weights, In-Context, and Copying Behaviour

To understand the effect of target-context similarity on the emergence of different forms of learning, we design two probes for in-weights and in-context learning, together with a third diagnostic for the propensity to copy context labels, and analyze these as training progresses.

Probes to detect different forms of learning

Let $\hat{y}_C = \text{argmax}_y P_\theta(y|C, x_*)$ denote the predicted output under set $C = [x_1, y_1, \ldots x_k, y_k]$ and $\hat{y}_\phi$ denote zero-shot prediction. Let $\mathcal{K}(x, x')$ denote similarity between two inputs and $\text{sim}(y, y')$ between outputs.

In-Weights Learning Probe. We quan-

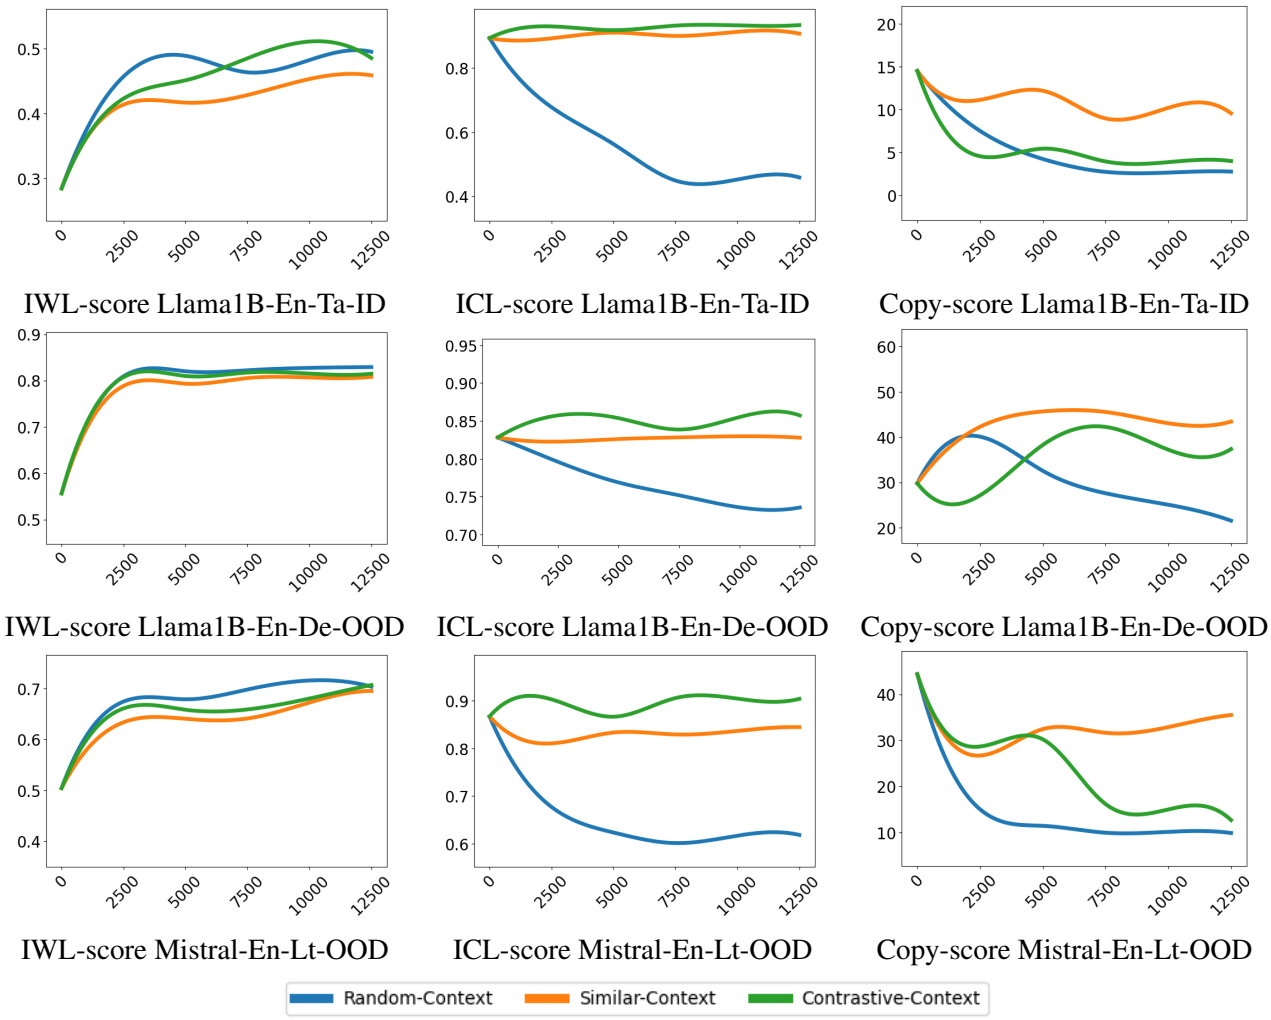

Figure 4: Emergence of different forms of learning in three training methods: Random-Context, Similar-Context, and Contrastive-Context. X-axis is training steps and Y-axis denotes scores of one of the three probes. Results of other model-task and datasets in Appendix Figure F. IWL-score of Similar-Context is lowest, ICL-score of Random-Context diminishes fast with training, Copy-score (a stress test under shuffled contexts, see text) of Similar-Context is higher. Contrastive-Context provides best retention of ICL and IWL capabilities without resorting to copying. Curves are cubic-spline interpolations through the evaluated checkpoints, shown for readability.

tify the in-weights learning as the similarity between predictions under random and empty context. IWL-score(P_θ) $=$ $\mathbb{E}_{\mathrm{x}_*}\, \mathbb{E}_{C\sim\text{Random}(D,k)}\, \text{sim}(\hat{\mathrm{y}}_C, \hat{\mathrm{y}}_\phi)$

In-Context Learning Probe. We quantify ICL capability by comparing prediction $\hat{\mathrm{y}}_C$ under a one-similar context C to the labels in context weighted by their similarity to x_*. One-similar context is obtained by injecting a paraphrase for x_* with gold y_* among remaining $k-1$ random examples, so $C = [\mathrm{x}_1, \mathrm{y}_1, \ldots, \mathrm{x}_{i^*} = \text{paraphrase}(\mathrm{x}_*), \mathrm{y}_{i^*} = \mathrm{y}_*, \ldots \mathrm{x}_k, \mathrm{y}_k]$. ICL-score$(P_\theta)$ $=$

$$\mathbb{E}_{\mathrm{x}_*}\, \mathbb{E}_{C\sim\text{One-similar}(D,\mathrm{x}_*,k)} \sum_{j=1}^{k} \frac{\mathcal{K}(\mathrm{x}_*, \mathrm{x}_j)\text{sim}(\hat{\mathrm{y}}_C, \mathrm{y}_j)}{\sum_i \mathcal{K}(\mathrm{x}_*, \mathrm{x}_i)}$$

Blind Copy Probe. A model with propensity to copy blindly from the context would output a $\hat{\mathrm{y}}$ similar to a label y_i in-context irrespective of x_*'s similarity to x_i. We quantify this by shuffling the one-similar context (above) so that no x_i and y_i are correctly matched. An example for $k = 3$ is: $C = [\mathrm{x}_1, \mathrm{y}_*, \mathrm{x}_2 = \text{paraphrase}(\mathrm{x}_*), \mathrm{y}_3, \mathrm{x}_3, \mathrm{y}_1]$. Pure ICL would output y_3 based on similarity of x_2 with x_* whereas a model that copies independent of x could output one of the other labels. Thus, we define copy score as maximum similarity to a label other than the one attached to x_{i^*} where i^* is the position of x_*'s paraphrase.

Copy-score(P_θ) $=$

$$\mathbb{E}_{\mathrm{x}_*}\, \mathbb{E}_{C\sim\text{Shuffle}(\text{One-similar}(D,\mathrm{x}_*,k))} \max_{i\in C, i\neq i^*} \text{sim}(\hat{\mathrm{y}}_C, \mathrm{y}_i)$$

We used cosine similarity of sentence embeddings from all-MiniLM-L6-v2 (Wang et al., 2020) for $\mathcal{K}$

and COMET for sim for the first two probes but BLEU for the copy score to measure the model's propensity to blindly copy the lexical tokens. The shuffled context is deliberately off-distribution: no model encounters it during training or deployment. We therefore treat this probe as a stress test that isolates whether a model's use of context is conditioned on input similarity, rather than as a measurement of deployment behaviour.

Main observations. In Figure 4 we show the emergence of different forms of learning as measured by these three probes for various model, language pair, and dataset combinations. More combinations appear in the Appendix F. We can make the following observations from these graphs: (1) The IWL-score (first column) of Random-Context, and Contrastive-Context steadily increases with training steps with Contrastive-Context lagging only slightly behind. However, for Similar–Context the IWL-score is distinctly lower, sometimes by a significant margin. This, coupled with the observation that Similar-Context provides lower accuracy in Figure 2 when target-context similarity is low, shows that Similar-Context is less effective to learn the task in-weights, and/or is distracted by irrelevant examples. (2) The ICL-score (second column) shows that for Random-Context ICL capability steadily drops with training, explaining why it provides almost the same accuracy across different grades of target-context similarity. For Contrastive-Context and Similar-Context, the ICL capability stays almost the same or increases. (3) The copy score (last column) decreases with training using Random-Context, and Contrastive-Context and gets substituted by either in-weights or in-context learning. In contrast, for Similar-Context the copy score is higher than for the other models. On real datasets, Similar-Context training may not consistently show increased copy tendencies because not every target will find all similar top-k examples, and the training data may naturally contain a mix of random and similar contexts. But in spite of the natural mixing, we observe Similar-Context to result in reduced IWL and more blind copying compared to Contrastive–Context. In Table 3 we present an example of blind copying even under normal deployment context.

3.4 Ablations

We present brief ablations on our method here. Appendix D has more details. **Sensitivity to ϵ, p values.** Contrastive-Context creates contrast using parameter p to control the split of random Vs similar contexts, and control the fraction of highly similar synthetic examples (ϵ). In Figure 5(a) we show the method is robust to alternative values ϵ, p between 0.1 and 0.5. **Quality of paraphrasing.** To study impact paraphrase quality, we evaluated with two alternative models `gemini-2.5-flash-lite` a strong model, and `Qwen2.5-7B-Instruct` as an open source, possibly weaker model compared to `gemini-2.0-flash-lite`. Figure 5(b) shows that Contrastive-Context remains stable across these three paraphrasing models. **Role of different types of contrasts.** We evaluate the role of different types of contrasts that Contrastive-Context creates in Figure 5(c). We set $\epsilon = 0$ to suppress paraphrase-augmentation and observe that test instances with high target- context similarity suffers. We set $p = 0$ to suppress the generation of similar real examples, which causes medium-scale accuracy to drop.

4 Related Work

Fine-tuning with in-context examples. Zero-shot fine-tuning is known to erode the ICL ability of pre-trained LLMs (Brown et al., 2020), an effect documented across instruction tuning (Duan et al., 2024; Goyal et al., 2025; Wang et al., 2026) and domain-adapted machine translation (Alves et al., 2023). Fine-tuning in ICL mode (IC-Train) is the standard remedy: Alves et al. (2023) show for MT that it preserves ICL ability that zero-shot fine-tuning loses, Zheng et al. (2024) report similar gains on English–German and English–Chinese, Bornschein et al. (2025) study IC-Train as an efficient adaptation mechanism, and Min et al. (2022) apply it in multi-task meta-training. All of this work selects in-context examples *at random*, and we showed that this preserves ICL far less reliably, and in the high-similarity regime can be *worse* than zero-shot fine-tuning.

Similarity-based example selection. More recent methods select in-context examples by similarity to the input. At inference time this is well established to improve ICL (Levy et al., 2023; Kothyari et al., 2025; Zebaze et al., 2025); Zheng et al. (2024) carry such selection into training. Gao and Das (2024) also use contrastive examples, but define contrast by the *correctness* of the context rather than relatedness to the target, and at prompting rather than fine-tuning time.

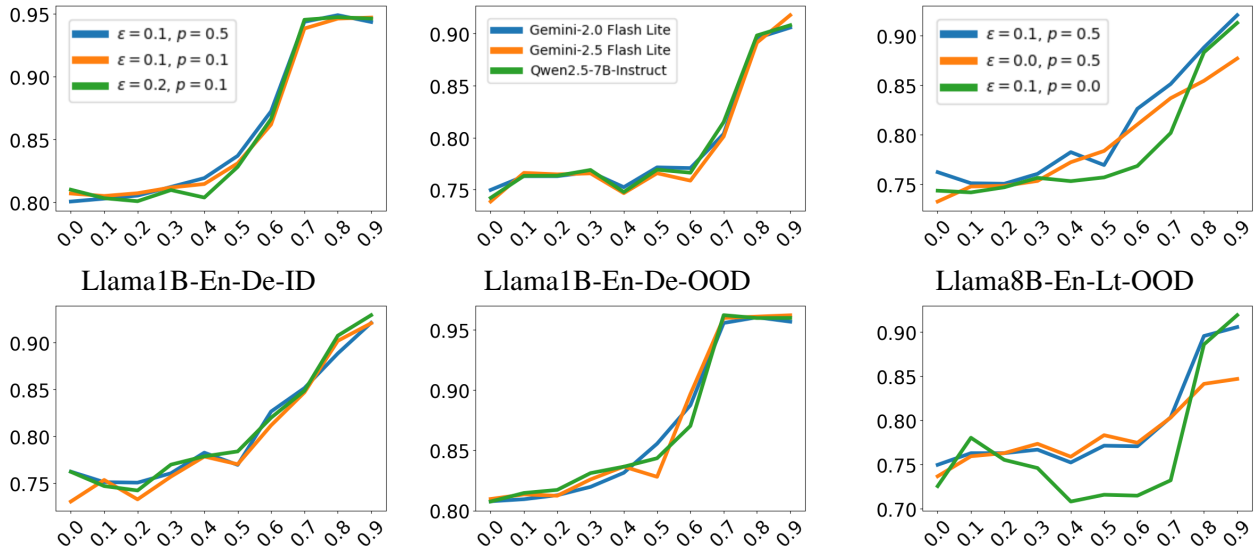

(a) Sensitivity to ϵ, p values (b) Quality of paraphrasing (c) Role of different types of contrasts

Figure 5: Ablations to show (a) the robustness of Contrastive-Context to variant non-zero ϵ, p values, (b) various paraphrasing models, (c) importance of various levels of similarity in the training data. X-axis is target-context similarity and Y-axis is accuracy.

Effect of context structure on model learning. Closest in spirit, Chan et al. (2022) show that training-data distribution governs whether ICL or in-weights learning (IWL) emerges: bursty distributions, where related examples recur in context, elicit ICL, non-bursty ones elicit IWL, and Zipfian data yields both. Singh et al. (2023) observe that ICL can be transient, decaying into IWL with continued training, and in-context repetition has been shown to accelerate ICL (Zucchet et al., 2025; Bratuli et al., 2025). We recast this in terms of graded target–context *similarity*, which generalizes burstiness and repetition to real NLP tasks where relatedness is continuous rather than binary. Our distinctive finding that contrast *within* a context matters does not follow from burstiness alone, and is what separates in-context learning from indiscriminate copying of context labels.

ICL–IWL tradeoffs under task mixtures. A substantial literature studies ICL emergence when training over a mixture of T tasks from a family, where the k in-context examples share the target's task f_τ (Reddy, 2024; Singh et al., 2024; Akyürek et al., 2024; Edelman et al., 2024; Park et al., 2025; Nguyen and Reddy, 2025; Wurgaft et al., 2025; Ku et al., 2025; Kim et al., 2025; Singh et al., 2025; Fu et al., 2024). These studies use synthetic regression, classification, or sequence-completion tasks (Garg et al., 2022; Zhang et al., 2023a; Olsson et al., 2022; Rajaraman et al., 2024; Shi et al., 2024; Agarwal and Sarawagi, 2025) with randomly drawn contexts. We instead fine-tune on a *single*, real sequence-to-sequence task, where the axis of variation is not which task an example belongs to but how related it is to the target.

5 Conclusion

We studied how the similarity between a target and its in-context examples during IC-Train shapes the model that results. Evaluated across the full spectrum of target-context relatedness, the two prevailing strategies prove accurate over complementary parts of that spectrum: random contexts yield a model that gains little from a related example even when one is placed in its context, while similarity-retrieved contexts weaken accuracy on targets without close neighbors. Probes for in-weights learning, in-context learning, and copying trace these profiles to distinct training dynamics. Both point to the same missing ingredient, and we showed that introducing contrast in similarity both across training batches and within a single context, yields models that remain accurate across the whole spectrum on machine translation, Text-to-SQL, and semantic parsing. For practitioners deploying a model alongside a growing labeled pool, the similarity structure of IC-Train data is a design choice worth making deliberately.

Limitations

Dependence on an external paraphraser. Contrastive-Context assumes access to a model that can produce meaning-preserving near-duplicates of a target input. For MT we paraphrase the English source, i.e. the high-resource side of the pair; whether comparable paraphrase quality is attainable when the input side is itself low-resource is untested. We ablate the choice of paraphraser (Figure 5), but not the regime in which paraphrases are semantically drifted or degenerate.

Task and language coverage. All our tasks are fixed, single-task sequence-to-sequence problems with relatively short inputs: translation into four languages (En→X only), Text-to-SQL, semantic parsing, and synthetic alignment. We do not evaluate classification, long-form generation, multi-turn dialogue, or reasoning-heavy tasks, where the notion of a "similar" example—and hence of a useful context—is less well defined. We also restrict contexts to $k \in [0, 5]$; the behaviour of the three context regimes under many-shot contexts, where retrieved examples naturally span a wider similarity spectrum, remains open.

Hyperparameters and compute overhead. We do not tune p and ϵ per task; we only show robustness over the range 0.1–0.5 on a subset of settings, and we give no principled procedure for choosing them as a function of pool size or task. We report wall-clock figures for these steps but do not present a compute-matched comparison against the baselines.

References

Harshvardhan Agarwal and Sunita Sarawagi. 2025. The missing alignment link of in-context learning on sequences.

Ekin Akyürek, Bailin Wang, Yoon Kim, and Jacob Andreas. 2024. In-context language learning: Architectures and algorithms. In *Forty-first International Conference on Machine Learning*.

Duarte Alves, Nuno Guerreiro, João Alves, José Pombal, Ricardo Rei, José de Souza, Pierre Colombo, and Andre Martins. 2023. Steering large language models for machine translation with finetuning and in-context learning. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 11127–11148, Singapore. Association for Computational Linguistics.

Jorg Bornschein, Clare Lyle, Yazhe Li, Amal Rannen-Triki, Xu Owen He, and Razvan Pascanu. 2025. Fine-tuned in-context learners for efficient adaptation. *Preprint*, arXiv:2512.19879.

Jelena Bratuli, Sudhanshu Mittal, David T. Hoffmann, Samuel Böhm, Robin Tibor Schirrmeister, Tonio Ball, Christian Rupprecht, and Thomas Brox. 2025. Unlocking in-context learning for natural datasets beyond language modelling. *Preprint*, arXiv:2501.06256.

Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc.

Stephanie C.Y. Chan, Adam Santoro, Andrew K. Lampinen, Jane X. Wang, Aaditya K. Singh, Pierre H. Richemond, James L. McClelland, and Felix Hill. 2022. Data distributional properties drive emergent in-context learning in transformers. In *Proceedings of the 36th International Conference on Neural Information Processing Systems*, NIPS '22, Red Hook, NY, USA. Curran Associates Inc.

Hanyu Duan, Yixuan Tang, Yi Yang, Ahmed Abbasi, and Kar Yan Tam. 2024. Exploring the relationship between in-context learning and instruction tuning. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 3197–3210, Miami, Florida, USA. Association for Computational Linguistics.

Ezra Edelman, Nikolaos Tsilivis, Benjamin L. Edelman, eran malach, and Surbhi Goel. 2024. The evolution of statistical induction heads: In-context learning markov chains. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.

Jingwen Fu, Tao Yang, Yuwang Wang, Yan Lu, and Nanning Zheng. 2024. Breaking through the learning plateaus of in-context learning in transformer. In *Forty-first International Conference on Machine Learning*.

Xiang Gao and Kamalika Das. 2024. Customizing language model responses with contrastive in-context learning. *Preprint*, arXiv:2401.17390.

Shivam Garg, Dimitris Tsipras, Percy Liang, and Gregory Valiant. 2022. What can transformers learn in-context? a case study of simple function classes. In *Advances in Neural Information Processing Systems*.

Sachin Goyal, Christina Baek, J Zico Kolter, and Aditi Raghunathan. 2025. Context-parametric inversion:

Why instruction finetuning may not actually improve context reliance. In *The Thirteenth International Conference on Learning Representations*.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, and et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Lélio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. 2023. Mistral 7b. *Preprint*, arXiv:2310.06825.

Jaeyeon Kim, Sehyun Kwon, Joo Young Choi, Jongho Park, Jaewoong Cho, Jason D. Lee, and Ernest K. Ryu. 2025. Task diversity shortens the icl plateau. *Preprint*, arXiv:2410.05448.

Mayank Kothyari, Sunita Sarawagi, Soumen Chakrabarti, Gaurav Arora, and Srujana Merugu. 2025. Diverse in-context example selection after decomposing programs and aligned utterances improves semantic parsing. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*.

Alexander Y. Ku, Thomas L. Griffiths, and Stephanie C. Y. Chan. 2025. Predictability shapes adaptation: An evolutionary perspective on modes of learning in transformers. *Preprint*, arXiv:2505.09855.

Anoop Kunchukuttan, Pratik Mehta, and Pushpak Bhattacharyya. 2018. The IIT Bombay English-Hindi parallel corpus. In *Proceedings of the Eleventh International Conference on Language Resources and Evaluation (LREC 2018)*, Miyazaki, Japan. European Language Resources Association (ELRA).

Itay Levy, Ben Bogin, and Jonathan Berant. 2023. Diverse demonstrations improve in-context compositional generalization. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2023, Toronto, Canada, July 9-14, 2023*, pages 1401–1422. Association for Computational Linguistics.

Haoyang Li, Jing Zhang, Cuiping Li, and Hong Chen. 2023. Resdsql: Decoupling schema linking and skeleton parsing for text-to-sql. In *AAAI*.

Haoyang Li, Jing Zhang, Hanbing Liu, Ju Fan, Xiaokang Zhang, Jun Zhu, Renjie Wei, Hongyan Pan, Cuiping Li, and Hong Chen. 2024a. Codes: Towards building open-source language models for text-to-sql. In *SIGMOD*.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. 2024b. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. *Advances in Neural Information Processing Systems*, 36.

Sewon Min, Mike Lewis, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2022. MetaICL: Learning to learn in context. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2791–2809, Seattle, United States. Association for Computational Linguistics.

Alex Nguyen and Gautam Reddy. 2025. Differential learning kinetics govern the transition from memorization to generalization during in-context learning. In *The Thirteenth International Conference on Learning Representations*.

NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, Anna Sun, Skyler Wang, Guillaume Wenzek, Al Youngblood, Bapi Akula, Loic Barrault, Gabriel Mejia Gonzalez, Prangthip Hansanti, and 20 others. 2024. Scaling neural machine translation to 200 languages. *Nature*, 630(8018):841–846.

Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova Dassarma, Tom Henighan, Benjamin Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Scott Johnston, Andy Jones, John Kernion, Liane Lovitt, and 7 others. 2022. In-context learning and induction heads. *ArXiv*, abs/2209.11895.

Core Francisco Park, Ekdeep Singh Lubana, and Hidenori Tanaka. 2025. Competition dynamics shape algorithmic phases of in-context learning. In *The Thirteenth International Conference on Learning Representations*.

Nived Rajaraman, Marco Bondaschi, Ashok Vardhan Makkuva, Kannan Ramchandran, and Michael Gastpar. 2024. Transformers on markov data: Constant depth suffices. In *ICML 2024 Workshop on Mechanistic Interpretability*.

Gowtham Ramesh, Sumanth Doddapaneni, Aravinth Bheemaraj, Mayank Jobanputra, Raghavan AK, Ajitesh Sharma, Sujit Sahoo, Harshita Diddee, Mahalakshmi J, Divyanshu Kakwani, Navneet Kumar, Aswin Pradeep, Srihari Nagaraj, Kumar Deepak, Vivek Raghavan, Anoop Kunchukuttan, Pratyush Kumar, and Mitesh Shantadevi Khapra. 2022. Samanantar: The Largest Publicly Available Parallel Corpora Collection for 11 Indic Languages. *Transactions of the Association for Computational Linguistics*, 10:145–162.

Gautam Reddy. 2024. The mechanistic basis of data dependence and abrupt learning in an in-context classification task. In *The Twelfth International Conference on Learning Representations*.

Zhenmei Shi, Junyi Wei, Zhuoyan Xu, and Yingyu Liang. 2024. Why larger language models do in-context learning differently? In *Forty-first International Conference on Machine Learning*.

Aaditya Singh, Stephanie Chan, Ted Moskovitz, Erin Grant, Andrew Saxe, and Felix Hill. 2023. The transient nature of emergent in-context learning in transformers. *Advances in neural information processing systems*, 36:27801–27819.

Aaditya K. Singh, Ted Moskovitz, Sara Dragutinovic, Felix Hill, Stephanie C. Y. Chan, and Andrew M. Saxe. 2025. Strategy coopetition explains the emergence and transience of in-context learning. *Preprint*, arXiv:2503.05631.

Aaditya K Singh, Ted Moskovitz, Felix Hill, Stephanie C.Y. Chan, and Andrew M Saxe. 2024. What needs to go right for an induction head? a mechanistic study of in-context learning circuits and their formation. In *Forty-first International Conference on Machine Learning*.

Jörg Tiedemann. 2012. Parallel data, tools and interfaces in opus. In *Proceedings of the Eight International Conference on Language Resources and Evaluation (LREC'12)*, Istanbul, Turkey. European Language Resources Association (ELRA).

Inacio Vieira, Will Allred, Séamus Lankford, Sheila Castilho, and Andy Way. 2024. How much data is enough data? fine-tuning large language models for in-house translation: Performance evaluation across multiple dataset sizes. In *Proceedings of the 16th Conference of the Association for Machine Translation in the Americas (Volume 1: Research Track)*, pages 236–249, Chicago, USA. Association for Machine Translation in the Americas.

Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. 2020. Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. *Preprint*, arXiv:2002.10957.

Yihan Wang, Andrew Bai, Nanyun Peng, and Cho-Jui Hsieh. 2026. On the loss of context awareness in general instruction fine-tuning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Daniel Wurgaft, Ekdeep Singh Lubana, Core Francisco Park, Hidenori Tanaka, Gautam Reddy, and Noah D. Goodman. 2025. In-context learning strategies emerge rationally. *Preprint*, arXiv:2506.17859.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, and 22 others. 2024. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Armel Randy Zebaze, Benoît Sagot, and Rachel Bawden. 2025. In-context example selection via similarity search improves low-resource machine translation. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 1222–1252, Albuquerque, New Mexico. Association for Computational Linguistics.

Ruiqi Zhang, Spencer Frei, and Peter L. Bartlett. 2023a. Trained transformers learn linear models in-context. *ArXiv*, abs/2306.09927.

Yusen Zhang, Jun Wang, Zhiguo Wang, and Rui Zhang. 2023b. Xsemplr: Cross-lingual semantic parsing in multiple natural languages and meaning representations. In *ACL*.

Jiawei Zheng, Hanghai Hong, Xiaoli Wang, Jing Su, Yong Liang, and Shikai Wu. 2024. Fine-tuning large language models for domain-specific machine translation. *ArXiv*, abs/2402.15061.

Nicolas Zucchet, Francesco D'Angelo, Andrew Kyle Lampinen, and Stephanie C.Y. Chan. 2025. The emergence of sparse attention: impact of data distribution and benefits of repetition. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Appendix Table of Contents

A Prompts for various tasks **12**
- A.1 Prompts for obtaining translations with in-context examples . . . 12
- A.2 Prompt to generate paraphrases . . . 13

B Datasets and Additional Tasks **14**
- B.1 Additional Setup Details . . . 14
- B.2 Low Resource Translation . . . 15
- B.3 Text-to-SQL . . . 15
- B.4 Synthetic Alignment Reasoning Task . . . 15

C Anecdotes showing Contrastive-Context outperform Similar-Context **17**

D Ablations **17**
- D.1 Experiments with non-zero ϵ and p . . . 17
- D.2 Importance of training with different similarity levels . . . 17
- D.3 Intra-Context Vs Inter-Context Contrast . . . 19
- D.4 Synthetic Paraphrases for Data Augmentation Vs Contrast . . . 20
- D.5 Experiments with various paraphrasing models . . . 20
- D.6 Effect of LoRA Target Modules . . . 20

E Additional Experiments Comparing Different Training Strategies **21**

F Additional Experiments on Learning Dynamics **24**
- F.1 Learning Dynamics (Unsmoothed) . . . 27

A Prompts for various tasks

A.1 Prompts for obtaining translations with in-context examples

- ```
  Translate the source text from English to German.
  Source: Redistribution is only possible if there is actually something produced
  ↪ to be redistributed.
  Target:
  ```
- ```
  Consider the following 4 translations from English to German.
  Example 1:
  Source: This is now to be given concrete shape in the proposal to give EU
  ↪ citizens the right of free movement and residence.
  Target: Jetzt soll sie in dem Vorschlag über das Recht der Unionsbürger auf
  ↪ Freizügigkeit und Aufenthalt konkret ausgestaltet werden.

  Example 2:
  Source: These data can only be seen as striking, alerting us to the urgent need
  ↪ for humanitarian aid to these countries to be properly directed to the area
  ↪ of health and the provision of basic medical care to highly deprived
  ↪ communities.
  Target: Diese Zahlen sind erschreckend und müssen für uns ein Alarmsignal sein,
  ↪ dass es dringend notwendig ist, die humanitäre Hilfe für diese Länder in
  ↪ angemessener Weise in den Bereich des Gesundheitswesens und der
  ↪ Bereitstellung von medizinischer Grundversorgung für die am meisten
  ↪ benachteiligten Gemeinschaften zu lenken.

  Example 3:
  ```

```
    Source: There are a number of actions we need to take in future in the area of
    ↪  budget control within the agencies.
    Target: Es gibt eine Reihe von Manahmen, die wir in Zukunft im Bereich der
    ↪  Haushaltskontrolle in den Agenturen ergreifen müssen.

    Example 4:
    Source: It is precisely in the area of the environment that people's
    ↪  dissatisfaction has been most deeply felt.
    Target: Gerade im Umweltbereich spürt man die Unzufriedenheit der Bevölkerung
    ↪  am deutlichsten.

    Translate the source text from English to German.
    Source: It is precisely in the area of budget policy that these three basic
    ↪  principles need to be given concrete expression.
    Target:
```

A.2 Prompt to generate paraphrases

```
Generate 5 diverse paraphrases for the given input English sentence.
Before paraphrasing, carefully analyze the original sentences and keywords.
Ensure significant variation in phrasing, structure, and wording while preserving
↪  meaning.
The paraphrases should be such that their translation into another language could
↪  still be the same.
Only modify the grammar and sentence structure such that all the paraphrases should
↪  be translatable to a common sentence in another language.

### Inputs:
Original Sentences:
{sentences}

### Example:
1.
#### Input:
Original Sentences:
[
    "Sentence 1": "In New York City, the streets were alive with the sound of
    ↪  honking taxis and chatter from the crowded sidewalks.",
    "Sentence 2": "\"I can't believe it's already time for the annual reunion,\"
    ↪  Emily said, her voice filled with excitement.",
    "Sentence 3": "The Amazon rainforest is home to countless species of plants and
    ↪  animals, many of which have yet to be discovered.",
]

#### Example Paraphrases:
{{
    "Sentence 1": [
        "The streets of New York City buzzed with the noise of honking taxis and
        ↪  voices from the bustling sidewalks.",
        "In the heart of New York City, honking taxis and lively chatter filled the
        ↪  streets.",
        "New York City's streets were vibrant with the sounds of honking taxis and
        ↪  the chatter of crowded sidewalks.",
```

```
        "In the bustling streets of New York City, honking taxis and lively
        ↪  conversations created a lively atmosphere.",
        "The sound of honking taxis and the chatter from crowded sidewalks made the
        ↪  streets of New York City come alive."
    ],
    "Sentence 2": [
        "Emily exclaimed, her excitement evident, \"I can't believe it's time for
        ↪  the annual reunion already.\"",
        "With excitement in her voice, Emily said, \"I can't believe the annual
        ↪  reunion is here already.\"",
        "Emily's voice was filled with excitement as she said, \"I can't believe
        ↪  it's already time for the annual reunion.\"",
        "With a voice full of excitement, Emily remarked, \"I can't believe the
        ↪  annual reunion is already upon us.\"",
        "Filled with excitement, Emily exclaimed, \"I can't believe it's already
        ↪  time for the annual reunion!\""
    ],
    "Sentence 3": [
        "The Amazon rainforest shelters an untold number of plant and animal
        ↪  species, many still waiting to be discovered.",
        "Countless species of plants and animals call the Amazon rainforest home,
        ↪  many of which remain undiscovered.",
        "Home to countless species of plants and animals, the Amazon rainforest
        ↪  holds many that are yet to be discovered.",
        "The Amazon rainforest is a habitat for innumerable species of plants and
        ↪  animals, many of which are still unknown.",
        "Many species of plants and animals, yet to be discovered, thrive in the
        ↪  Amazon rainforest."
    ],
}}

### Output Instructions:
- Generate exactly 5 paraphrases.
- Ensure correct JSON output syntax like the example
- Always escape quotes within a paraphrase (like in example 2)
```

B Datasets and Additional Tasks

B.1 Additional Setup Details

For training, we adopted LoRA with a rank of 16, a scaling factor of $\alpha = 32$, and a dropout rate of 0.05. ϵ is set to 0.1 and p is set to 0.5. Only the attention projection matrices (except the output projection matrix) were trained. We also trained all the attention projection matrices along the MLP layer matrices as an ablation. The results for this ablation can be found in Appendix D.6. Training was performed with a batch size of 2 using Adam with a learning rate of 2×10^{-4} and a linear decay learning rate scheduler. To obtain paraphrases for Contrastive-Context, we employed the `gemini-2.0-flash-lite` model.

For the experiments in the main text, the training data size is 50000, and for Contrastive-Context, ϵ is set to 0.1. So, only 5000 paraphrases have to be generated. These are prompted to the paraphrasing model in groups of 10; hence, 500 API calls are required. This takes around 1.2 hours. This can be sped up greatly with multithreading.

We use 50000 training instances and 5000 validation instances for all models. Generating these instances by retrieving similar examples from a BM25 index takes approximately 25 minutes using 32 parallel processes.

B.2 Low Resource Translation

	Train		ID Test		OOD Test	
LP	Source	Size	Source	Size	Source	Size
En-De	Europarl	50000	Flores+	1012	EMEA	500
En-Hi	Samanantar	50000	Flores+	1012	Judicial	500
En-Ta	Samanantar	50000	Flores+	1012	Tanzil	500
En-Lt	Europarl	50000	Flores+	1012	EMEA	500

Table 1: Datasets Used for MT. Flores+ refers to Flores-200. **Sources:** Europarl, Tanzil, and EMEA from Tiedemann (2012); Samanantar from Ramesh et al. (2022); Flores+ from NLLB Team et al. (2024); Judicial from Kunchukuttan et al. (2018).

Dataset	Link
Europarl	Helsinki-NLP/europarl
Tanzil	Helsinki-NLP/tanzil
EMEA	Helsinki-NLP/emea
Samanantar	ai4bharat/samanantar
Flores+	openlanguagedata/flores_plus
Judicial	cfilt/iitb-english-hindi

Table 2: Links to datasets used

B.3 Text-to-SQL

The paraphrases for this task are obtained using OpenAI-O3-mini. Given an NL query, SQL pair (x, y), we prompted the LLM to generate first a paraphrased SQL $\tilde{\mathrm{y}}$ that differs from y only in the mention of literal constants, and then generating the corresponding NL question $\tilde{\mathrm{x}}$ describing the SQL $\tilde{\mathrm{y}}$. An example paraphrased pair appears below:

x: Is molecule TR183 known to be a carcinogen?
y: SELECT T.label FROM molecule AS T WHERE T.molecule_id = 'TR183
.
$\tilde{\mathrm{x}}$: Molecule KAC16 is carcinogenic. Yes or No?
$\tilde{\mathrm{y}}$: SELECT T.label FROM molecule AS T WHERE T.molecule_id = 'KAC16

The prompt to the LLM comprises of generic natural language instructions, followed by a description of the schema and metadata of the database queried, followed by the in-context examples, and then the current test question x_*. Following standard practice, instead of providing the entire database metadata, we filter a subset of schema information (relevant table and column names) using an open source schema filtering tool [1] (Li et al., 2023), (Li et al., 2024a).

B.4 Synthetic Alignment Reasoning Task

We use the dataset as discussed by (Agarwal and Sarawagi, 2025) for reasoning about the alignment between two synthetic sequences. This task first defines a vocabulary $\mathcal{V}$ of symbols, each symbol $\sigma \in \mathcal{V}$ is associated with a probabilistic finite state automata (PFA) to generate sequences of length c over elements over its own vocabulary. The input sequence x contains m symbols chosen from $\mathcal{V}$, the output sequence y is obtained by a fixed (unknown) permutation of x, and then sampling a length c sequence from the FSA for each element in the permuted x. More details can be found in (Agarwal and Sarawagi, 2025). In order to correctly generate the output for a sequence, the model needs to reason about the input-output alignment and then the PFA of each token in $\mathcal{V}$. Examples appear below:

For this task we choose $|\mathcal{V}| = 25$ and ran under two settings of sequence lengths ($m = 4,\ c = 4$) and ($m = 6,\ c = 6$) with the alignment as a permutation. The training is done on 200 instances for 5

[1] https://github.com/RUCKBReasoning/text2sql-schema-filter

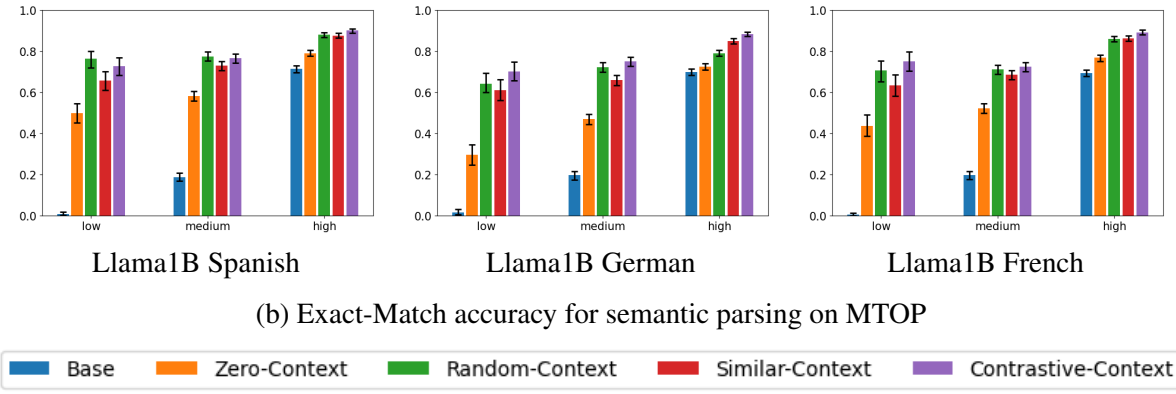

(b) Exact-Match accuracy for semantic parsing on MTOP

Figure 6: Accuracy (Y-axis) against three levels of similarity (X-axis). The similarity ranges here are - (a) Low: $0-33$, Medium: $33-67$, High: $67-100$, (b) Low: $0-0.33$, Medium: $0.33-0.67$, High: $0.67-1$. Error bars show 95% confidence intervals.

Prompt	Prompt string: $x^1 : y^1 \;\; x^2 : y^2 \;\; x^3 : y^3 \;\; x_*$
Example $m = 3, c = 2$	`ACB: rijjpr CAB: jjriwp ABC: rtprjh BCA:`
Example $m = 2, c = 4$	`AC: ririjhjh CA: jjhhriir AB: rttrprrp BA:`

epochs. In the original paper, each instance sampled its own PFA to simulate an infinite mixture. In this paper, since our goal is to develop ICL-IWL mixture for a fixed task, the PFAs per symbol were fixed throughout the entire process. We sampled in-context examples at different similarity levels as follows: the Random-Context data is sampled using a vocabulary of size 25, and the Similar-Context data is created by using a vocabulary of size 12. For the Contrastive-Context data, the paraphrases are generated by taking the vocabulary as the set of letters used in the target example, x_*. Hence, the paraphrase is simply a shuffling of the letters in x_*. The results are presented in Figure 7.

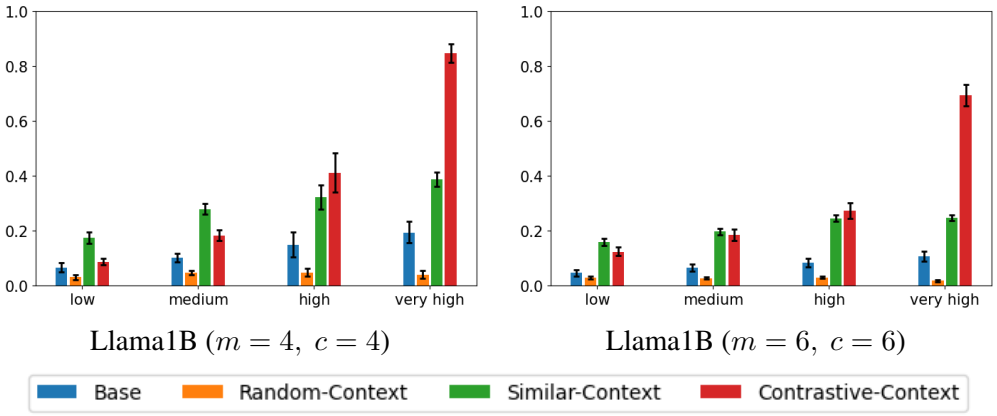

Figure 7: X-axis: Level of maximum Jaccard similarity between the sets of letters used in the target and the context examples. Y-axis: Average maximum accepted length of respective subsequence in y for the PFAs of each letter in x_*. Random-Context performs worse than the baseline irrespective of the similarity level, probably because it has learnt the PFAs before learning the alignment, and thus it chooses a PFA at random for every letter in x_*. On the other hand, the baseline model performs better because it blindly copies an ICL target example sequence corresponding to an ICL context example having some symbols in the same position as x_*. Although Contrastive-Context performs second only to Similar-Context in the low and medium similarity regions, Contrastive-Context outperforms all the models in the high and very high similarity regions.

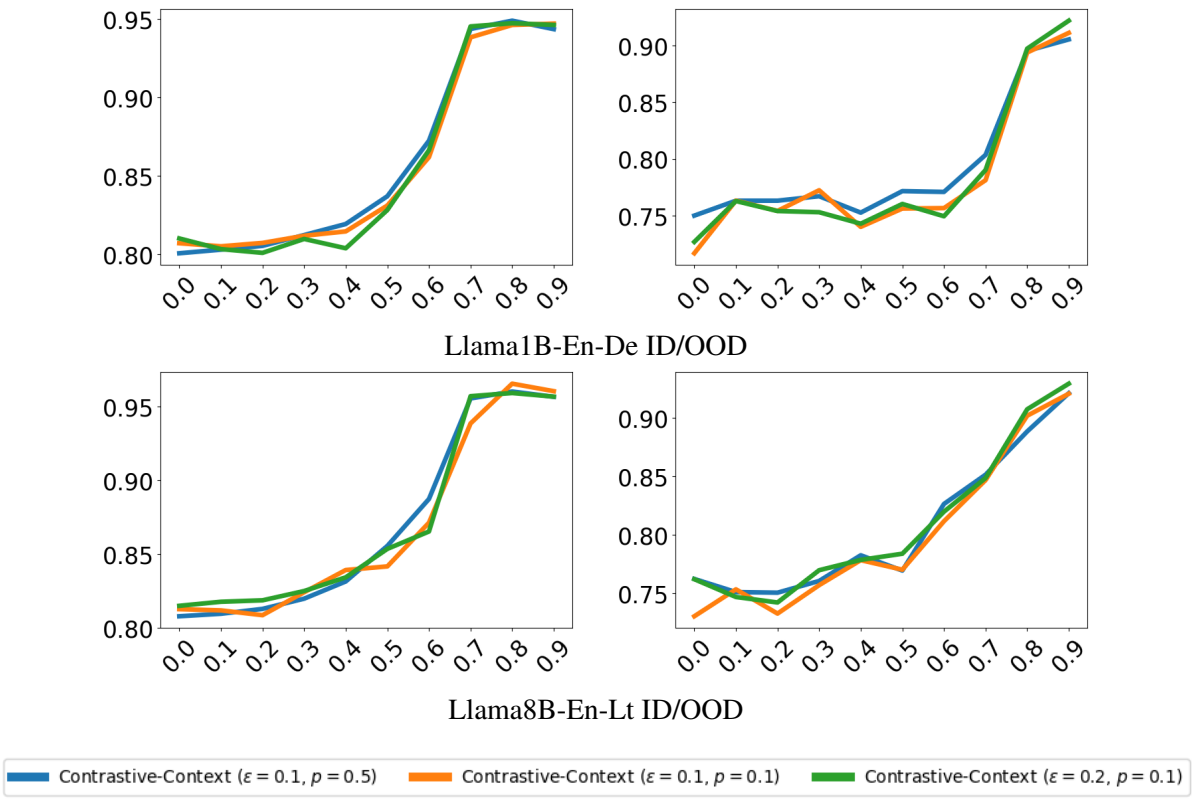

Figure 8: Contrastive-Context is robust to specific choices of ϵ and p as long as sufficient contrast is maintained. X-axis is target-example similarity and Y-axis is accuracy.

C Anecdotes showing Contrastive-Context outperform Similar-Context

Here we evaluate the Similar-Context and Contrastive-Context models on Random-Context data, where no demonstration is guaranteed to resemble the query. Table 3 shows two representative prompts. In both, Similar-Context reproduces one of the in-context targets verbatim instead of translating the query, while Contrastive-Context produces a translation under the same prompt. In (b) the copied demonstration shares only the token *Iraq* with the query, suggesting that a shallow surface cue is enough to trigger the copy.

D Ablations

D.1 Experiments with non-zero ϵ and p

We present robustness of Contrastive-Context to alternative choice of ϵ and p parameters in Figure 8.

D.2 Importance of training with different similarity levels

Contrastive-Context includes in-context examples at three levels of similarity — randomly selected examples which are of low similarity, naturally occurring top-k most similar examples from the training pool which are typically of medium similarity levels, and highly similar examples obtained by paraphrasing the target. We establish the necessity of employing these multiple similarity levels during training by creating two ablations: first, we remove the highly similar paraphrased examples by setting $\epsilon = 0$, and second, we remove the natural similar examples by setting $p = 0$.

Figure 9 shows evaluation on test instances created with mixed similarity levels as described in Section 3.1 (Evaluation Setup). We observe that test instances with high target-context similarity suffer on models trained with $\epsilon = 0$. This is evidenced by the dip in accuracy of the **orange** line compared to the original Contrastive-Context (blue line) in the right quarter of the X-axis). Likewise, the model trained without naturally occurring Top-K pairs ($p = 0$), performs poorly on test instances where IC examples

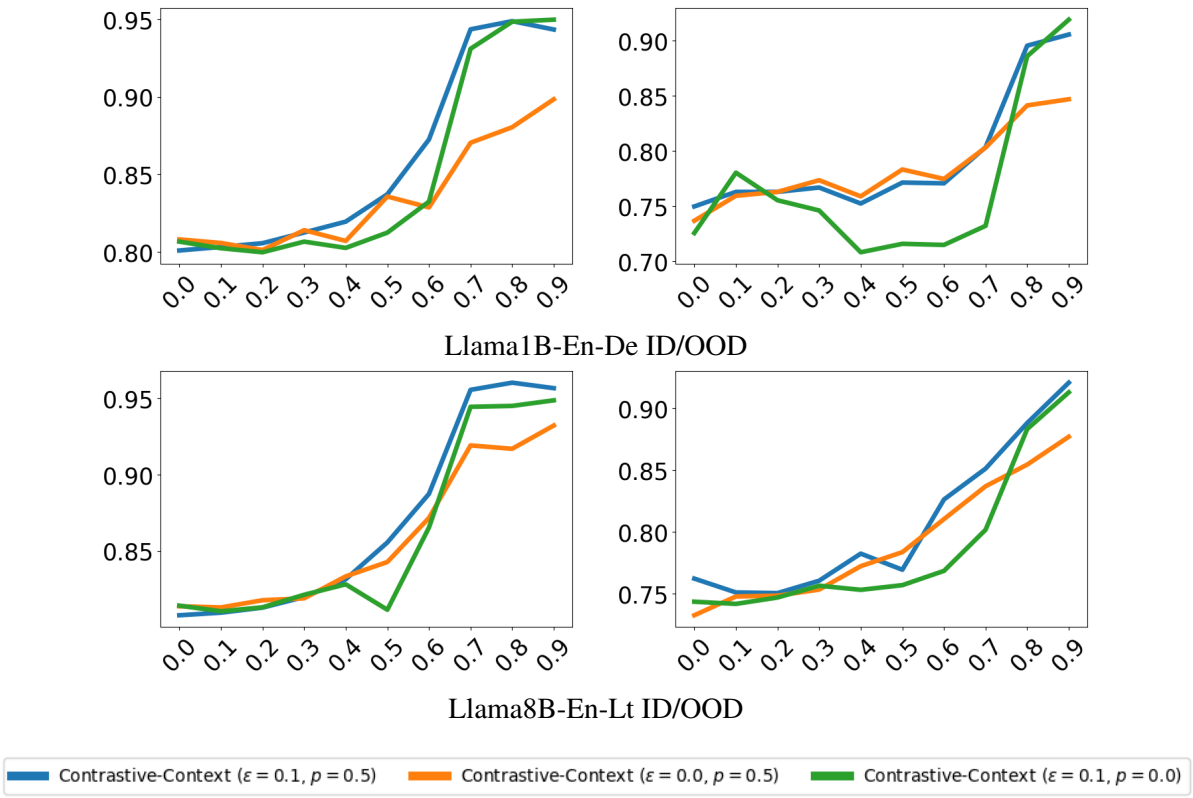

Figure 9: Ablations on Contrastive-Context to establish the importance of training with different similarity levels: X-axis is target-example similarity and Y-axis is accuracy. Removing highly similar examples obtained via paraphrasing ($\epsilon = 0$) causes test accuracy in the high similarity range to suffer. Removing natural similar examples ($p = 0$) causes accuracy in the medium range to suffer.

are at medium levels of similarity with the target. This is seen by the dip in accuracy of the **green** line in the middle part of the X axis.

Overall, we observe the trend that it is important for IC-Train to be trained at multiple similarity levels in order to provide the best accuracy under all levels of relatedness of the test example with the context.

D.3 Intra-Context Vs Inter-Context Contrast

Contrastive-Context creates contrasts in similarity levels both amongst examples within a context, and across contexts. Here we present ablations to understand the importance of contrasts within a context by comparing with a baseline that mixes Random-Context and Similar-Context to preserve only inter context contrast. We present results in Figure 10. On these tasks, the Random-Context and Similar-Context mixture (**green** line) are as good as Contrastive-Context (**blue** line) in the low and medium similarity regime. In the high similarity regime, because of the absence of paraphrases, this mixture loses out. When we remove the paraphrases from Contrastive-Context (**orange** line) the two methods are equivalent on this task. However, conceptually, a model that mixes Random-Context and Similar-

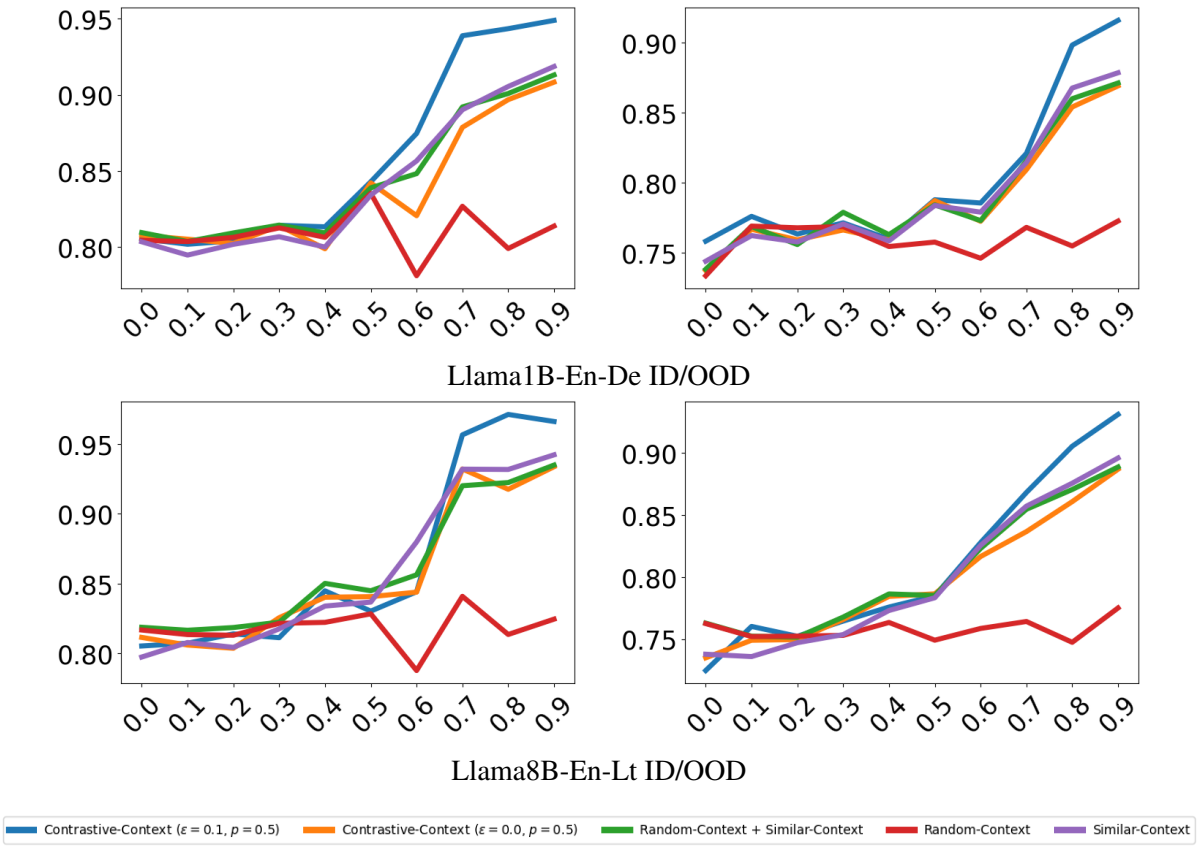

Figure 10: Comparing a Random-Context+Similar-Context mixture with Contrastive-Context that additionally also creates contrasts within a context, and other baselines Random-Context and Similar-Context. X-axis is target-example similarity and Y-axis is accuracy.

Context could learn to swing between ignoring the context (IWL) and blindly copying from the context based on aggregated similarity with the context. This forms a IWL+Copy mixture, which might not perform well when tested with only a subset of context examples as relevant. Contrast within a context promotes true in-context learning where similarity between the x-s determines which y-s are copied. In real-life limited data regimes, the top-K similar examples may differ in their similarity to the target, and thus Random-Context+Similar-context may behave like Contrastive-Context like we observed above.

D.4 Synthetic Paraphrases for Data Augmentation Vs Contrast

An interesting question is whether the observed accuracy boost with synthetic paraphrased examples is due to increased contrast, or vanilla data augmentation. To answer this question, we added the synthetic paraphrases for all target training examples to the training pool as augmented data. We then invoked the Similar-Context method on this augmented dataset. A comparison appears in Figure 11. We observe that with such data augmentation Similar-Context does perform well on test data in the high similarity range but shows a huge drop in accuracy in the low similarity range. This is because Similar-Context goes for the Top-K most similar examples, and all training instances contain a paraphrase in the context causing no incentive for IWL to develop. Contrastive-Context's use of paraphrases in a contrastive setting is important to develop IWL+ICL mixtures.

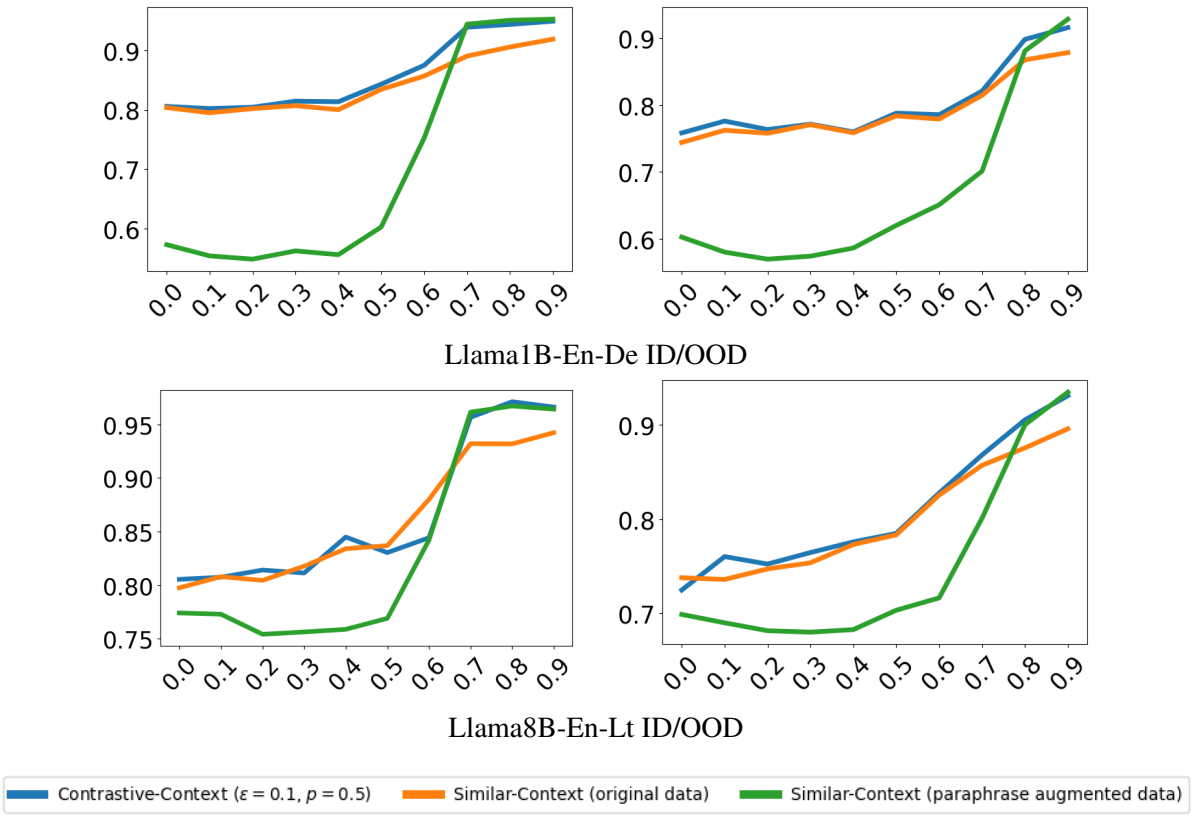

Figure 11: Comparing Similar-Context on original data and on data augmented with paraphrases. X-axis is similarity of test target instance with in-context examples, and Y axis is accuracy. Contrastive-Context uses paraphrases selectively to create contrast, whereas Similar-Context greedily selects Top-K most similar examples. While performance of Similar-Context improves in the high similarity range, it suffers in the low similarity range because IWL does not develop with highly similar context.

D.5 Experiments with various paraphrasing models

Although our proposed method relies on an external paraphrasing model (specifically, `gemini-2.0-flash-lite`), we demonstrate that the choice of the paraphraser has negligible impact on overall performance. To validate this, we conduct additional experiments using two alternative models `gemini-2.5-flash-lite` as an example of a strong model, and `Qwen2.5-7B-Instruct` as an open source, possibly weaker model. The results in Figure 12 shows that Contrastive-Context remains stable regardless of whether the paraphrasing model is open-source or proprietary.

D.6 Effect of LoRA Target Modules

We investigate the effect of applying LoRA to different subsets of transformer projections. In our main experiments, we keep the target modules as only the attention projection matrices (except the output

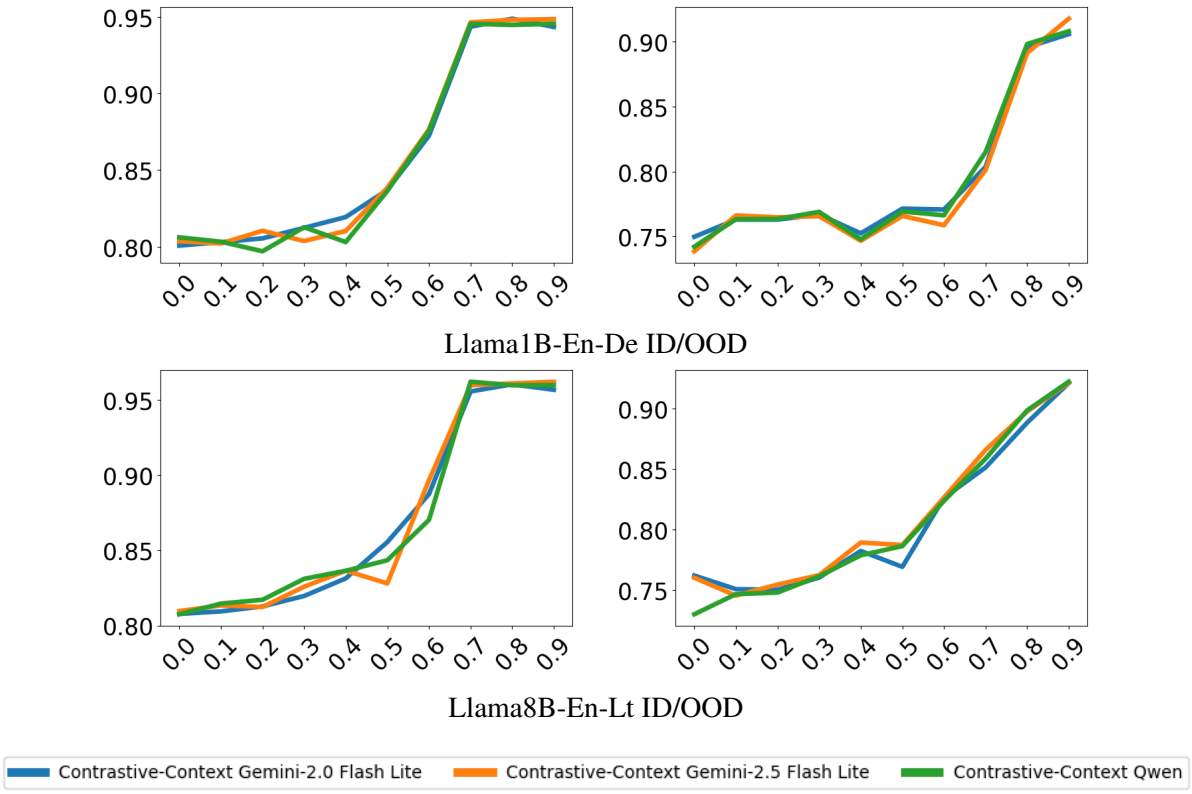

Figure 12: Comparing Contrastive-Context with highly similar examples created with various paraphrasing models. X-axis is target-example similarity and Y-axis is accuracy on the test data. Contrastive-Context is robust to the choice of the paraphrasing model.

projection matrix), but as an ablation, we also experiment by keeping all the attention projection matrices and the MLP layer matrices trainable. The comparison is shown in Figure 13. We observe that with additional trainable parameters, the overall performance increases, but the general trend of results across training regimes remains unchanged.

E Additional Experiments Comparing Different Training Strategies

We present results of more model-language pair combinations that could not be fit in Figure 2 of the main paper in Figure 14.

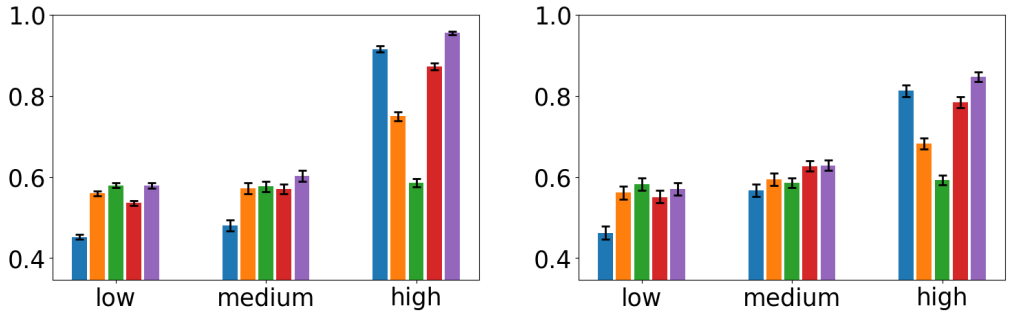

Llama1B-En-Lt ID/OOD for training only attention projection weights

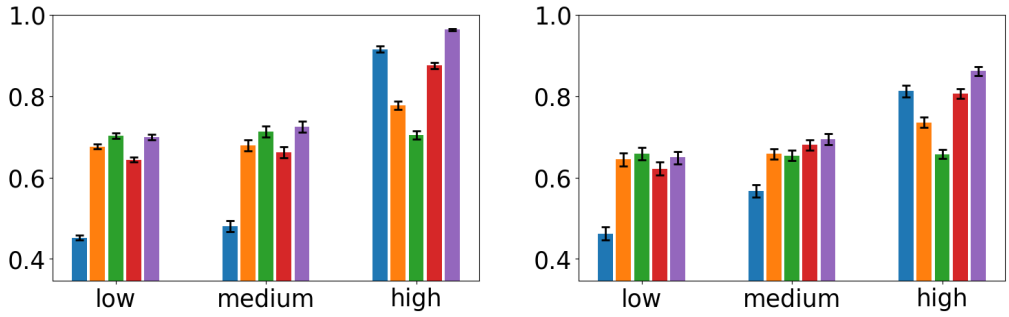

Llama1B-En-Lt ID/OOD for training attention projection weights + MLP layer weights

Figure 13: Comparing training only attention projection weights (except output projection) with training attention projection weights + MLP layer weights. X-axis: Level of maximum similarity of target with in-context examples. The similarity ranges here are - Low: $0-0.33$, Medium: $0.33-0.67$, High: $0.67-1$. Y-axis: Accuracy (COMET score).

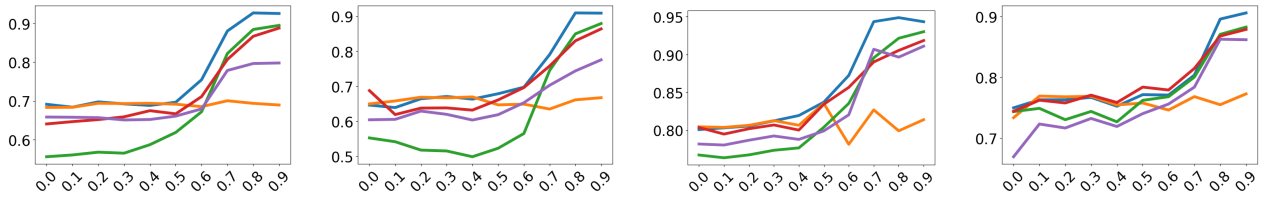

Llama1B-En-Hi ID/OOD

Llama1B-En-De ID/OOD

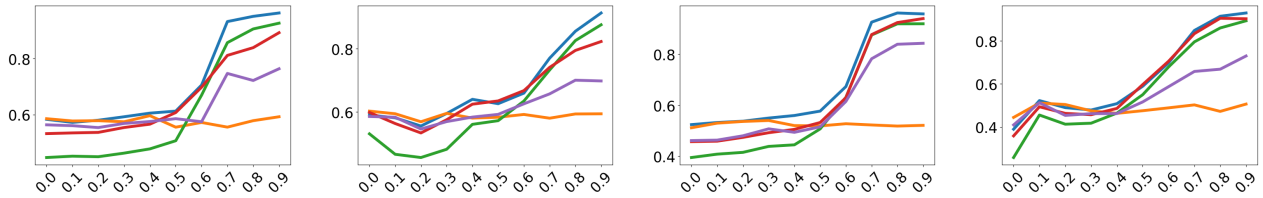

Llama1B-En-Lt ID/OOD

Llama1B-En-Ta ID/OOD

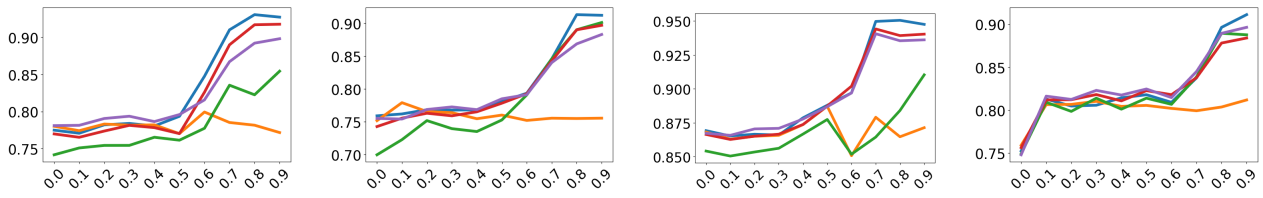

Llama8B-En-Hi ID/OOD

Llama8B-En-De ID/OOD

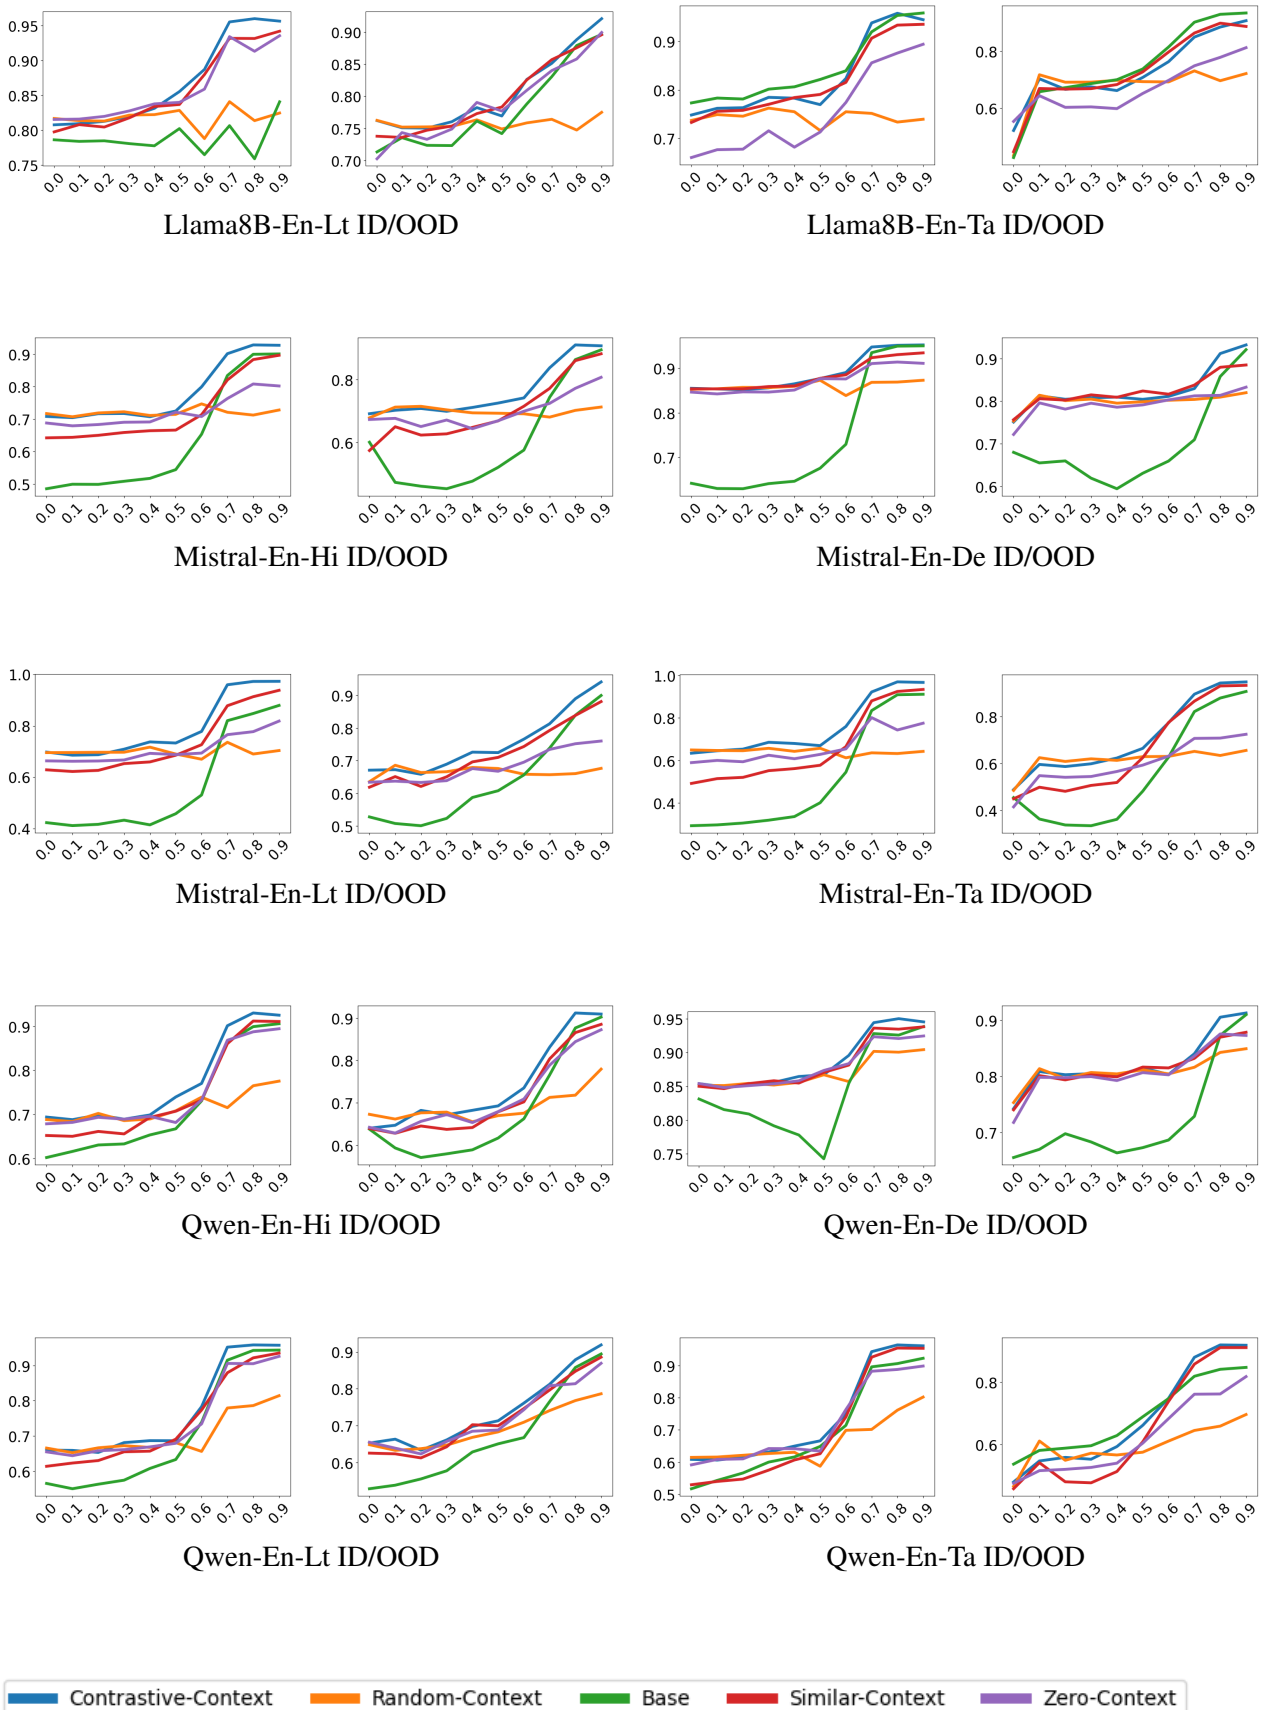

Figure 14: All 32 plots for four different models tested on four different language pairs with ID and OOD distribution. X-axis is target-example similarity and Y-axis is accuracy.

F Additional Experiments on Learning Dynamics

We present results of more model-language pair combinations that could not be fit in Figure 4 of the main paper in Figure 15.

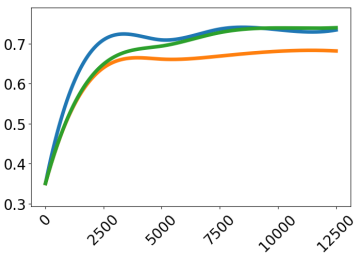

IWL score: 1B-En-Hi-ID

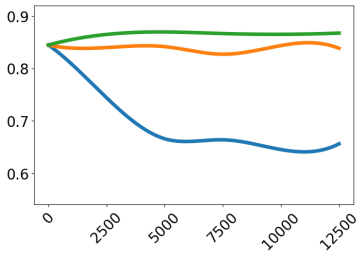

ICL score: 1B-En-Hi-ID

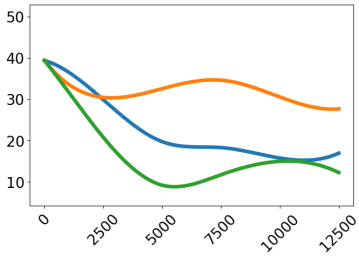

Copy score: 1B-En-Hi-ID

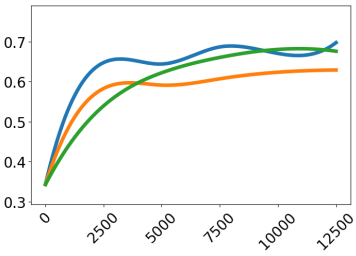

IWL score: 1B-En-Hi-OOD

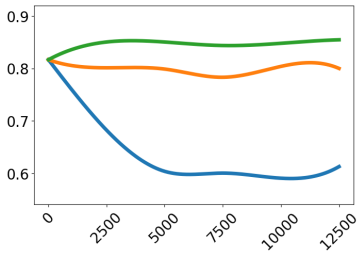

ICL score: 1B-En-Hi-OOD

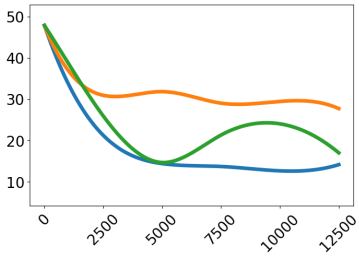

Copy score: 1B-En-Hi-OOD

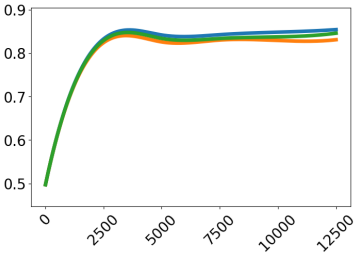

IWL score: 1B-En-De-ID

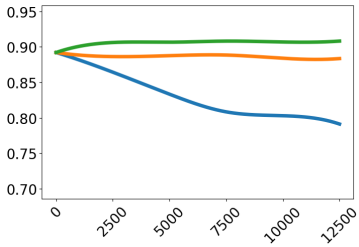

ICL score: 1B-En-De-ID

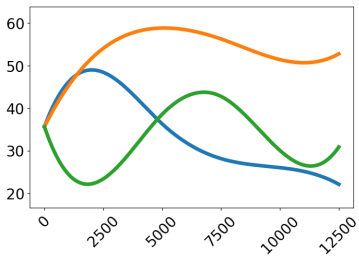

Copy score: 1B-En-De-ID

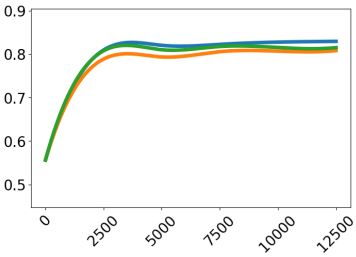

IWL score: 1B-En-De-OOD

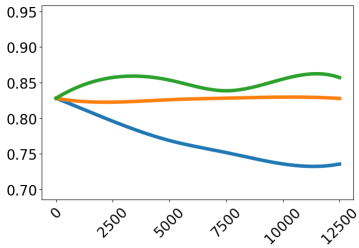

ICL score: 1B-En-De-OOD

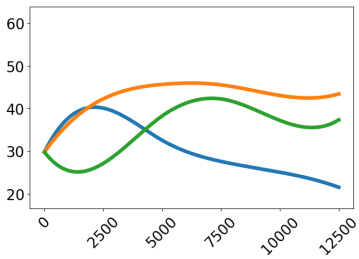

Copy score: 1B-En-De-OOD

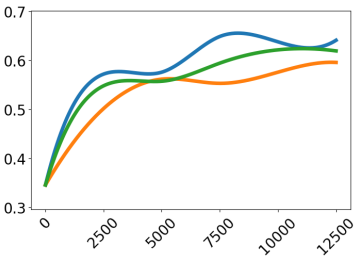

IWL score: 1B-En-Lt-ID

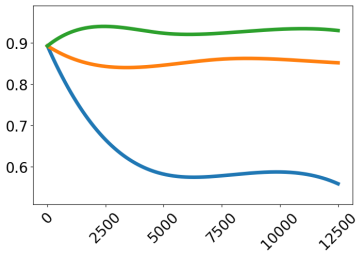

ICL score: 1B-En-Lt-ID

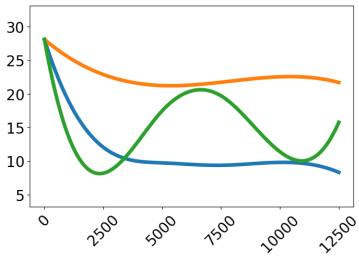

Copy score: 1B-En-Lt-ID

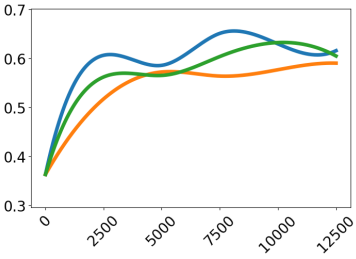

IWL score: 1B-En-Lt-OOD

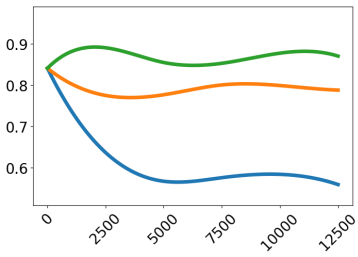

ICL score: 1B-En-Lt-OOD

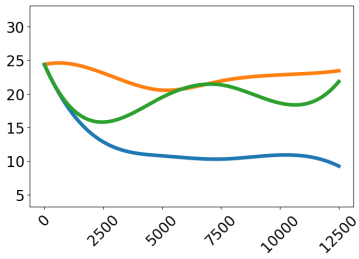

Copy score: 1B-En-Lt-OOD

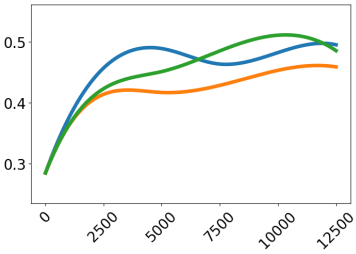

IWL score: 1B-En-Ta-ID

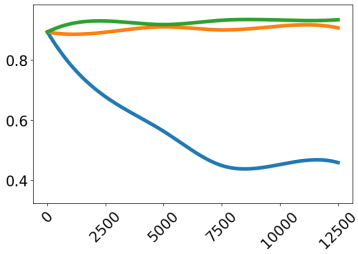

ICL score: 1B-En-Ta-ID

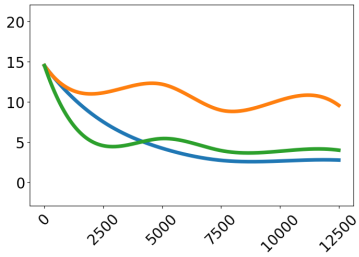

Copy score: 1B-En-Ta-ID

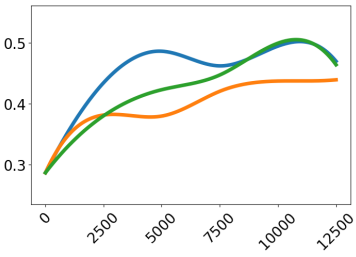

IWL score: 1B-En-Ta-OOD

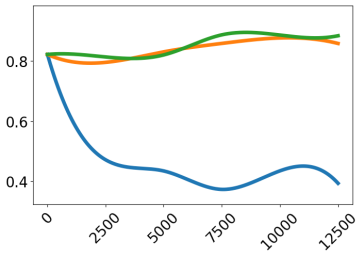

ICL score: 1B-En-Ta-OOD

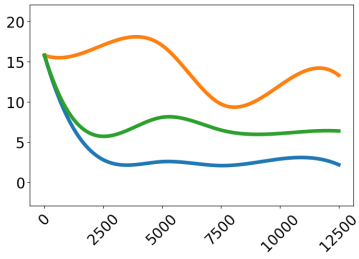

Copy score: 1B-En-Ta-OOD

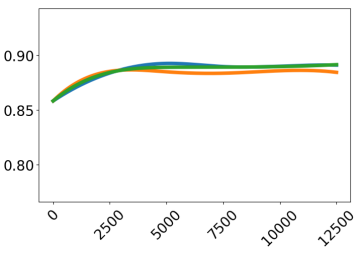

IWL score: 8B-En-De-ID

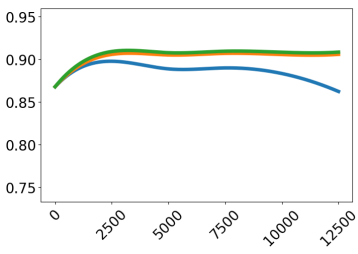

ICL score: 8B-En-De-ID

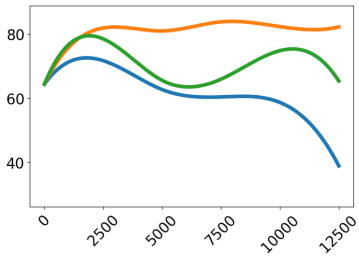

Copy score: 8B-En-De-ID

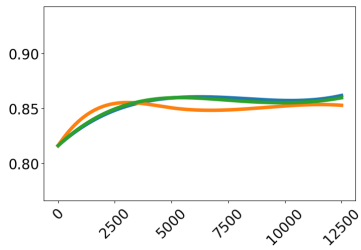

IWL score: 8B-En-De-OOD

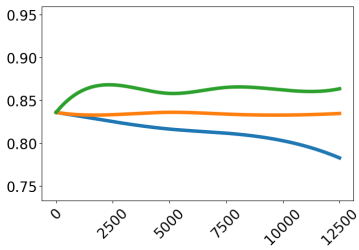

ICL score: 8B-En-De-OOD

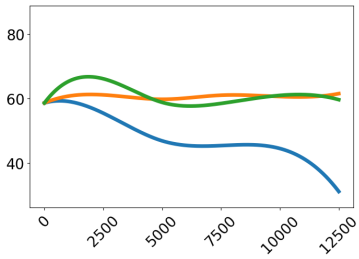

Copy score: 8B-En-De-OOD

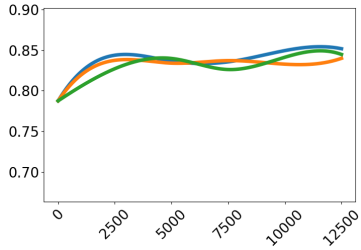

IWL score: 8B-En-Lt-ID

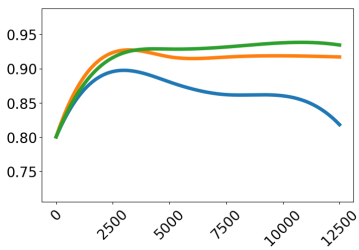

ICL score: 8B-En-Lt-ID

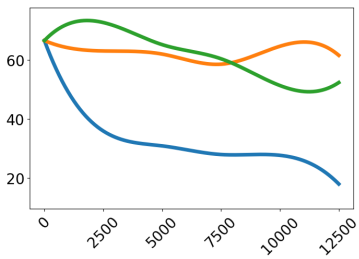

Copy score: 8B-En-Lt-ID

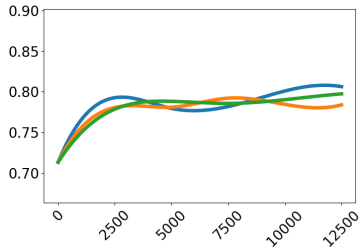

IWL score: 8B-En-Lt-OOD

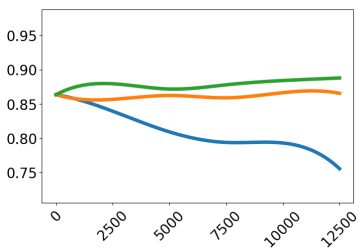

ICL score: 8B-En-Lt-OOD

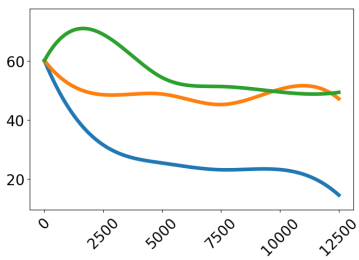

Copy score: 8B-En-Lt-OOD

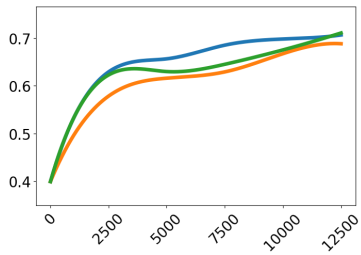

IWL score: Mistral-En-Lt-ID

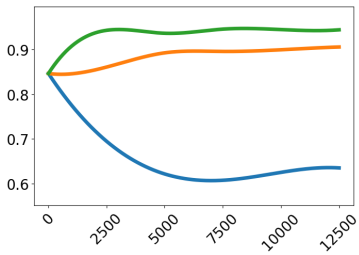

ICL score: Mistral-En-Lt-ID

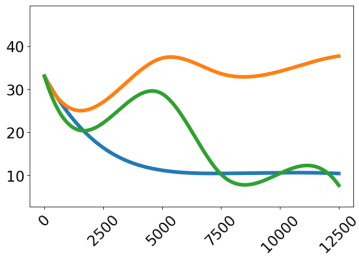

Copy score: Mistral-En-Lt-ID

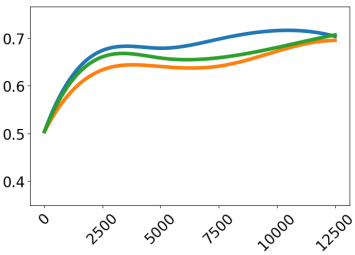

IWL score: Mistral-En-Lt-OOD

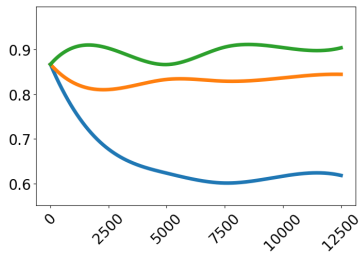

ICL score: Mistral-En-Lt-OOD

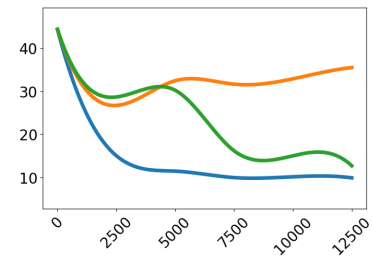

Copy score: Mistral-En-Lt-OOD

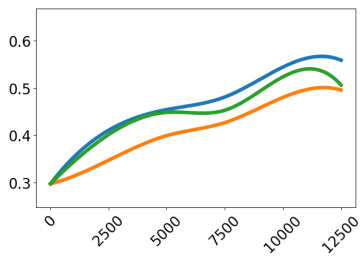

IWL score: Mistral-En-Ta-ID

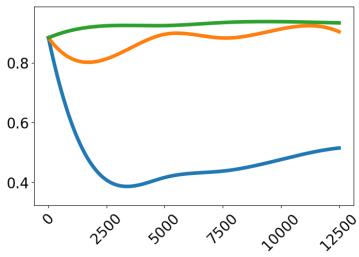

ICL score: Mistral-En-Ta-ID

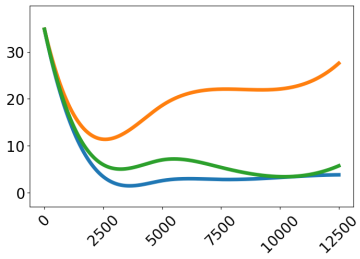

Copy score: Mistral-En-Ta-ID

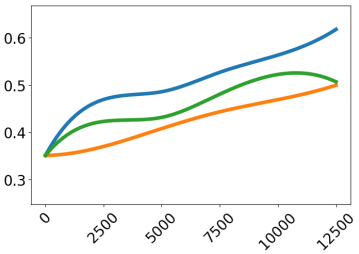

IWL score: Mistral-En-Ta-OOD

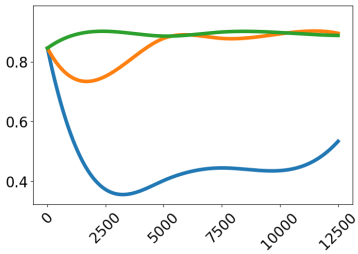

ICL score: Mistral-En-Ta-OOD

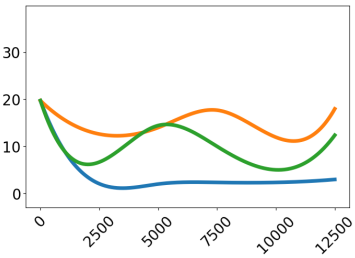

Copy score: Mistral-En-Ta-OOD

Figure 15: Learning dynamics plots for other models and datasets. X-axis is training steps and Y-axis denotes scores of one of the three probes.

F.1 Learning Dynamics (Unsmoothed)

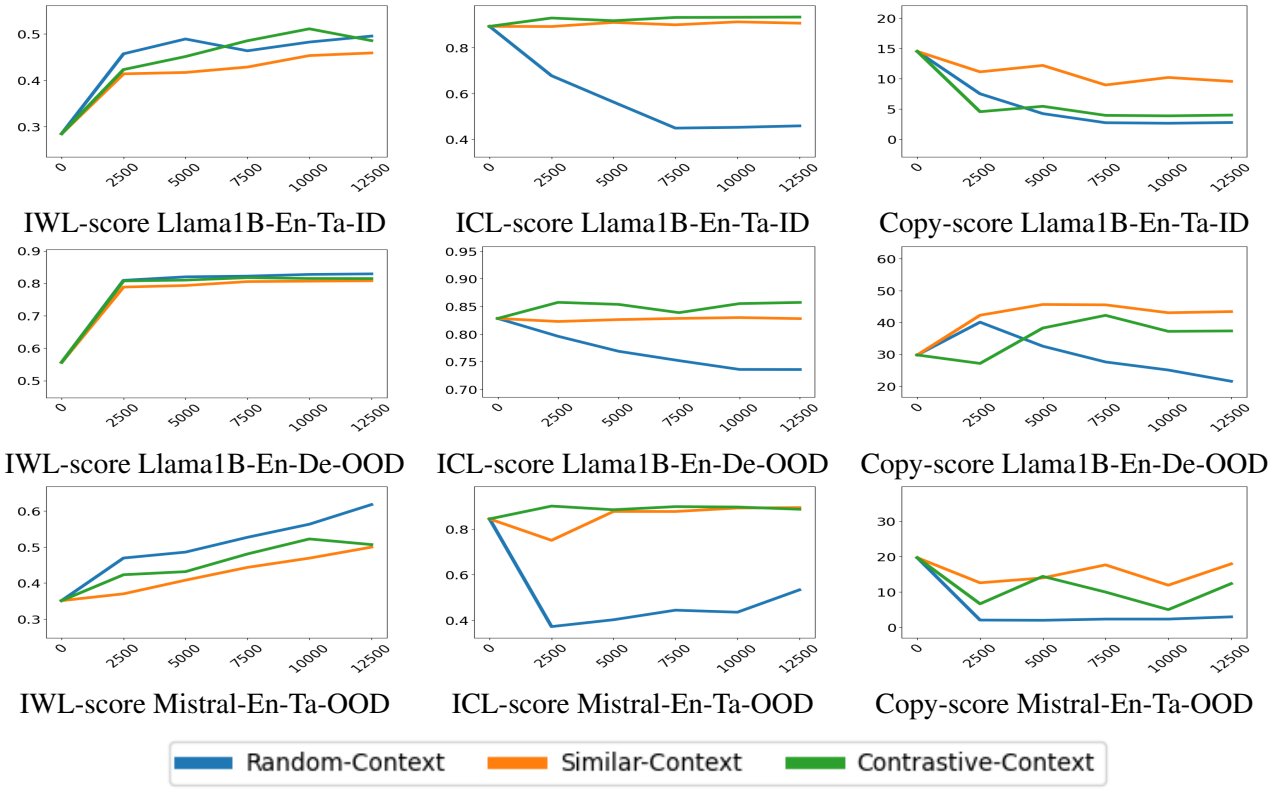

Figure 16: Unsmoothed learning dynamics plots for a few models and datasets. X-axis is training steps and Y-axis denotes scores of one of the three probes.

(a) Qwen · OOD · En–Lt	
x_1	*Each tablet contains 121.3 mg lactose.*
y_1	*Kiekvienoje tablet yra 121, 3 mg laktoz*
x_2	(hallucinations), mania (mental condition characterised by episodes of overactivity, elation or irritability), paranoia, suicidal thoughts
y_2	Galite netekti riebal Js kojose, rankose ir veide, padaugti riebal pilve ir kituose vidaus
x_3	Nitisinone is a competitive inhibitor of 4-hydroxyphenylpyruvate dioxygenase, an enzyme which precedes fumarylacetoacetate hydrolase in the tyrosine catabolic pathway.
y_3	Nitizinonas yra konkuruojantis 4 - hidroksifenilpiruvato dioksigenazs (fermento, esanio prie fumarilacetoacetato hidrolaz tirozino kataboliniame kelyje) inhibitorius.
x_4	Humalog 100 U/ ml, solution for injection in cartridge
y_4	Humalog 100 V/ ml, injekcinis tirpalas utaise
x_5	Nespo 10 micrograms solution for injection in a pre-filled pen (SureClick) Nespo 15 micrograms solution for injection in a pre-filled pen (SureClick) . . .
y_5	Nespo 10 mikrogram injekcinis tirpalas upildytomis virktimo priemonmis (SureClick) Nespo 15 mikrogram injekcinis tirpalas upildytomis virktimo priemonmis (SureClick) . . .
Query	The list below contains adverse reactions seen in clinical trials.
Similar-Context	*Kiekvienoje tablet yra 121, 3 mg laktoz.*
Contrastive-Context	emiau pateikta neigiam reakcij, matom klinikiniuose bandymuose, sraas.
Reference	Toliau pateikiamas nepageidaujam reakcij, pastebt klinikini tyrim metu, sraas.
(b) Llama-1B · ID · En–Ta	
x_1	Tours are cheaper for larger groups, so if you're by yourself or with just one friend, try to meet other people and form a group of four to six for a better per-person rate.
y_1	பெரிய குழுக்களுக்கு சுற்றுப்பயணங்கள் மலிவானதாகும், எனவே, உங்களுக்கோ அல்லது ஒரு நண்பருடனோ இருந்தால், மற்றவர்களைச் சந்திக்க முயற்சி செய்யுங்கள், மேலும் ஒரு நபருக்கு ஒரு சிறந்த விகிதத்திற்காக நான்கு முதல் ஆறு பேர் கொண்ட குழுவை உருவாக்கவும்.
x_2	The valves opening allowed a pressure release for the system and oil flowed on a pad to a tank that can hold 55,000 barrels (2.3 million gallons).
y_2	வால்வுகளைத் திறப்பதென்பது, அமைப்பு அழுத்தத்தில் இருந்து விடுபடுவதை அனுமதித்தது மற்றும் 55,000 பேரல்களை (2.3 மில்லியன் கேலன்கள்) வைத்திருக்கக்கூடிய ஒரு டேங்குக்கான திண்டில் எண்ணெய் பாய்ந்தது.
x_3	*Iraq's Abu Ghraib prison has been set alight during a riot.*
y_3	ஈராக்கின் அபு கிரைப் சிறை ஒரு கலவரத்தின் போது தீ வைக்கப்பட்டது.
x_4	Fifteen of these rocks are attributed to the meteorite shower last July.
y_4	கடந்த ஜூலை மாதம் விண்கற்கள் பொழிந்ததற்கு இவற்றின் பதினைந்து பாறைகள் காரணமாக இருந்தது.
x_5	The ACMA also found that despite the video being streamed on the Internet, Big Brother had not breached online content censorship laws as the media had not been stored on Big Brother's website.
y_5	இணையத்தில் வீடியோ ஒளிபரப்பப்பட்டபோதிலும், பிக் பிரதர் இணையதளத்தில் மீடியா சேமிக்கப்படாததால், பிக் பிரதர் ஆன்லைன் உள்ளடக்க தணிக்கைச் சட்டங்களை மீறவில்லை என்பதையும் ACMA கண்டறிந்தது.
Query	The Iraq Study Group presented its report at 12.00 GMT today.
Similar-Context	ஈராக்கின் அபு கிரைப் சிறை ஒரு கலவரத்தின் போது தீ வைக்கப்பட்டது.
Contrastive-Context	ஈராக் செய்தியாளர்கள் இன்று காலை 12.00 மணிக்கு அறிக்கையை வெளியிட்டனர்.
Reference	ஈராக் ஆய்வுக் குழு, தனது அறிக்கையை இன்று 12.00 GMT-க்கு வழங்கியது.

Table 3: Two Random-Context prompts shown in full, with the outputs they elicit. Each prompt contains five demonstrations (x_i, y_i). *Red italics* mark the demonstration whose target Similar-Context reproduces verbatim as its translation; blue marks the output of Contrastive-Context, which translates the query in both cases. The very long fifth demonstration in (a) is truncated here for space; the model saw it in full.