arXiv:2508.15610v3 [cs.AI] 3 Feb 2026

Transduction is All You Need for Structured Data Workflows

Alfio Massimiliano Gliozzo *gliozzo@us.ibm.com*
IBM

Naweed Khan *naweed.khan@ibm.com*
IBM

Christodoulos Constantinides *christodoulos.constantinides@ibm.com*
IBM

Nandana Mihindukulasooriya *nandana@ibm.com*
IBM

Nahuel Defosse *nahuel.defosse@ibm.com*
IBM

Gaetano Rossiello *gaetano.rossiello@ibm.com*
IBM

Junkyu Lee *junkyu.lee@ibm.com*
IBM

Abstract

This paper introduces `Agentics`, a functional agentic AI framework for building LLM-based structured data workflow pipelines. Designed for both research and practical applications, `Agentics` offers a new data-centric paradigm in which agents are embedded within data types, enabling logical transduction between structured states. This design shifts the focus toward principled data modeling, providing a declarative language where data types are directly exposed to large language models and the data values are composed through transductions between input and output types. We present a range of structured data workflow tasks and empirical evidence demonstrating the effectiveness of this approach, including data wrangling, text-to-SQL semantic parsing, and domain-specific multiple-choice question answering, and data-driven scientific discovery tasks.

1 Introduction

Large language models (LLMs) have demonstrated remarkable capabilities in natural language understanding, reasoning, and tool use. Recent advances in LLM-based agent systems equipped with human-level text generation and conversational abilities have opened promising directions in software engineering, scientific research, and a wide range of tasks that can be automated (Hosseini & Seilani, 2025; Agashe et al., 2025). In the emerging paradigm of agentic AI, LLMs are integrated with external tools, structured knowledge sources, and memory modules to form specialized agents. These agents, each designed with distinct functions and behaviors, collaborate to solve complex tasks (Acharya et al., 2025; Moshkovich & Zeltyn, 2025; Huang & Huang, 2025; Han et al., 2024).

Despite growing interest in agentic AI, current systems remain poorly suited for structured data workflows, where inputs and outputs are governed by explicit schemas and semantics (Hopkins et al., 2022). Embedding structured data into natural language often fails in enterprise use cases such as analytics and data transformation (Sarirete et al., 2022; Heck, 2024; Putri & Athoillah, 2024), due to the lack of compositional guarantees, which leads to brittle workflows, cascading errors, and limited reproducibility.

```python
class ProductReview(BaseModel):
    reviewer: str
    text: str
    stars: int
```

```python
class SentimentSummary(BaseModel):
    sentiment: Literal["positive", "neutral", "
        negative"]
    reason: str
```

```
{"reviewer": "Alice", "text": "Excellent product
    quality and fast delivery!", "stars": 5},
{"reviewer": "Bob", "text": "It's okay, but the
    package was damaged", "stars": 3},
{"reviewer": "Carol", "text": "Terrible experience,
    broken after one use!", "stars": 1}
```

```
{"sentiment": "positive", "reason": "Excellent
    quality and fast delivery"},
{"sentiment": "neutral", "reason": "Okay product, but
    package issues"},
{"sentiment": "negative", "reason": "Broke after one
    use"}
```

Figure 1: Logical Transduction Applied to Sentiment Summary

While existing frameworks (LangChain, 2025; CrewAIInc, 2025; Pydantic, 2025a; Dibia et al., 2024; Khattab et al., 2024) offer modular agent composition and tool integration, they fall short of addressing the foundational challenge. Namely, how to endow agentic systems with algebraic structure that ensures robustness, modularity, and interpretability. As a result, these pipelines remain fragile, i.e., lacking formal semantics and struggling with structured data integration.

To overcome these limitations, we introduce `Agentics`, a framework for agentic AI grounded in *logical transduction algebra*, a formalism for representing and composing transformations between structured inputs, intermediate states, and outputs. `Agentics` provides a unifying programming model for generative structured data workflows, treating each step as an asynchronous transduction rather than a prompt-chained interaction. This abstraction enables modularity, parallelism, and schema-constrained transduction, addressing the fragility and lack of formal semantics in existing agentic pipelines.

At the core of `Agentics` lies the notion of logical transduction, namely, a typed transformation that maps an input object of one schema into an output object of another. The key distinction is that `Agentics` treats agents as stateless transducers operating over well-defined data types, hence shifting away from chat-based multi-agent architectures. Unlike conversational agents that rely on serialized dependencies and multi-turn dialogue, `Agentics` agents support fully asynchronous execution.

By design, `Agentics` exposes a programmatic interface to LLMs in which all input and output data are represented as typed objects, ensuring schema validation and constraint checking. Interestingly, such structured types are particularly well handled by LLMs, and align with function-calling patterns, making them particularly well-suited for reliable and interpretable inference(Rossiello et al., 2021; 2023). Figure 1 illustrates a simple example of logical transduction. A `ProductReview` object containing a reviewer's name, review text, and rating is transduced into a new object that includes a sentiment label and a rationale. The LLM generates these new fields based on the structured input and the schema of the target type, no additional prompt engineering is required. Beyond individual transductions, `Agentics` introduces an asynchronous *map-reduce style* programming model for asynchronous workflow composition, enabling scalable, and controllable pipelines. In `Agentics` programming model, every step is a typed transformation, enabling reproducibility and adaptability in real-world applications.

In this paper, we focus on proposing a new programming model for agentic AI, rather than introducing novel algorithms or engineered solutions for improving task-specific performance. Nevertheless, our evaluation shows that `Agentics` achieves competitive or improved results due to the structured prompting compared with baselines. Our emphasis is on the advantages of shifting toward a data-centric paradigm, demonstrating how this approach simplifies the development of LLM-based workflows while enhancing compositionality, scalability, and execution efficiency.

Section 2 summarizes related works and highlights the key difference between `Agentics` and existing frameworks. Section 3 develops logical transduction algebra and an asynchronous programming model, Section 4 develops the technical implementation of the proposed framework as a Python library, Section 5 shows experiment results on a wide variety of tasks that are closely related to generative structured data workflow tasks, such as data-driven knowledge discovery (Majumder et al., 2025), schema-matching (Parciak et al., 2024), text-to-SQL semantic parsing (Hendrix et al., 1978; Androutsopoulos et al., 1995; Li et al., 2023), and schema rich domain-sensitive multiple choice question answering tasks (Constantinides et al., 2025) [1] .

[1] Due to the space limitation, we provide details of the implementation of each task and additional tasks in the Appendix.

2 Related Work

For decades, AI has been framed as an effort to emulate human intelligence. The Turing Test exemplifies this anthropomorphic ideal: a machine is deemed intelligent if its conversation is indistinguishable from a human's. The rise of LLMs has amplified this framing. By enabling rapid prototyping of intelligent systems through natural language prompts, developing AI agents has become more intuitive and accessible. Consequently, many agentic AI frameworks adopt agentic metaphors such as memory, planning, and tool use, often realized via chat-based interfaces.

This approach has driven remarkable progress in consumer applications. Yet, as tasks demand greater semantic precision, prompt-centric methods often prove brittle, opaque, and hard to scale, especially in structured data environments where reproducibility and accuracy are paramount. Agentic AI frameworks typically position the agent as the locus of intelligence, with data as passive input. While effective for open-ended tasks, this model struggles in deterministic workflows. In enterprise settings where querying, transformation, and integration of structured data are common, conversational agents can introduce error propagation and unpredictable behavior.

To address the limitations of prompt-centric agentic systems, recent research has introduced techniques such as guardrails (Ouyang et al., 2022; Dong et al., 2024; Zhang et al., 2024), self-reflection (Shinn et al., 2023; Asai et al., 2024), and correction strategies (Madaan et al., 2023; Pan et al., 2024). While these methods enhance reliability, most frameworks still depend on free-form text or loosely structured prompts that remain fragile and difficult to verify, especially in tasks requiring high semantic precision. In practice, techniques that offer more reliable interfaces such as structured decoding based on predefined schemas (Rossiello et al., 2023) or type-safe libraries (Pydantic, 2025b) have gained traction. These approaches are increasingly supported by modern software stacks built around LLMs.

The shift toward structured data computation is reflected in several emerging frameworks. `Pydantic-AI` (Pydantic, 2025a) emphasizes type-safe agentic programming and serves as agent framework for building structured AI applications using Pydantic types and multi-agent systems. `LangGraph` (LangChain, 2025) enables orchestration of stateful agents over finite state machines, supporting complex control flows. `CrewAI` (CrewAIInc, 2025) demonstrates strong performance in multi-agent coordination and message passing between agents and tools. `DSPy` (Khattab et al., 2024) pioneers declarative abstractions for prompt engineering and optimization, tightly coupled with structured templates. While these frameworks have evolved toward conversational multi-agent coordination in networked environments, `Agentics` proposes a shift toward computation centered on data semantics and type-driven transformations—enabling more robust, scalable, and interpretable workflows for generative structured data tasks.

3 Logical Transduction Algebra

`Agentics` leverages asynchronous, parallel LLM inference to support enterprise-scale workflows over structured data. To ground this capability, we introduce *Logical Transduction Algebra* (LTA): a typed, compositional calculus for building, analyzing, and optimizing LLM-powered pipelines. Our work is closely related to relational algebra (Codd, 1970) and the MapReduce programming model (Dean & Ghemawat, 2008). Here, we present an abridged version of the formal Logical Transduction Algebra in the main paper. The full details are provided in the Appendix.

A *logical transduction* is a semantically grounded transformation from an object x of type X to an object y of type Y such that each field of y is logically justified by information in x under the constraints of the source/target schemas. Concretely, schemas are realized as `Pydantic` types in Python library implementation, which allows type-checking and constraints make intermediate states explicit and auditable, while aligning naturally with LLM function-calling behaviors. Logical transduction is executed between any two Agentics (AGs).

Definition: Agentics (AG) Let Θ be the universe of types. A type $T \in \Theta$ is a finite set of named slots $T = \{(s_i, T_{s_i})\}$ with $T_{s_i} \in \Theta$. An *Agentic structure* AG bundles a schema and a list of instances:

$$AG := \{\, s_{\text{atype}} : \Theta,\ \ s_{\text{states}} : \text{List}[\, s_{\text{atype}} \,] \,\}.$$

We write $AG[X]$ for an agentic structure with schema X and $\mathbf{x} = AG[X]$ for a particular instance list. Concatenation of state lists endows instances of a fixed $AG[X]$ with a monoid structure, giving a simple but useful algebra over batches.

Definition: Transduction operator ($\ll$) The basic operator of LTA is the left-shift $\ll$, which maps a source object into the target schema:

$$\mathbf{y} \;:=\; AG[Y] \ll x \quad \text{where} \quad \mathbf{y}.s_{\text{states}} = \{\, y \;:\; y \text{ satisfies } Y \text{ and is logically inferred from } x \,\}.$$

Operationally, $\ll$ renders typed inputs into prompts, invokes an LLM (optionally with tools/RAG/few-shot), and parses/validates the result into the target type Y.

Definition: Prompt function (P) A prompt function $P : \text{List}[T] \to \texttt{str}$ serializes typed states to text, bridging structured data and LLM inputs. Zero-shot transduction applies P per state:

$$\mathbf{y}[i] \;=\; AG[Y] \ll P(\mathbf{x}[i]),$$

The default prompt function of AG is the pydantic model_dump() method with returns a json dictionary representing the state.

Lemma: Properties of LTA Let the *transduction context* , i.e. the LLM, decoding settings, tools, and few-shot used by the AG, fixed. Then the following conditions applies:

- *Conditional determinism*: Re-invoking $\ll$ on the same x under the same context yields the same y, enabling reproducibility.
- *Statelessness*: y depends only on x and the context, not on other inputs, enabling asynchronous parallel execution.
- *Compositionality*: If $\mathbf{y} = AG[Y] \ll \mathbf{x}$ and $\mathbf{z} = AG[Z] \ll \mathbf{y}$, then $\mathbf{z} = AG[Z] \ll AG[Y] \ll \mathbf{x}$, giving functional-style pipeline composition.

AType Operators In LTA, operations among agentic types (ATypes) provide the foundation for composing and reasoning about structured workflows. Each type X is defined as a set of named slots (s_i, T_{s_i}), and standard set operations are applied component-wise: the union $X \cup Y$ collects all slots present in either schema, the intersection $X \cap Y$ keeps only shared slots, the difference $X \setminus Y$ removes slots from X that also appear in Y, and the Cartesian product $X \times Y$ builds composite types pairing slots from both. At the level of instances, Agentic structures AG[X] (a schema plus a list of states) form a monoid under concatenation of state lists, with the empty instance as identity.

Asynchronous MapReduce. To scale beyond single-step transduction, LTA provides two higher-order operators:

$$\texttt{amap} : (AG[X], f) \to AG[Y], \qquad \texttt{areduce} : (AG[X], g) \to AG[Y],$$

where $f : X \to \text{List}[Y]$ is applied independently to each x_i (filter/transform/fan-out), and $g : \text{List}[X] \to Y$ aggregates many states (summaries, rankings, joins). Because $\ll$ is stateless, `aMap` execute the function f, which might embed transductions, in asynchronous fashion. In contrast, `aReduce` functions takes all the states of the input agentic at once, and therefore cannot be asyncronously executed.

In short, LTA provides a formalism to make assumptions explicit; $\ll$ provides a uniform *contract* between stages; `amap`/`areduce` expose parallel structure and enable hierarchical summaries; and conditional determinism/statelessness support reproducible, high-throughput execution. Together, LTA turns LLM pipelines from brittle prompt chains into modular, optimized, and auditable programs.

4 Technical Implementation

`Agentics` framework is distributed as a Python library. As briefly introduced in the logical transduction algebra, we define the Agentic structure AG as a container for a list of typed objects. Within this framework, LLM inference is conceptualized as logical transduction, i.e., the process of inducing one object from anothers based on a predefined type schema. The framework is built upon two essential components: the Agentic structure, implemented as the meta-class `AG`, and the `PydanticTransducer`, built on top of `Pydantic` to enable type-safe object transformation. In this section, we describe these core elements and present examples that illustrate the *Agentic programming model* in practice.

4.1 Example Usage of Metaclass AG

The metaclass `AG` is a lightweight yet expressive container for typed data and its execution context. It provides the following three essential capabilities. First, it binds a *typed schema*, which we call `AType`, a `pydantic.BaseModel` subclass. Second, it holds a *list of states*, where each state is a validated instance of `AType`. Third, it carries an *execution context of transducer*, such as `llm`, tools, prompt template, batch size, and decoding parameters.

Metaclass AG The metaclass `AG` allows structured states to be natively represented in Python, while remaining agnostic to the specific LLM providers (e.g., OpenAI, Google DeepMind, Anthropic, Meta AI, WatsonX, etc). The following pseudo code shows an example of `AG`.

```
class AG(BaseModel):
    atype: Type[BaseModel]               # target schema (Pydantic model)
    states: List[BaseModel] = []         # instances of 'atype'
    # execution context
    llm: Any = None                      # LLM client/handle
    tools: Optional[List[Tool]] = []     # optional tool registry
    prompt_fn: Optional[Callable] = lambda x: x.model_dump()
    batch_size: int = 20                 # async batch size
```

Given an `AType` and a list of typed instances, an instance of `AG` carries the states and manages their transformation either through logical transduction or via asynchronous functions in a map/reduce-style program.

Pydantic Models as Schemas An `AType` can be any subclass of `Pydantic BaseModel`. For instance, a news schema `StockNews` may be defined as shown below, and an `AG` instance `news` can be instantiated with the states of type `StockNews`.

```
class StockNews(BaseModel):
    ticker: str    # e.g., "AAPL"
    content: str   # raw news article text
    date: datetime # publication time

news = AG(atype=StockNews, states=[StockNews(ticker="AAPL", content="Apple shares surged after
    record iPhone sales were announced.",date="2025-01-15")])
```

Note that an `AG` instance behaves like a list enriched with a type system.

```
news = AG.from_csv("apple_news.csv")
for article in news:    # iterate states like a list
    print(article.ticker)
news.add_attribute("sentiment", str)
news.rebind_atype(StockNews) # rebind to new schema with added attribute.
```

Logical Transduction ($\ll$) `Agentics` framework overloads the $\ll$ operator to express the *logical transduction* between Agentic structures. The left operand is the *target* `AG`, which becomes populated by transducing the states of the right operand. The following example shows an application of logical transduction to sentiment summary by populating the `SentimentScore` instances from `news`.

```
class SentimentScore(BaseModel):
    ticker: str
    sentiment: float   # normalized score in [-1, 1]
    label: str         # "pos", "neutral", "neg"
output = await (AG(atype=SentimentScore) << news)
#output[0]: SentimentScore(ticker="AAPL", sentiment=0.8, label="pos")
```

Map/Reduce Paradigm The metaclass `AG` applies asynchronous mapping to its list of states, executing either the logical transduction or a (possibly asynchronous) function f via `AG.amap(f)` on each `state` independently. Since transductions are stateless, `amap` evaluations over multiple states can run asynchronously in parallel. In other words,

the mapping can be batched for efficiency. In contrast, `AG.areduce(f)` aggregates a collection of states into a single (or small set of) output(s). Because it consumes the entire list of states, it operates synchronously and is not parallelizable.

Example Workflow By chaining logical transduction (≪), `amap`, and `areduce`, a complex workflow can be expressed declaratively. For example, a sentiment-driven stock ranking can be built as follows. This combination of typed schemas, logical transduction, and asynchronous map–reduce execution yields workflows that are composable, interpretable, and highly scalable.

```
# (1) Gather news for each ticker
news_ag = await AG(atype=StockNews).amap(fetch_news_for_ticker)
# (2) Extract sentiment per article
scores_ag = await AG(atype=SentimentScore) << news_ag
# (3) Aggregate to stock-level sentiment
stock_sentiment = scores_ag.areduce(group_by_ticker_mean)
# (4) Rank portfolio by sentiment
portfolio_ranking = stock_sentiment.areduce(rank_portfolio)
```

4.2 PydanticTransducer

Logical transduction triggers creation and execution of a `PydanticTransducer`, which is a stateless agent whose role is to generate a valid instance of the target `AType` given text input. The textual input can represent virtually any concept accessible to LLMs, while the structured output ensures reliability in downstream tasks. Because agents in `Agentics` avoids shared conversational memory, it naturally supports asynchronous execution and efficient scale-out.

To ensure high throughput and responsiveness across datasets of varying sizes, the transduction operator processes data in configurable batches, typically ranging from 8 to 32 items. As shown in the figure on the right, the execution time decreases as the batch size increases from 1, with performance gains saturating around a batch size of 32. This trend reflects near-linear improvements in processing speed across all tasks, highlighting the efficiency of batched execution. The underlying asynchronous map/reduce programming model —grounded in a logical transduction algebra, which we call `Agentic` programming model enables scalable and robust computation, adapting seamlessly to diverse data workloads.

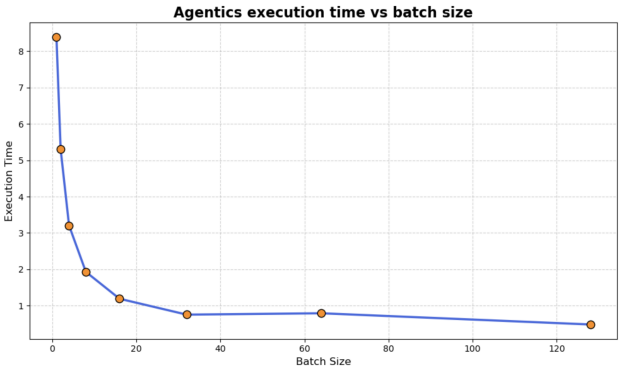

Figure 2: Average time (sec) per question

5 Experiments

In this paper, we evaluate `Agentics`, a programming framework for building generative data workflows with a focus on computational efficiency, scalability, ease of design, and accuracy. While prior work on agentic AI frameworks has emphasized capabilities such as planning and tool use, we focus on the quality and structure of agentic workflow pipelines, particularly in data-centric tasks. We test the following hypotheses: (1) `Agentics` supports a data-centric paradigm through declarative data modeling via type schemas, which decouples logical agent workflows from chat-centric paradigms. This enables intuitive and functional composition of pipelines, (2) the structured prompts induced from declarative data models reduce the burden of manual prompt engineering while maintaining or improving task performance. They are effective in tasks that benefit from clear contexts.

5.1 Data Workflow Tasks

We evaluate canonical structured data workflow tasks such as schema matching with healthcare data (Parciak et al., 2024), data imputation with the Buy and Restaurant datasets (Narayan et al., 2022), and text-to-SQL benchmarks (Li et al., 2023; Zheng et al., 2024).

5.1.1 Schema Matching

Schema matching is a canonical data workflow task that identifies mappings between semantically identical elements in two relational schemas. Recent works such as (Parciak et al., 2024) have explored the use of LLMs for the schema matching task. The schema matching benchmark by Parciak et al. (2024) contains 9 datasets, each consisting of a source table from the MIMIC-IV dataset Johnson et al. (2023) and a target table from the OHDSI OMOP Common Data Model OHDSI (2019). Specifically, the mapping problem involves identifying which source column can be matched to a column in the target table. Table 1 shows the list of datasets with their name of the table from MIMIC-IV and OMOP schemas, and also shows the number of columns in each table. The number of matching tables (GT Matches) are shown in the right most column.

Dataset	MIMIC		OMOP		Candidates	GT Matches
	table name	columns	table name	columns		
AdCO	admissions	16	condition_occurrence	16	256	2
AdVD	admissions	16	visit_detail	19	304	5
AdVO	admissions	16	visit_occurrence	17	272	8
DiCO	diagnoses_icd	5	condition_occurrence	16	80	2
LaMe	labevents	10	measurement	20	200	10
PaPe	patients	6	person	18	108	5
PrDE	prescriptions	17	drug_exposure	23	391	6
SeVD	services	5	visit_detail	19	95	5
TrVD	transfers	7	visit_detail	19	133	6
Total					**1839**	**49**

Table 1: Schema Matching Dataset. The table shows the names of tables from MIMIC-IV and OMOP schemas, and the number of matching tables.

Data Model The schema of the source and target table follows the following data model with table and column names and descriptions. `Attributes` defines a list of `Attribute` states.

```
class Attribute(BaseModel):
    relation_name: Optional[str] = Field(None, description="table name")
    relation_description: Optional[str] = Field(None, description="table description")
    attribute_name: Optional[str] = Field(None, description="column name")
    attribute_description: Optional[str] = Field(None, description="column description")
class Attributes(BaseModel):
    attributes: Optional[list[Attribute]] = Field(None, description="list of Attributes")
```

Logical Transduction Given a source table with M columns and a target table with N columns, we implemented the mapping task as described in (Parciak et al., 2024) in two variations: *1-to-1* and *1-to-N*. In the *1-to-1* setting, each prompt contains a pair of columns from the source and target tables, and the LLM is tasked with determining whether they are semantically equivalent. This setting requires MN LLM calls. In the *1-to-N* setting, each prompt contains a column from the source table and N columns from the target table, requiring only M LLM calls. The LLM evaluates all possible pairs and identifies all semantically equivalent matches in a single inference.

The following pseudocode shows the `Agentics` program that creates `AG` instances for the source and target tables. To parallelize the mapping, `mimic_omop` creates a product of the two `AG` objects and adds an `invertible` attribute to transduce the truth assessment of the mappings.

```
mimic_data = AG.from_states(mimic_states, Attribute)
omop_data = AG.from_states(omop_states, Attributes)
mimic_omop = mimic_data.product(omop_data)
mimic_omop = mimic_omop.add_attribute(slot_name = "invertible", slot_type = list[bool])
mimic_omop = await mimic_omop.self_transduction(mimic_omop.fields, ["invertible"])
```

F1-Scores We have implemented the schema matching task in `Agentics` and ensembled results from two open-source models, GPT-OSS and LLaMA-4. Ensembling the two pipelines from two models takes the disjunction of their assessments, which may decrease accuracy but increase recall.

Table 2 shows detailed results from the experiment. The mean F1-score over 9 datasets is summarized as follows. For the *1-to-1* setting, the GPT-3.5 baseline achieves 0.241, while Agentics achieves **0.325**. For the *1-to-N* setting, the GPT-3.5 baseline achieves **0.398**, and Agentics achieves 0.382, slightly behind. Note that we achieved significantly better performance with smaller-parameter models in the *1-to-1* setting, and slightly worse results in the *1-to-N* setting.

Datasets	1-to-1		1-to-N	
	GPT 3.5 Baseline	Agentics gpt-oss & llama4	GPT 3.5 Baseline	Agentics gpt-oss & llama4
AdCO	0.000	0.500	0.133	0.250
AdVD	0.000	0.330	0.083	0.250
AdVO	0.235	0.400	0.320	0.143
DiCO	0.667	0.500	0.800	0.667
LaMe	0.471	0.364	0.500	0.500
PaPe	0.571	0.000	0.500	0.667
PrDE	0.222	0.000	0.417	0.211
SeVD	0.000	0.500	0.400	0.400
TrVD	0.000	0.333	0.429	0.600
Mean	**0.241**	**0.325**	**0.398**	**0.382**

Table 2: Schema matching F1 scores. GPT 3.5 baseline results are from Parciak et al. (2024).

5.1.2 Text-to-SQL

We evaluated text-to-SQL pipelines composed of multiple components on challenging text-to-SQL benchmarks. The dev set of BIRD-bench (Li et al., 2023) contains 1,534 questions across 11 databases, and the dev set of the Archer dataset (Zheng et al., 2024) contains 104 questions across 10 databases.[2] In our evaluation, we focus on the data-centric design aspects of `Agentic` text-to-SQL pipelines. Composing various components such as randomly generated few-shot examples, schema linker outputs, keywords from topic models, and sub-questions with optimized prompts collectively improves the execution match results by **10.33%** over the baseline performance of Llama-3.3-70B.

Asynchronous Pipeline The text-to-SQL task can introduce the following data model for each problem in the dataset, where a natural language `question` and the SQL `ddl` scripts for creating the database tables are the input fields provided, and the remaining fields are generated by passing through asynchronous pipelines. The `enrichment` field annotates semantic meanings for the tables and columns, `sql_query` is the generated SQL from LLMs, and `execution_result` is the resulting data frame obtained by executing the `sql_query`.

```
class Text2SQLTask(BaseModel):
    question: str = Field(description="The input natural language question")
    ddl: str = Field(description="The database schema in DDL")
    enrichment: Optional[DB] = Field(description="Additional database enrichments")
    few_shots: Optional[Problem] = Field(description="question, sql pair")
    sql_query: Optional[str] = Field(description="SQL query generated")
    execution_result: Optional[List[Dict[str, str]]] = Field(description="resulting table")
```

[2]State-of-the-art performance on both benchmarks is available on the official leaderboards. The best dev performance is 76.14% on BIRD-bench (`https://bird-bench.github.io/`) and 38.46% on Archer (`https://sig4kg.github.io/archer-bench/`).

The asynchronous pipeline concatenates the logical transductions in order, as shown in the pseudo-code below. A series of logical transductions generates the empty fields in the `Text2SQLTask` state object. Note that the overall execution of all state objects is performed concurrently, either per transduction step or per line of the pseudo-code.

```
text2sql("enrichment.keywords")<<text2sql("enrichment.description")<<text2sql("ddl_schema")
text2sql("enrichment.subquestions") << text2sql("question")
text2sql("enrichment.linked_schema") << text2sql("enrichment.ddl_schema", "enrichment.subquestions", "question"
    )
text2sql("few_shots.question", "few_shots.sql_query")<<text2sql("ddl_schema", "enrichment.description")
text2sql("few_shots") = text2sql.filter(valid_sql, "few_shots")
text2sql("sql_query") << text2sql("question", "enrichment", "few_shots")
text2sql("execution_result") = text2sql.amap(execute_sql_query)
```

Brid-Dev Result We evaluated execution accuracy on BIRD-dev using a simple prompt and a composite workflow that includes additional logical transductions with randomly generated few-shot examples (**FS**), keyword enrichment (**KW**), sub-question enrichment (**SQ**), schema linking enrichment (**SL**), and an optimized prompt template (**OP**).

Table 3 highlights the performance gains on average by including additional transductions on top of the base prompt. We note that every individual transduction (i.e. FS / KW / SQ / SL / OP) does not in fact improve average performance on all models. However, techniques such as schema linking and optimized prompts yield greater improvements on a few models than the other approaches. Having a simple modeling framework, like `Agentics`, allows for more sophisticated augmentations and feedback loops to improve the creation of such programs.

	Model		
Method	llama-3-3 70b-instruct	mistral-large	llama-4-maverick 17b-128e-instruct-fp8
P	50.51 ± 0.71	47.20 ± 0.78	53.88 ± 1.37
P+FS	50.32 ± 0.52 (-0.19)	45.63 ± 0.9 (-1.57)	53.47 ± 0.56 (-0.41)
P+KW	50.32 ± 0.32 (-0.19)	45.09 ± 0.81 (-2.11)	53.03 ± 0.82 (-0.85)
P+SQ	51.58 ± 0.58 $(+1.07)$	46.15 ± 0.44 (-1.05)	52.33 ± 0.62 (-1.55)
P+SL	52.75 ± 0.07 $(+2.24)$	49.54 ± 1.23 $(+2.34)$	50.46 ± 0.05 (-3.42)
P+OP	54.90 ± 0.53 $(+4.39)$	48.93 ± 0.5 $(+1.73)$	53.15 ± 0.65 (-0.74)
P+FS+KW+SQ+SL+OP	60.84 ± 0.7 $(+10.33)$	55.41 ± 1.2 $(+8.21)$	56.94 ± 0.23 $(+3.06)$

Table 3: Execution accuracy on BIRD-dev, testing a prompt **P** with the simplified task vs. additional transductions in the composite workflow. **SL** replaces the full DDL schema with a linked schema, **KW** includes keyword topic modeling, **FS** randomly generates sql-validated few-shot question-query pairs, **SQ** extracts sub-questions, and **OP** optimizes the prompt template.

Across all models, the composite workflow consistently improved performance. Llama-3.3-70B showed the largest gain, improving from $50.51\% \pm 0.71$ to $60.84\% \pm 0.53$, a **10.33%** increase. Mistral-Large improved by **8.21%**, and Llama-4-maverick-17B saw a more modest gain of **3.06%**. These results demonstrate that structured prompting and modular transductions in Agentics can significantly enhance execution accuracy. Inspecting the impact of each individual component, **SL** and **OP** contributed 2.24% and 4.49% individually, while other components did not improve accuracy on their own. However, combining all components results in a higher gain of 10.33%, exceeding the sum of individual improvements.

Archer-Dev Result The Archer benchmark presents challenges that require LLMs to perform commonsense, arithmetic, and hypothetical reasoning in order to correctly generate SQL expressions from natural language questions. In our experiments, we adopt a simplified pipeline that bypasses intermediate components and directly generates the `sql_query` from the `ddl` and optional `commonsense_knowledge` hints.

Table 4 summarizes the evaluation result on Archer dataset.

Model	With Commonsense	Without Commonsense
GPT-OSS-120B	$35\% \pm 2$	$28\% \pm 2$
Llama-3.3-70B	$15\% \pm 1$	$15\% \pm 1$
Llama-4-17B	$30\% \pm 2$	$25\% \pm 3.2$

Table 4: Execution accuracy on Archer-dev. We evaluted 104 dev problems with and without commonsense knowledge provided in the dataset.

We evaluate three LLMs, GPT-OSS-120B, Llama-3.3-70B, and Llama-4-17B under two conditions, with and without commonsense knowledge. Based on the average execution match over 10 trials, GPT-OSS-120B shows a clear improvement when commonsense knowledge is incorporated, increasing from 28% to 35%. Llama-3.3-70B performs identically in both settings 15%, and Llama-4-17B also improved execution match performance from 25% to 30%. For reference, the best reported performance on the Archer dev set is 38.46% using the GPT-o1 model. Our results with GPT-OSS-120B demonstrate comparable performance, highlighting its effectiveness despite the simplified pipeline and smaller parameter models.

5.1.3 Data Imputation

The data imputation task is also a canonical example in generative structured data workflows. We consider the task of filling in missing entries with plausible substitutions for a given tabular record containing one or more missing field values. In `Agentics`, the input record can be defined as a semantic data type. The missing attribute is then transduced from the known attributes using logical transduction. We evaluate data imputation implemented in `Agentics` on the Buy and Restaurant datasets Narayan et al. (2022). In the Buy dataset, given a product name and its description, the model is tasked with predicting the manufacturer as the missing value. In the Restaurant dataset, the goal is to infer the type of restaurant based on its name, address, and phone number.

The results of our framework in zero- and few-shot settings are presented in Table 5.

Dataset	Model	zero-shot	few-shot
Buy	gpt-oss	70.77	70.77
	llama-3.3	49.23	58.46
	llama-4	72.31	75.38
Restaurant	gpt-oss	79.07	80.23
	llama-3.3	67.44	48.84
	llama-4	66.28	47.67

Table 5: Accuracy results for different models and few-shot settings on Buy and Restaurant data imputation tasks.

The zero-shot accuracy performance of the GPT-OSS-120B and Llama-4-17B models is 70.77% and 72.31% on the Buy dataset, and 79.07% and 66.28% on the Restaurant dataset, respectively. These results are comparable to the baseline performance of GPT-3-175B, which achieves 84.5% on Buy and 70.9% on Restaurant.

5.2 Domain-Specific Multi-Choice Question Answer

5.2.1 Dataset

We evaluate the *FailureSensorIQ* dataset, a recently proposed domain-specific multiple-choice QA benchmark designed to assess understanding of sensor relationships and failure modes (Constantinides et al., 2025). We demonstrate the stable performance of structured prompting in `Agentics` across both the original single-correct MCQA and its perturbation variants. The single-correct MCQA set contains 2,667 questions spanning various industrial assets, with questions around identifying the right sensor that can detect a given failure mode for a given asset, or identifying the right failure mode that a given asset and sensor can detect. This requires a nuanced understanding of sensor behavior, failure propagation, and asset-specific operational logic, and performing logical deductions across different knowledge about the asset. An example query is given as follows.

```
{
  "Question": "For electric motor, if a failure event rotor windings fault occurs, which sensor out of the
       choices is the most relevant sensor regarding the occurrence of the failure event?",
  "Options":["A. partial discharge", "B. resistance",  "C. oil debris", "D. current", "E. voltage"]
}
```

The perturbations are designed to be knowledge-invariant. The *simple perturbation* involves renaming option letters, while the *complex perturbation* combines option letter renaming with question rephrasing.

5.2.2 Methods

Baseline results are obtained from the leaderboard that evaluates the standardized prompts with at most three trials for invalid responses.

Our approach leverages the `Agentics` framework to perform schema-constrained transduction from structured input to structured output. Each input instance is represented using a subset of fields from the FailureSensorIQ schema—specifically, the question, asset name, option ids, options, and subject—which are sufficient to ground the reasoning process in both linguistic and domain-specific context.

To improve inference efficiency, `Agentics` supports *parallel batch execution* via the `amap` operation. This distributes multiple structured prompts across concurrent model invocations, significantly reducing total runtime. Unlike sequential prompting, which processes one question at a time, batch transduction enables scalable evaluation and deployment.

Logical Self Transduction Each problem in *FailureSensorIQ* dataset has multiple fields, answer `options`, industrial `asset_name`, `relevancy` context, `question_type`, `subject` of the question, and the `answer` to the question. The structured prompting compose the input prompt from the data model and the Agentics framework performs self-transduction from all other fields to the `answer` with additional instruction as follows.

```
fsiq_benchmark = await fsiq_benchmark.self_transduction(
    input_fields=["question", "options", "option_ids",  "asset_name", "relevancy", "question_type", "subject"],
    output_fields=["answer"],
    instructions=("Read the input questions, all possible answers, and background task information. This is a
         multiple choice test, where one of the options is true and the others are false. Select the answer with
          the highest likelihood of being correct, and return it along with a confidence score and a verbal
         assessment explaining your judgment."))
```

5.2.3 Experiment Setting

We evaluate four models ranging from 8B to 405B parameters: **Qwen3-8B** (Yang et al., 2025), **Llama-3.3-70B-Instruct** (Touvron et al., 2023), **Mistral-Large-Instruct-2407** (Jiang et al., 2024), **Llama-3-405B-Instruct** (Touvron et al., 2023). Models are tested using both the original FailureSensorIQ baseline pipeline and the `Agentics` framework. The baseline uses loosely formatted natural language prompts and retries up to three times if the output is invalid. In contrast, `Agentics` uses structured prompting and schema-constrained decoding. To measure execution time, we host `Qwen3-8B` on a dedicated node with an `A100 80GB GPU` running `VLLM`. Other models are accessed via cloud computing platform. We vary batch sizes to assess scalability and throughput.

5.2.4 Experiment Results

Model	# Params	Baseline	`Agentics`
Qwen3-8B	8B	45.86	60.18 (+14.32)
Llama-3.3-70B	70B	41.69	50.73 (+9.04)
Mistral-Large	123B	50.09	58.41 (+8.32)
Llama-3-405B	405B	51.26	52.90 (+1.64)

Table 6: Accuracy (%) of on 2,667 FailureSensorIQ instances.

Accuracy Table 6 shows the accuracy comparison. Agentics consistently improved performance over the baseline prompt performance for all models. Qwen3-8B showed the largest gain, improving from 45.86% to 60.18%. Llama-3.3-70B and Mistral-Large saw gains of **9.04%** and **8.32%**, respectively. The largest model, Llama-3-405B, showed a

modest improvement of **1.64%**, Notably, `Qwen3-8B` achieves a major improvement of **14.32%**, placing it just behind `openai-o1`, 60.4%. These results suggest that prompting through logical transduction helps unlock latent reasoning capabilities, even in models with limited parameter counts.

Robustness `Agentics` demonstrates robustness against knowledge-invariant perturbations. We follow `FailureSensorIQ's` perturbation study using the `Agentics` framework to study if there are any robustness benefits that comes with the structured workflows that our framework offers. We experiment with the following knowledge invariant perturbations: (1) Option letter renaming; changing the option letters from of A., B., C., to other letters like P., Q., R. We'll call this "Simple" Perturbation. (2) Option letter renaming and question rephrasing done by an LLM. We'll call this "Complex" Perturbation. We use the same perturbed datasets from the original paper.

In the original `FailureSensorIQ` experiments, all models experienced significant drops in performance, ranging from 5% to 20% . However, `Agentics` showed minimal change, ranging from just 0.08% to 0.19% under the simple perturbation, and from 2.21% to 2.44% under the complex perturbation.

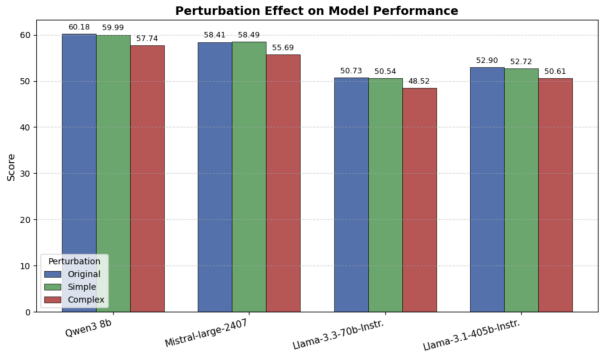

Figure 3: Domain Specific MCQA Perturbation Results.

5.3 Data Driven Scientific Discovery

The DiscoveryBench (Majumder et al., 2025) dataset is a relatively recent benchmark designed to evaluate the full pipeline of data-driven scientific discovery. It contains 264 real-world discovery tasks spanning six scientific domains, including sociology, engineering, and various experimental sciences. Each domain is defined by structured data, accompanying metadata, and a natural-language discovery goal. To solve the DiscoveryBench tasks, it is crucial to interpret data semantically, identify salient variables, generate hypotheses, and compose multi-step reasoning workflows.

Table 7: Baseline Performance Comparison on HMS Metrics

Agents	**HMS**
NoDataGuess (GPT-4o)	0.0
CodeGen (GPT-4o)	15.5
React (GPT-4o)	15.4
DataVoyager (GPT-4o)	15.4
Agentics (GPT-OSS120B)	21.4
Reflexion (Oracle)	24.5

Since the input of this dataset is in CSV files with numerical fields, many of the state of the art approaches rely on code generation and execution, namely the agentic AI systems produce Python code and execute it on the numeric heavy dataset, interpret the result, and refine the hypothesis. Similar to other structured data workflows, failure to constrain the type semantics results in brittle pipelines and limiting the scalability.

In `Agentics` framework, we follow the schema constrained composition of the data transformations with logical transduction and asynchronous MapReduce. Each row of the raw dataset first maps to a typed object that extracts

evidence via asynchronous Map transductions. Then, typed intermediate states are aggregated through asynchronous Reduce transduction that synthesizes into a final state with the Answer type.

We evaluate the hypothesis match score (HMS) defined in the paper. The resulting pipeline processes all 285 discovery tasks under one hour with batch processing 20 parallel calls. Table 7 shows the final scores. An agent using GPTOSS-120B implemented in `Agentics` shows an HMS score of 21.4 in the DB-REAL category, which is higher than other agents in the baseline using GPT-4o.

5.4 Prompt Optimization

Automatic prompt optimization (APO) is an essential task because LLM performance is highly sensitive to prompt structure, tone, and formatting. The prompt function in Definition 9 plays a central role in logical transduction, which can be conceptualized as a negotiation of meaning between agents. Among various APO approaches (Ramnath et al., 2025; Li et al., 2025), logical transduction algebra naturally supports OPRO-style methods (Yang et al., 2023; Ye et al., 2024; Opsahl-Ong et al., 2024), which follow a local search procedure based on a cycle of generate-select-evaluate. In the `Agentics` framework, the candidate prompt templates are generated by logical transduction from the meta-type that describes the prompt optimization task. Additionally, logical transduction algebra naturally parallelizes local search, such as generating and evaluating candidate prompts.

5.4.1 Dataset

We evaluated GSM8K (Cobbe et al., 2021) and FailureSensorIQ (Constantinides et al., 2025) benchmarks. The former enables direct comparison with existing APO techniques. For GSM8K, we use the first 500 training examples for prompt optimization and evaluate the final performance on the full test set. The latter benchmark allows us to assess additional gains beyond the default prompting provided by logical transduction. We randomize the dataset, selecting 500 examples for training and fixing 1,000 examples for the test set. In APO, the training set is used to construct demonstrations for candidate generation and to evaluate candidates during optimization.

5.4.2 Methods

We designed the meta-prompt for prompt generation by analyzing those from OPRO (Yang et al., 2023) and PE2 (Ye et al., 2024)[3]. Our focus is on the impact of local search hyperparameters on final test performance and the overall running time. Following Yang et al. (2023), the optimization meta-prompt includes 3 demonstration tasks and the top 8 candidates, along with their evaluated scores in ascending order. In our experiments, we vary the number of parallel candidate generations from 1 to 8, and the batch size for asynchronous LLM API calls from 1 to 20.

5.4.3 Optimization Scope

Following the common experiment settings (Yang et al., 2023; Ye et al., 2024; Zheng et al., 2025; Spiess et al., 2025), we optimize the prompt template for the zero-shot chain of thought technique (Kojima et al., 2022) by specifying the role, goal, expected output types, and the description of input task as well. This can be flexibly incorporated into APO, as candidate generation follows logical transduction from field descriptions to a prompt template. We initialize all prompt components as empty strings and iteratively refine them.

To understand the impact of prompt components, we manually constructed a prompt using the parameters and evaluated it with the LLaMA-3.3-70B model. The default prompt, which only shows the input and output field names, achieved just 5% accuracy. Adding an expected output description increased accuracy to 66%, and including all prompt parameters further improved it to 67%. These results suggest that optimizing both output expectations and imperative phrasing is essential for effective prompt design.

5.4.4 Performance Improvement

We present experiment results showing the improvement of test scores on both the GSM8K and MCQA datasets.

[3]Relevant open-source implementations are `https://github.com/psunlpgroup/GreaterPrompt`, `https://github.com/stanfordnlp/dspy`, `https://github.com/google-deepmind/opro`.

- The optimized prompt templates improve the test scores to 85 for Llama-3.3-70B and 91 for Qwen3-8B, which is consistent with findings reported in the literature. For the FailureSensorIQ dataset, the test score of Llama-3.3-70B was further improved to 54%.
- The plots with batch size 1 indicate that the Llama-3.3-70B model discovered better prompts when using a larger batch size. In contrast, the Qwen3-8B model identified a good prompt with a smaller batch size. However, it's important to note that the Qwen3 series models have been trained on data derived from the GSM8K dataset as well as other math-related datasets. This prior exposure may contribute to their relatively strong performance on GSM8K, even with smaller batch sizes or less prompt optimization effort, and should be considered when interpreting the results.

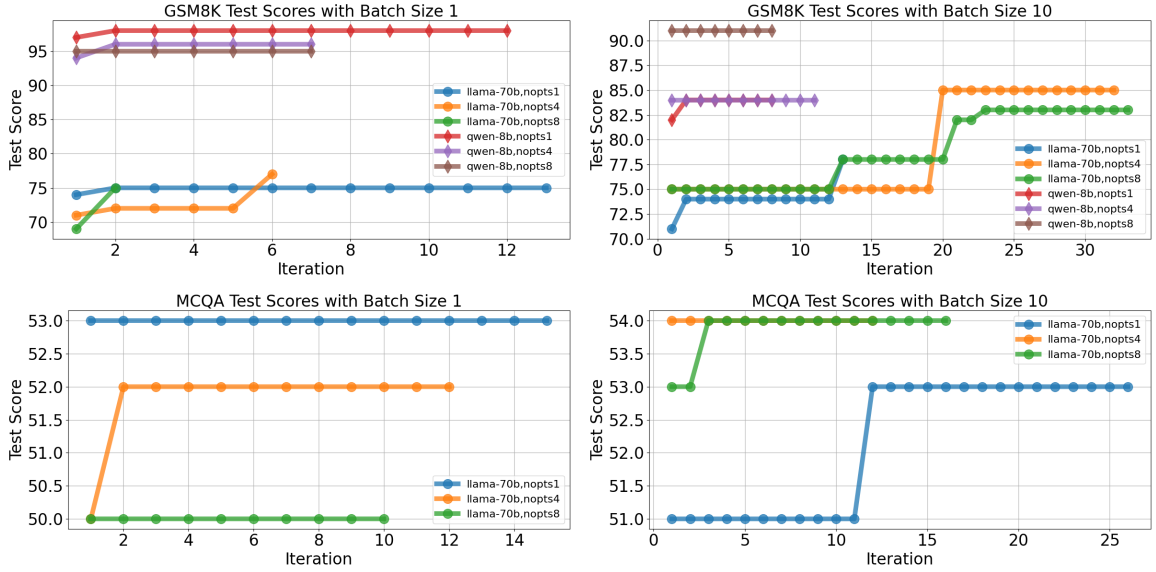

Figure 4: Improvement of test score over iterations: The x-axis represents the number of iterations, and the y-axis shows the test score evaluated using the best prompt template found up to that iteration.

5.4.5 Running Time

Figure 5 shows the average running time per iteration during prompt optimization for the GSM8K amd MCQA dataset.

- In GSM8K, we observe that the improvement in running time is most significant at a batch size of 4, after which the gap gradually decreases. Beyond a batch size of 10, the running time saturates or increases due to the overhead caused by invalid outputs in the batch results.
- The running time results from the MCQA dataset follow similar trends to those observed with the GSM8K dataset. For smaller models like Qwen3-8B, if the model fails to follow instructions and produce output in the expected structured format, the transduction step often fails to return a valid JSON object. This leads to increased running time due to the additional overhead of the error recovery process.

6 Conclusion

We present a principled framework for agentic AI, grounded in a novel logical transduction algebra and a scalable asynchronous programming model. This framework redefines how agents interact with data through a declarative, type-driven approach, enabling robust and efficient execution across diverse tasks. Despite its strengths, the current framework has several limitations. First, it primarily focuses on type-driven transduction, which may not generalize well to tasks requiring richer contextual understanding or instruction-following behavior. Many real-world tasks involve

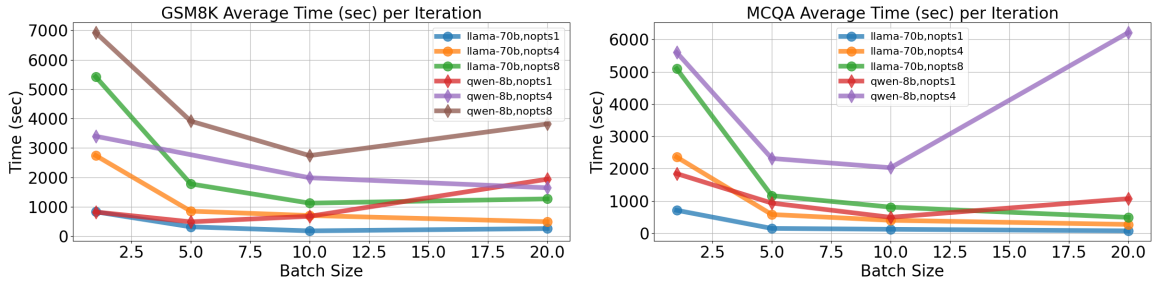

Figure 5: Average running time per iteration: The x-axis represents the batch size of the asynchronous execution, and the y-axis shows the average running time in seconds.

implicit signals or external context that go beyond type annotations. Second, the integration of tool usage remains underexplored. Future work will explore several promising directions. One is transduction that incorporates instruction or retrieval. Another is enhanced tool integration, enabling agents to invoke external tools within the transduction pipeline. Additionally, extending the framework to support interoperability with other agentic AI frameworks could unlock broader capabilities of agentic systems.

References

Deepak Bhaskar Acharya, Karthigeyan Kuppan, and B Divya. Agentic ai: Autonomous intelligence for complex goals–a comprehensive survey. *IEEe Access*, 2025.

Saaket Agashe, Jiuzhou Han, Shuyu Gan, Jiachen Yang, Ang Li, and Xin Eric Wang. Agent s: An open agentic framework that uses computers like a human. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=lIVRgt4nLv.

Ion Androutsopoulos, Graeme D Ritchie, and Peter Thanisch. Natural language interfaces to databases–an introduction. *Natural language engineering*, 1(1):29–81, 1995.

Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=hSyW5go0v8.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Edgar F Codd. A relational model of data for large shared data banks. *Communications of the ACM*, 13(6):377–387, 1970.

Christodoulos Constantinides, Dhaval C Patel, Shuxin Lin, Claudio Guerrero, SUNIL DAGAJIRAO PATIL, and Jayant Kalagnanam. FailuresensorIQ: A multi-choice QA dataset for understanding sensor relationships and failure modes. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2025. URL https://openreview.net/forum?id=9KfkMAy2ut.

CrewAIInc. CrewAI Documentation. http://docs.crewai.com/, 2025. Accessed: 27-Jan-2026.

Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.

Victor Dibia, Jingya Chen, Gagan Bansal, Suff Syed, Adam Fourney, Erkang Zhu, Chi Wang, and Saleema Amershi. Autogen studio: A no-code developer tool for building and debugging multi-agent systems. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pp. 72–79, 2024.

Yi Dong, Ronghui Mu, Gaojie Jin, Yi Qi, Jinwei Hu, Xingyu Zhao, Jie Meng, Wenjie Ruan, and Xiaowei Huang. Position: building guardrails for large language models requires systematic design. In *Proceedings of the 41st International Conference on Machine Learning*, ICML'24. JMLR.org, 2024.

Shanshan Han, Qifan Zhang, Yuhang Yao, Weizhao Jin, Zhaozhuo Xu, and Chaoyang He. Llm multi-agent systems: Challenges and open problems. *CoRR*, 2024.

Petra Heck. What about the data? a mapping study on data engineering for ai systems. In *Proceedings of the IEEE/ACM 3rd International Conference on AI Engineering-Software Engineering for AI*, pp. 43–52, 2024.

Gary G Hendrix, Earl D Sacerdoti, Daniel Sagalowicz, and Jonathan Slocum. Developing a natural language interface to complex data. *ACM Transactions on Database Systems (TODS)*, 3(2):105–147, 1978.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *International Conference on Learning Representations*, 2021. URL https://openreview.net/forum?id=d7KBjmI3GmQ.

Danielle Hopkins, Debra J Rickwood, David J Hallford, and Clare Watsford. Structured data vs. unstructured data in machine learning prediction models for suicidal behaviors: A systematic review and meta-analysis. *Frontiers in digital health*, 4:945006, 2022.

Soodeh Hosseini and Hossein Seilani. The role of agentic ai in shaping a smart future: A systematic review. *Array*, pp. 100399, 2025.

Ken Huang and Jerry Huang. Ai agent tools and frameworks. In *Agentic AI: Theories and Practices*, pp. 23–50. Springer, 2025.

Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.

Alistair EW Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Ayad Shammout, Steven Horng, Tom J Pollard, Sicheng Hao, Benjamin Moody, Brian Gow, et al. Mimic-iv, a freely accessible electronic health record dataset. *Scientific data*, 10(1):1, 2023.

Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Saiful Haq, Ashutosh Sharma, Thomas T Joshi, Hanna Moazam, Heather Miller, et al. Dspy: Compiling declarative language model calls into state-of-the-art pipelines. In *The Twelfth International Conference on Learning Representations*, 2024.

Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35:22199–22213, 2022.

LangChain. LangGraph Documentation. https://docs.langchain.com/oss/python/langgraph/overview, 2025. Accessed: 27-Jan-2026.

Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhua Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. *Advances in Neural Information Processing Systems*, 36:42330–42357, 2023.

Wenwu Li, Xiangfeng Wang, Wenhao Li, and Bo Jin. A survey of automatic prompt engineering: An optimization perspective. *arXiv preprint arXiv:2502.11560*, 2025.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems*, 36:46534–46594, 2023.

Bodhisattwa Prasad Majumder, Harshit Surana, Dhruv Agarwal, Bhavana Dalvi Mishra, Abhijeetsingh Meena, Aryan Prakhar, Tirth Vora, Tushar Khot, Ashish Sabharwal, and Peter Clark. Discoverybench: Towards data-driven discovery with large language models. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://openreview.net/forum?id=vyflgpwfJW.

Dany Moshkovich and Sergey Zeltyn. Taming uncertainty via automation: Observing, analyzing, and optimizing agentic ai systems. *arXiv preprint arXiv:2507.11277*, 2025.

Avanika Narayan, Ines Chami, Laurel Orr, and Christopher Ré. Can foundation models wrangle your data? *Proc. VLDB Endow.*, 16(4):738–746, December 2022. ISSN 2150-8097. doi: 10.14778/3574245.3574258. URL https://doi.org/10.14778/3574245.3574258.

OHDSI. *The Book of OHDSI: Observational Health Data Sciences and Informatics*. OHDSI, 2019.

Krista Opsahl-Ong, Michael Ryan, Josh Purtell, David Broman, Christopher Potts, Matei Zaharia, and Omar Khattab. Optimizing instructions and demonstrations for multi-stage language model programs. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 9340–9366, 2024.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.

Liangming Pan, Michael Saxon, Wenda Xu, Deepak Nathani, Xinyi Wang, and William Yang Wang. Automatically correcting large language models: Surveying the landscape of diverse automated correction strategies. *Transactions of the Association for Computational Linguistics*, 12:484–506, 2024.

Marcel Parciak, Brecht Vandevoort, Frank Neven, Liesbet M Peeters, and Stijn Vansummeren. Schema matching with large language models: an experimental study. *Proceedings of the VLDB Endowment. ISSN*, 2150:8097, 2024.

Rani Kurnia Putri and Muhammad Athoillah. Artificial intelligence and machine learning in digital transformation: Exploring the role of ai and ml in reshaping businesses and information systems. In *Advances in Digital Transformation-Rise of Ultra-Smart Fully Automated Cyberspace*. IntechOpen, 2024.

Pydantic. GenAI Agent Framework, the Pydantic way. https://ai.pydantic.dev/, 2025a. Accessed: 27-Jan-2026.

Pydantic. Pydantic Valiadation. https://docs.pydantic.dev/latest/, 2025b. Accessed: 27-Jan-2026.

Kiran Ramnath, Kang Zhou, Sheng Guan, Soumya Smruti Mishra, Xuan Qi, Zhengyuan Shen, Shuai Wang, Sangmin Woo, Sullam Jeoung, Yawei Wang, Haozhu Wang, Han Ding, Yuzhe Lu, Zhichao Xu, Yun Zhou, Balasubramaniam Srinivasan, Qiaojing Yan, Yueyan Chen, Haibo Ding, Panpan Xu, and Lin Lee Cheong. A systematic survey of automatic prompt optimization techniques. In Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng (eds.), *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 33078–33110, Suzhou, China, November 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.1681. URL https://aclanthology.org/2025.emnlp-main.1681/.

Gaetano Rossiello, Nandana Mihindukulasooriya, Ibrahim Abdelaziz, Mihaela Bornea, Alfio Gliozzo, Tahira Naseem, and Pavan Kapanipathi. Generative relation linking for question answering over knowledge bases. In *International Semantic Web Conference*, pp. 321–337. Springer, 2021.

Gaetano Rossiello, Md Faisal Mahbub Chowdhury, Nandana Mihindukulasooriya, Owen Cornec, and Alfio Massimiliano Gliozzo. Knowgl: Knowledge generation and linking from text. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 37, pp. 16476–16478, 2023.

Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach (4th Edition)*. Pearson, 2020. ISBN 9780134610993. URL http://aima.cs.berkeley.edu/.

Akila Sarirete, Zain Balfagih, Tayeb Brahimi, Miltiadis D Lytras, and Anna Visvizi. Artificial intelligence and machine learning research: towards digital transformation at a global scale. *Journal of Ambient Intelligence and Humanized Computing*, 13(7):3319–3321, 2022.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36:8634–8652, 2023.

Claudio Spiess, Mandana Vaziri, Louis Mandel, and Martin Hirzel. Autopdl: Automatic prompt optimization for llm agents. *arXiv preprint arXiv:2504.04365*, 2025.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. In *The Twelfth International Conference on Learning Representations*, 2023.

Qinyuan Ye, Mohamed Ahmed, Reid Pryzant, and Fereshte Khani. Prompt engineering a prompt engineer. In *Findings of the Association for Computational Linguistics: ACL 2024*, pp. 355–385, 2024.

Zhexin Zhang, Leqi Lei, Lindong Wu, Rui Sun, Yongkang Huang, Chong Long, Xiao Liu, Xuanyu Lei, Jie Tang, and Minlie Huang. Safetybench: Evaluating the safety of large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15537–15553, 2024.

Danna Zheng, Mirella Lapata, and Jeff Pan. Archer: A human-labeled text-to-sql dataset with arithmetic, commonsense and hypothetical reasoning. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 94–111, 2024.

Wenliang Zheng, Sarkar Snigdha Sarathi Das, Yusen Zhang, and Rui Zhang. Greaterprompt: A unified, customizable, and high-performing open-source toolkit for prompt optimization. *arXiv preprint arXiv:2504.03975*, 2025.

A Logical Transduction Algebra

A.1 Formalization

We presented an abridged version of the formal Logical Transduction Algebra in the main paper. Our work is closely related to relational algebra (Codd, 1970) and the MapReduce programming model (Dean & Ghemawat, 2008). This enables the composition of data transformation pipelines and supports an efficient programming model that leverages the stateless and asynchronous nature of LLM inference.

A.1.1 Algebraic Structures

Types and Agentic Structure We define types and meta-types, collectively referred to as the Agentic Structure (AG), and establish a sound algebra over the types and states within it.

Definition 1 (Types)**.** Let Θ denote the universe of all possible *types*, $\Theta = \{X, Y, Z, T, \dots\}$, where each type $T \in \Theta$ is a collection of named fields (s_i, T_{s_i}):

$$T := \{(s_1, T_{s_1}), (s_2, T_{s_2}), \dots, (s_n, T_{s_n})\},$$

with each s_i representing a string-valued slot name, and each $T_{s_i} \in \Theta$ denoting the corresponding type of that slot.

Definition 2. Given two types X and Y, we define standard set operations component-wise:

$$\begin{aligned}
X \cup Y &= \{(s_i, T_{s_i}) \mid (s_i, T_{s_i}) \in X \text{ or } (s_i, T_{s_i}) \in Y\},\\
X \cap Y &= \{(s_i, T_{s_i}) \mid (s_i, T_{s_i}) \in X \text{ and } (s_i, T_{s_i}) \in Y\},\\
X \setminus Y &= \{(s_i, T_{s_i}) \mid (s_i, T_{s_i}) \in X \text{ and } (s_i, T_{s_i}) \notin Y\},\\
X \times Y &= \{((s_i, T_{s_i}), (s_j, T_{s_j})) \mid (s_i, T_{s_i}) \in X, (s_j, T_{s_j}) \in Y\}
\end{aligned}$$

Definition 3 (Agentic Structure AG)**.** An *Agentic structure* AG is a meta-type that bundles a type schema s_{atype} [4] and a corresponding list of instances, referred to as **states** s_{states}:

$$AG := \begin{Bmatrix} s_{\text{atype}} : \Theta, \\ s_{\text{states}} : \text{List}[\Theta] \end{Bmatrix}$$

Notation conventions: Types are denoted by uppercase letters. Instances of types are denoted by lowercase letters, with $t : T$ indicating that t is an instance of type T. Lists are written in boldface, so $\mathbf{t} : T$ represents a list of instances of type T. We use the shorthand $AG[X]$ to denote an Agentic structure with $s_{\text{atype}} = X$. A boldface lowercase symbol, such as $\mathbf{x} = AG[X]$, represents an instance of $AG[X]$. We also overload the notation to access the list of states: $x_i = \mathbf{x}[i] = \mathbf{x}.s_{\text{states}}[i]$ refers to the i-th state of the Agentic instance $\mathbf{x}$.

In Logical Transduction Algebra, we focus on structured data and its transformation around agents encapsulating LLMs. The algebraic structure of composing two Agentic instances of the same type can be shown as follows.

Proposition 1 (Monoid of Agentic Instances)**.** Let $AG[X]$ be an Agentic structure and let ξ be the set of all instances of $AG[X]$. Define a binary operation $\circ$ on ξ such that for any $\mathbf{x}_1, \mathbf{x}_2 \in \xi$, their composition $\mathbf{x} = \mathbf{x}_1 \circ \mathbf{x}_2$ is an Agentic instance whose state list is the concatenation: $\mathbf{x}.s_{\text{states}} := \mathbf{x}_1.s_{\text{states}} \circ \mathbf{x}_2.s_{\text{states}}$. Then, the pair $(\xi, \circ)$ forms a *monoid*, where the identity element is the Agentic instance with an empty state list: $\mathbf{e}.s_{\text{states}} := []$.

Proof. We verify the three monoid properties:

Closure: Let $\mathbf{x}_1, \mathbf{x}_2 \in \xi$. Then $\mathbf{x} = \mathbf{x}_1 \circ \mathbf{x}_2$ has a state list formed by concatenating two valid state lists, which is itself valid. Hence, $\mathbf{x} \in \xi$.

Associativity: For any $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3 \in \xi$,

$$\begin{aligned}
((\mathbf{x}_1 \circ \mathbf{x}_2) \circ \mathbf{x}_3).s_{\text{states}} &= (\mathbf{x}_1.s_{\text{states}} \circ \mathbf{x}_2.s_{\text{states}}) \circ \mathbf{x}_3.s_{\text{states}}\\
&= \mathbf{x}_1.s_{\text{states}} \circ (\mathbf{x}_2.s_{\text{states}} \circ \mathbf{x}_3.s_{\text{states}})\\
&= (\mathbf{x}_1 \circ (\mathbf{x}_2 \circ \mathbf{x}_3)).s_{\text{states}}.
\end{aligned}$$

[4] Given two types X and Y, the standard set operations such as union, intersection, complement, and product can be defined component-wise.

Identity: Let $\mathbf{e} \in \xi$ be the Agentic instance with an empty state list. Then for any $\mathbf{x} \in \xi$,

$$(\mathbf{e} \circ \mathbf{x}).s_{\text{states}} = [] \circ \mathbf{x}.s_{\text{states}} = \mathbf{x}.s_{\text{states}},$$
$$(\mathbf{x} \circ \mathbf{e}).s_{\text{states}} = \mathbf{x}.s_{\text{states}} \circ [] = \mathbf{x}.s_{\text{states}}.$$

Thus, $(\xi, \circ)$ satisfies closure, associativity, and identity, and is therefore a monoid. □

The standard operators follow standard algebraic principles such as the product, $\mathbf{x} \times \mathbf{y}$ for $\mathbf{x} \in AG[X]$ and $\mathbf{y} \in AG[Y]$, the equivalence, $\mathbf{x_1} \sim \mathbf{x_2}$ for $\mathbf{x_1}, \mathbf{x_2} \in AG[X]$, and the quotient, $\mathbf{z}/\mathbf{y}$ for $\mathbf{z} \in AG[X \times Y]$ and $\mathbf{y} \in AG[Y]$.

Product of the Agentic Structure Next, we define the product of Agentic structures, a construction that plays a foundational role in modeling and executing complex, multi-dimensional data workflows. By combining two Agentic structures into a single product structure, we can represent composite types—such as paired entities, coupled processes, or input-output relationships—within a unified algebraic framework. This formulation ensures that operations applied to states remain well-defined, type-safe, and composable, preserving the monoidal properties of each component. The product structure is especially valuable in scenarios involving joint reasoning, parallel transformations, or structured transductions across heterogeneous data streams.

Definition 4 (Product of Agentic Structures)**.** Let $AG[X]$ and $AG[Y]$ be two Agentic structures defined over distinct types X and Y, respectively. We define their product as a new Agentic structure $AG[T]$, where the type T is the Cartesian product of the two types:

$$T : X \times Y.$$

Given instances $\mathbf{x} : AG[X]$ and $\mathbf{y} : AG[Y]$, we define an instance $\mathbf{t} : AG[T]$ such that:

$$\mathbf{t}.s_{\text{states}} = (\mathbf{x}.s_{\text{states}}, \mathbf{y}.s_{\text{states}}),$$

i.e., the state list of $\mathbf{t}$ is the pair of state lists from $\mathbf{x}$ and $\mathbf{y}$.

Proposition 2 (Monoid of Agentic Product)**.** Let ξ_X and ξ_Y be the set of all instances of $AG[X]$ and $AG[Y]$, respectively, and let ξ_T be the set of all instances of $AG[T]$.

Define a binary operation $\circ$ on ξ_T as follows:

$$(\mathbf{x}_1, \mathbf{y}_1) \circ (\mathbf{x}_2, \mathbf{y}_2) := (\mathbf{x}_1 \circ \mathbf{x}_2, \mathbf{y}_1 \circ \mathbf{y}_2),$$

where $\circ$ on each component denotes concatenation of state lists:

$$(\mathbf{x}_1 \circ \mathbf{x}_2).s_{\text{states}} := \mathbf{x}_1.s_{\text{states}} \circ \mathbf{x}_2.s_{\text{states}},$$

and similarly for $\mathbf{y}_1 \circ \mathbf{y}_2$.

Then, the structure $(\xi_T, \circ)$ forms a *monoid*, with the identity element given by the pair of Agentic instances with empty state lists:

$$\mathbf{e}_T := (\mathbf{e}_X, \mathbf{e}_Y), \quad \text{where } \mathbf{e}_X.s_{\text{states}} = [], \quad \mathbf{e}_Y.s_{\text{states}} = [].$$

Proof. We verify the three monoid properties for $(\xi_T, \circ)$.

Closure: Let $(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2) \in \xi_T$. Then their composition is:

$$(\mathbf{x}_1 \circ \mathbf{x}_2, \mathbf{y}_1 \circ \mathbf{y}_2),$$

where each component is a valid Agentic instance due to closure in $(\xi_X, \circ)$ and $(\xi_Y, \circ)$. Hence, the result is in ξ_T.

Associativity: Let $(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2), (\mathbf{x}_3, \mathbf{y}_3) \in \xi_T$. Then:

$$\begin{aligned}((\mathbf{x}_1, \mathbf{y}_1) \circ (\mathbf{x}_2, \mathbf{y}_2)) \circ (\mathbf{x}_3, \mathbf{y}_3) &= (\mathbf{x}_1 \circ \mathbf{x}_2, \mathbf{y}_1 \circ \mathbf{y}_2) \circ (\mathbf{x}_3, \mathbf{y}_3) \\ &= ((\mathbf{x}_1 \circ \mathbf{x}_2) \circ \mathbf{x}_3, (\mathbf{y}_1 \circ \mathbf{y}_2) \circ \mathbf{y}_3) \\ &= (\mathbf{x}_1 \circ (\mathbf{x}_2 \circ \mathbf{x}_3), \mathbf{y}_1 \circ (\mathbf{y}_2 \circ \mathbf{y}_3)) \\ &= (\mathbf{x}_1, \mathbf{y}_1) \circ ((\mathbf{x}_2, \mathbf{y}_2) \circ (\mathbf{x}_3, \mathbf{y}_3)),\end{aligned}$$

using associativity in each component.

Identity: Let $\mathbf{e}_T := (\mathbf{e}_X, \mathbf{e}_Y)$, where $\mathbf{e}_X$ and $\mathbf{e}_Y$ are identity elements in ξ_X and ξ_Y, respectively. Then for any $(\mathbf{x}, \mathbf{y}) \in \xi_T$:

$$(\mathbf{e}_X, \mathbf{e}_Y) \circ (\mathbf{x}, \mathbf{y}) = (\mathbf{e}_X \circ \mathbf{x}, \mathbf{e}_Y \circ \mathbf{y}) = (\mathbf{x}, \mathbf{y}),$$
$$(\mathbf{x}, \mathbf{y}) \circ (\mathbf{e}_X, \mathbf{e}_Y) = (\mathbf{x} \circ \mathbf{e}_X, \mathbf{y} \circ \mathbf{e}_Y) = (\mathbf{x}, \mathbf{y}).$$

Hence, $\mathbf{e}_T$ is the identity element. □

Quotient of the Agentic Structure To complement the expressiveness of product structures, the quotient of Agentic structures provides a principled mechanism for abstraction and generalization. By defining an equivalence relation over Agentic instances—such as grouping together states that differ only in irrelevant or redundant dimensions—we can collapse fine-grained distinctions into coarser, semantically meaningful categories. This is especially useful in scenarios involving behavioral equivalence, or clustering of similar agentic behaviors. The quotient structure enables reasoning at a higher level of abstraction while preserving the algebraic properties of the original system. In distributed settings, it supports compression, deduplication, and aggregation of stateful computations.

Definition 5 (Equivalence Relation on Agentic Instances)**.** Let ξ_X be the set of Agentic instances over type X.

An equivalence relation $\sim$ on ξ_X is defined by a relation $\mathcal{R}$ on state lists $\mathbf{s} : X$ such that for any $\mathbf{x}, \mathbf{y} \in \xi_X$,

$$\mathbf{x} \sim \mathbf{y} \iff \mathcal{R}(\mathbf{x}.s_{\text{states}}, \mathbf{y}.s_{\text{states}}),$$

where $\mathcal{R}$ satisfies the following properties:

- *Reflexivity:* $\mathcal{R}(\mathbf{s}, \mathbf{s})$ for all state lists $\mathbf{s}$.
- *Symmetry:* If $\mathcal{R}(\mathbf{s}_1, \mathbf{s}_2)$, then $\mathcal{R}(\mathbf{s}_2, \mathbf{s}_1)$.
- *Transitivity:* If $\mathcal{R}(\mathbf{s}_1, \mathbf{s}_2)$ and $\mathcal{R}(\mathbf{s}_2, \mathbf{s}_3)$, then $\mathcal{R}(\mathbf{s}_1, \mathbf{s}_3)$.

The specific form of $\mathcal{R}$ depends on the semantics of the Agentic structure. A common choice is *statewise equivalence*, defined below.

Definition 6 (Statewise Equivalence of Agentic Instances)**.** Let ξ_X be the set of Agentic instances over type X.

Define an equivalence relation $\sim$ on ξ_X such that for any $\mathbf{x}, \mathbf{y} \in \xi_X$,

$$\mathbf{x} \sim \mathbf{y} \iff \mathbf{x}.s_{\text{states}} \equiv \mathbf{y}.s_{\text{states}},$$

where $\equiv$ denotes elementwise equivalence of state lists. That is,

$$\mathbf{x}.s_{\text{states}} = [x_1, x_2, \ldots, x_n], \quad \mathbf{y}.s_{\text{states}} = [y_1, y_2, \ldots, y_n],$$

and for all $i = 1, \ldots, n$, we have $x_i \approx y_i$ under a given equivalence relation $\approx$ on X.

The relation $\approx$ on X may be defined in various ways, such as:

- *Syntactic equality:* $x_i = y_i$.
- *Observational equivalence:* $f(x_i) = f(y_i)$ for some observable function $f : X \to O$.
- *Abstract equivalence:* x_i and y_i belong to the same equivalence class under a domain-specific partition of X.

This relation groups Agentic instances whose state trajectories are equivalent up to the equivalence of individual states.

Definition 7 (Quotient of Agentic Structure)**.** Let $AG[X]$ be an Agentic structure over type X, and let $\sim$ be an equivalence relation on the set of Agentic instances ξ_X.

The *quotient Agentic structure*, denoted $AG[X/\sim]$, is defined as follows:

- The type $X/\sim$ is the set of equivalence classes of X under the induced relation $\approx$ on individual states.
- The set of instances $\xi_{X/\sim}$ consists of equivalence classes $\langle \mathbf{x} \rangle$ of Agentic instances $\mathbf{x} \in \xi_X$, where

$$\langle \mathbf{x} \rangle := \{\mathbf{y} \in \xi_X \mid \mathbf{y} \sim \mathbf{x}\}.$$

- The state list of an equivalence class $\langle \mathbf{x} \rangle$ is defined as:

$$\langle \mathbf{x} \rangle.s_{\text{states}} := \{\mathbf{y}.s_{\text{states}} \mid \mathbf{y} \sim \mathbf{x}\}.$$

This structure abstracts over individual Agentic instances by identifying those whose states are equivalent.

Proposition 3 (Monoid Structure on Quotient Agentic Structure)**.** Let $(\xi_X, \circ)$ be a monoid of Agentic instances over type X, and let $\sim$ be a congruence relation on ξ_X, i.e., for all $\mathbf{x}_1 \sim \mathbf{x}_2$ and $\mathbf{y}_1 \sim \mathbf{y}_2$, we have:

$$\mathbf{x}_1 \circ \mathbf{y}_1 \sim \mathbf{x}_2 \circ \mathbf{y}_2.$$

Then, the quotient structure $(\xi_X/\sim, \circ)$ forms a monoid, where:

- Elements are equivalence classes $\langle \mathbf{x} \rangle$.
- The operation is defined by:

$$\langle \mathbf{x} \rangle \circ \langle \mathbf{y} \rangle := \langle \mathbf{x} \circ \mathbf{y} \rangle.$$

- The identity is $\langle \mathbf{e} \rangle$, where $\mathbf{e}$ is the identity in $(\xi_X, \circ)$.

Proof. We verify the monoid properties on the quotient structure:

Well-definedness: If $\mathbf{x}_1 \sim \mathbf{x}_2$ and $\mathbf{y}_1 \sim \mathbf{y}_2$, then by congruence:

$$\mathbf{x}_1 \circ \mathbf{y}_1 \sim \mathbf{x}_2 \circ \mathbf{y}_2,$$

so $\langle \mathbf{x}_1 \circ \mathbf{y}_1 \rangle = \langle \mathbf{x}_2 \circ \mathbf{y}_2 \rangle$.

Associativity: Follows from associativity in ξ_X:

$$\langle \mathbf{x} \rangle \circ (\langle \mathbf{y} \rangle \circ \langle \mathbf{z} \rangle) = \langle \mathbf{x} \circ (\mathbf{y} \circ \mathbf{z}) \rangle = \langle (\mathbf{x} \circ \mathbf{y}) \circ \mathbf{z} \rangle = (\langle \mathbf{x} \rangle \circ \langle \mathbf{y} \rangle) \circ \langle \mathbf{z} \rangle.$$

Identity: For any $\langle \mathbf{x} \rangle$,

$$\langle \mathbf{e} \rangle \circ \langle \mathbf{x} \rangle = \langle \mathbf{e} \circ \mathbf{x} \rangle = \langle \mathbf{x} \rangle, \quad \langle \mathbf{x} \rangle \circ \langle \mathbf{e} \rangle = \langle \mathbf{x} \circ \mathbf{e} \rangle = \langle \mathbf{x} \rangle.$$

□

Example 1 (Quotient of a Product Agentic Structure)**.** Let $AG[X]$ and $AG[Y]$ be Agentic structures over types X and Y, respectively. Their product $AG[T]$ is defined by

$$s_{\text{atype}} = X \times Y, \quad s_{\text{states}} = (\mathbf{x}.s_{\text{states}}, \mathbf{y}.s_{\text{states}})$$

for instances $\mathbf{x} \in AG[X]$ and $\mathbf{y} \in AG[Y]$.

We define an equivalence relation $\sim$ on ξ_T such that for any $\mathbf{t}_1, \mathbf{t}_2 \in \xi_T$,

$$\mathbf{t}_1 \sim \mathbf{t}_2 \iff \forall i,\ x_i^{(1)} \approx x_i^{(2)},$$

where $x_i^{(1)}$ and $x_i^{(2)}$ are the first components of the i-th state in $\mathbf{t}_1$ and $\mathbf{t}_2$, respectively, and $\approx$ is an equivalence relation on X.

As a concrete example, let the types and the equivalence $\approx$ be defined as:

- $X = \{\text{Red}, \text{Green}, \text{Blue}\}$ (colors),

- $Y = \{\text{Circle}, \text{Square}\}$ (shapes),
- Define $\approx$ on X by:

$$\text{Red} \approx \text{Green}, \quad \text{Blue} \not\approx \text{Red}, \quad \text{Blue} \not\approx \text{Green}.$$

Consider two Agentic instances:

$$\mathbf{t}_1.s_{\text{states}} = [(\text{Red}, \text{Circle}), (\text{Green}, \text{Square})],$$

$$\mathbf{t}_2.s_{\text{states}} = [(\text{Green}, \text{Circle}), (\text{Red}, \text{Square})].$$

Then $\mathbf{t}_1 \sim \mathbf{t}_2$ because:

$$\text{Red} \approx \text{Green}, \quad \text{Green} \approx \text{Red}.$$

Note that the shape components (second elements) are not constrained by the equivalence relation.

The quotient Agentic structure $AG[T/\sim]$ consists of equivalence classes $\langle \mathbf{t} \rangle$ of Agentic instances under $\sim$, where:

$$\langle \mathbf{t} \rangle := \{\mathbf{t}' \in \xi_T \mid \mathbf{t}' \sim \mathbf{t}\}.$$

This structure abstracts over differences in the first component of the state tuples according to $\approx$, while preserving the full state list structure.

A.1.2 The Transduction Operator

Equipped with Agentic structures that form a monoid, we obtain a sound abstraction for composing the data workflows in a functional programming style. This foundation enables the introduction of the *logical transduction operator*, which utilizes LLMs as transductive inference engines. We now define a series of logical transduction operators, organized by increasing levels of complexity. These operators are designed with explicit consideration of types and Agentic structures.

Transduction Operator Overloading

Definition 8 (Transduction)**.** Given an information object x and a target Agentic structure $AG[Y]$, the *transduction* of x into $AG[Y]$ is defined as:

$$\mathbf{y} := Y \ll x = \bigcup_i (y_i, T_{s_i}),$$

where each $y_i \in T_{s_i}$ is a value assigned to slot s_i of type T_{s_i}, logically inferred from x. Here, the operator $\ll$ denotes a logical transduction process, implemented via an LLM, that maps x to a structured output conforming to the type Y.

To support the definition of logical transduction operators, we introduce a generic function that renders typed objects into textual representations suitable for LLM input.

Definition 9 (Prompt Function)**.** Given a type $T \in \Theta$, a *prompt function* P is a mapping that renders a list of states $\mathbf{t} : T$ into an information object, leveraging the string-valued slot names associated with T. Formally, $P : \text{List}[T] \to \texttt{str}$.

Prompt functions serve as a bridge between structured data and natural language, enabling logical transduction operators to interface with LLMs by converting typed instances into semantically meaningful prompts.

We now define two specific forms of logical transduction: zero-shot and few-shot.

Definition 10 (Zero-Shot Logical Transduction)**.** Let $\mathbf{x} = AG[X]$ and $\mathbf{y} = AG[Y]$ be Agentic structures over types X and Y, respectively. A *zero-shot logical transduction* from $\mathbf{x}$ to $\mathbf{y}$ is defined component-wise as:

$$\mathbf{y}[i] = Y \ll P(\mathbf{x}[i]),$$

where $P : X \to \texttt{str}$ is a prompt function that renders each instance $\mathbf{x}[i]$ into a textual prompt.

Next, we also show more overloaded transduction operators in the case of the few-shot transduction, $\mathbf{y}[i] = Y \ll (P(\mathbf{x}[i]) \oplus FS(\mathbf{x}, \mathbf{y}))$ with a few-shot function $FS(\mathbf{x}, \mathbf{y}) := P((\mathbf{x}', \mathbf{y}'))$, and a syntactic sugar such as self-transduction.

Definition 11 (Few-Shot Logical Transduction)**.** Let $\mathbf{x} = AG[X]$ and $\mathbf{y} = AG[Y]$ be Agentic structures over types X and Y, respectively. A *few-shot logical transduction* from $\mathbf{x}$ to $\mathbf{y}$ is defined for all indices i such that $\mathbf{y}[i] = \emptyset$ as:

$$\mathbf{y}[i] = Y \ll (P(\mathbf{x}[i]) \oplus FS(\mathbf{x}, \mathbf{y}))\,,$$

where:

- $P : X \to \texttt{str}$ is a prompt function that renders an instance $\mathbf{x}[i]$ into a textual prompt.
- $\oplus$ denotes prompt concatenation.
- $FS(\mathbf{x}, \mathbf{y})$ is the *few-shot context*, defined as:

$$FS(\mathbf{x}, \mathbf{y}) := P((\mathbf{x}', \mathbf{y}'))\,,$$

  where $(\mathbf{x}', \mathbf{y}')$ is the projection of $(\mathbf{x}, \mathbf{y})$ onto the subset of indices for which $\mathbf{y}[i] \neq \emptyset$.

Finally, we introduce self-transduction as syntactic sugar within the programming model for logical transductions.

Definition 12 (Self Transduction)**.** Let $\mathbf{x} \in AG[X]$ be an Agentic structure, and let $Y, Z \subset X$ be two disjoint subsets of types. A *self transduction* is a function that produces a modified Agentic structure $\mathbf{x}' \in AG[X]$, defined as:

$$\mathbf{x}' = \mathbf{x} \ll_{Y,Z} := \mathbf{x} \cup (\mathbf{x}[Y] \ll \mathbf{x}[Z])\,,$$

where $\mathbf{x}[Y]$ denotes the *rebind operator*, which extracts an $AG[Y]$ from $\mathbf{x}$ by retaining only the slots in Y that overlap with X.

Properties of Transduction Operator Next, we formalize key properties of the transduction operator. These properties are foundational for enabling scalable, parallel, and composable computation.

Proposition 4 (Conditional Determinism)**.** Let σ denote a fixed transduction context, which may include components such as a few-shot context, additional instructions, external tools, or memory. Let the LLM configuration—comprising model weights, temperature, and decoding strategy—also be fixed. Then, for any input x_i, the transduction $y_i := Y \ll x_i$ is deterministic.

Proof. When the model parameters and transduction context σ are fixed, and the LLM is configured with deterministic settings (e.g., temperature set to zero and caching enabled), the output y_i is uniquely determined by the input x_i and the context σ. Therefore, the transduction process is deterministic under these conditions. □

Proposition 5 (Statelessness)**.** Logical transduction operators are stateless. The output of a transduction $y_i := Y \ll x_i$ depends only on x_i and the transduction context, and not on any prior or future transductions.

Proof. By definition, the transduction operator $\ll$ does not rely on conversational memory or sequential state. Each y_i is computed independently from x_i and σ, enabling asynchronous evaluation. □

Proposition 6 (Compositionality)**.** Let $AG[X]$, $AG[Y]$, and $AG[Z]$ be Agentic structures over types X, Y, and Z, respectively. Suppose $\mathbf{y}' = \mathbf{y} \ll \mathbf{x}$, and $\mathbf{z}' = \mathbf{z} \ll \mathbf{y}'$. Then the composite transduction holds, $\mathbf{z}' = \mathbf{z} \ll \mathbf{y} \ll \mathbf{x}$.

Proof. Each transduction step applies $\ll$ component-wise to the state list of the input Agentic structure. Since the output of $\mathbf{y} \ll \mathbf{x}$ is an Agentic structure $AG[Y]$, it can be used as input to the next transduction. Thus, the composition is well-defined and yields $\mathbf{z}'$. □

Implications for Distributed and Concurrent Computing These properties make the transduction operator $\ll$ particularly well-suited for distributed and concurrent computing paradigms. The *Conditional Determinism* ensures reproducibility and traceability in distributed pipelines, the *Statelessness* enables parallel execution, allowing transductions to be mapped across shards of data without coordination or shared state, and the *Compositionality* supports modular pipeline construction, akin to functional composition in MapReduce, where intermediate Agentic structures can be chained and reused.

A.1.3 Asynchronous MapReduce

The programming model of the `Agentics` supports asynchronous execution of mapping and reduction operations over Agentic structures, enabling scalable and composable data workflows. We formalize these operations as `amap` and `areduce`, which extend the MapReduce by Dean & Ghemawat (2008).

Definition 13 (Asynchronous Map (`amap`))**.** Let $AG[X]$ be an Agentic structure over type X, and let $f : X \to \text{List}[Y]$ be an asynchronous mapping function. Then the asynchronous map operator is defined as:

$$\texttt{amap} : (AG[X], f) \to AG[Y],$$

where the output Agentic structure $\mathbf{y} = \texttt{amap}(\mathbf{x}, f)$ satisfies: $\mathbf{y}.s_{\text{states}} = \bigcup_i f(x_i)$, and the union preserving the original order of inputs.

The function f may return an empty list by removing x_i from the output acting as a filter, map each x_i to a single output acting as a transformer, or map each x_i to multiple outputs acting as fan-out. Note that `amap` operator is executed asynchronously across all input states, enabling parallelism and scalability in distributed environments.

The `amap` operator is executed asynchronously across all input states, enabling parallelism and scalability in distributed environments.

Definition 14 (Asynchronous Reduce (`areduce`))**.** Let $AG[X]$ be an Agentic structure and let $f : \text{List}[X] \to Y$ be an asynchronous reduction function. Then the asynchronous reduce operator is defined as:

$$\texttt{areduce} : (AG[X], f) \to AG[Y],$$

where the output Agentic structure $\mathbf{y} = \texttt{areduce}(\mathbf{x}, f)$ satisfies: $\mathbf{y}.s_{\text{states}} = f(\mathbf{x})$.

Unlike `amap`, which applies f to each state individually, `areduce` applies f to the entire states $\mathbf{x}$ at once. This is useful for summarization or aggregation, such as generating a report or computing statistics over the full dataset. Since LLMs have limited context windows, applying `areduce` to a large dataset may be intractable. In such cases, scalable strategies such as hierarchical or batched reduction can be employed by applying `areduce` to random subsets and merging the results.

Composability with Logical Transduction `amap` and `areduce` can be composed with the logical transduction operator $\ll$ to build expressive and modular workflows. For example,

$$\mathbf{y} = \texttt{amap}(\mathbf{x},\, x \mapsto Y \ll x), \quad \mathbf{z} = \texttt{areduce}(\mathbf{y},\, f),$$

where the transduction $Y \ll x$ is embedded within the mapping function. As we can see, Agentic structure enables distributed, asynchronous, and semantically typed computation over structured data.

B Computing Infrastructure

In the experiments, we benchmark open-weight instruct tuned models ranging from larger or smaller parameter version of GPT-OSS-120B, Llama-3.3, Llama-4, Qwen-3, and Mistral. For the experiment that measures the running time, we host LLMs in local vLLM server with four A-100-80GB GPUs for Llama-3-3-70B model, and one A-100-80GB GPU for other 8B parameter models.

- In the Text-to-SQL experiment, we used GPT-OSS-120B, Llama-3.3-70B, Mistral-Large, and Llama-4-Maverick-17B models. These models were also run in a cloud computing environment.

- In the schema matching and data imputation experiments, we used GPT-OSS-120B, and Llama-4-Maverick-17B models. These models were also run in a cloud computing environment
- In the Domain Specific MCQA experiments, we used instruction-tuned models such as Qwen3-8B, Llama-3.3-70B, Mistral-Large, and Llama-3-405B. We locally hosted the Qwen3-8B model to measure running time, while the other three models were used in a cloud computing environment.
- In automatic prompt optimization experiment, we locally hosted Qwen3-8B and Llama-3.3-70B models.

The following shows the parameters for hosing Qwen3-8B and Llama-3.3-70B models with vLLM.

```
GPUS=4
CPUS=16
MEM=200GB
MODEL="meta-llama/Llama-3.3-70B-Instruct"
LEN=16000
  vllm serve ${MODEL} \
  --max-model-len ${LEN} \
  --tensor-parallel-size ${GPUS} \
  --gpu-memory-utilization 0.9

GPUS=1
CPUS=8
MEM=64GB
LEN=8000
MODEL="Qwen/Qwen3-8B"
vllm serve ${MODEL} \
  --max-model-len ${LEN} \
  --tensor-parallel-size ${GPUS} \
  --gpu-memory-utilization 0.9
```

Listing 1: vLLM parameters and computing resources

C Design Patterns and Use Cases

The `Agentics` framework provides a concrete realization of the logical transduction algebra. We elaborate on various design patterns and use cases for domain-specific multi-choice QA, text-to-SQL pipelines, clustering, and automatic prompt optimization. This section demonstrates how the `Agentics` framework supports a wide range of generative structured data workflows through reusable design patterns. Each use case highlights a different aspect of the `Agentics` programming model, showcasing its flexibility, scalability, and composability.

C.1 Semantic Parsing Text-to-SQL

Text-to-SQL is an essential task for broadening the accessibility of structured data interaction, allowing users to query databases without needing to understand the underlying decisions made by data engineers. Loosely considered a translation task, questions are posed to a database and first translated into SQL queries before being executed and answers retrieved and answers retrieved. In practice, this task involves multiple stages of reasoning, whereby one has to interact with the schema of the structured data, as well as understand and decompose the question into its constituent parts.

C.1.1 Data Models

`Agentics` supports this workflow by chaining multiple transduction steps and integrating them with traditional Python logic. First, let's consider setting up the Pydantic types with the required task information.

```
class Text2SQLTask(BaseModel):
    question: str = Field(description="The input natural language question.")
    ddl: str = Field(description="The database schema in DDL format (e.g., CREATE TABLE statements).")
    sql_query: Optional[str] = Field(description="The SQL query to be generated from the question.")
    execution_result: Optional[List[Dict[str, str]]] = Field(description="The resulting table from executing
        the SQL query.")
```

Listing 2: Data model for the simplified Text2SQL task

From the above Pydantic types, the simplified task is to map the question and DDL to the sql_query. However, we can break down this complex operation into declarative data modeling steps to improve task performance, including:

- enriching the database so that there are additional fields like description of the schema and business descriptions
- decompose the user question into constituent parts input that can be also a part of optimization!
- optimize the prompt template using the final input fields.

```python
class Text2SQLTask(BaseModel):
    question: str = Field(description="The input natural language question.")
    ddl: str = Field(description="The database schema in DDL format (e.g., CREATE TABLE statements).")
    enrichment: Optional[DB] = Field(description="Additional database enrichments  that applies to all problem instances")
    sql_query: Optional[str] = Field(description="The SQL query to be generated from the question.")
    execution_result: Optional[List[Dict[str, str]]] = Field(description="The resulting table from executing the SQL query.")

class DB(BaseModel):
    description: Optional[str] = Field(description="A Description of the business purpose of the db, what use cases it is good for how what type of information it contain")
    keywords: Optional[list[str]] = Field(description="A list of keywords describing the content of the database. Produce Keywords that are: Domain-Relevant: Reflects the thematic area (e.g., education, healthcare, finance). Purpose-Oriented: Indicates the type of insights the database supports (e.g., performance tracking, demographic analysis). Unambiguous: Avoids generic or overly broad terms. Interoperable: Aligns with standard taxonomies when possible (e.g., DataCite or UNSDG classification). Examples of Strong Keywords: student_outcomes, climate_metrics, financial_forecasting, public_health_indicators, supply_chain_kpis")
    few_shots: Optional[list[QuestionSQLPair]] = Field(description="A selection of the generated question-sql pair to be used as examples of how to generate a sql from a question.")
    subquestions: Optional[list[str]] = Field(description="a list of subquestions inside question")
    schema_link: Optional[str] = Field(description="the output of the schema linker in the form of DDL script showing relevant table, column, and values for the question")

class QuestionSQLPair(BaseModel):
    question: Optional[str]
    sql_query: Optional[str]
```

Listing 3: Data model for the compositional Text2SQL task

C.1.2 Meta-Prompts and Prompt Templates

We define the initial prompt that performs the main text-to-SQL task. The following shows the prompt template used as input to the experiment, which may be optionally modified by the automated prompt optimization flag.

```
prompt_template = "
Translate the input natural language **question** to a valid SQLite query that can be executed on the following database in **dbs**.
Do your best to apply the following rules when generating SQL.
- Cleary understand the **question** and given database description in **dbs**.
- Database description is given in the **dbs** with the following fields.
- The database description under **dbs** contains DDL scripts or natural language description of the database, tables, columns, and values.
- Only SELECT statements are allowed, do not produce any DDL or DML.
- When writing `SELECT <column>`, only include the columns specifically mentioned in the question.
- Use **evidence** to find correct column names or the values of the columns or other expressions.
- If you see 'None' or `None` in the [Value examples] for a column, prioritize using `JOIN <table>` or `WHERE <column> IS NOT NULL` to handle potential missing data effectively.
- Use `WHERE <column> IS NOT NULL` in `WHERE` if you are sorting with `<column>`.
- Use alias in `SELECT` is consistently in the expressions.
- Use `WHERE <alias> IS NOT NULL` in `WHERE` if you are sorting with `<alias>`.

Input is provided under SOURCE.
"{input_spec_str}"

Generate Output as JSON decodable format
"{{"
generated_sql_query: a valid SQLite query that translates nautral language question
"}}"\n
"
```

Listing 4: The Prompt Template and Meta-Prompt for Automatic Prompt Optimization

C.1.3 Main Algorithm

Next, we show how to define the compositional text-to-SQL pipeline in `Agentics`. We see that the entire pipeline can be implemented through compositions of logical transductions defined over the data models.

```
text2sql("enrichment.keywords") << text2sql("enrichment.description") << text2sql("ddl_schema")
text2sql("few_shots.question", "few_shots.sql_query") << text2sql("ddl_schema", "enrichment.description")  # synthetic pair
text2sql("few_shots") = text2sql("few_shots") + k_shot  # additional augmentations
text2sql("few_shots") = text2sql.filter(valid_sql, "few_shots")  # executes sql
text2sql("enrichment.subquestions") << text2sql("question")
text2sql("enrichment.linked_schema") << text2sql("enrichment.ddl_schema", "enrichment.subquestions", "question")
input_fields = ["question", "enrichment"]
text2sql.instructions = optimize(prompt_template, input_fields)
text2sql("sql_query") << text2sql(*input_fields)
text2sql("execution_result") = text2sql.amap(execute_sql_query)
```

Listing 5: Pseudo-code for the Compositional Text-to-SQL Workflow

C.2 Domain-Specific Multiple Choice Question Answering

Domain-specific Multiple Choice Question Answering (MCQA) tasks present unique challenges for LLMs, particularly when grounded in technical domains that are unfamiliar or underrepresented in pretraining corpora.

C.2.1 FailureSensorIQ Benchmark

In this subsection, we demonstrate how the `Agentics` framework supports structured reasoning and robust performance on the *FailureSensorIQ* benchmark—a dataset designed to evaluate LLMs' understanding of failure modes and sensor relationships in Industry 4.0 (Constantinides et al., 2025).

Unlike widely used QA datasets such as MMLU (Hendrycks et al., 2021), FailureSensorIQ introduces a novel domain with no prior exposure to the models under evaluation. This makes it a strong testbed for assessing generalization and reasoning capabilities in high-stakes, real-world industrial contexts. The benchmark includes 8,296 questions across 10 assets, with both single- and multi-answer formats. Despite the presence of strong reasoning models, the best-performing `openai-o1` achieves only 60.4 precent accuracy on single-answer questions, underscoring the dataset's difficulty.

C.2.2 Schema-Guided LLM Reasoning

The `Agentics` approach to MCQA leverages self-transduction and schema-driven prompting using `Pydantic` models. This structured prompt format contrasts with the loosely formatted natural language prompts used in baseline evaluations. By explicitly encoding the input-output schema (e.g., JSON fields for question, options, and selected answers), `Agentics` reduces decoding errors and enforces type safety. This is particularly beneficial in multi-answer settings, where ambiguity in output formatting can lead to evaluation mismatches and degraded performance.

We observe that this structured prompting pattern not only improves accuracy but also enhances robustness to perturbations. For example, when question phrasing or distractor options are altered, schema-constrained decoding helps maintain consistent model behavior. This suggests that `Agentics`' structured approach offers a degree of perturbation resilience, addressing one of the key weaknesses identified in the original FailureSensorIQ benchmark.

In addition to accuracy improvements, the `Agentics` framework enables parallel batch execution of MCQA tasks, significantly reducing inference time. This is achieved through the amap operation, which distributes structured prompts across multiple model invocations. Compared to sequential prompting, this design pattern yields substantial speedups, making it practical for large-scale evaluation and deployment.

C.2.3 Data Model

The data model for domain-specific MCQA in `Agentics` is defined using `Pydantic`, which enforces structural constraints and type safety during both prompt construction and response decoding. This schema-guided approach ensures that the model's outputs conform to expected formats, reducing parsing errors and improving evaluation reliability.

The `FailureSensorIQ` class encapsulates the core components of each QA instance, including the question text, list of options, associated metadata (e.g., asset name, question type), and the model-generated answer. The nested `Answer` class captures the model's selected answer, a numerical confidence score, and a free-text explanation that provides insight into the model's reasoning process.

```python
class Answer(BaseModel):
    answer_letter: str = Field(description="The selected answer letter")
    confidence: float = Field(description="Confidence score")
    assessment: str = Field(description="Rationale for the answer")

class FailureSensorIQ(BaseModel):
    id: int = Field(description="Unique identifier for the question")
    question: str = Field(description="The question text")
    options: List[str] = Field(description="List of answer options")
    option_ids: List[str] = Field(description="List of option identifiers")
    asset_name: str = Field(description="Name of the industrial asset")
    relevancy: str = Field(description="Relevancy context or metadata")
    question_type: str = Field(description="Type of question")
    subject: str = Field(description="Subject or topic of the question")
    system_answer: Answer = Field(description="Model-generated answer object")
```

Listing 6: Data Model for Domain-Specific Multiple Choice QA

C.2.4 Main Algorithm

The workflow begins by instantiating an AG object with the `FailureSensorIQ` schema and a specified batch size. Each example from the dataset is parsed into a structured `FailureSensorIQ` instance and appended to the agent's internal state.

The core operation is the self transduction, which performs schema-guided inference over the specified input fields—such as question, options, and asset name—and generates structured outputs in the system answer field. The transduction is guided by a natural language instruction that defines the task: selecting the most plausible answer from a set of options, along with a confidence score and rationale.

```python
# Initialize the \texttt{Agentics} benchmark with the FailureSensorIQ schema and batch size
fsiq_benchmark = AG(atype=FailureSensorIQ, batch_size=40)

# Load dataset and populate agent states
dataset = load_dataset("cc4718/FailureSensorIQ")
for example in dataset:
    fsiq_benchmark.states.append(FailureSensorIQ(**example))

# Run self-transduction with structured input and output fields
fsiq_benchmark = await fsiq_benchmark.self_transduction(
    input_fields=[
        "question", "options", "option_ids",
        "asset_name", "relevancy", "question_type", "subject"
    ],
    output_fields=["system_answer"],
    instructions=(
        "Read the input questions, all possible answers, and background task information. "
        "This is a multiple choice test, where one of the options is true and the others are false. "
        "Select the answer with the highest likelihood of being correct, and return it along with "
        "a confidence score and a verbal assessment explaining your judgment."
    )
)
```

Listing 7: Pseudo Code for Domain-Specific Multiple Choice QA

C.3 Prompt Optimization

Prompt optimization is a critical component in leveraging large language models (LLMs) for complex tasks. The performance of LLMs is highly sensitive to variations in prompt structure, tone, and even the positioning of textual components. Minor changes—such as rephrasing imperative sentences or reordering blocks—can significantly impact the model's output.

Prompts typically follow a structured format, often divided into system and user sections. The system prompt provides general context, instructions, and expectations for the output, while the user prompt contains task-specific information. Common elements include task descriptions, constraints, few-shot examples, input/output format specifications, and guiding phrases like "Let's think step by step". These components are frequently organized using structured formats such as Markdown or JSON schemas.

C.3.1 General Framework for Prompt Optimization

Ramnath et al. (2025) and Li et al. (2025) have summarized existing prompt optimization techniques into a generic prompt optimization framework. Algorithm 1 presented in Ramnath et al. (2025), formalizes the process of prompt optimization as follows. Given a task model M_{task}, an initial prompt $\rho \in V$, the goal of an prompt optimization system M_{PO} is to obtain the best performing prompt-template ρ^{opt} under a metric $f \in F$ and eval-set D_{val} that maximizes expected performance:

$$\rho^{\text{opt}} := \arg\max_{\rho \in V} \mathbb{E}_{x \sim D_{\text{val}}} \left[f\left(M_{\text{task}}(\rho \oplus x) \right) \right].$$

Since the objective function is not tractable due to the combinatorial nature of discrete token-sequence search spaces, the optimization process typically follows a generate-select-evaluate cycle, akin to local search algorithms (Russell & Norvig, 2020). Most methods begin with a predefined prompt template that specifies the structure and content to be included. An initial prompt may be constructed manually or generated automatically using an LLM.

Given a specific task, the dataset is usually partitioned into training and validation sets. The training set is used to optimize the prompt, while the validation set is employed for tuning hyperparameters. Additionally, a held-out test set is used to evaluate final performance. Candidate prompts are generated, filtered based on performance metrics, and refined over successive iterations, with incremental improvements guided by feedback from model outputs.

Algorithm 1: Prompt Optimization Framework

$P_0 := \{\rho_1, \rho_2, \ldots, \rho_k\}$ ▷ Initial seed prompts
$D_{val} := \{(x_1, y_1)\}_{i=1}^{n}$ ▷ Validation set
$f_1, \ldots, f_m \in F$ ▷ Inference evaluation
for $t = 1, 2, \ldots, N$ **do** ▷ Iteration depth
    $G_t := M_{PO}(P, D_{val}, F)$ ▷ Generate prompt candidates with M_{PO}
    $P_t := Select(G_t, D_{val}, F)$ ▷ Filter and retain candidates
    if $f_{convergence} \leq \epsilon$ **then** ▷ Optionally check for early convergence
        exit
return $\arg\max_{\rho \in P_N} E_{x \sim D_{val}} [f(M_{task}(\rho \oplus x))]$

C.3.2 Parallelizing Prompt Optimization with Transduction Algebra

Within the `Agentics` framework, the prompt function defined in Definition 9 maps type information into structured information objects for transduction. Prompt optimization in this context is naturally expressed using transduction algebra. Since candidate prompts can be generated by logical transduction, we adopt meta-prompt-based optimization strategies similar to those proposed in (Yang et al., 2023; Ye et al., 2024; Opsahl-Ong et al., 2024). Our approach emphasizes two key aspects:

- The `Agentics` framework supports a declarative style of prompting, where prompts are constructed to encode rich contextual and type-level information rather than procedural instructions.
- The optimization process follows the generate-select-evaluate cycle described in Algorithm 1. Importantly, the transduction algebra enables this optimization loop to be expressed in a functional and parallelizable manner. Prompt candidates can be generated and evaluated independently, allowing for efficient execution. This abstraction not only improves scalability but also decouples the optimization logic from the underlying execution strategy.

In summary, the `Agentics` framework provides a principled and extensible foundation for prompt optimization. It integrates declarative prompt construction, transduction algebra, and parallel search strategies into a unified system that supports both expressiveness and scalability. Next, we present a functional design pattern for implementing prompt optimization using transduction algebra, abstracting away procedural details common in existing approaches. In Section 5.4, our experiments demonstrate that declarative context optimization improves performance and that parallelization yields substantial runtime gains.

C.3.3 Data Models

We begin by defining two data models using Pydantic: `OptimizationTask` and `GSM8K`. The `OptimizationTask` schema captures the components of a prompt template—such as role, goal, expected output, and imperative instructions—along with a score field for evaluation. The `GSM8K` schema represents the target task, including the question, ground-truth answer, model-generated reasoning, and correctness flag.

```python
class OptimizationTask(BaseModel):
    # a list of demo tasks used to guide prompt generation
    demos: list[Any] = Field(description="optimization demo tasks")
    # role description to be embedded in the prompt (e.g., 'You are a math tutor')
    role: str = Field(description="New role instruction")
    # goal statement describing what the prompt aims to achieve
    goal: str = Field(description="New goal instruction")
    # criteria for what constitutes a good or acceptable output
    expected_output: str = Field(description="New expected_output instruction")
    # imperative phrase (e.g., 'Let's think step by step')
    imperative: str = Field(description="New imperative")
    # evaluation score assigned to the prompt after testing on validation data
    score: int = Field(description="evaluation score")

class GSM8K(BaseModel):
    question: str = Field(description="a grade school math question.")
    answer: str = Field(description="the ground-truth answer")
    response_think: str = Field(description="the step by step reasoning")
    response_answer: str = Field(description="the final answer")
    # boolean flag indicating whether the response answer is correct
    correct: bool
```

Listing 8: Data Models for GSM8K Prompt Optimization

C.3.4 Meta Prompt and Optimized Prompts

The meta-prompt (`OPT_META_INSTRUCTION`) guides the generation of new prompt templates by describing the structure and expectations for the optimizer. It includes historical context from previous iterations and instructs the model to avoid redundancy. The user prompt template (`USER_PROMPT_TEMPLATE`) is instantiated with the optimized parameters and used to evaluate candidate prompts on the validation set.

The following shows the meta-prompt used for optimizing the prompt template for the GSM8K dataset.

```
OPT_META_INSTRUCTION = "Your proposed prompt template will be used in the following way.
* You are "role" -- this role must be suitable for solving the demo task.
* Your personal goal is: "goal" -- the goal achieves the outputs given inputs.
* This is the expected criteria for your final answer "expected_output" -- this constrains the output format.
* You can add a short imperative instruction "imperative" -- this comes after the input of the task.

[[Several demo tasks of input and outputs will be provided when you solve problem.]]

[[The previous optimized prompt templates with scores appear from the worst to the best.]]
{optimization_history}

* Given the previous optimization results, don't generate duplicate or similar prompt templates.
* Generate prompt template that achieves the best score, and succint and concise instructions.
"


USER_PROMPT_TEMPLATE = "
You are {role}.
Your personal goal is: {goal}.
This is the expected criteria for your final answer: {expected_output}.

solve the following task.
{question}

{imperative}
"
```

Listing 9: Meta Prompt and Template for GSM8K Prompt Optimization

In the following, we show the prompts returned by APO, including both system and user prompts. The system prompt consists of three components: role, goal, and expected_output. In the user prompt, an imperative statement appears after each question. A score of 89 was evaluated on the validation set, using 100 problems sampled from the training set, which is higher than the test score of 85.

```
{
"role": "Elite Mathematical Problem Solver",
```

```
"goal": "To swiftly and accurately determine the precise numerical solution by meticulously dissecting complex
    problem statements, identifying pivotal information, and applying a comprehensive array of advanced
    mathematical concepts, formulas, and logical reasoning to achieve an optimal solution, ensuring efficiency,
     precision, and reliability in all calculations, while providing accurate and relevant answers",
"expected_output": "A single, exact numeric value that directly addresses and solves the given mathematical
    problem, ensuring all calculations are correctly executed, based on the information provided, and presented
     in a clear, concise manner, with strict adherence to mathematical rules and consideration of all given
    data",
"imperative": "Thoroughly analyze the problem statement, extract key information, apply pertinent mathematical
    principles and logical reasoning to derive the accurate numerical answer, and provide the final answer in
    the required numeric format, while validating the solution through re-evaluation of calculations and
    verification of the accuracy of the obtained result, and ensuring precision, accuracy, and reliability in
    all steps of the calculation process.",
"score": 89
}
```

Listing 10: Optimized Prompt for GSM8K using Llama-3.3-70B Model

The following result is obtained from the Qwen3-8B model. A score of 98 was evaluated on the validation set, using 100 problems sampled from the training set. This score is higher than the test score of 92.

```
{
"role": "Math Problem Solver",
"goal": "Solve complex problems by decomposing them into sequential steps, applying arithmetic operations (
    percentages, fractions, averages), ensuring unit consistency, and presenting the final answer in a boxed
    format.",
"expected_output": "A single numeric value boxed (e.g., \\boxed{64}) representing the solution, derived through
     precise step-by-step calculations with attention to percentages, fractions, and averages.",
"imperative": "Analyze the problem statement, execute calculations step-by-step with focus on percentages,
    fractions, and unit conversions, then box the final numeric result.",
"score": 98
}
```

Listing 11: Optimized Prompt for GSM8K using Qwen3-8B Model

The following shows the meta prompt for optimizing the prompt template for MCQA dataset.

```
OPT_META_INSTRUCTION = "Your proposed prompt template will be used in the following way.
* You are "role" -- this role must be suitable for solving the demo task.
* Your personal goal is: "goal" -- the goal achieves the outputs given inputs.
* This is the expected criteria for your final answer "expected_output" -- this constrains the output format.
* Extract"task_context" from demo tasks to explain the problem context -- this comes before the input of the
    task.
* You can add a short imperative instruction "imperative" -- this comes after the input of the task.


[[Several demo tasks of input and outputs will be provided when you solve problem.]]

[[The previous optimized prompt templates with scores appear from the worst to the best.]]
{optimization_history}

* Given the previous optimization results, don't generate duplicate or similar prompt templates.
* Generate prompt template that achieves the best score, and succint and concise instructions.
"

# Template used to instantiate a user prompt from the optimized parameters.
# This is applied to each validation example.
USER_PROMPT_TEMPLATE = "
You are {role}.
Your personal goal is: {goal}.
This is the expected criteria for your final answer: {expected_output}.

This is the general task context.
{task_context}

solve the following task.
{question}
{options}
{option_ids}
{asset_name}
{relevancy}
{question_type}
{subject}

{imperative}
"
```

Listing 12: Meta Prompt and Template for MCQA Prompt Optimization

In the following, we show the system and user prompts returned by APO. The system prompt includes three components: role, goal, and expected_output. In the user prompt, task_context appears before each question, and an imperative statement follows each question. The test score was 54.

```
{
"role": "Industrial Asset Diagnostician",
"goal": "To accurately identify the most relevant sensors for detecting specific failure modes in various industrial assets and determine the least relevant failure events for abnormal sensor readings, optimizing asset performance and reducing downtime",
"expected_output": "A concise description of the most relevant sensor for monitoring a specific failure mode or the least relevant failure event that does not significantly contribute to detecting a particular failure mode in the given asset, including the option id or description",
"task_context": "The task involves analyzing the relationship between sensor readings, failure modes, and industrial assets such as steam turbines, aero gas turbines, electric generators, and compressors, to determine the most relevant sensors for monitoring specific failure modes and the least relevant failure events for abnormal sensor readings", "imperative": "Analyze the provided asset, sensor readings, and failure modes to identify the most relevant sensor for monitoring a specific failure mode or the least relevant failure event for an abnormal sensor reading, considering the relationship between failure modes, sensors, and assets",
"score": 54}
```

Listing 13: Optimized Prompt for MCQA using Llama-3.3-70B Model

C.3.5 Main Algorithm

The main optimization loop follows the generate-select-evaluate cycle described in Algorithm 1. It begins by preparing the training and validation sets using the `Agentics` (`AG[GSM8K]`). Demo tasks are extracted and transformed into `OptimizationTask` instances.

In each iteration, candidate prompt templates are generated via self-transduction using the meta-prompt. These templates are applied to the validation set using the user prompt format. The responses are evaluated using the grading function defined in `GSM8K.grade`, and scores are assigned. The best-performing prompts are retained using a filtering function (`keep_best_k`). This loop continues until convergence or a maximum number of iterations is reached. The use of transduction and asynchronous execution enables parallel evaluation of prompt candidates, improving scalability and runtime efficiency.

```
# load GSM8K training data into \texttt{Agentics} abstraction
trainset = AG.from_json("gsm8k_train.jsonl")
# truncate to a subset for training
trainset = trainset.truncate_states(num_trains)
# create demo tasks from training examples
demosets = create_optimization_demos(trainset, num_demos)
# convert demo tasks into OptimizationTask instances
optimization_tasks = OptimizationTask.create_optimization_tasks(demosets)
# prepare validation set from remaining examples
validationset = trainset.truncate_states(num_trains, num_trains + num_devs)

# initialize optimizer AG[OptimizationTask] with demo tasks
optimizer = AG.from_states(optimization_tasks, atype=OptimizationTask)

# set default parameters and prompt configuration
set_default_params(optimizer)
optimizer.prompt_template = "{{"demo tasks":{demos}}}"
optimizer.prompt_params = {"role": "Prompt optimizer.", "goal": "Propose diverse prompt templates that achieves high performance for the demo task given as input.", "backstory": "Understand the problem domain given the demo task example and propose what answer should be generated.", "expected_output": "the outputs are role, goal, and the expected output description, and imperative sentence for solving provided tasks."}

# initialize list to store optimized prompt tasks
optimized_tasks = []
for iter_ind in range(max_iter):
    # generate candidate prompt templates using meta-prompt and transduction
    # transduction from demos to prompt parameters
    optimizer.instructions = OPT_META_INSTRUCTION.format(
        optimization_history = get_history_string(optimized_tasks))
    optimizer = asyncio.run(optimizer.self_transduction(
        ["demos"],
        ["role", "goal", "expected_output", "imperative"]))

    # apply candidate prompts to validation set using user prompt format
    opt_eval = optimizer * validationset
    opt_eval.prompt_template = USER_PROMPT_TEMPLATE
    opt_eval = asyncio.run(opt_eval.self_transduction(
        ["role", "goal", "expected_output", "imperative", "question"],
        ["response_think", "response_answer"]))

    # evaluate responses using GSM8K grading function
    executed_tasks = opt_eval / validationset
    for ind, exectask in enumerate(executed_tasks):
        exectask = asyncio.run(exectask.amap(GSM8K.grade))
        setattr(optimizer[ind], "score", summary["score"])

    # retain top-performing prompts for next iteration
    optimized_tasks.extend(optimizer.states)
    optimized_tasks = keep_best_k(optimized_tasks)
```

Listing 14: Pseudo Code for GSM8K Prompt Optimization

C.4 Comparing with Other Frameworks

In this section, we show code examples to compare the programming models of Agentics, DSPy, and Langgraph.

- DSPy is a very popular framework and pioneered the programming model framework for LLMs. DSPy uses Pydantic models (Signatures) to define data types and offers a PyTorch-like programming experience through Modules and Adaptors. We can write a program with a computational graph-like design for the workflows inside the Module.

- LangGraph focuses on stateful agent orchestration using graph-based control flows. It models workflows as nodes and edges, where each node represents an agent or tool and edges define message-passing dependencies.

- Agentics offers a principled algebra for typed, composable workflows and fine-grained asynchronous execution. While DSPy supports asynchronous calls, its primary focus remains on declarative prompt optimization and compilation strategies. DSPy typically achieves asynchrony by compiling entire programs asynchronously or using an async alternative to forward. In contrast, Agentics enables functional-style asynchronous execution at the level of individual transformations through logical transduction and asynchronous map/reduce functions.

C.4.1 Overall Comparison

Instead of arguing that one framework is superior, we aim to highlight the different design intents and goals that Agentics, DSPy, and LangGraph were created to achieve. Each framework reflects a distinct philosophy.

- LangGraph and the LangChain ecosystem offer a mature codebase with many useful libraries for agent orchestration, tracing, and state management. In fact, Agentics leverages prompt functions provided by LangGraph.

- DSPy pioneered declarative programming for LLM workflows and introduced the mechanism for constructing prompts from Pydantic types (Signatures), which remains a core innovation.

- Agentics, by contrast, is designed for massive asynchronous LLM API calls that transduce strings formatted from one type to another. It directly manages workloads through the AG meta-class, which orchestrates asynchronous transductions at a fine-grained level. This differs from DSPy, which organizes logic imperatively in Modules, and from LangGraph, which requires defining states and graphs for transformations.

Due to this design, Agentics programs tend to be simpler and more concise, with most effort focused on specifying types and composing pipelines, as shown in multiple examples in the paper. Error handlings are also easier because the AG meta-class serves as a single control point for Pydantic transducers.

In DSPy, execution is distributed across multiple Modules, and in LangGraph, across graph nodes. When a desired program can be implemented using DSPy's predefined Modules, the overhead in programming is negligible. Similarly, LangGraph offers powerful features like state tracing, which is difficult in Agentics framework.

In enterprise structured data workflows, the primary requirement is often the ability to process large batches of transformations quickly and efficiently under strict schema constraints.

Next, we will present code examples comparing the composability and programming models of Agentics, DSPy, and LangGraph. The Python scripts were first written by hand and then added comments and unified the styles using the Gemini-2.5 model.

C.4.2 Text-to-SQL in Agentics

All three implementations handle data types in a similar way using Pydantic models, and we provide their asynchronous versions for fairness. In the Agentics example, the AG class loads the BIRD-dev dataset and performs a single logical transduction from input fields to the output SQL query. Internally, the AG class maintains a list of typed states, and the PydanticTransducer generates transductions by applying prompt instructions to each state. These transductions are executed asynchronously, enabling fine-grained parallelism and efficient error recovery.

```python
from pydantic import BaseModel, Field
from typing import Optional, List, Union
from agentics.core.agentics import Agentics as AG
import asyncio
import os
import json

# Pydantic Data Model (State Definition)
class TextToSQLItem(BaseModel):
    """
    Defines the structure for a single Text-to-SQL example.
    This acts as the state for the Agentics self-transduction.
    """
    question: str = Field(..., description="Natural language question for SQL generation.")
    evidence: List[str] = Field([], description="Hints/Notes for writing the SQL query.")
    ddl_schema: str = Field(..., description="DDL script for database tables.")
    generated_sql_query: Optional[str] = Field(None, description="The generated SQL query.")


async def generate_sql_query(
    dataset: AG[TextToSQLItem],
    input_fields: List[str]
) -> AG[TextToSQLItem]:
    """
    Uses Agentics self-transduction to translate the input fields (question, evidence, ddl_schema)
    into the desired output field (generated_sql_query).
    """
    output_fields = ["generated_sql_query"]

    # Execute the self-transduction (the LLM call)
    # The instructions set on the dataset will be used here.
    dataset = await dataset.self_transduction(
        input_fields,
        output_fields,
        instructions=dataset.instructions # Pass instructions explicitly for clarity
    )
    return dataset


if __name__ == "__main__":
    # Load the data into an Agentics Dataset object
    text2sql_dataset = AG.from_json(
        "bird_dev_extended.jsonl",
        atype=TextToSQLItem,
        jsonl=True
    )

    # Define Instructions
    # These instructions guide the LLM during the self-transduction step.
    text2sql_dataset.instructions = """
    You are an expert SQL writer. Your task is to translate the input natural language
    **question** to a valid SQLite query based on the provided **ddl_schema**.
    Use the **evidence** as hints. Your final output must *only* contain the SQL query.
    """

    # Run the Transduction
    input_fields_to_use = ["question", "evidence", "ddl_schema"]
    final_dataset = asyncio.run(
        generate_sql_query(
            text2sql_dataset,
            input_fields_to_use
        )
    )
```

Listing 15: Code for Text-to-SQL in Agentics

C.4.3 Text-to-SQL in DSPy

A DSPy program is structured around three core components. Signatures, which define data types, Adapters, which encapsulate LLM API calls, and Modules, which implement program logic. Similar to PyTorch, the forward method in a DSPy Module is where the workflow logic resides. This may include calling functions or chaining multiple LLM calls through other Modules. This design supports modularity and composability. However, asynchronous execution in

DSPy applies at the level of the forward method, meaning that fine-grained parallelism requires implementing each step as a separate Module.

```python
import dspy
import asyncio
from typing import List
from dspy.datasets import DataLoader
from dspy import Signature, InputField, OutputField, Predict, Module, settings
from dspy.adapters import ChatAdapter

# Placeholder for the missing schema retrieval function
def get_schema_description(db_id: str) -> str:
    """Mock function to simulate getting the DB schema for a given ID."""
    return f"""

# Placeholder for the missing SQL cleanup function
def cleanup_sql(sql_str: str) -> str:
    """Simple cleanup to remove common LLM markdown wraps."""
    sql_str = sql_str.strip()
    if sql_str.lower().startswith("```sql"):
        sql_str = sql_str[len("```sql"):].strip()
    if sql_str.endswith("```"):
        sql_str = sql_str[:-len("```")].strip()
    return sql_str.replace('\n', ' ')


# SIGNATURE
class TextToSQLBasicSignature(Signature):
    """You are an expert SQL writer and database analyst. Your task is to translate natural language QUESTION
       into a valid SQL query that can be executed against a given database schema."""
    DB_SCHEMA: str = InputField(desc="DB table definitions showing names, columns, and relations.")
    NOTE: str = InputField(desc="Hints for writing down the SQL query.")
    QUESTION: str = InputField(desc="Natural language question to convert to SQL.")
    SQL: str = OutputField(desc="Syntactically correct SQL with correct column names using suitable tables and
        rows. Your response should *only* contain the SQL query. Do not provide explanations or any other text.
        ")

# DSPy ADAPTER
# This adapter is primarily used to post-process the LLM's output (cleanup).
class SQLChatAdapter(ChatAdapter):
    def parse(self, signature, completion):
        # Use the base class parsing to get the output fields (e.g., {"SQL": "..."})
        fields = super().parse(signature, completion)

        # Apply the cleanup function to the SQL output field
        if "SQL" in fields:
            fields["SQL"] = cleanup_sql(fields["SQL"])
        return fields

# DSPy MODULE
class TextToSQLPrompter(Module):
    def __init__(self):
        super().__init__()
        self.program = Predict(TextToSQLBasicSignature)

    async def aforward(self, db_id: str, question: str, evidence: List[str]):
        """
        The native asynchronous forward pass for the module.
        Uses .acall() on the internal predictor.
        """
        table_descriptions = get_schema_description(db_id)
        evidence_str = "\n".join(evidence)

        # Use .acall() and await the result
        prediction = await self.program.acall(
            DB_SCHEMA=table_descriptions,
            NOTE=evidence_str,
            QUESTION=question
        )
        return prediction

async def run_batch_aforward(model: TextToSQLPrompter, dev_set: List[dspy.Example]):
    """Runs a batch of DSPy predictions concurrently using the module's aforward."""
    print(f"Starting concurrent execution for {len(dev_set)} questions...")

    tasks = []

    for example in dev_set:
        # Get the input dictionary for the module
        inputs = example.inputs()

        # Create an asynchronous task using the module's aforward method
        async def aforward(model, inputs, original_q):
            # Call the custom module's aforward
            prediction = await model.aforward(**inputs)
            return {
                "question": original_q,
                "generated_sql": prediction.SQL,
                "error": None
            }
        tasks.append(aforward(model, inputs, example.question))
```

```
    # Run all tasks concurrently and wait for all to complete
    results = await asyncio.gather(*tasks)
    return results


if __name__ == "__main__":
    # Example setup using OpenAI
    lm = dspy.OpenAI(model="llm model id",)
    dspy.settings.configure(lm=lm, adapter=SQLChatAdapter())

    # Load Data
    data_loader = DataLoader()
    dev_set = data_loader.from_json(file_path="bird_dev_extended.jsonl",
                                    fields=["db_id", "question", "evidence"])
    # Instantiate the Module and Evaluate
    text2sql = TextToSQLPrompter()

    # Run the full batch asynchronously using aforward
    final_results = asyncio.run(run_batch_aforward(text2sql, dev_set))
```

Listing 16: Code for Text-to-SQL in DSPy

C.4.4 Text-to-SQL in Langgraph

The LangGraph example demonstrates how an application is defined as a state graph, starting with a Graph State and adding nodes for specific operations, such as generating SQL from a question using an LLM. Compared to DSPy, LangGraph expresses program logic declaratively through graph nodes and edges, which is intuitive for multi-step workflows. Asynchronous execution is supported by adding asynchronous nodes. However, parallelism is controlled at the graph level, not at the level of individual nodes. In practice, the run-batch function loops over problems and invokes the application per state, meaning fine-grained parallelism is limited. By contrast, Agentics enables pipeline-level control where each transformation (analogous to a node) can execute asynchronously and independently, offering greater flexibility and scalability for structured workflows.

```
import operator
import asyncio
from typing import TypedDict, Annotated, List, Dict

# LangGraph and LangChain Imports
from langgraph.graph import StateGraph, END, START
from langchain_core.prompts import ChatPromptTemplate
from langchain_openai import ChatOpenAI

# --- Graph State Definition ---
class TextToSQLState(TypedDict):
    """
    Represents the state of our graph.
    """
    db_schema: str
    question: str
    sql_query: Annotated[str, operator.replace]
    error: Annotated[str, operator.replace]


# --- Prompt Template ---
SQL_PROMPT_TEMPLATE = """You are an expert SQL writer and database analyst. Your task is to translate natural
     language QUESTION into a valid SQL query that can be executed against a given database schema.

DB_SCHEMA:
{db_schema}

QUESTION:
{question}

NOTE:
You must only output the syntactically correct SQL query. Do not provide explanations or any other text.
"""
sql_prompt = ChatPromptTemplate.from_template(SQL_PROMPT_TEMPLATE)


# --- Asynchronous Node for SQL Generation (The core change) ---
async def generate_sql_async(state: TextToSQLState) -> TextToSQLState:
    """
    Generates a SQL query from the question and schema using an LLM asynchronously.

    NOTE: The function is now 'async def' and uses 'await llm.ainvoke'.
    """
    # Initialize your chosen LLM (LangChain clients are often thread-safe)
    llm = ChatOpenAI(model="gpt-4o", temperature=0)

    # Format the prompt
    formatted_prompt = sql_prompt.format(
```

```
        db_schema=state["db_schema"],
        question=state["question"]
    )

    # Run the LLM asynchronously
    try:
        response = await llm.ainvoke(formatted_prompt)
        sql_query = response.content.strip()

        # Simple cleanup (removing markdown fences)
        if sql_query.lower().startswith("```sql"):
            sql_query = sql_query[len("```sql"):].strip()
        if sql_query.endswith("```"):
            sql_query = sql_query[:-len("```")].strip()

        # print(f"Generated SQL for '{state['question'][:30]}...': {sql_query}") # Uncomment for debugging

        # Update the state with the generated SQL
        return {"sql_query": sql_query, "error": ""}

    except Exception as e:
        print(f"Error during SQL generation for: {state['question'][:30]}...: {e}")
        # Update state with the error
        return {"sql_query": "", "error": f"Generation failed: {e}"}


# --- Graph Construction ---
# Build the Graph
workflow = StateGraph(TextToSQLState)

# Add the single Text-to-SQL node (using the async function)
workflow.add_node("sql_generator", generate_sql_async)

# Define the single, linear path
workflow.add_edge(START, "sql_generator")
workflow.add_edge("sql_generator", END)

# Compile the graph
app = workflow.compile()


# --- Asynchronous Batch Runner ---
async def run_batch(problems, app):
    """
    Runs the LangGraph app concurrently for a list of questions.
    """
    tasks = []

    for problem in problems:
        # Prepare the initial state for the current question
        initial_state = {
            "db_schema": problem["db_schema"],
            "question": problem["question"],
            "sql_query": "",
            "error": "",
        }
        # Create an asynchronous task for each question using app.ainvoke()
        tasks.append(app.ainvoke(initial_state))

    # Run all tasks concurrently and wait for them to complete
    concurrent_results = await asyncio.gather(*tasks)

    # Process the results
    final_results = []
    for state in concurrent_results:
        final_results.append({
            "question": state['question'],
            "sql_query": state['sql_query'],
            "error": state['error']
        })

    print("Concurrent execution finished.")
    return final_results


if __name__ == "__main__":
    # The asyncio.run() function is used to execute the top-level async function
    # assume that we loaded list of dictionaries storeing text-to-sql problems.
    final_results = asyncio.run(run_batch(problems, app))
```

Listing 17: Code for Text-to-SQL in Langgraph